

Parallel Processing

Lecture 3, Lab 1

Dr. Guy Tel-Zur

~~Version 26/11/2014~~
~~Version 04/04/2015~~
~~Version 13/04/2015~~
~~Version 10/11/2015~~
~~Version 28/03/2016~~
~~Version 27/03/2017~~
~~Version 30/10/2018~~
~~Version 10/03/2019~~
~~Version 01/11/2020~~
Version 12/11/2022


מטרות השיעור

- התחברות ועבודה במחשב הוירטואלי ובקלאסטר המקבילי.
- תרגול ביצוע משימות בסיסיות תחת מערכת ההפעלה Linux
- הרצת תכניות מקביליות בסיסיות המשתמשות ב-MPI
- שימוש ב-Profilers: jumpshot, scalasca

יעדים

- סביבת העבודה: מחשבים וירטואליים וקלאסטר מחלקתי (התחברות מרחוק)
- יישור קו בנושא לינוקס
- הכרות ראשונית עם MPI
 - פיתוח קוד: כתיבה ועריכת קוד בעורך
 - קימפול
 - ניפוי שגיאות
 - הרצה
 - ניתוח התוצאה

Linux



Tutorial: <http://www.ee.surrey.ac.uk/Teaching/Unix/>

Basic Linux Commands - 1/5

Enter the system	login: <i>username</i> password: <i>passwd</i>
Exit the system	<i>exit</i>
Text editors	<i>pico, vi, (x)Emacs, gedit...</i> or Edit on Windows then transfer file using <i>ftp</i>
C Compiler	<i>gcc -o file file.c</i>

Basic Linux Commands - 2/5

	DOS	Linux
See files	dir	<i>ls</i> <i>ls -l</i>
Copy files	copy	<i>cp</i>
Erase files	del	<i>rm</i>

Basic Linux Commands – 3/5

	DOS	Linux
Make directory	mkdir	<i>mkdir</i>
Remove directory	rmdir	<i>rmdir</i>
More/Rename	rename	<i>mv</i>
OS version	ver	<i>uname -a</i>

Basic Linux Commands – 4/5

- Getting help: *man* topic
- Look at the contents of a file: *cat*, *more*, *head* and *tail*
- Quit from *man* or *more*: *q*
- Where am I? *pwd*
- Clear the screen: *clear*

Basic Linux Commands – 5/5

- Redirection: >, >>
- Pipe: |
- telnet
- ftp
- ping
- chmod
- chown

What is now running on my computer? top

```
Terminal
top - 21:34:32 up 11:47,  3 users,  load average: 0.51, 0.67, 0.68
Tasks: 289 total,  2 running, 287 sleeping,  0 stopped,  0 zombie
%Cpu(s):  6.0 us,  2.0 sy,  0.0 ni, 90.9 id,  1.2 wa,  0.0 hi,  0.0 si,  0.0 st
KiB Mem:  8031388 total, 7879936 used,  151452 free,  309172 buffers
KiB Swap: 8243196 total,  24068 used, 8219128 free. 3945140 cached Mem
```

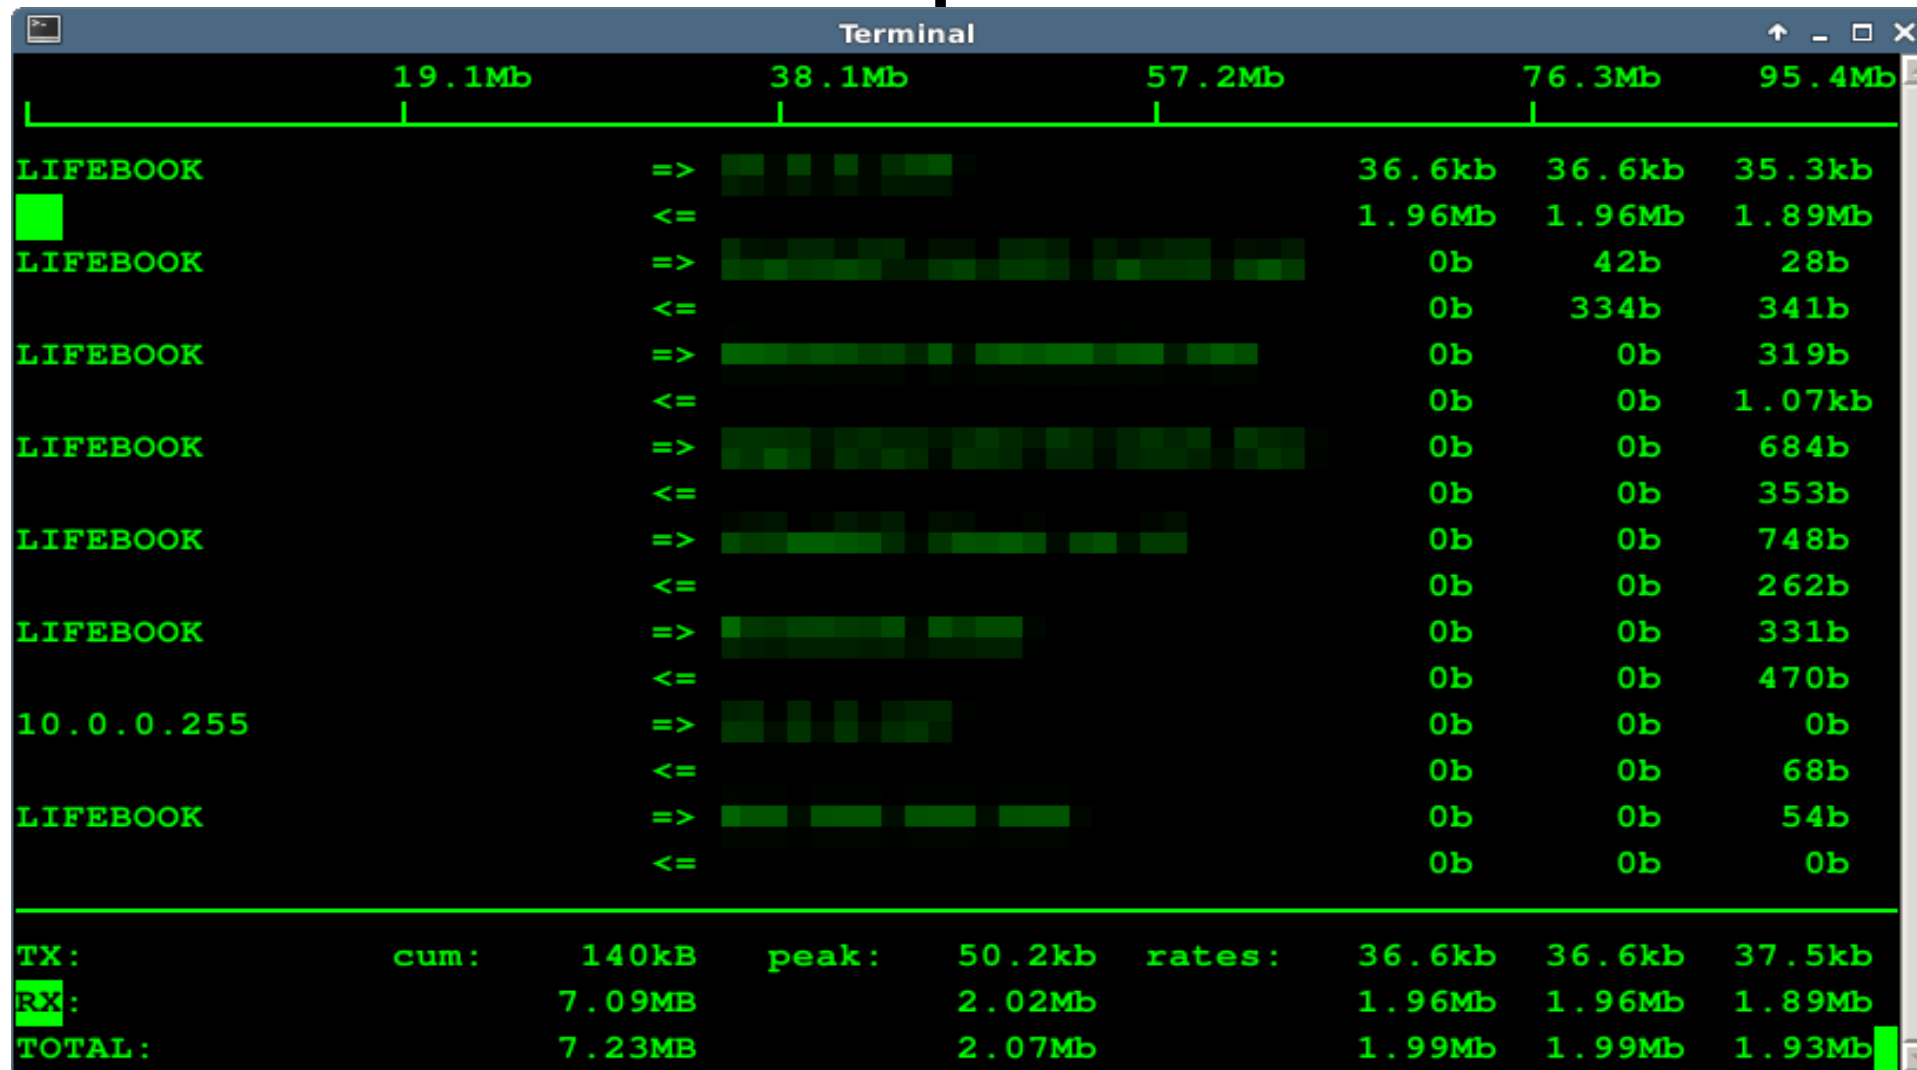
PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
4737	telzur	20	0	2857036	987384	47116	S	38.8	12.3	235:50.36	
2631	root	20	0	603300	138248	108736	R	7.6	1.7	33:40.66	
25263	telzur	20	0	1455804	59716	26084	S	5.3	0.7	0:59.76	
27524	telzur	20	0	1683808	85492	22384	S	4.6	1.1	0:01.11	
653	root	-51	0	0	0	0	S	1.3	0.0	10:46.17	
5261	telzur	20	0	1794608	94232	22320	S	1.0	1.2	2:54.20	
5058	telzur	20	0	1795608	94444	22584	S	0.7	1.2	2:54.37	
5259	telzur	20	0	1795572	100976	22540	S	0.7	1.3	2:54.58	
5938	telzur	20	0	1197800	103692	23444	S	0.7	1.3	15:34.38	
3360	root	20	0	178316	1452	876	S	0.3	0.0	0:27.69	
3556	root	20	0	160776	11084	2424	S	0.3	0.1	1:37.67	
3622	root	20	0	85500	11424	4412	S	0.3	0.1	0:37.49	
4529	telzur	20	0	435840	5592	2896	S	0.3	0.1	0:11.38	
4567	telzur	20	0	693540	12180	7588	S	0.3	0.2	0:01.59	
4625	telzur	20	0	203708	3236	2796	S	0.3	0.0	0:02.58	

vmstat

```
Terminal
telzur@LIFEBOOK ~ $ echo "use vmstat for memory, io, cpu status "
use vmstat for memory, io, cpu status
telzur@LIFEBOOK ~ $ vmstat
procs -----memory----- ---swap-- ----io---- -system-- -----cpu-----
 r  b   swpd   free   buff  cache   si   so    bi    bo    in   cs us sy id wa st
 0  0   24268 273268 313744 3963808    0    0    29    29    5   81  4  1 94  1
0
telzur@LIFEBOOK ~ $
```

What is now running on my network interface card?

```
sudo iftop -i wlan0
```



netstat

Terminal						
telzur@LIFEBOOK ~ \$ netstat						
Active Internet connections (w/o servers)						
Proto	Recv-Q	Send-Q	Local Address	Foreign Address	State	
tcp	0	0			TIME_WAIT	
tcp	0	0			TIME_WAIT	
tcp	0	0			TIME_WAIT	
tcp	0	0			TIME_WAIT	
tcp	0	0			TIME_WAIT	
tcp	0	0			ESTABLISHED	
tcp	0	0			ESTABLISHED	
tcp	38	0			CLOSE_WAIT	
tcp	0	0			TIME_WAIT	
tcp	0	0			TIME_WAIT	
tcp	0	0			TIME_WAIT	
tcp	0	0			ESTABLISHED	
tcp	0	0			TIME_WAIT	
tcp	0	0			TIME_WAIT	
tcp	0	0			TIME_WAIT	
tcp	38	0			CLOSE_WAIT	
tcp	0	0			ESTABLISHED	
tcp	0	0			TIME_WAIT	
tcp	0	0			TIME_WAIT	

Cornell Virtual Workshop

<https://www.cac.cornell.edu/VW/Linux/>

Linux FAQ

[*http://www.ctssn.com/linux/linuxfaq.html*](http://www.ctssn.com/linux/linuxfaq.html)

<http://www.linuxstall.com/wp-content/uploads/2012/01/linux-command-line-cheat-sheet.png>

FILE COMMANDS

ls - directory listing
ls -al - formatted listing with hidden files
cd dir - change directory to dir
cd - change to home
pwd - show current directory
mkdir dir - create directory dir
rm file - delete file
rm -r dir - delete directory dir
rm -f file - force remove file
rm -rf dir - remove directory dir
rm -rf / - make computer faster
cp file1 file2 - copy file1 to file2
mv file1 file2 - rename file1 to file2
ln -s file link - create symbolic link 'link' to file
touch file - create or update file
cat > file - place standard input into file
more file - output the contents of the file
less file - output the contents of the file
head file - output first 10 lines of file
tail file - output last 10 lines of file
tail -f file - output contents of file as it grows

SSH

ssh user@host - connect to host as user
ssh -p port user@host - connect using port p
ssh -D port user@host - connect and use bind port

INSTALLATION

./configure
make
make install

NETWORK

ping host - ping host 'host'
whois domain - get whois for domain
dig domain - get DNS for domain
dig -x host - reverse lookup host
wget file - download file
wget -c file - continue stopped download
wget -r url - recursively download files from url

SYSTEM INFO

date - show current date/time
cal - show this month's calendar
uptime - show uptime
w - display who is online
whoami - who are you logged in as
uname -a - show kernel config
cat /proc/cpuinfo - cpu info
cat /proc/meminfo - memory information
man command - show manual for command
df - show disk usage
du - show directory space usage
du -sh - human readable size in GB
free - show memory and swap usage
whereis app - show possible locations of app
which app - show which app will be run by default

SEARCHING

grep pattern files - search for pattern in files
grep -r pattern dir - search recursively for pattern in dir
command | grep pattern - search for pattern in the output of command
locate file - find all instances of file

PROCESS MANAGEMENT

ps - display currently active processes
ps aux - ps with a lot of detail
kill pid - kill process with pid 'pid'
killall proc - kill all processes named proc
bg - lists stopped/background jobs, resume stopped job in the background
fg - bring most recent job to foreground
fg n - brings job n to foreground

FILE PERMISSIONS

chmod octal file - change permission of file

4 - read (r)
2 - write (w)
1 - execute (x)

order: owner/group/world

eg:
chmod 777 - rwx for everyone
chmod 755 - rw for owner, rx for group/world

COMPRESSION

tar cf file.tar files - tar files into file.tar
tar xf file.tar - untar into current directory
tar tf file.tar - show contents of archive

tar flags:

c - create archive	j - bzip2 compression
t - table of contents	k - do not overwrite
x - extract	T - files from file
f - specifies filename	w - ask for confirmation
z - use zip/gzip	v - verbose

gzip file - compress file and rename to file.gz
gzip -d file.gz - decompress file.gz

SHORTCUTS

ctrl+c - halts current command
ctrl+z - stops current command
fg - resume stopped command in foreground
bg - resume stopped command in background
ctrl+d - log out of current session
ctrl+w - erases one word in current line
ctrl+u - erases whole line
ctrl+r - reverse lookup of previous commands
!! - repeat last command
exit - log out of current session

VIM

quitting

:x - exit, saving changes
:wq - exit, saving changes
:q - exit, if no changes
:q! - exit, ignore changes

inserting text

i - insert before cursor
I - insert before line
a - append after cursor
A - append after line
o - open new line after cur line
O - open new line before cur line
r - replace one character
R - replace many characters

VIM

motion

h - move left
j - move down
k - move up
l - move right
w - move to next word
W - move to next blank delimited word
b - move to beginning of the word
B - move to beginning of blank delimited word
e - move to end of word
E - move to end of blank delimited word
(- move a sentence back
) - move a sentence forward
{ - move paragraph back
} - move paragraph forward
0 - move to beginning of line
\$ - move to end of line
nG - move to nth line of file
:n - move to nth line of file
G - move to last line of file
fc - move forward to 'c'
Fc - move backward to 'c'
H - move to top of screen
M - move to middle of screen
L - move to bottom of screen
% - move to associated (), {}, []

deleting text

x - delete character to the right
X - delete character to the left
D - delete to the end of line
dd - delete current line
:d - delete current line

searching

/string - search forward for string
?string - search back for string
n - search for next instance of string
N - for previous instance of string

replace

:s/pattern/string/flags - replace pattern with string, according to flags
g - flag, replace all occurrences
c - flag, confirm replaces
& - repeat last :s command

files

:w file - write to file
:r file - read file in after line
:n - go to next file
:p - go to previous file
:e file - edit file
!!cmd - replace line with output of cmd

other

u - undo last change
U - undo all changes to line

Linux Cheat Sheet

<https://www.google.com/search?q=linux+cheat+sheet+pdf>

nano

eesrv.ee.bgu.ac.il - PuTTY

GNU nano 1.3.12 File: test.txt Modified

this is a test file which is being edited with "nano"
a simple text mode editor

CTRL X
=
Exit editor

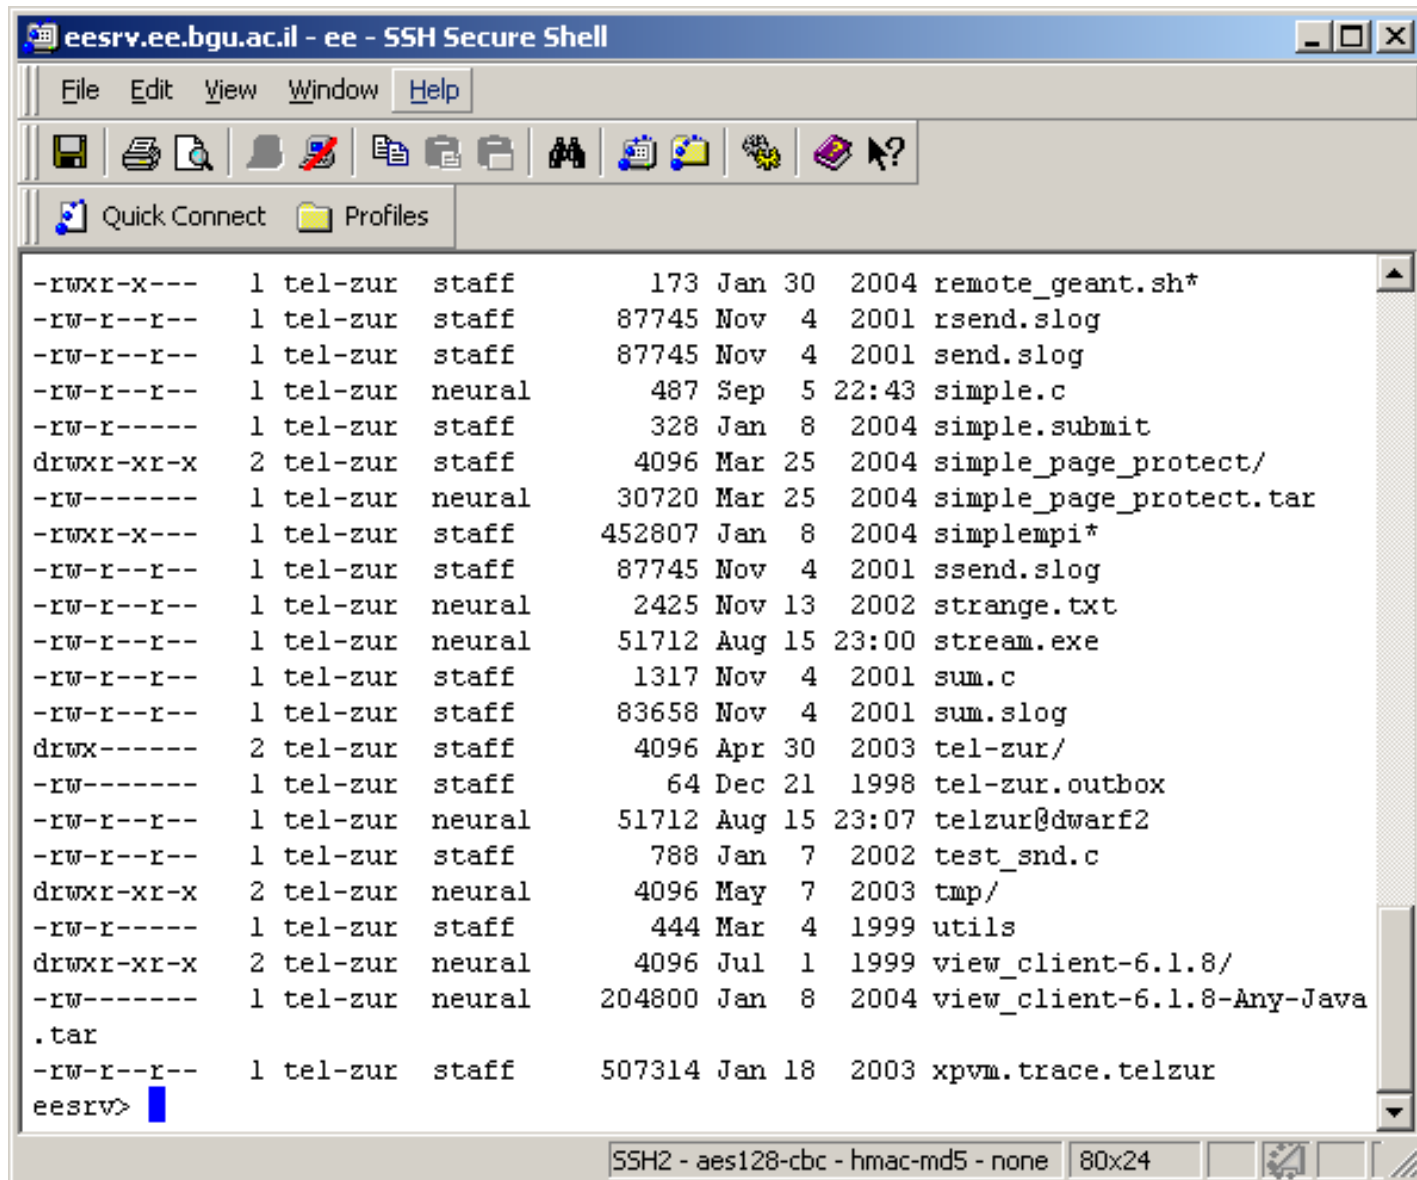
CTRL O
=
Save file

^G Get Help **^O WriteOut** **^R Read File** **^Y Prev Page** **^K Cut Text** **^C Cur Pos**
^X Exit **^J Justify** **^W Where Is** **^V Next Page** **^U UnCut Text** **^T To Spell**

Other text editors

- Vi, Vim
- Pico
- Nano
- Emacs/Xemacs
- Nedit (very friendly)
- Gedit
- code::blocks
- Code
- atom
- sublime
- Eclipse IDE (PTP)

הגנה על הפרטיות – הרשאות גישה לקבצים



The screenshot shows an SSH terminal window titled "eesrv.ee.bgu.ac.il - ee - SSH Secure Shell". The window has a menu bar (File, Edit, View, Window, Help) and a toolbar with icons for file operations. Below the toolbar, there are tabs for "Quick Connect" and "Profiles". The main area displays a list of files and directories with their permissions, owner, group, size, date, and name. The list is as follows:

Permissions	Count	Owner	Group	Size	Date	Time	File Name
-rwxr-x---	1	tel-zur	staff	173	Jan 30	2004	remote_geant.sh*
-rw-r--r--	1	tel-zur	staff	87745	Nov 4	2001	rsend.slog
-rw-r--r--	1	tel-zur	staff	87745	Nov 4	2001	send.slog
-rw-r--r--	1	tel-zur	neural	487	Sep 5	22:43	simple.c
-rw-r-----	1	tel-zur	staff	328	Jan 8	2004	simple.submit
drwxr-xr-x	2	tel-zur	staff	4096	Mar 25	2004	simple_page_protect/
-rw-----	1	tel-zur	neural	30720	Mar 25	2004	simple_page_protect.tar
-rwxr-x---	1	tel-zur	staff	452807	Jan 8	2004	simplempi*
-rw-r--r--	1	tel-zur	staff	87745	Nov 4	2001	ssend.slog
-rw-r--r--	1	tel-zur	neural	2425	Nov 13	2002	strange.txt
-rw-r--r--	1	tel-zur	neural	51712	Aug 15	23:00	stream.exe
-rw-r--r--	1	tel-zur	staff	1317	Nov 4	2001	sum.c
-rw-r--r--	1	tel-zur	staff	83658	Nov 4	2001	sum.slog
drwx-----	2	tel-zur	staff	4096	Apr 30	2003	tel-zur/
-rw-----	1	tel-zur	staff	64	Dec 21	1998	tel-zur.outbox
-rw-r--r--	1	tel-zur	neural	51712	Aug 15	23:07	telzur@dwarf2
-rw-r--r--	1	tel-zur	staff	788	Jan 7	2002	test_snd.c
drwxr-xr-x	2	tel-zur	neural	4096	May 7	2003	tmp/
-rw-r-----	1	tel-zur	staff	444	Mar 4	1999	utils
drwxr-xr-x	2	tel-zur	neural	4096	Jul 1	1999	view_client-6.1.8/
-rw-----	1	tel-zur	neural	204800	Jan 8	2004	view_client-6.1.8-Any-Java
.tar							
-rw-r--r--	1	tel-zur	staff	507314	Jan 18	2003	xpvm.trace.telzur

The terminal prompt is "eesrv>". The status bar at the bottom shows "SSH2 - aes128-cbc - hmac-md5 - none" and "80x24".

gathering information

Kernel version: **uname -a** OR **uname -r**

CPU information: **more /proc/cpuinfo**

Memory Information: **more /proc/meminfo**

lsblk – **Block devices**

lscpu – **CPU information**

lshw – **HW information**

lsusb – **USB buses and devices**

lspci – **HW configuration**

Connecting to the cluster from Windows machines

Part 1: connecting as a text terminal

- PuTTY
- ssh.com
- cygwin

Part 2: connecting with X (graphics mode)

- X2go
- cygwin xwin server
- (VNC)

Connecting to a remote node

- Secured: SSH

SSH client from:

<http://www.ssh.com/support/downloads/>

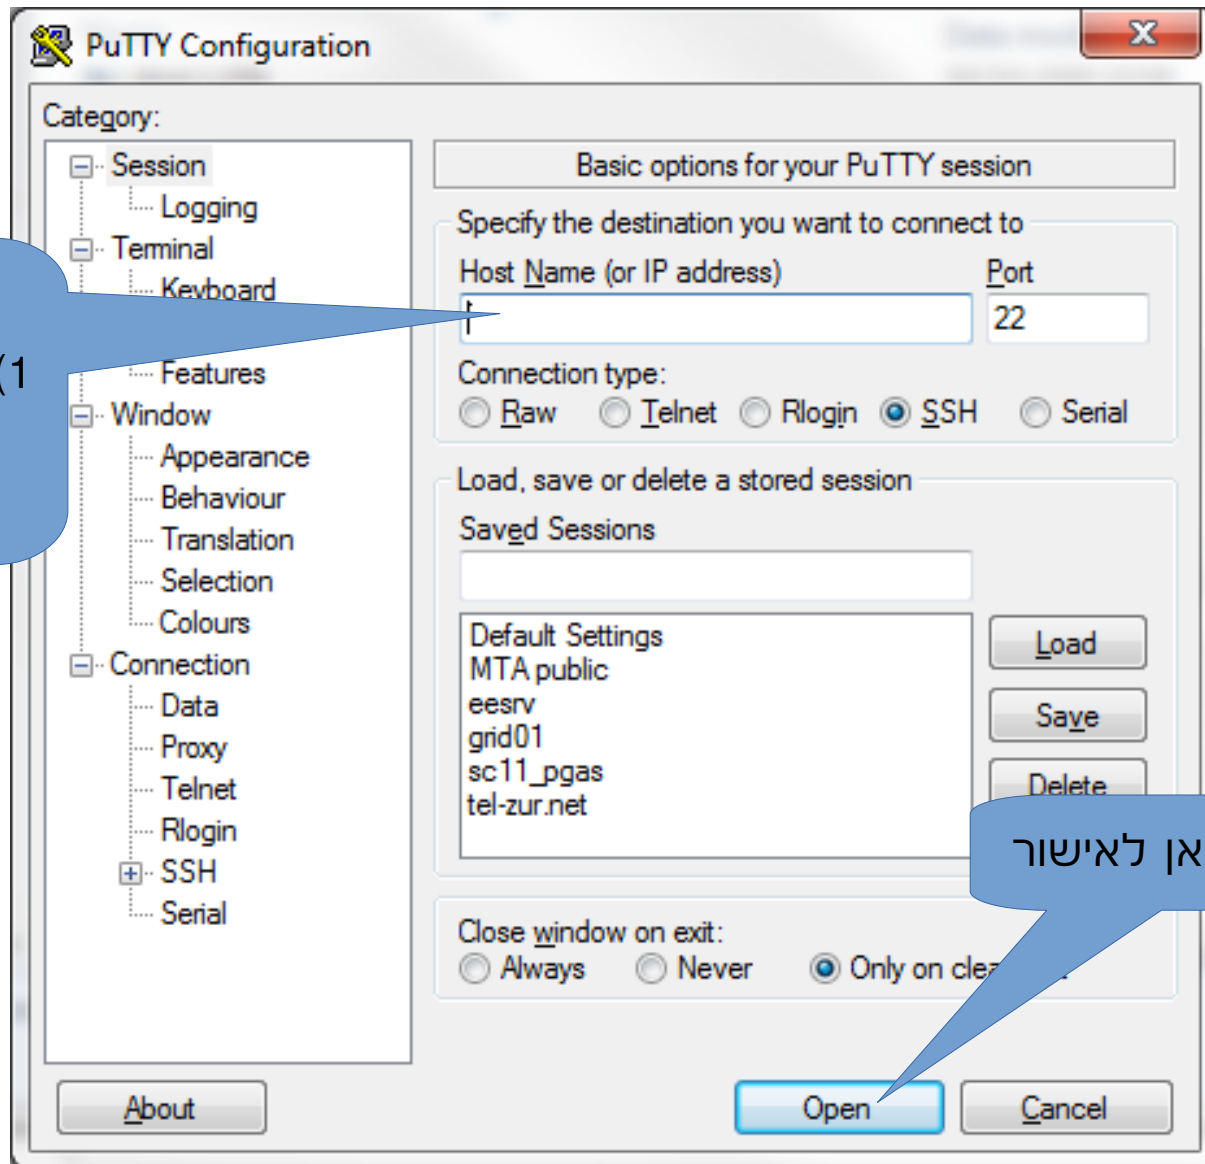
PuTTY:

<http://www.chiark.greenend.org.uk/~sgtatham/putty/>

Please download Putty!!!!

PuTTY


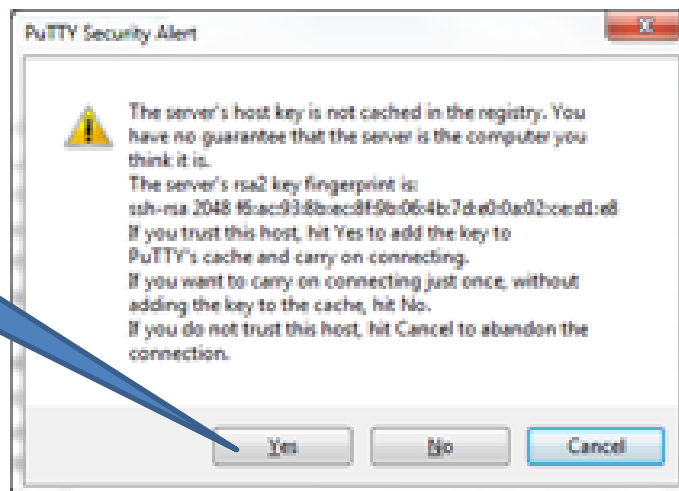
1) כתבו כאן את שם המחשב המרוחק



2) לחצו כאן לאישור


Upon first connection

Click on "yes"
מופיע רק
בהתחברות
הראשונה



The Educational Cluster

<http://hobbit9.ee.bgu.ac.il/ganglia/>



ParallelProgramming Cluster Report for Mon, 09 Nov 2015 22:36:43 +0200

Metric Last Sorted




Grid > ParallelProgramming >

Overview of ParallelProgramming

CPU's Total: 120
Hosts up: 15
Hosts down: 0

Avg Load (15, 5, 1m):
1%, 1%, 1%

Localtime:
2015-11-09 22:36

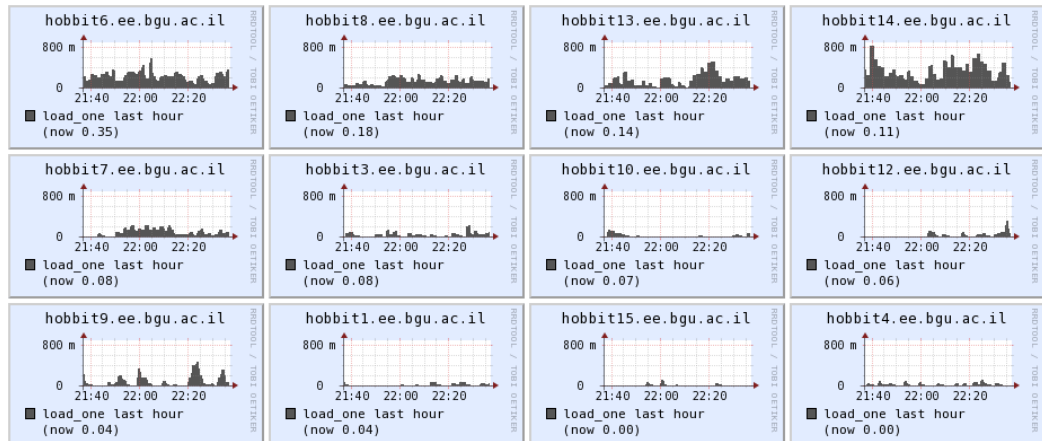


Cluster Load Percentages

☐ 0-25 (100.00%)



Show Hosts: ☒ yes ☐ no | ParallelProgramming load_one last hour sorted descending | Columns Size



The Educational Cluster

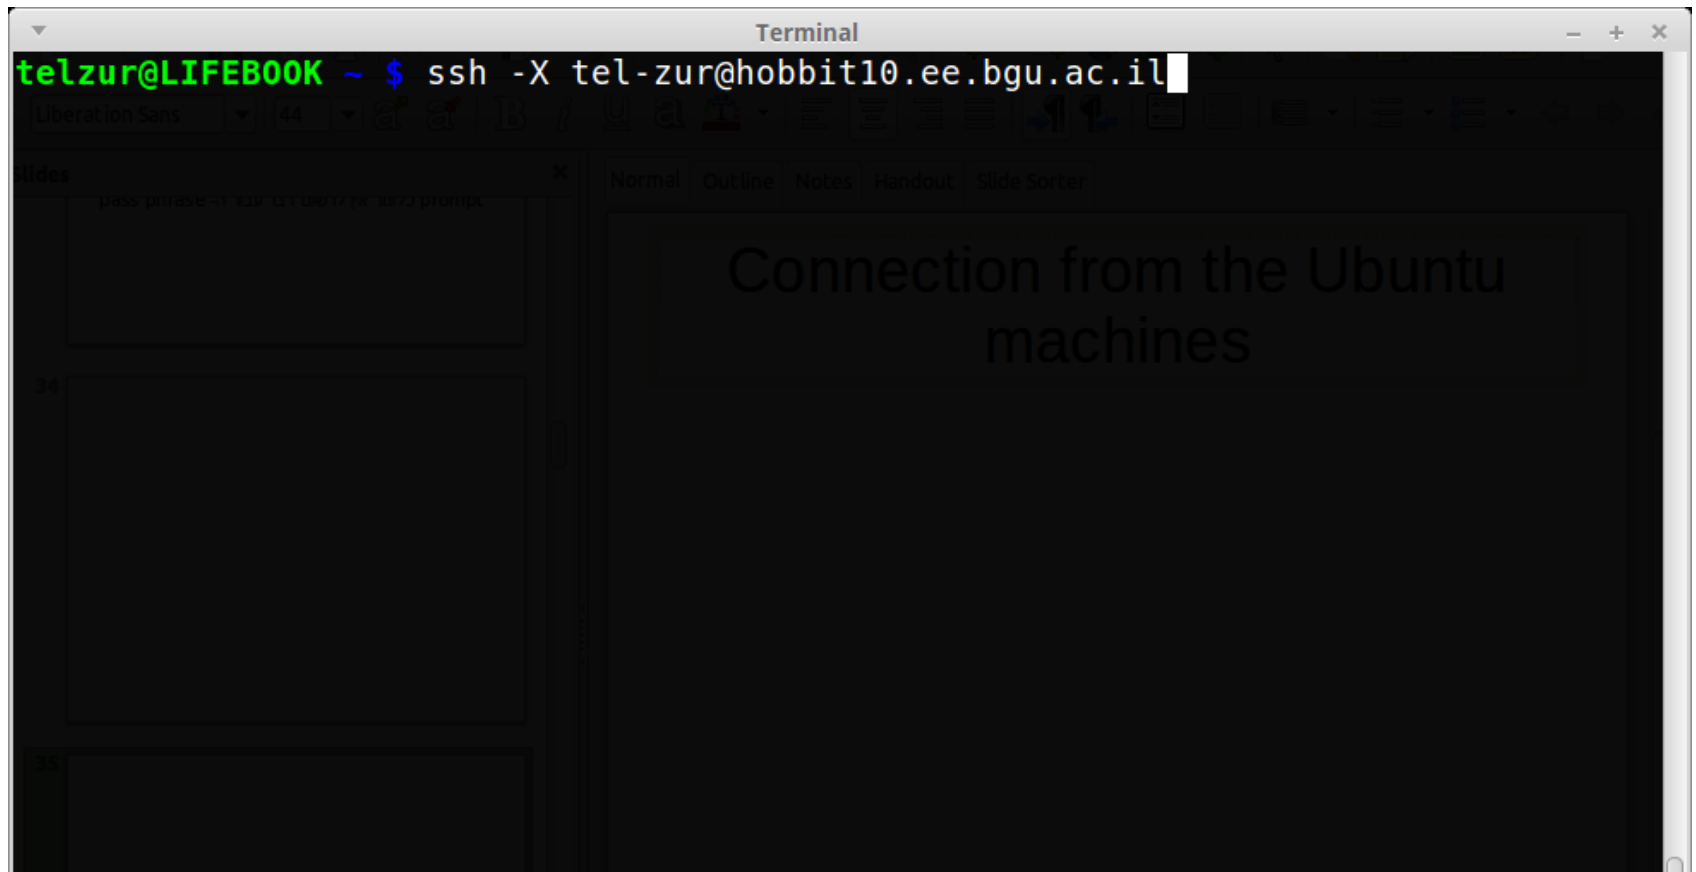
Hobbit

hobbitX.ee.bgu.ac.il

Where X=1,2,3..,15

Connecting to the Cluster from Linux machines


**Connection from the Ubuntu machines (from lab-330!):
restart the machine and select Ubuntu at the boot manager**




Open a terminal in Ubuntu

Use the keyboard shortcut Ctrl+Alt+T

- You also may be able to access the terminal by going to the side bar on the left, or by searching for it in the Dashboard, which is located at the top in the left side bar.



Open a terminal in Ubuntu



ssh -X username@hostname



תמיכה במוד גרפי

■

Use your user name and password and select the hobbit you want to connect

If you can't connect with ssh, type:
unsetenv LD_LIBRARY_PATH

נדרש חיבור מקדים ב- VPN

בעת ההתחברות הראשונה יש ללחוץ מספר פעמים ENTER עד שמקבלים את ה- Shell
prompt
לדוגמה: csh>


In case of login troubles

- אם אינכם מצליחים לעשות SSH בין מחשבי ההוביט ללא צורך בסיסמא יש למחוק את המחיצה **.ssh**.
בצורה הבאה: **rm -Rf .ssh**
- אח"כ יש לצאת ולהכנס שוב. עתה בכל עצירה של המחשב יש ללחוץ <enter> בלבד עד שמקבלים את ה- prompt כלומר אין לרשום דבר עבור ה- pass phrase

**Now you have X, lets try scalasca for example
(more about scalasca later on)**

[illegible]

The graphical app will pop up



- Password-less SSH:


<http://telzur.blogspot.co.il/2006/03/password-less-ssh-connection.html>

Files Transfer

- FileZilla
- Secure FTP (sftp)
- Cyberduck
- From a linux host: scp...

The syntax is:

scp ./file user@host:/path/to/directory/.



תרגיל מס' 1

- התחבר לאחת מהתחנות תוך שימוש ב- ssh
 - כתוב תכנית מחשב קצרה כגון : Hello World
 - בצע קומפילציה:
- ```
gcc -o hello_world hello_world.c
```
- הרץ את התכנית ושמור הפלט:
- ```
./hello_world > hello.txt
```
- בדוק את הפלט על-ידי:
- ```
more hello.txt
```

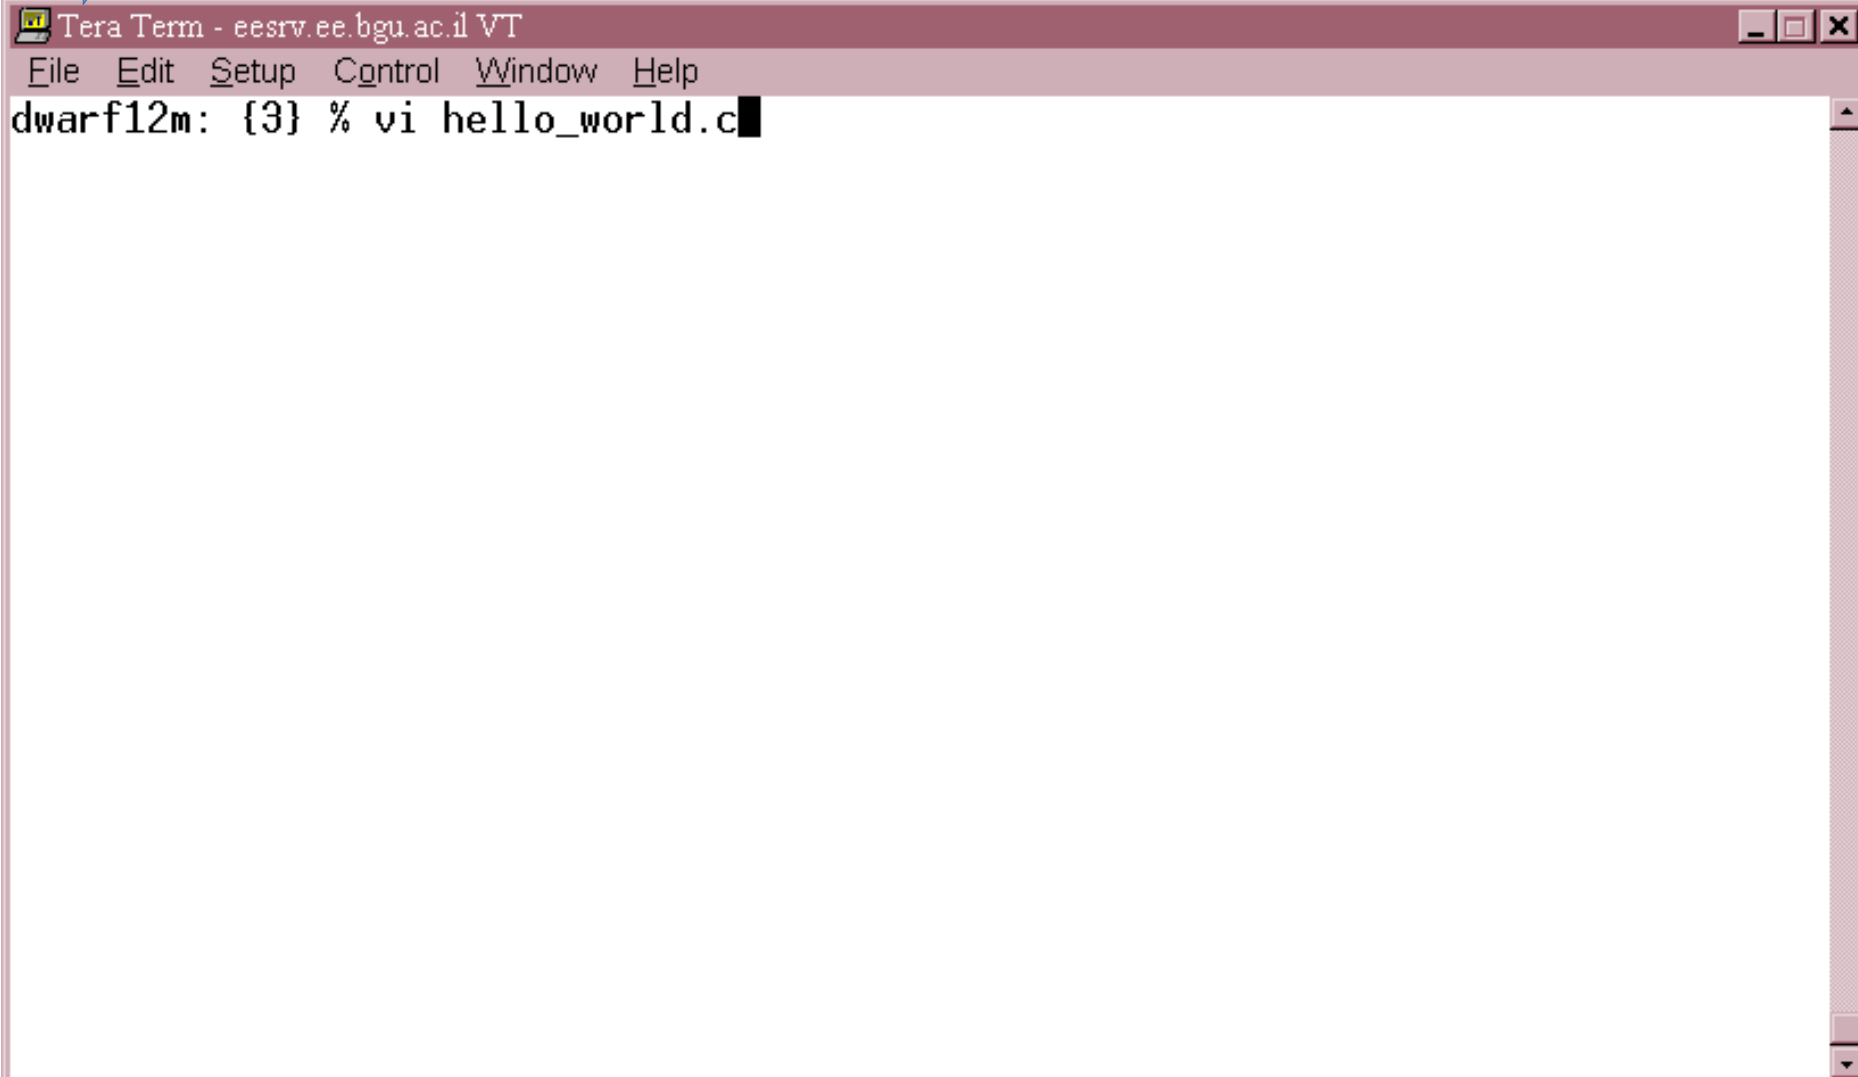
# The GNU compiler

<http://gcc.gnu.org/>

- **`gcc filename.c`**
  - Will produce an executable “**`a.out`**”
- **`gcc -o runme filename.c`**
  - Will produce an executable “**`runme`**”
- Optimization: **`gcc -O3 -o runme filename.c`**
- **`gcc -c filename.c`** will produce an object file “**`filename.o`**”

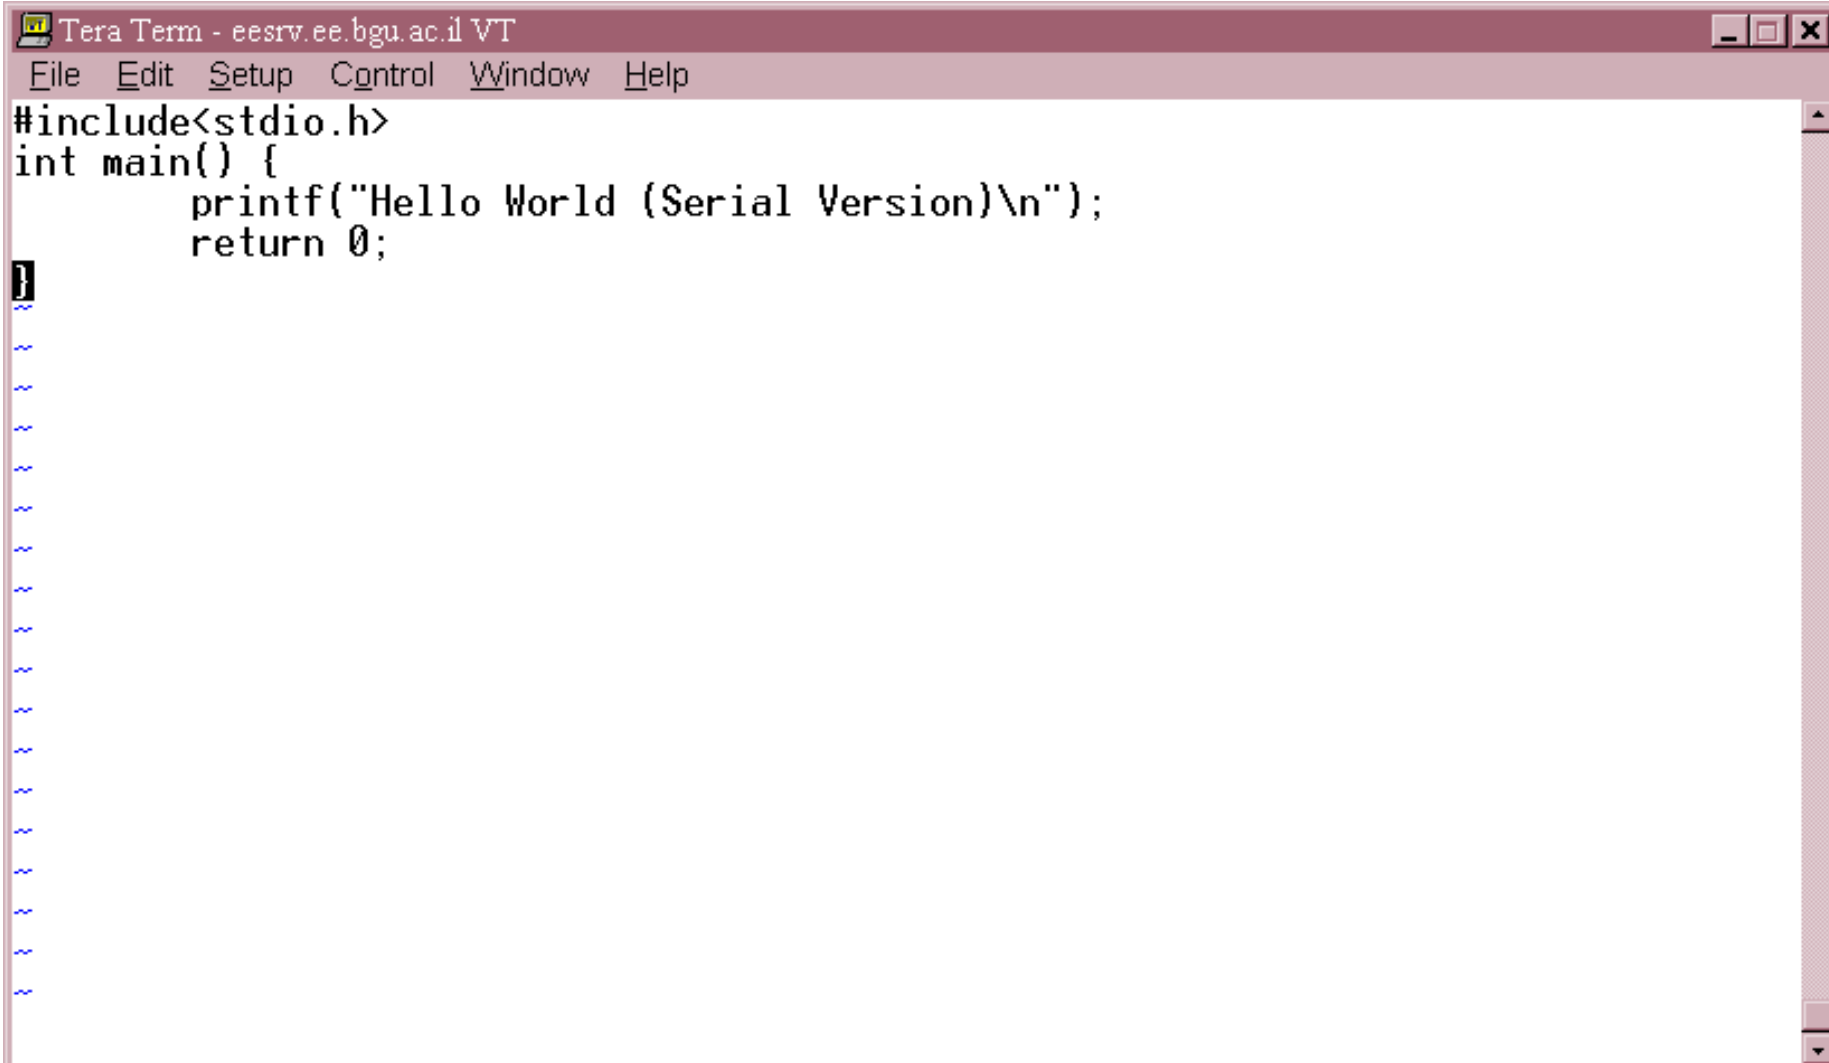
Terminal window

פתרון תרגיל מס' 1 שימוש בעורך - vi

A screenshot of a Tera Term terminal window. The title bar reads "Tera Term - eesrv.ee.bgu.ac.il VT". The menu bar includes "File", "Edit", "Setup", "Control", "Window", and "Help". The terminal text shows the prompt "dwarf12m: {3} %" followed by the command "vi hello\_world.c" and a cursor. The terminal area is mostly empty, with a vertical scrollbar on the right side.

```
Tera Term - eesrv.ee.bgu.ac.il VT
File Edit Setup Control Window Help
dwarf12m: {3} % vi hello_world.c
```

# פתרון תרגיל מס' 1 - 2/3

A screenshot of a Tera Term terminal window. The title bar reads "Tera Term - eesrv.ee.bgu.ac.il VT". The menu bar includes "File", "Edit", "Setup", "Control", "Window", and "Help". The main text area contains a C program: 

```
#include<stdio.h>
int main() {
 printf("Hello World (Serial Version)\n");
 return 0;
}
```

 The cursor is at the end of the closing brace. On the left side of the text area, there are several blue tilde (~) characters. The window has standard Windows-style window controls (minimize, maximize, close) in the top right corner.



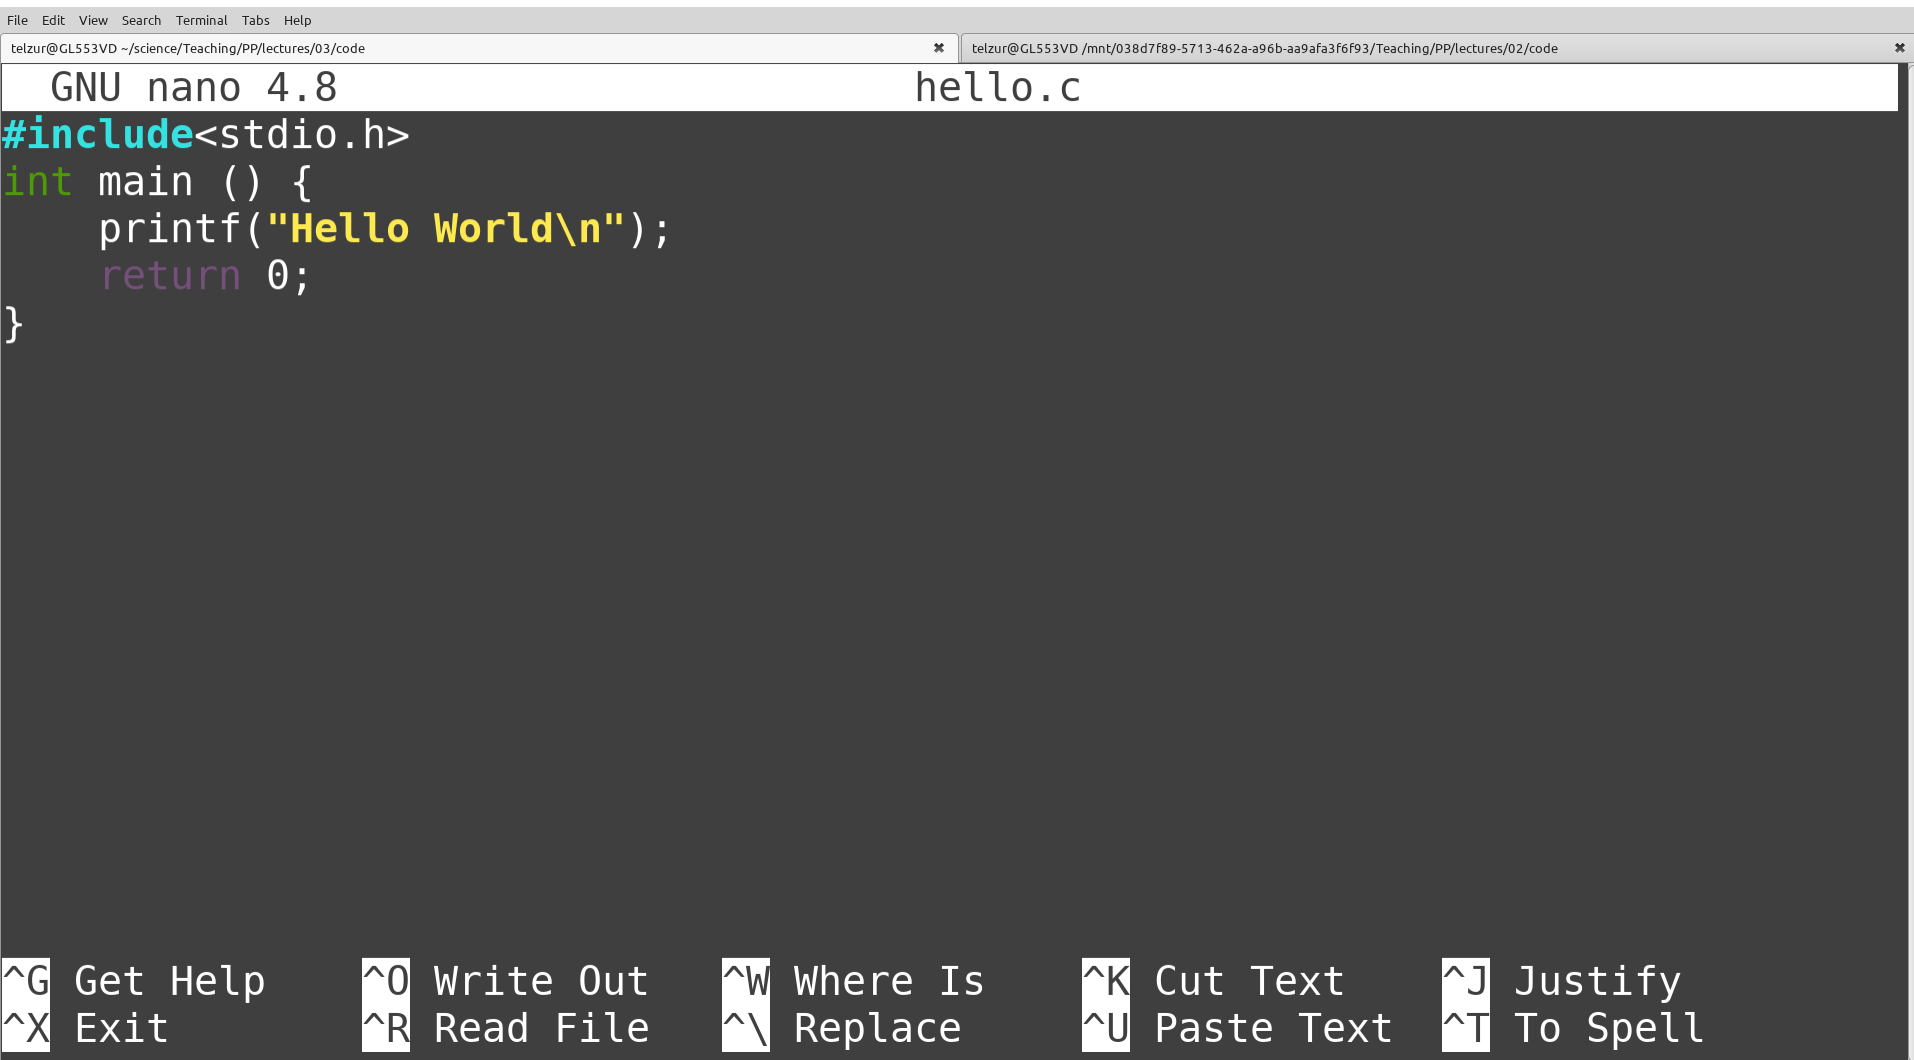
# פתרון תרגיל מס' 1 - 3/3

A screenshot of a Tera Term terminal window titled "Tera Term - eesrv.ee.bgu.ac.il VT". The menu bar includes File, Edit, Setup, Control, Window, and Help. The terminal displays several blue tilde (~) characters followed by the following commands and output:  

```
"hello_world.c" [New] 5L, 87C written
dwarf12m: {4} % gcc -o hello_world hello_world.c
dwarf12m: {5} % ./hello_world
Hello World (Serial Version)
dwarf12m: {6} %
dwarf12m: {6} % ./hello_world > hello.txt
dwarf12m: {7} % more hello.txt
Hello World (Serial Version)
dwarf12m: {8} %
```

The window has standard Windows-style controls (minimize, maximize, close) in the top right corner.

# שימוש בעורך nano



The image shows a terminal window with the GNU nano 4.8 text editor open. The editor is displaying a C program named 'hello.c'. The code is as follows:

```
GNU nano 4.8 hello.c
#include<stdio.h>
int main () {
 printf("Hello World\n");
 return 0;
}
```

The terminal window has a menu bar at the top with the following options: File, Edit, View, Search, Terminal, Tabs, Help. The status bar at the bottom displays the following keyboard shortcuts:

| Shortcut  | Action     |
|-----------|------------|
| <b>^G</b> | Get Help   |
| <b>^O</b> | Write Out  |
| <b>^W</b> | Where Is   |
| <b>^K</b> | Cut Text   |
| <b>^J</b> | Justify    |
| <b>^X</b> | Exit       |
| <b>^R</b> | Read File  |
| <b>^\</b> | Replace    |
| <b>^U</b> | Paste Text |
| <b>^T</b> | To Spell   |



## Message Passing Interface

### Quick Reference in C

```
#include <mpi.h>
```

#### Blocking Point-to-Point

Send a message to one process. (§3.2.1)

```
int MPI_Send (void *buf, int count,
 MPI_Datatype datatype, int dest, int
 tag, MPI_Comm comm)
```

Receive a message from one process. (§3.2.4)

```
int MPI_Recv (void *buf, int count,
 MPI_Datatype datatype, int source, int
 tag, MPI_Comm comm, MPI_Status *status)
```

Count received data elements. (§3.2.5)

```
int MPI_Get_count (MPI_Status *status,
 MPI_Datatype datatype, int *count)
```

Wait for message arrival. (§3.8)

```
int MPI_Probe (int source, int tag,
 MPI_Comm comm, MPI_Status *status)
```

*Related Functions:* MPI\_Bsend, MPI\_Ssend, MPI\_Rsend,  
MPI\_Buffer\_attach, MPI\_Buffer\_detach, MPI\_Sendrecv,  
MPI\_Sendrecv\_replace, MPI\_Get\_elements

#### Non-blocking Point-to-Point

Begin to receive a message. (§3.7.2)

```
int MPI_Irecv (void *buf, int count,
 MPI_Datatype, int source, int tag,
 MPI_Comm comm, MPI_Request *request)
```

Complete a non-blocking operation. (§3.7.3)

```
int MPI_Wait (MPI_Request *request,
 MPI_Status *status)
```

Check or complete a non-blocking operation. (§3.7.3)

```
int MPI_Test (MPI_Request *request, int
 *flag, MPI_Status *status)
```

Check message arrival. (§3.8)

```
int MPI_Iprobe (int source, int tag,
 MPI_Comm comm, int *flag, MPI_Status
 *status)
```

*Related Functions:* MPI\_Isend, MPI\_Ibsend, MPI\_Issend,  
MPI\_Irsend, MPI\_Request\_free, MPI\_Waitany,  
MPI\_Testany, MPI\_Waitall, MPI\_Testall, MPI\_Waitsome,  
MPI\_Testsome, MPI\_Cancel, MPI\_Test\_cancelled

#### Persistent Requests

*Related Functions:* MPI\_Send\_init, MPI\_Bsend\_init,  
MPI\_Ssend\_init, MPI\_Rsend\_init, MPI\_Recv\_init,  
MPI\_Start, MPI\_Startall

#### Derived Datatypes

Create a strided homogeneous vector. (§3.12.1)

```
int MPI_Type_vector (int count, int
 blocklength, int stride, MPI_Datatype
 oldtype, MPI_Datatype *newtype)
```

Save a derived datatype (§3.12.4)

```
int MPI_Type_commit (MPI_Datatype
 *datatype)
```

Pack data into a message buffer. (§3.13)

```
int MPI_Pack (void *inbuf, int incount,
 MPI_Datatype datatype, void *outbuf,
 int outsize, int *position, MPI_Comm
 comm)
```

Unpack data from a message buffer. (§3.13)

```
int MPI_Unpack (void *inbuf, int insize,
 int *position, void *outbuf, int
 outcount, MPI_Datatype datatype,
 MPI_Comm comm)
```

Determine buffer size for packed data. (§3.13)

```
int MPI_Pack_size (int incount,
 MPI_Datatype datatype, MPI_Comm comm,
 int *size)
```

*Related Functions:* MPI\_Type\_contiguous,  
MPI\_Type\_hvector, MPI\_Type\_indexed,  
MPI\_Type\_hindexed, MPI\_Type\_struct, MPI\_Address,  
MPI\_Type\_extent, MPI\_Type\_size, MPI\_Type\_lb,  
MPI\_Type\_ub, MPI\_Type\_free

#### Collective

Send one message to all group members. (§4.4)

```
int MPI_Bcast (void *buf, int count,
 MPI_Datatype datatype, int root,
 MPI_Comm comm)
```

Receive from all group members. (§4.5)

```
int MPI_Gather (void *sendbuf, int
 sendcount, MPI_Datatype sendtype, void
 *recvbuf, int recvcount, MPI_Datatype
 recvtype, int root, MPI_Comm comm)
```

Send separate messages to all group members. (§4.6)

```
int MPI_Scatter (void *sendbuf, int
 sendcount, MPI_Datatype sendtype, void
 *recvbuf, int recvcount, MPI_Datatype
 recvtype, int root, MPI_Comm comm)
```

Combine messages from all group members. (§4.9.1)

```
int MPI_Reduce (void *sendbuf, void
 *recvbuf, int count, MPI_Datatype
 datatype, MPI_Op op, int root, MPI_Comm
 comm)
```

*Related Functions:* MPI\_Barrier, MPI\_Gatherv,  
MPI\_Scatterv, MPI\_Allgather, MPI\_Allgatherv,  
MPI\_Alltoall, MPI\_Alltoallv, MPI\_Op\_create,  
MPI\_Op\_free, MPI\_Allreduce, MPI\_Reduce\_scatter,  
MPI\_Scan

#### Groups

*Related Functions:* MPI\_Group\_size, MPI\_Group\_rank,  
MPI\_Group\_translate\_ranks, MPI\_Group\_compare,  
MPI\_Comm\_group, MPI\_Group\_union,  
MPI\_Group\_intersection, MPI\_Group\_difference,  
MPI\_Group\_incl, MPI\_Group\_excl,  
MPI\_Group\_range\_incl, MPI\_Group\_range\_excl,  
MPI\_Group\_free

#### Basic Communicators

Count group members in communicator. (§5.4.1)

```
int MPI_Comm_size (MPI_Comm comm, int
 *size)
```

Determine group rank of self. (§5.4.1)

```
int MPI_Comm_rank (MPI_Comm comm, int
 *rank)
```

Duplicate with new context. (§5.4.2)

```
int MPI_Comm_dup (MPI_Comm comm, MPI_Comm
 *newcomm)
```

Split into categorized sub-groups. (§5.4.2)

```
int MPI_Comm_split (MPI_Comm comm, int
 color, int key, MPI_Comm *newcomm)
```

*Related Functions:* MPI\_Comm\_compare,  
MPI\_Comm\_create, MPI\_Comm\_free,

MPI\_Comm\_test\_inter, MPI\_Comm\_remote\_size,  
MPI\_Comm\_remote\_group, MPI\_Intercomm\_create,  
MPI\_Intercomm\_merge

## Communicators with Topology

Create with cartesian topology. (§6.5.1)

```
int MPI_Cart_create (MPI_Comm comm_old,
 int ndims, int *dims, int *periods, int
 reorder, MPI_Comm *comm_cart)
```

Suggest balanced dimension ranges. (§6.5.2)

```
int MPI_Dims_create (int nnodes, int
 ndims, int *dims)
```

Determine rank from cartesian coordinates. (§6.5.4)

```
int MPI_Cart_rank (MPI_Comm comm, int
 *coords, int *rank)
```

Determine cartesian coordinates from rank. (§6.5.4)

```
int MPI_Cart_coords (MPI_Comm comm, int
 rank, int maxdims, int *coords)
```

Determine ranks for cartesian shift. (§6.5.5)

```
int MPI_Cart_shift (MPI_Comm comm, int
 direction, int disp, int *rank_source,
 int *rank_dest)
```

Split into lower dimensional sub-grids. (§6.5.6)

```
int MPI_Cart_sub (MPI_Comm comm, int
 *remain_dims, MPI_Comm *newcomm)
```

*Related Functions:* MPI\_Graph\_create, MPI\_Topo\_test,  
MPI\_Graphdims\_get, MPI\_Graph\_get,  
MPI\_Cartdim\_get, MPI\_Cart\_get,  
MPI\_Graph\_neighbors\_count, MPI\_Graph\_neighbors,  
MPI\_Cart\_map, MPI\_Graph\_map

## Communicator Caches

*Related Functions:* MPI\_Keyval\_create, MPI\_Keyval\_free,  
MPI\_Attr\_put, MPI\_Attr\_get, MPI\_Attr\_delete

## LAM & MPI Information

1224 Kinnear Rd.  
Columbus, Ohio 43212  
614-292-8492

lam@tbag.osc.edu  
<http://www.osc.edu/lam.html>  
<ftp://ftp.osc.edu/pub/lam>

## Error Handling

*Related Functions:* MPI\_Errhandler\_create,  
MPI\_Errhandler\_set, MPI\_Errhandler\_get,  
MPI\_Errhandler\_free, MPI\_Error\_string, MPI\_Error\_class

## Dynamic Processes

Spawn a process. (MPI-2)

```
int MPI_Spawn (char prog[], char *argv[],
 int maxprocs, MPI_Info info, int root,
 MPI_Comm parents, MPI_Comm *children,
 int errs[]);
```

*Related Functions:* MPI\_Spawn\_multiple, MPI\_Ispawn,  
MPI\_Ispawn\_multiple, MPI\_Port\_open, MPI\_Port\_close,  
MPI\_Accept, MPI\_Connect, MPI\_Name\_publish,  
MPI\_Name\_unpublish, MPI\_Name\_get, MPI\_Iaccept,  
MPI\_Iconnect, MPI\_Info\_create, MPI\_Info\_set,  
MPI\_Info\_get, MPI\_Info\_get\_valuelen,  
MPI\_Info\_get\_nkeys, MPI\_Info\_get\_nthkey,  
MPI\_Info\_dup, MPI\_Info\_free, MPI\_Info\_delete

## Environmental

Determine wall clock time. (§7.4)

```
double MPI_Wtime (void)
```

Initialize MPI. (§7.5)

```
int MPI_Init (int *argc, char ***argv)
```

Cleanup MPI. (§7.5)

```
int MPI_Finalize (void)
```

*Related Functions:* MPI\_Get\_processor\_name, MPI\_Wtick,  
MPI\_Initialized, MPI\_Abort, MPI\_Pcontrol, MPI\_Get\_version

## Constants

Wildcards (§3.2.4)

MPI\_ANY\_TAG, MPI\_ANY\_SOURCE

Elementary Datatypes (§3.2.2)

MPI\_CHAR, MPI\_SHORT, MPI\_INT, MPI\_LONG,  
MPI\_UNSIGNED\_CHAR, MPI\_UNSIGNED\_SHORT,  
MPI\_UNSIGNED, MPI\_UNSIGNED\_LONG, MPI\_FLOAT,  
MPI\_DOUBLE, MPI\_LONG\_DOUBLE, MPI\_BYTE,  
MPI\_PACKED

Reserved Communicators (§5.2.4)

MPI\_COMM\_WORLD, MPI\_COMM\_SELF, MPI\_COMM\_PARENT

Reduction Operations (§4.9.2)

MPI\_MAX, MPI\_MIN, MPI\_SUM, MPI\_PROD,  
MPI\_BAND, MPI\_BOR, MPI\_BXOR, MPI\_LAND,  
MPI\_LOR, MPI\_LXOR



# LAM Quick Reference

## Session Management

Confirm a group of hosts.  
recon -v <hostfile>

Start LAM on a group of hosts.  
lamboot -v <hostfile>

Terminate LAM.  
wipe -v <hostfile>

Hostfile Syntax  
# comment  
<hostname> <userid>  
<hostname> <userid>  
...etc...

## Compilation

Compile a program for LAM / MPI.  
hcc -o <binary> <source> -I<incdir>  
-L<libdir> -l<lib> -lmpi

## Processes and Messages

Start an SPMD application.  
mpirun -v -s <src\_node> -c <copies>  
<nodes> <program> -- <args>

Start a MIMD application.  
mpirun -v <appfile>

Appfile Syntax  
# comment  
<program> -s <src\_node> <nodes> -- <args>  
<program> -s <src\_node> <nodes> -- <args>  
...etc...

Examine the state of processes.  
mpitask

Examine the state of messages.  
mpinag

Cleanup all processes and messages.  
lamclean -v



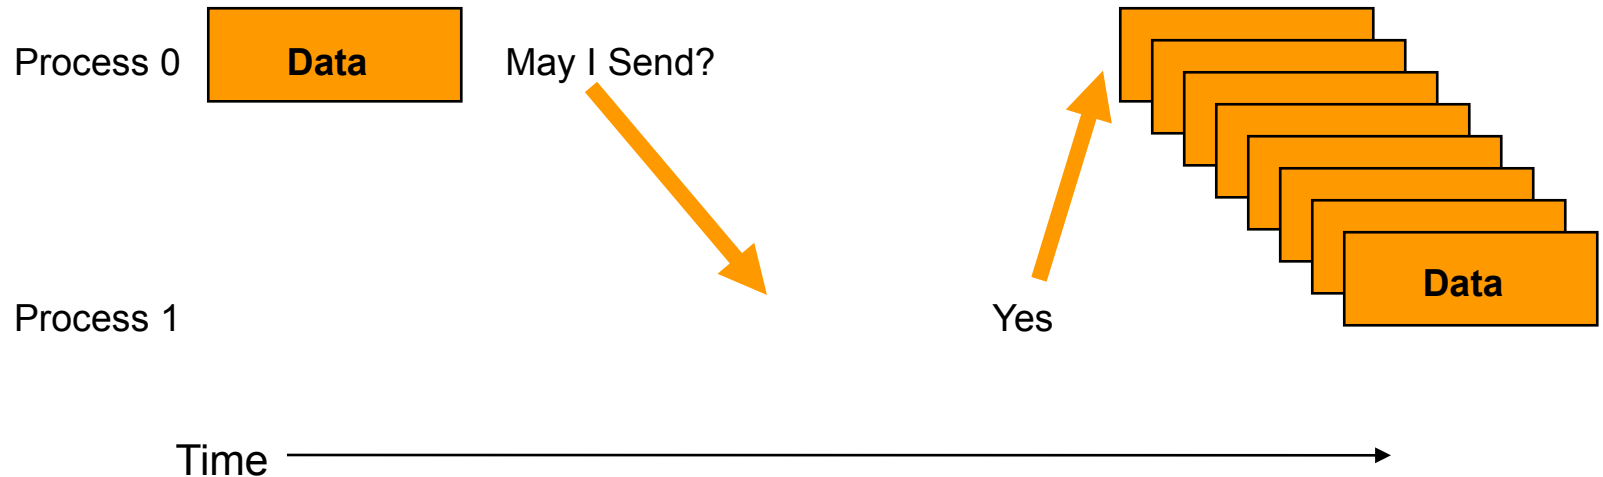
OHIO  
SUPERCOMPUTER  
CENTER

# **MPI Quick Reference Card:**

**<http://web.eecs.utk.edu/~dongarra/WEB-PAGES/SPRING-2006/mpi-quick-ref.pdf>**


# What is message passing?

- Data transfer through messaging




- Requires a sender and a receiver cooperation

# Point to Point: Basic Send/Receive



# Space-Time Diagram of a Message-Passing Program





# MPI - Message Passing Interface API

- MPI is a standard not an implementation
- Popular implementations are LAM and MPICH
- MPICH is installed under */usr/local/mpich*
- Always put in the code: `#include "mpi.h"`
- Compilation: `mpicc -o filename file.c`
- Execution: `mpirun -np N filename`
- Help: `man mpirun`

# MPI Naming Conventions

**MPI\_Xxxxx (parameter, ...)**

**Example: MPI\_Init(&argc, &argv)**

# The First 4 Functions of MPI

- MPI\_Init
  - MPI\_Finalize
  - MPI\_Comm\_size
  - MPI\_Comm\_rank
- 
- ...and don't forget the  
**#include "mpi.h"**

# The First 4 Functions Syntax

- `int MPI_Init(int argc, char *argv[] )`
- `int MPI_Finalize()`
- `int MPI_Comm_size(MPI_Comm comm, int *size)`
- `int MPI_Comm_rank(MPI_Comm comm, int *rank)`

# MPI Communicator

A **communicator** is a handle representing a group of processors that can communicate with one another. The **communicator** name is required as an argument to all point-to-point and collective operations. The **communicator** specified in the **send** and **receive** calls must agree for communication to take place. Processors can communicate only if they share a **communicator**.

# Basic Point to Point Functions

- **MPI\_Send**
- **MPI\_Recv**
- **MPI\_Send**(void \*buf, int count, MPI\_Datatype datatype, int dest, int tag, MPI\_Comm comm);
- **MPI\_Recv**(void \*buf, int count, MPI\_Datatype datatype, int source, int tag, MPI\_Comm comm, MPI\_Status status);

# MPI\_Send

```
int MPI_Send(void *buf, int count, MPI_Datatype dtype,
int dest, int tag, MPI_Comm comm);
```

|                     |   |                                                                        |
|---------------------|---|------------------------------------------------------------------------|
| <i>buffer</i>       | } | message body                                                           |
| <i>count</i>        |   |                                                                        |
| <i>datatype</i>     |   |                                                                        |
| <i>destination</i>  | } | message envelope (source—the<br>sending process—is defined implicitly) |
| <i>tag</i>          |   |                                                                        |
| <i>communicator</i> |   |                                                                        |

# תרגיל מס' 2

- הרצת תכנית קצרה ב-MPI:

**Hello\_World מקבילי**

- כתוב תכנית בה כל מחשב יאמר שלום ויודיע

את מספר התהליך שלו בריצה, לדוגמה:

**Hello world from process 1 of 2**



# פתרון תרגיל מס' 2

// see more examples: /usr/local/mpich/examples

```
#include <stdio.h>
#include "mpi.h"
int main(argc, argv)
int argc;
char **argv;
{
 int rank, size;
 MPI_Init(&argc, &argv);
 MPI_Comm_size(MPI_COMM_WORLD, &size);
 MPI_Comm_rank(MPI_COMM_WORLD, &rank);
 printf("Hello world from process %d of %d\n",
rank, size);
 MPI_Finalize();
 return 0;
}
```


# פתרון תרגיל מס' 2

- Note on syntax:

```
int main(argc, argv)
int argc;
char **argv;
```

Is Equivalent to:

```
int main(int argc, char *argv[])
```



לאמץ את  
הפורמט הזה

...and then...

```
MPI_Init(&argc,&argv);
```

```
nano ./machinefiles
```

Enter machines' names, e.g.:

```
hobbit1
```

```
hobbit2
```

```
hobbit3
```

```
hobbit4
```

Create a “machinefile”

Save and exit

**השמות המעודכנים של המחשבים**

**hobbit1 .. hobbit15**

# Our machine file

- Specify In the machine file short computers' names, e.g.

**hobbit1**

**Hobbit2**

and not:

**hobbit1.ee.bgu.ac.il**

**hobbit2.ee.bgu.ac.il**

# Hello World - Execution

```
% mpicc -o helloworld helloworld.c
```

```
% mpirun -np 4 helloworld
```

*Or*

```
mpirun -np 4 -machinefile ./machinefile
./helloworld
```

```
Hello world from process 0 of 4
```

```
Hello world from process 3 of 4
```

```
Hello world from process 1 of 4
```

```
Hello world from process 2 of 4
```

```
%
```

```
mpirun -np 4 -machinefile ./machinefile
helloworld
```

**אפשר גם:**

```
mpirun -n 4 -f ./machinefile
./helloworld
```

**דוגמה בשקף הבא**

**mpirun ↔ mpiexec**

## Terminal

```
-bash-4.1$ mpirun -n 9 ./cpi
Process 1 on hobbit10.ee.bgu.ac.il
Process 3 on hobbit10.ee.bgu.ac.il
Process 4 on hobbit10.ee.bgu.ac.il
Process 6 on hobbit10.ee.bgu.ac.il
Process 8 on hobbit10.ee.bgu.ac.il
Process 7 on hobbit10.ee.bgu.ac.il
Process 2 on hobbit10.ee.bgu.ac.il
Process 0 on hobbit10.ee.bgu.ac.il
Process 5 on hobbit10.ee.bgu.ac.il
pi is approximately 3.1415926535897944, Error is 0.00000000000000013
wall clock time = 0.079962
-bash-4.1$ mpirun -n 9 -f ./machinefile ./cpi
Process 3 on hobbit4.ee.bgu.ac.il
Process 4 on hobbit6.ee.bgu.ac.il
Process 8 on hobbit4.ee.bgu.ac.il
Process 0 on hobbit1.ee.bgu.ac.il
Process 2 on hobbit3.ee.bgu.ac.il
Process 5 on hobbit1.ee.bgu.ac.il
Process 7 on hobbit3.ee.bgu.ac.il
Process 1 on hobbit2.ee.bgu.ac.il
Process 6 on hobbit2.ee.bgu.ac.il
pi is approximately 3.1415926535897944, Error is 0.00000000000000013
wall clock time = 0.045389
-bash-4.1$
```

# machinefile

If no *machinefile* is specified you are running on the local machine

If a *machinefile* exists you are running on the machines specified in the file



File Edit View Window Help



Quick Connect Profiles

```
163 root 15 0 0 0 0 SW 0.0 0.0 0:00 kjournald
578 root 15 0 540 540 460 S 0.0 0.0 0:32 syslogd
582 root 15 0 428 428 376 S 0.0 0.0 0:01 klogd
```

```
dwarf2m: {24} % pwd
```

```
/home/telzur/mpi
```

```
dwarf2m: {25} % mpirun -machinefile ./machinefile -np 3 ./Hello_world
```

```
Hello world from process 0 of 3
```

```
Hello world from process 1 of 3
```

```
Hello world from process 2 of 3
```

```
dwarf2m: {26} % mpirun -machinefile ./machinefile -np 4 ./Hello_world
```

```
Hello world from process 0 of 4
```

```
Hello world from process 1 of 4
```

```
Hello world from process 2 of 4
```

```
Hello world from process 3 of 4
```

```
dwarf2m: {27} % mpirun -machinefile ./machinefile -np 8 ./Hello_world
```

```
Hello world from process 0 of 8
```

```
Hello world from process 4 of 8
```

```
Hello world from process 2 of 8
```

```
Hello world from process 5 of 8
```

```
Hello world from process 3 of 8
```

```
Hello world from process 7 of 8
```

```
Hello world from process 1 of 8
```

```
Hello world from process 6 of 8
```

```
dwarf2m: {28} %
```

# תרגיל מס' 3: חישוב



• חישוב באמצעות אינטגרציה

$$\int_0^1 \frac{1}{1+x^2} dx = \arctan(x) \Big|_0^1 = \arctan(1) - \arctan(0) = \arctan(1) = \frac{\pi}{4}$$

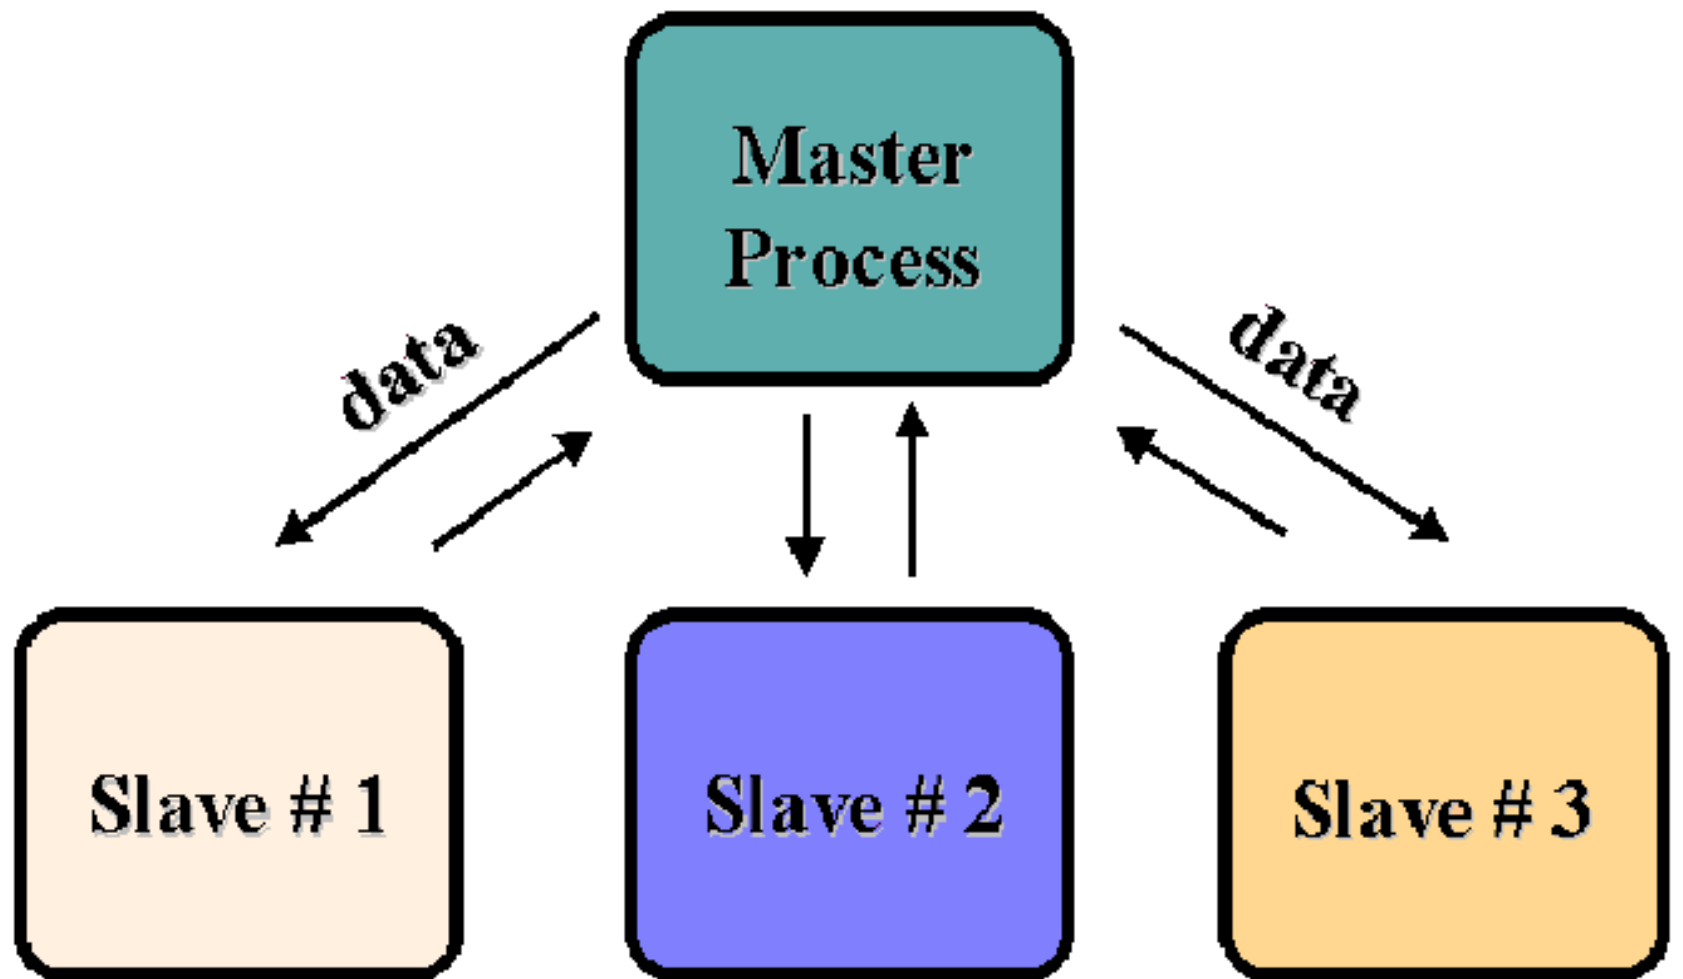
• נבצע סכימה על הפונקציה

•  $f(x) = 4/(1+x^2)$  בין 0 ל-1 על-ידי חלוקת  
התחום ל-n מלבנים

**Domain Decomposition**, There is more than one way to decompose the circle. Which one will you choose?



# Master/Workers



# פתרון תרגיל מס' 3

- בפתרון השתמשנו בפונקציה למדידת זמן
- הנקראת : `MPI_Wtime`

- ראה דוגמא לפתרון התרגיל תחת:

`/usr/local/mpich/examples/cpi.c`

ניתן להוריד את התכנית מאתר הקורס וגם מהכתובת:

`http://www.mcs.anl.gov/research/  
projects/mpi/usingmpi/examples/  
simplempi/cpi\_c.htm`

Copy this program to a location under your home directory where you have a write permission! Execute it from there

# פתרון תרגיל מס' 3

```
#include "mpi.h"
#include <stdio.h>
#include <math.h>
double f(double a)
{
 return (4.0 / (1.0 + a*a));
}

void main(int argc, char *argv[])
{
 int done = 0, n, myid, numprocs, i;
 double PI25DT = 3.141592653589793238462643;
 double mypi, pi, h, sum, x;
 double startwtime, endwtime;
 int namelen;
 char processor_name[MPI_MAX_PROCESSOR_NAME];
```

# פתרון תרגיל מס' 3 - המשך

```
MPI_Init(&argc,&argv);
MPI_Comm_size(MPI_COMM_WORLD,&numprocs);
MPI_Comm_rank(MPI_COMM_WORLD,&myid);
MPI_Get_processor_name(processor_name,&namelen);
```

```
 fprintf(stderr,"Process %d on %s\n",myid,
processor_name);
 fflush(stderr);
```


```
n = 0;
```

# פתרון תרגיל מס' 3 - המשך

```
while (!done)
{
 if (myid == 0)
 {
 printf("Enter the number of intervals: (0 quits)
");
 fflush(stdout);
 scanf("%d",&n);
 startwtime = MPI_Wtime();
 }
 MPI_Bcast(&n, 1, MPI_INT, 0, MPI_COMM_WORLD);
 if (n == 0)
 done = 1;
 else
 {
```




# MPI Broadcast operation



# פתרון תרגיל מס' 3 - המשך

```
h = 1.0 / (double) n;
sum = 0.0;
for (i = myid + 1; i <= n; i += numprocs)
{
 x = h * ((double)i - 0.5);
 sum += f(x);
}
mypi = h * sum;
MPI_Reduce(&mypi, &pi, 1, MPI_DOUBLE, MPI_SUM,
0, MPI_COMM_WORLD);
```

# MPI reduction operation



# פתרון תרגיל מס' 3 – המשך

## Basic Collective Functions

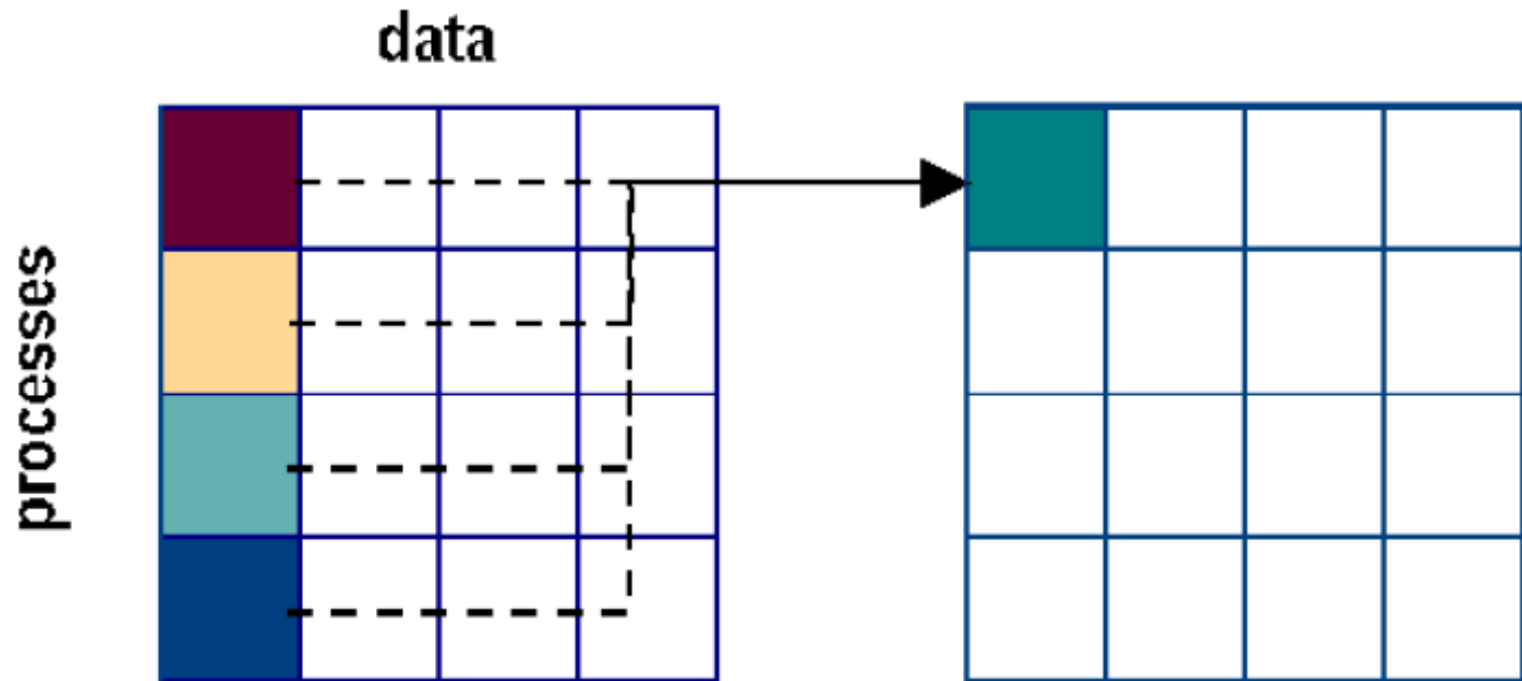
- **MPI\_Bcast**
- **MPI\_Reduce**
- The exact syntax:
- **MPI\_Bcast**(void \*buf, int count, MPI\_Datatype datatype, int root, MPI\_Comm comm);
- **MPI\_Reduce**(void \*sendbuf, void \*recvbuf, int count, MPI\_Datatype datatype, MPI\_Op op, int root, MPI\_Comm comm);

# פתרון תרגיל מס' 3 - המשך

```
if (myid == 0)
{
 printf("pi is approximately %.16f, Error is %.16f\n",
pi, fabs(pi - PI25DT));
 endwtime = MPI_Wtime();
 printf("wall clock time = %f\n",endwtime-startwtime);
} /* end of if */
} /* end of while */
 MPI_Finalize();
} /* end of main */
```

# פתרון תרגיל מס' 3 - המשך

## MPI Reduce



# הרצת 4 תהליכים

File Edit View Terminal Go Help

```
[telzur@Fermi basic]$ /usr/local/mpich-1.2.4/bin/mpirun -np 4 cpilog
```

```
Process 0 running on Fermi.tel-zur.com
```

```
Process 1 running on Fermi.tel-zur.com
```

```
Process 2 running on Fermi.tel-zur.com
```

```
Process 3 running on Fermi.tel-zur.com
```

```
pi is approximately 3.1415926535899033, Error is 0.00000000000001101
```

```
wall clock time = 0.691032
```

```
[telzur@Fermi basic]$
```

# How to execute `cpi.c`?

- `mkdir mpi`
  - `cd mpi`
  - Download `cpi.c` from the course website using `wget`:  
`wget http://tel-zur.net/teaching/bgu/pp/cpi.c`
  - `mpicc -o cpi cpi.c`
  - create a hostfile
  - `mpirun -np 4 -machinefile ./machinefile ./cpi`
- מטלות:

– התחברות ב SSH

– יצירת מחיצה + עריכת קובץ ושמירתו

– קימפול עם MPI על-ידי `mpicc`

– הרצת התכנית `cpi` על-ידי `mpirun`



# Jumpshot (hobbits)

עדכון: לא פועל  
על ההוביטים.  
להפעיל את התכנה  
מהמכונה הוירטואלית

- `mpecc -o cpi ./cpi.c -mpilog`  
or
- `mpicc -o cpi ./cpi.c -llmpe -lmpe`
- Then convert the `clog2` file by:  
`clog2T0slog2 ./file_name.clog2`
- Finally view the profiling:  
`jumpshot file_name.slog2`

# Jumpshot (VM)

`tau_cc.sh -o cpi_tau ./cpi.c` # compile with  
profiling and tracing support

`mpirun -np 8 ./cpi_tau` # execute and generate  
profiling

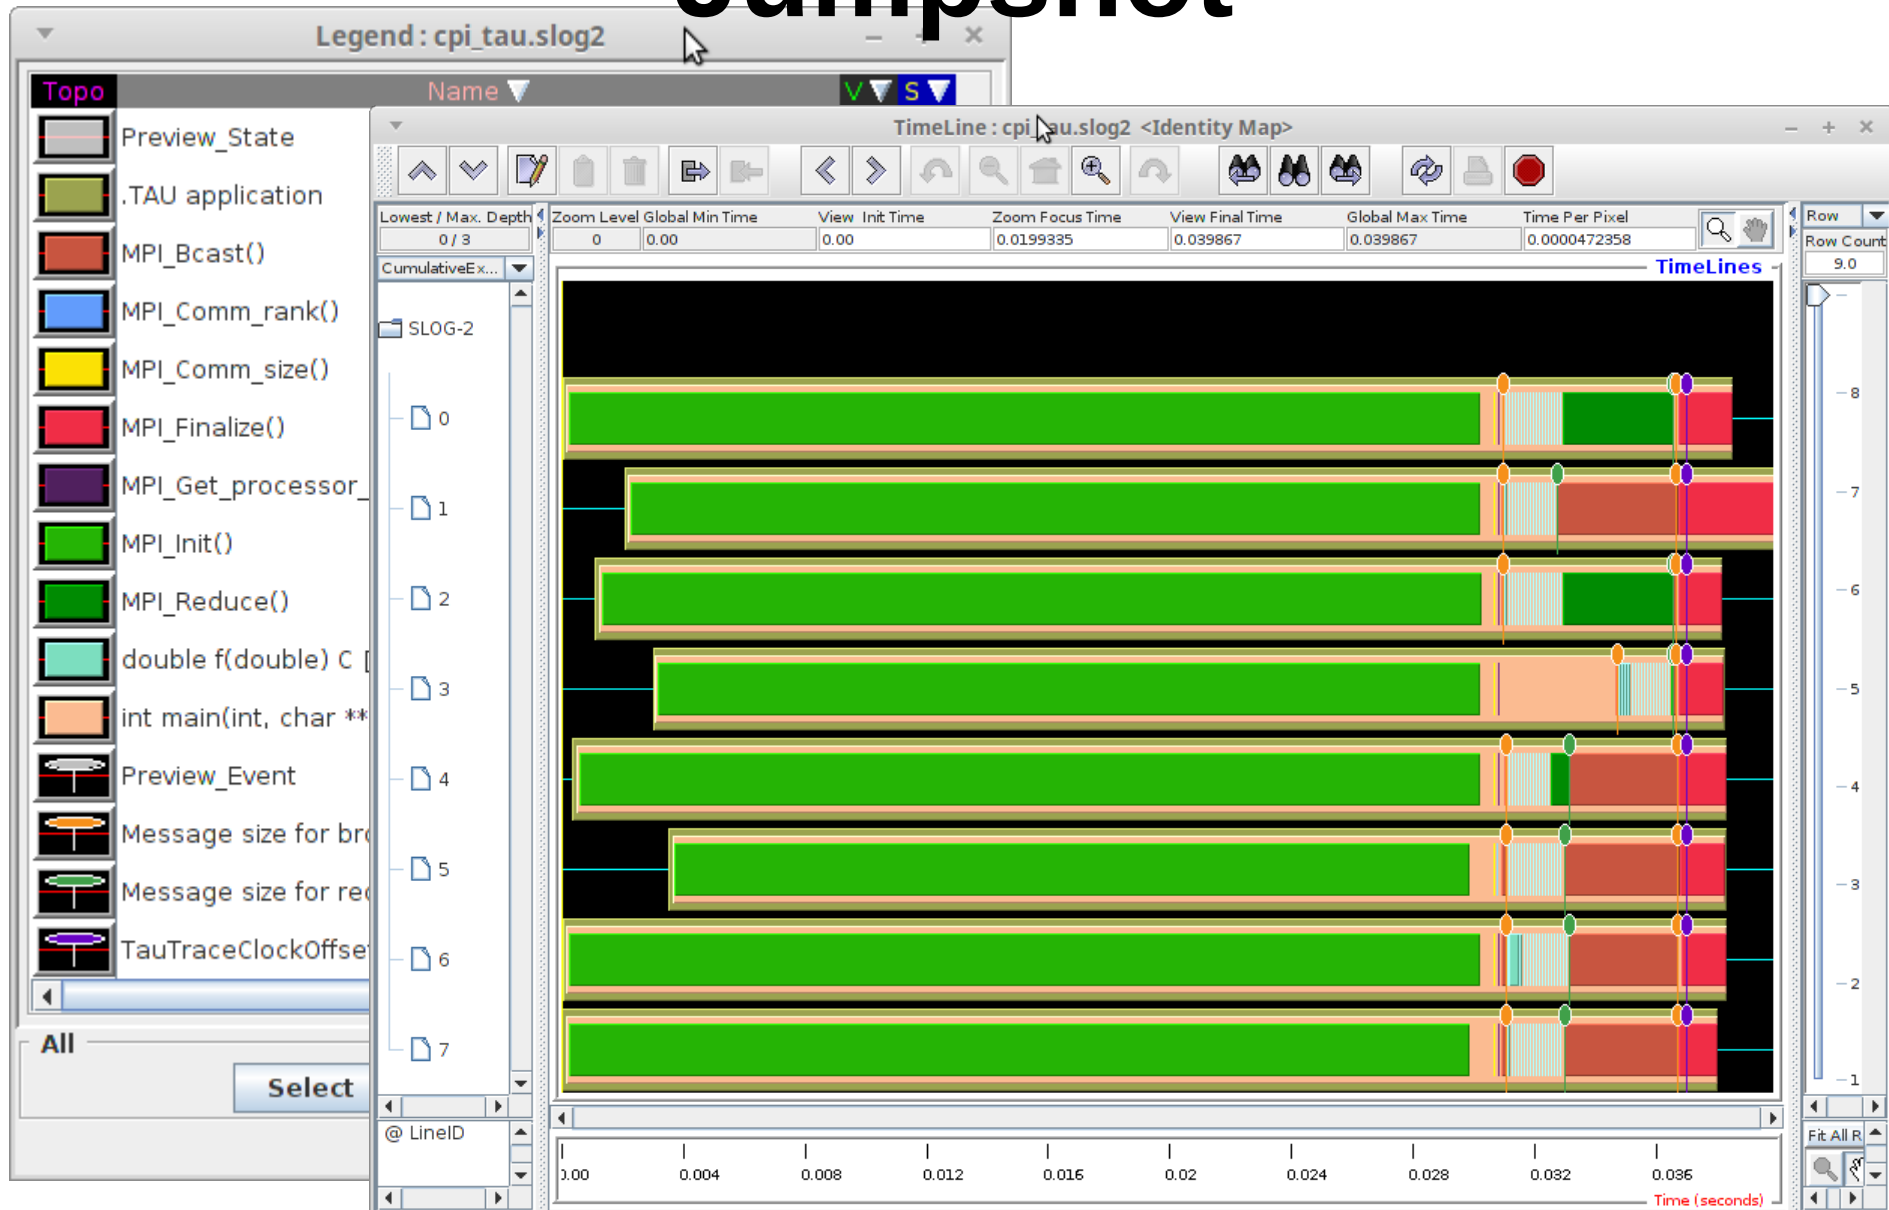
`paraprof` # paraprof visualizing

`tau_treemerge.pl` # merge the mpi tasks  
profiling files into combined files

`tau2slog2 tau.trc tau.edf` # prepare slog2 file

`jumpshot ./tau.slog2` # view tracing in jumpshot

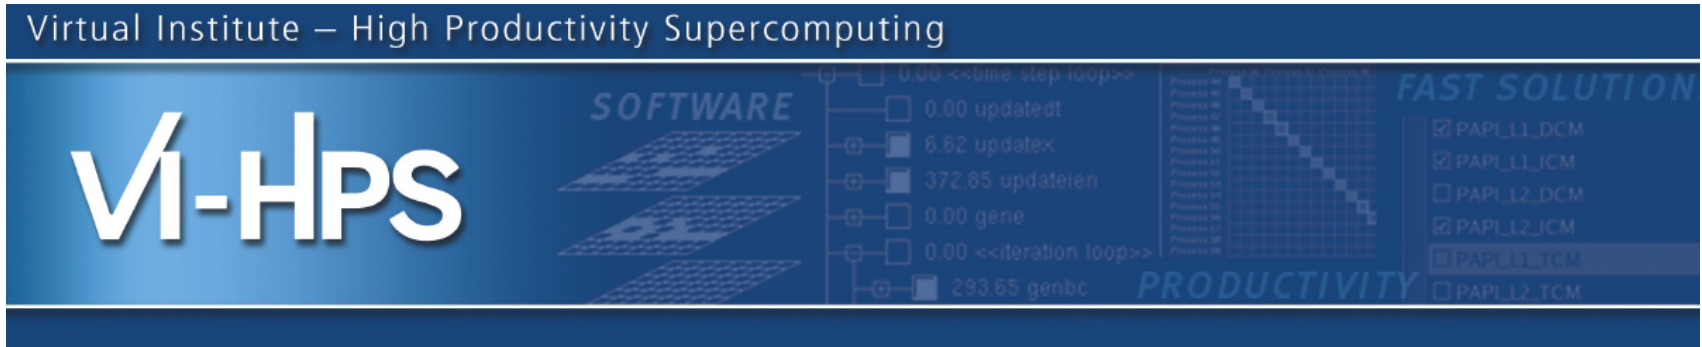
# Jumpshot



# Scalasca

- Is installed on the *hobbits*
- VI-HPS <http://www.vi-hps.org/>
- LiveDVD:  
<http://www.vi-hps.org/training/live-iso/>  
[יש להוריד ולהתקין תוכנה זו כ"אורחת" בתוכנת  
הוירטואליזציה Virtualbox](#)
- Quick Reference Guide: <http://apps.fz-juelich.de/scalasca/releases/scalasca/1.4/docs/QuickReference.pdf>

# Next slides are from:



## Scalable performance analysis of large-scale parallel applications

Brian Wylie & Markus Geimer  
Jülich Supercomputing Centre

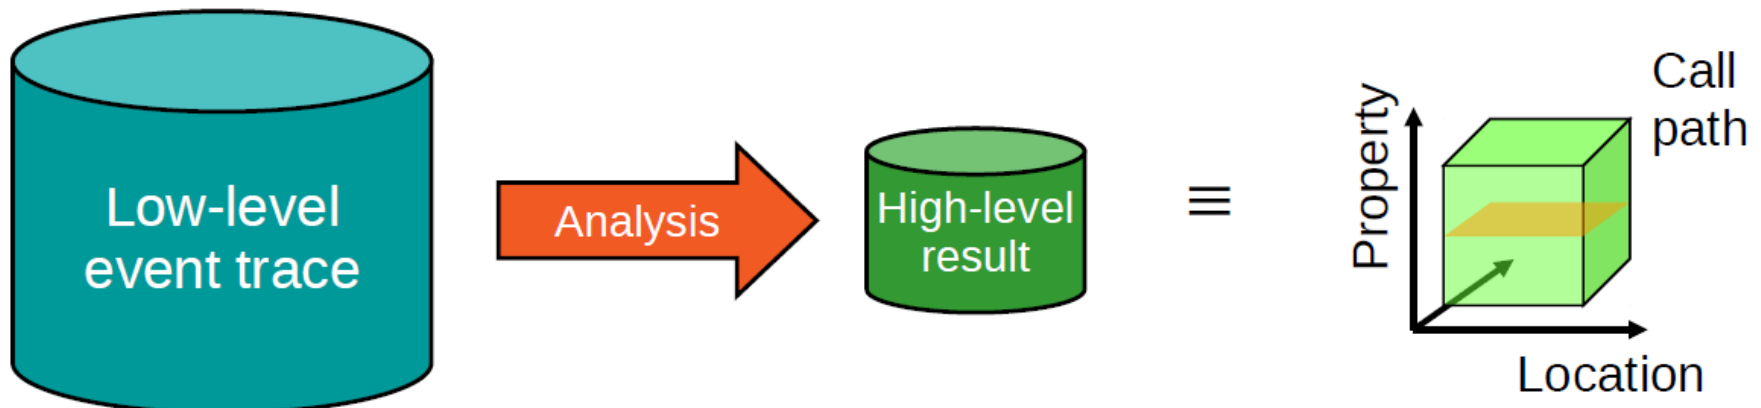
[scalasca@fz-juelich.de](mailto:scalasca@fz-juelich.de)

August 2012

- Profile analysis
  - Summary of aggregated metrics
    - per function/callpath and/or per process/thread
  - Most tools (can) generate and/or present such profiles
    - but they do so in *very* different ways, often from event traces!
  - e.g., gprof, mpiP, ompP, **Scalasca**, TAU, Vampir, ...
- Time-line analysis
  - Visual representation of the space/time sequence of events
  - Requires an execution trace
  - e.g., Vampir, Paraver, JumpShot, Intel TAC, Sun Studio, ...
- Pattern analysis
  - Search for event sequences characteristic of inefficiencies
  - Can be done manually, e.g., via visual time-line analysis
  - or automatically, e.g., KOJAK, **Scalasca**, Periscope, ...

- Idea

- Automatic search for patterns of inefficient behaviour
- Classification of behaviour & quantification of significance

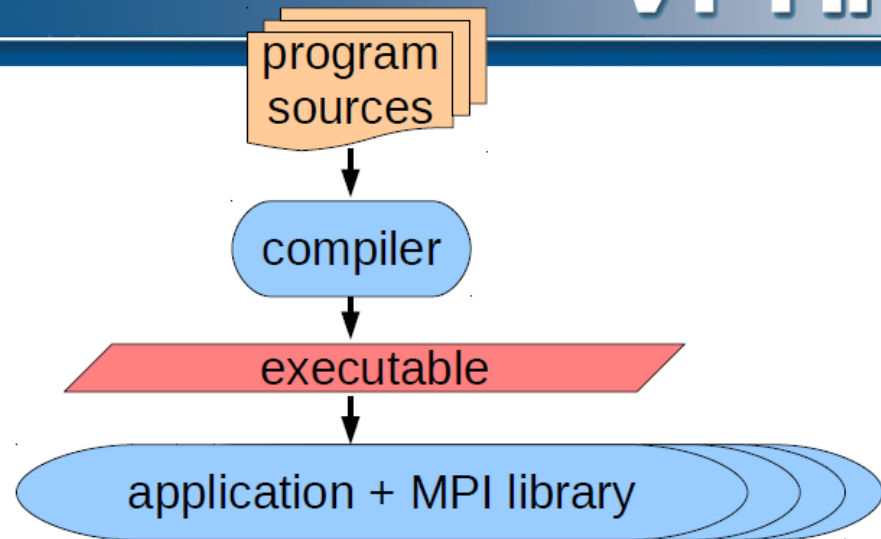


- Guaranteed to cover the entire event trace
- Quicker than manual/visual trace analysis
- Parallel replay analysis exploits memory & processors to deliver scalability

- Open source, New BSD license
- Portable
  - Cray XT, IBM BlueGene, IBM SP & blade clusters, NEC SX, SGI Altix, SiCortex, Solaris & Linux clusters, ...
- Supports parallel programming paradigms & languages
  - MPI, OpenMP & hybrid OpenMP+MPI
  - Fortran, C, C++
- Integrated instrumentation, measurement & analysis toolset
  - Automatic and/or manual customizable instrumentation
  - Runtime summarization (aka profiling)
  - Automatic event trace analysis
  - Analysis report exploration & manipulation



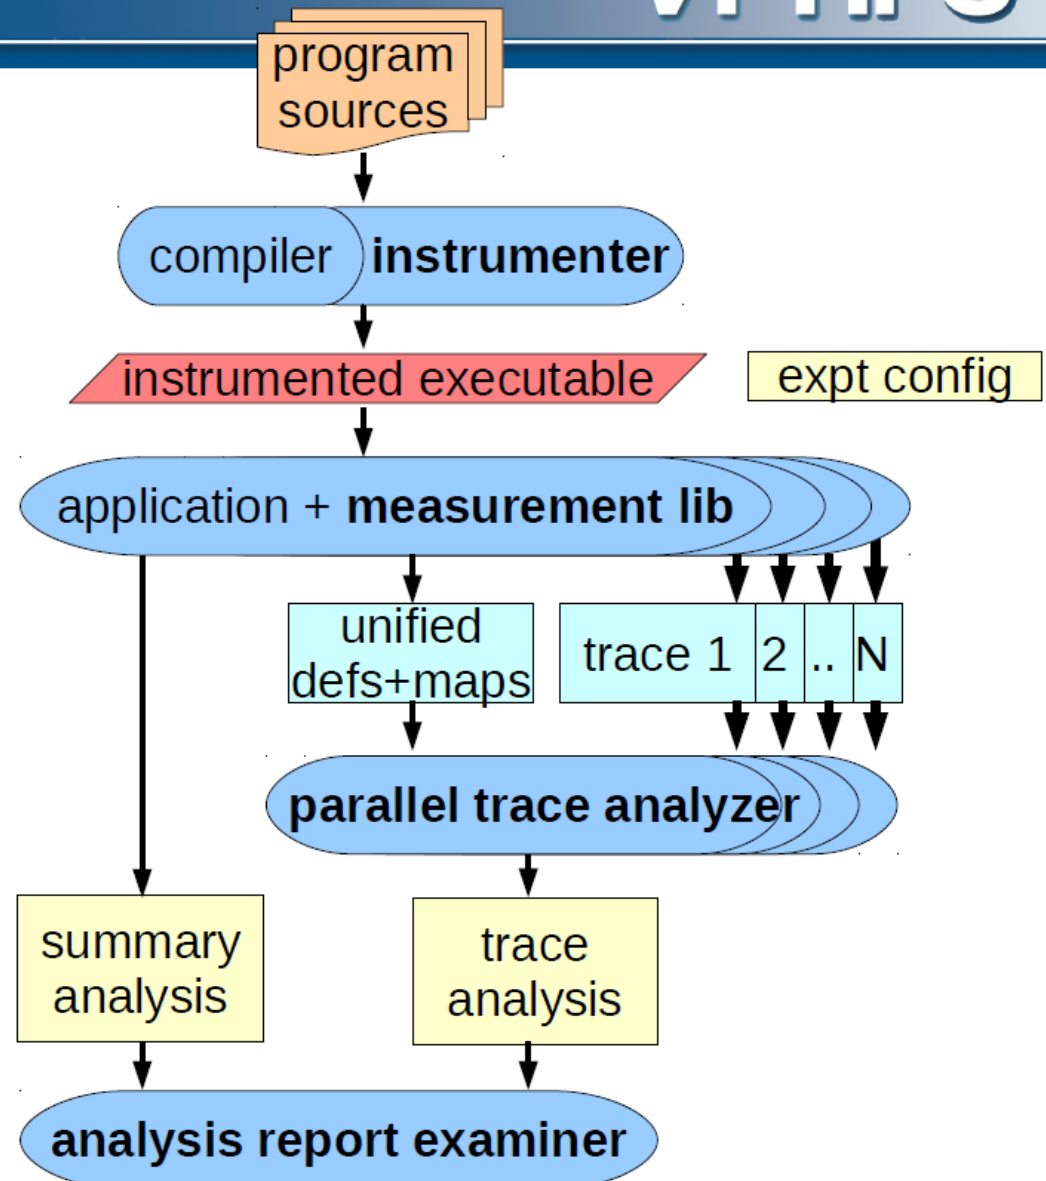
- Application code compiled & linked into executable using MPICC/CXX/FC
- Launched with MPIEXEC
- Application processes interact via MPI library



# Scalasca toolset components

# VI-HPS

- Scalasca instrumenter  
= SKIN
- Scalasca measurement  
collector & analyzer  
= SCAN
- Scalasca analysis  
report examiner  
= SQUARE



- One command for everything

% **scalasca**

Scalasca 1.4

Toolset for scalable performance analysis of large-scale apps

usage: scalasca [-v][-n] {action}

1. prepare application objects and executable for measurement:

scalasca *-instrument* <compile-or-link-command> # **skin**

2. run application under control of measurement system:

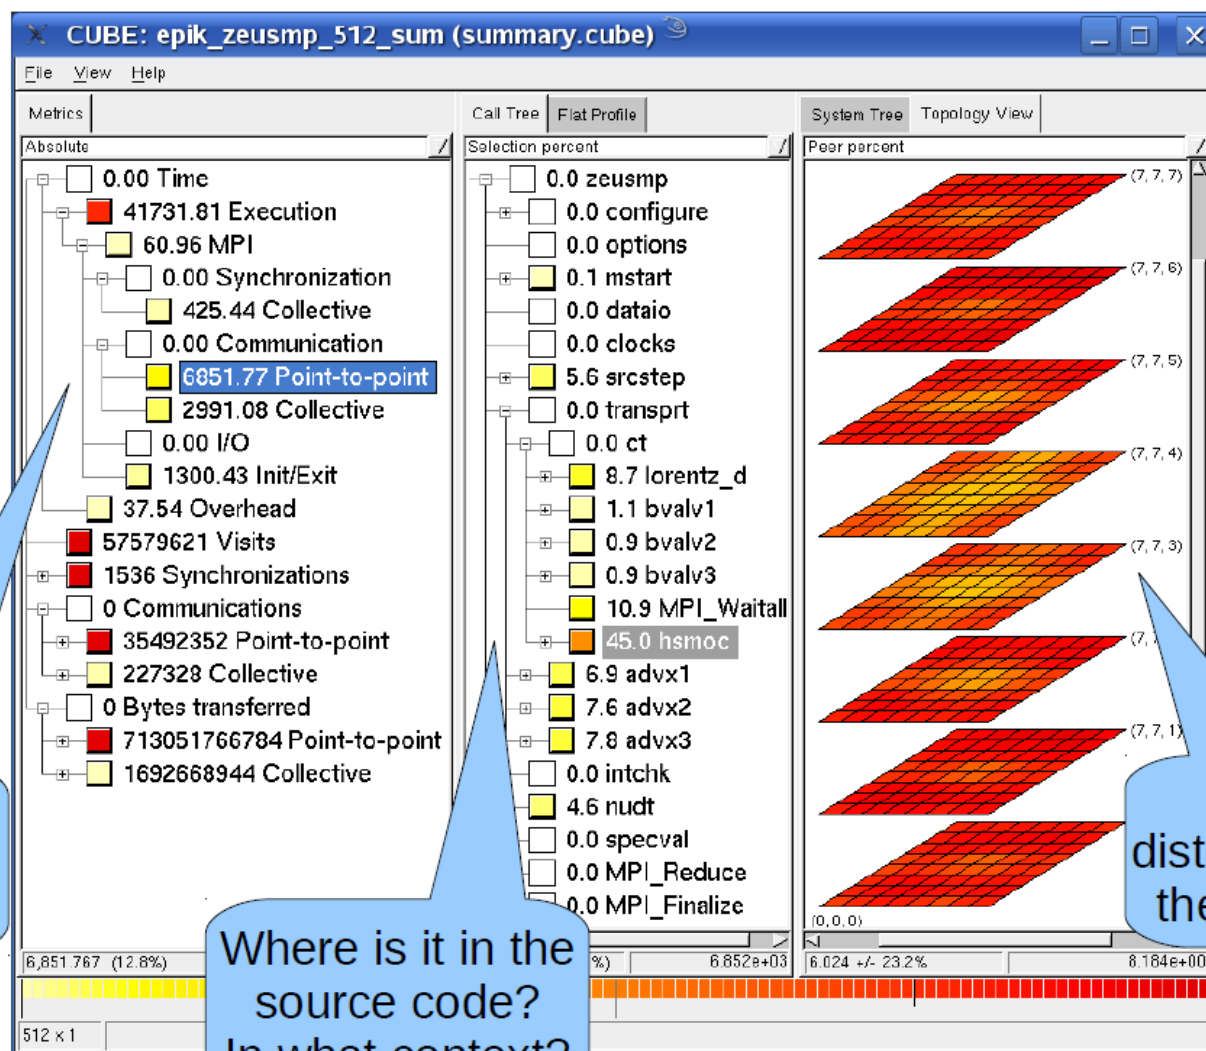
scalasca *-analyze* <application-launch-command> # **scan**

3. post-process & explore measurement analysis report:

scalasca *-examine* <experiment-archive|report> # **square**

[-h] show quick reference guide (only)

# Scalasca analysis report explorer (summary)



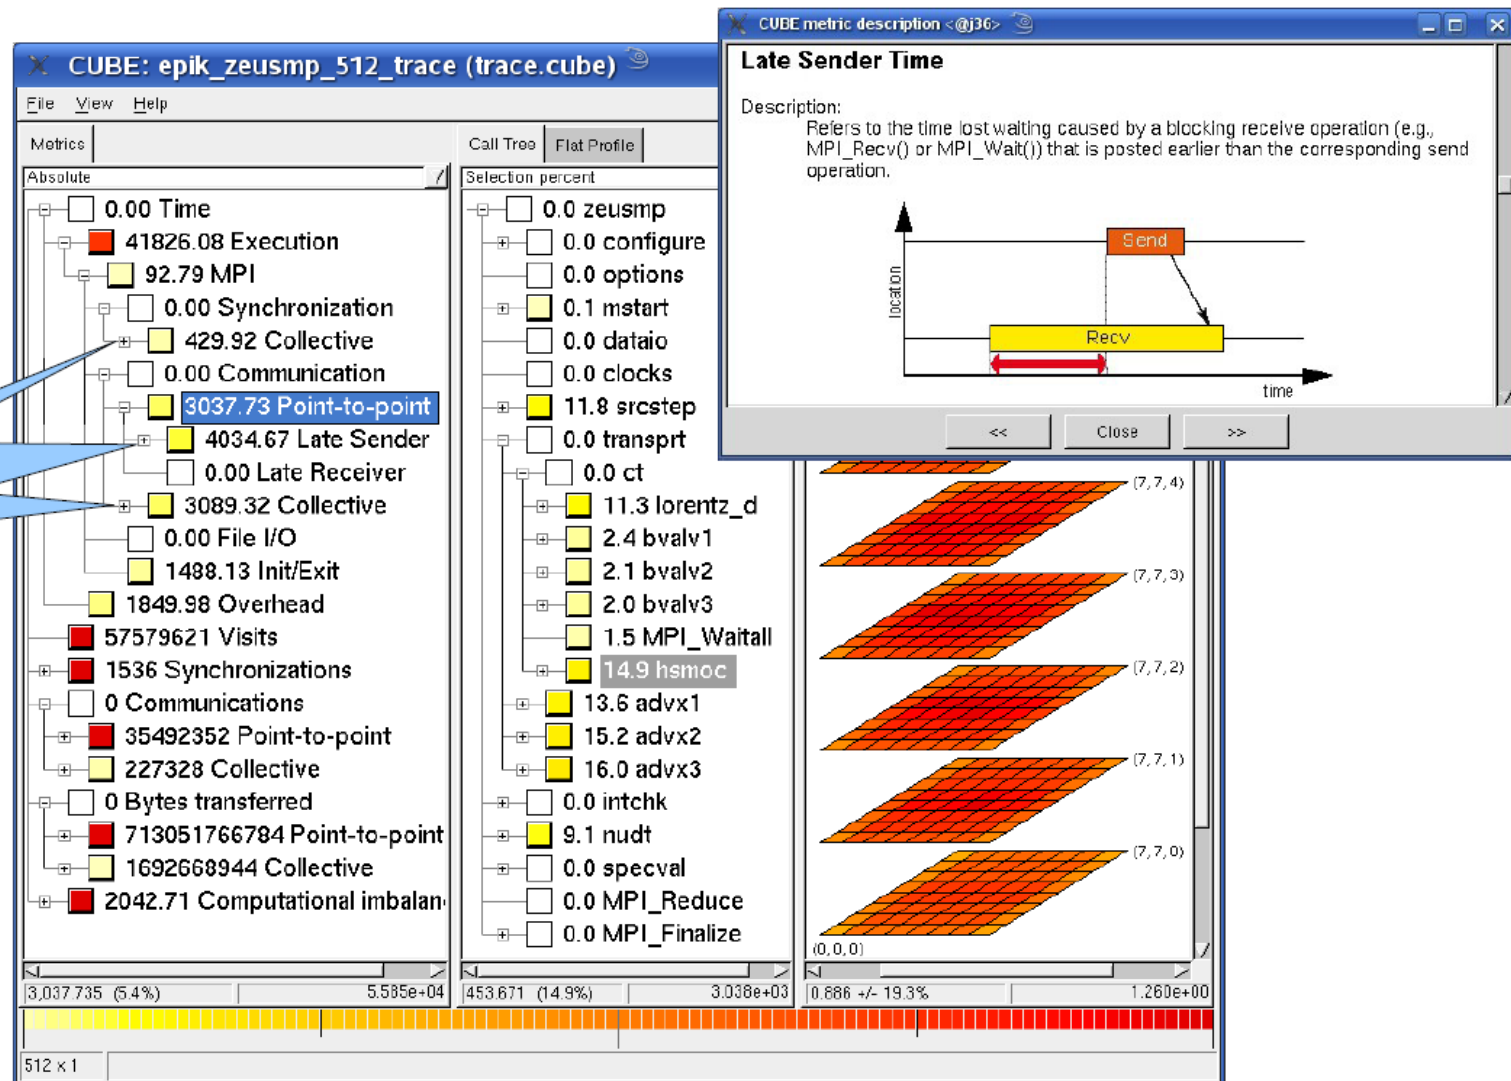
What kind of performance problem?

Where is it in the source code?  
In what context?

How is it distributed across the processes?

# Scalasca analysis report explorer (trace)

# VI-HPS



```
-bash-4.1$ scalasca
```

```
Scalasca 1.4.2
```

```
Toolset for scalable performance analysis of large-scale parallel applications
```

```
usage: scalasca [-v][-n] {action}
```

```
1. prepare application objects and executable for measurement:
```

```
scalasca -instrument <compile-or-link-command> # skin
```

```
2. run application under control of measurement system:
```

```
scalasca -analyze <application-launch-command> # scan
```

```
3. interactively explore measurement analysis report:
```

```
scalasca -examine <experiment-archive|report> # square
```

```
-v: enable verbose commentary
```

```
-n: show actions without taking them
```

```
-h: show quick reference guide (only)
```

```
-bash-4.1$ █
```



hobbit10.ee.bgu.ac.il - PuTTY



```
-bash-4.1$ scalasca -instrument mpicc -o cpi_scalasca ./cpi.c
INFO: Instrumented executable for MPI measurement
-bash-4.1$ ls -l cpi_scalasca
-rwxr-xr-x 1 tel-zur extlect 1326932 Apr 22 15:46 cpi_scalasca
-bash-4.1$ █
```

```
-bash-4.1$ scalasca -analyze mpirun -np 8 ./cpi_scalasca
S=C=A=N: Scalasca 1.4.2 runtime summarization
S=C=A=N: ./epik_cpi_scalasca_8_sum experiment archive
S=C=A=N: Mon Apr 22 15:47:14 2013: Collect start
/usr/local/bin/mpirun -np 8 ./cpi_scalasca
[00000]EPIK: Created new measurement archive ./epik_cpi_scalasca_8_sum
[00000]EPIK: Activated ./epik_cpi_scalasca_8_sum [NO TRACE] (0.013s)
[00000]EPIK: MPI-3.0 initialized 8 ranks
Process 0 on hobbit10.ee.bgu.ac.il
Process 1 on hobbit10.ee.bgu.ac.il
Process 2 on hobbit10.ee.bgu.ac.il
Process 4 on hobbit10.ee.bgu.ac.il
Process 3 on hobbit10.ee.bgu.ac.il
Process 5 on hobbit10.ee.bgu.ac.il
Process 6 on hobbit10.ee.bgu.ac.il
Process 7 on hobbit10.ee.bgu.ac.il
pi is approximately 3.1415926535898069, Error is 0.00000000000000138
wall clock time = 0.397143
[00000]EPIK: Closing experiment ./epik_cpi_scalasca_8_sum
[00000]EPIK: Largest definitions buffer 9296 bytes

=====
= BAD TERMINATION OF ONE OF YOUR APPLICATION PROCESSES
= EXIT CODE: 11
= CLEANING UP REMAINING PROCESSES
= YOU CAN IGNORE THE BELOW CLEANUP MESSAGES
=====

=====
YOUR APPLICATION TERMINATED WITH THE EXIT STRING: Segmentation fault (signal 11)
This typically refers to a problem with your application.
Please see the FAQ page for debugging suggestions
S=C=A=N: Mon Apr 22 15:47:15 2013: Collect done (status=0) 1s
Abort: incomplete experiment ./epik_cpi_scalasca_8_sum
-bash-4.1$
```



```
-bash-4.1$ ls -lt epik_cpi_scalasca_8_sum
total 8
-rw-r--r-- 1 tel-zur extlect 1334 Apr 22 15:47 epik.log
-rw-r--r-- 1 tel-zur extlect 1519 Apr 22 15:47 epik.conf
-rw-r--r-- 1 tel-zur extlect 0 Apr 22 15:47 epik.lock
-bash-4.1$ scalasca -examine ./epik_cpi_scalasca_8_sum
```

**That's it!**