

Parallel Processing

Xeon Phi

Guy Tel-Zur
4/1/2016

Agenda

- Introduction to Xeon Phi

<http://www.colfax-intl.com/nd/resources/Files/Stanford-Conference-2015-Scientific-Computing-with-Intel-Xeon-Phi-Coprocessors.pdf>

- Code modernization
 - Part 1: nbody
 - Part 2: Optimizations techniques

Tests were performed on `xphisrv2.cs.bgu.ac.il`

- Cpu 2x Intel(R) Xeon(R) CPU E5-2620 v2 @ 2.10GHz
- 6 cores each x 2 (hyperthreading), total 24 cores
- MIC: 5110P (x 2)



Intel® Xeon® Processor E5-2620 v2 (15M Cache, 2.10 GHz)

Specifications

Essentials

Performance

Memory Specifications

Expansion Options

Package Specifications

Advanced Technologies

Intel® Data Protection Technology

Intel® Platform Protection Technology

Compatible Products

Ordering / sSpecs / Steppings

Download Drivers

Specifications

Essentials

Status	Launched
Launch Date	Q3'13
Processor Number	E5-2620V2
Intel® Smart Cache	15 MB
Intel® QPI Speed	7.2 GT/s
# of QPI Links	2
Instruction Set	64-bit
Instruction Set Extensions	AVX
Embedded Options Available	No
Lithography	22 nm
Scalability	2S Only
VID Voltage Range	0.65–1.30V
Recommended Customer Price	BOX: \$410.00 TRAY: \$406.00
Datasheet	Link
Additional Information URL	Link

Performance

# of Cores	6
# of Threads	12
Processor Base Frequency	2.1 GHz
Max Turbo Frequency	2.6 GHz
TDP	80 W

Compare



Related Products

Intel® Xeon® Processor E5
v2 Family

Intel® Xeon® Processor
E5-2600 v2 Product
Family

Products formerly Ivy
Bridge EP

Quick Links

[Export Full Specifications >](#)

[Additional Information >](#)

[Search Distributors >](#)

[Support Overview >](#)


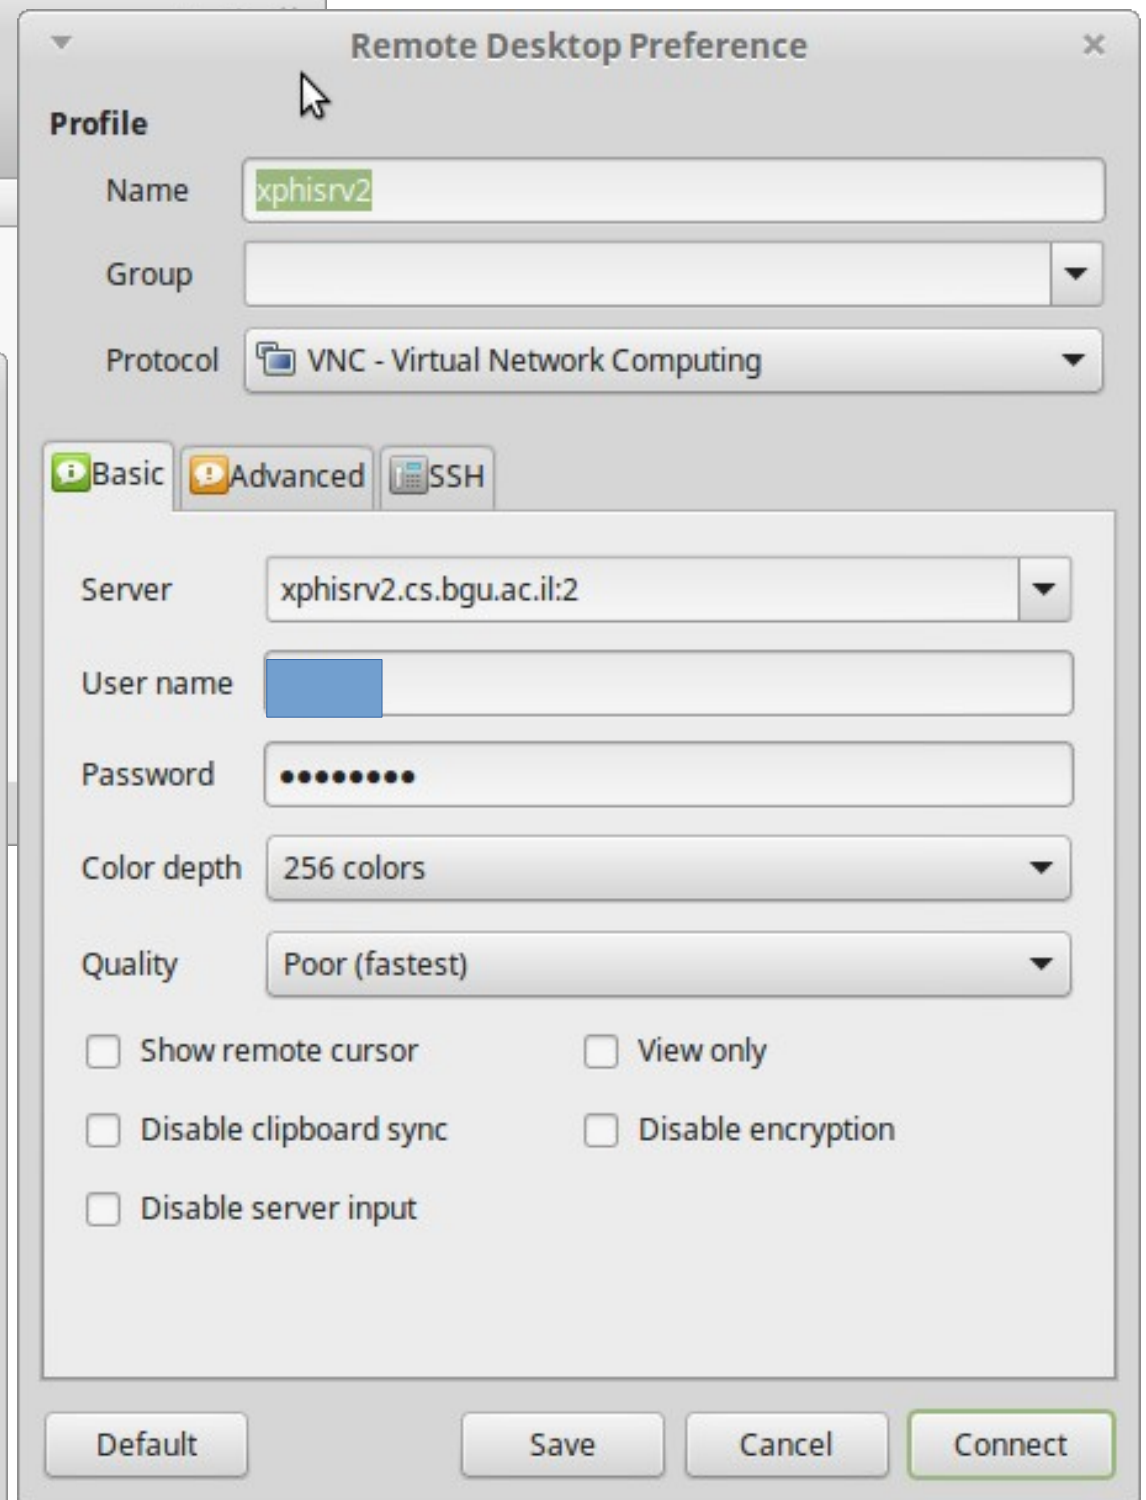
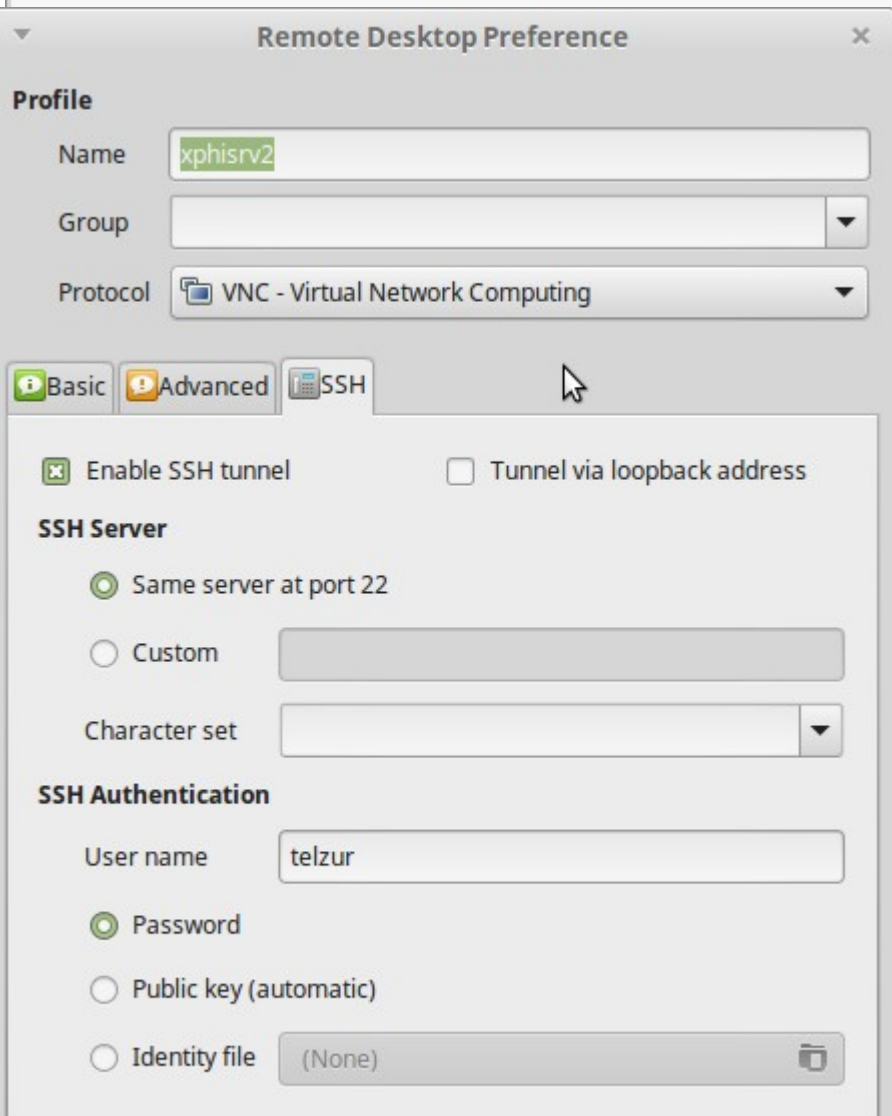

[Search all of intel.com >](#)

PCN/MDDS Information

SR1AN

930059: [PCN](#) | [MDDS](#)

931259: [PCN](#) | [MDDS](#)



Part 1: nbody

Colfax material:

<http://colfaxresearch.com/test-driving-intel-xeon-phi-coprocessors-with-a-basic-n-body-simulation/>

Paper:

<http://colfaxresearch.com/download/530/>

Example code is based on:

mini-nbody: A simple N-body Code

<https://github.com/harrism/mini-nbody>

By Mark Harris

GitHub

Search GitHub



Mark Harris

harrism



NVIDIA



Larnook, NSW Australia



<http://www.markmark.net>

nbody

- Unzip mini-nbody-master
- Baseline: put into comment // rate
- icc compile: `icc -std=c99 -qopenmp -o nbody ./nbody.c -lm`
- gcc compile: `gcc -std=c99 -fopenmp -o nbody_gcc ./nbody.c -lm`
- Working dir on xphisrv2:
`/home/telzur/Teaching/PP/XeonPhi`

Baseline version

- `export OMP_NUM_THREADS=24`

Tests on the CPUs

gcc versus icc (no optimizations)

```
[telzur@xphisrv2 mini-nbody-master]$ time ./nbody_gcc  
Iteration 1: 2.005 seconds  
Iteration 2: 1.991 seconds  
Iteration 3: 1.990 seconds  
Iteration 4: 1.990 seconds  
Iteration 5: 1.990 seconds  
Iteration 6: 1.990 seconds  
Iteration 7: 1.990 seconds  
Iteration 8: 1.989 seconds  
Iteration 9: 1.990 seconds  
Iteration 10: 1.991 seconds  
30000 Bodies: average 0.452 Billion Interactions / second
```

```
real 0m19.927s  
user 6m18.485s  
sys 0m0.010s
```

```
[telzur@xphisrv2 mini-nbody-master]$ time ./nbody_icc  
Iteration 1: 0.230 seconds  
Iteration 2: 0.202 seconds  
Iteration 3: 0.201 seconds  
Iteration 4: 0.202 seconds  
Iteration 5: 0.206 seconds  
Iteration 6: 0.202 seconds  
Iteration 7: 0.201 seconds  
Iteration 8: 0.201 seconds  
Iteration 9: 0.201 seconds  
Iteration 10: 0.204 seconds  
30000 Bodies: average 4.450 Billion Interactions / second
```

```
real 0m2.058s  
user 0m40.860s  
sys 0m0.026s
```

Test on the mic: nbody-orig

24 threads

```
telzur@xphisrv2-mic0:~/nbody$ time ./nbody-orig
```

Iteration 1: 0.733 seconds

Iteration 2: 0.691 seconds

Iteration 3: 0.691 seconds

Iteration 4: 0.688 seconds

Iteration 5: 0.686 seconds

Iteration 6: 0.689 seconds

Iteration 7: 0.689 seconds

Iteration 8: 0.693 seconds

Iteration 9: 0.689 seconds

Iteration 10: 0.693 seconds

30000 Bodies: average 1.304 Billion Interactions / second

real 0m6.980s

user 2m45.660s

sys 0m0.060s

240 threads

```
telzur@xphisrv2-mic0:~/nbody$ time ./nbody-orig
```

Iteration 1: 0.357 seconds

Iteration 2: 0.098 seconds

Iteration 3: 0.099 seconds

Iteration 4: 0.100 seconds

Iteration 5: 0.100 seconds

Iteration 6: 0.100 seconds

Iteration 7: 0.100 seconds

Iteration 8: 0.100 seconds

Iteration 9: 0.100 seconds

Iteration 10: 0.100 seconds

30000 Bodies: average 9.050 Billion Interactions / second

real 0m1.343s

user 3m59.170s

sys 0m2.300s

mic0, /home/telzur/nbody

Test on the mic: nbody-soa

24 threads

```
telzur@xphisrv2-mic0:~/nbody$ time ./nbody-soa
```

Iteration 1: 0.191 seconds

Iteration 2: 0.145 seconds

Iteration 3: 0.145 seconds

Iteration 4: 0.145 seconds

Iteration 5: 0.145 seconds

Iteration 6: 0.145 seconds

Iteration 7: 0.145 seconds

Iteration 8: 0.145 seconds

Iteration 9: 0.145 seconds

Iteration 10: 0.145 seconds

30000 Bodies: average 6.217 Billion Interactions / second

real 0m1.532s

user 0m34.830s

sys 0m0.050s

240 threads

```
telzur@xphisrv2-mic0:~/nbody$ time ./nbody-soa
```

Iteration 1: 0.274 seconds

Iteration 2: 0.026 seconds

Iteration 3: 0.027 seconds

Iteration 4: 0.027 seconds

Iteration 5: 0.027 seconds

Iteration 6: 0.027 seconds

Iteration 7: 0.027 seconds

Iteration 8: 0.026 seconds

Iteration 9: 0.027 seconds

Iteration 10: 0.027 seconds

30000 Bodies: average 33.747 Billion Interactions / second

real 0m0.607s

user 1m4.960s

sys 0m2.380s

diff ./nbody.c ./mic/nbody-soa.c

```
XPHT> diff ./nbody.c ./mic/nbody-soa.c
5a6
>
8c9
< typedef struct { float x, y, z, vx, vy, vz; } Body;
---
> typedef struct { float *x, *y, *z, *vx, *vy, *vz; } BodySystem;
16c17,18
< void bodyForce(Body *p, float dt, int n) {
---
>
> void bodyForce(BodySystem p, float dt, int n) {
22,24c24,26
<     float dx = p[j].x - p[i].x;
<     float dy = p[j].y - p[i].y;
<     float dz = p[j].z - p[i].z;
---
>     float dy = p.y[j] - p.y[i];
>     float dz = p.z[j] - p.z[i];
>     float dx = p.x[j] - p.x[i];
32c34
<     p[i].vx += dt*Fx; p[i].vy += dt*Fy; p[i].vz += dt*Fz;
---
>     p.vx[i] += dt*Fx; p.vy[i] += dt*Fy; p.vz[i] += dt*Fz;
44c46
<     int bytes = nBodies*sizeof(Body);
---
>     int bytes = 6*nBodies*sizeof(float);
46c48,50
<     Body *p = (Body*)buf;
---
>     BodySystem p;
>     p.x = buf+0*nBodies; p.y = buf+1*nBodies; p.z = buf+2*nBodies;
>     p.vx = buf+3*nBodies; p.vy = buf+4*nBodies; p.vz = buf+5*nBodies;
58,60c62,64
<     p[i].x += p[i].vx*dt;
<     p[i].y += p[i].vy*dt;
<     p[i].z += p[i].vz*dt;
---
>     p.x[i] += p.vx[i]*dt;
>     p.y[i] += p.vy[i]*dt;
>     p.z[i] += p.vz[i]*dt;
XPHT>
```

Test on the mic: nbody-block

24 threads

```
telzur@xphisrv2-mic0:~/nbody$ time ./nbody-block
```

Iteration 1: 0.196 seconds

Iteration 2: 0.152 seconds

Iteration 3: 0.152 seconds

Iteration 4: 0.152 seconds

Iteration 5: 0.152 seconds

Iteration 6: 0.152 seconds

Iteration 7: 0.152 seconds

Iteration 8: 0.152 seconds

Iteration 9: 0.152 seconds

Iteration 10: 0.152 seconds

30000 Bodies: average 5.923 Billion Interactions / second

real 0m1.601s

user 0m36.540s

sys 0m0.070s

240 threads

```
telzur@xphisrv2-mic0:~/nbody$ time ./nbody-block
```

Iteration 1: 0.280 seconds

Iteration 2: 0.027 seconds

Iteration 3: 0.027 seconds

Iteration 4: 0.027 seconds

Iteration 5: 0.027 seconds

Iteration 6: 0.027 seconds

Iteration 7: 0.027 seconds

Iteration 8: 0.027 seconds

Iteration 9: 0.027 seconds

Iteration 10: 0.027 seconds

30000 Bodies: average 33.335 Billion Interactions / second

real 0m0.614s

user 1m5.250s

sys 0m2.220s

```

XPHI>
XPHI> diff ./mic/nbody-soa.c ./mic/nbody-block.c
6c6
<
---
> #define CACHELINE 64 // size of cache line [bytes]
18,31c18,38
< void bodyForce(BodySystem p, float dt, int n) {
<   #pragma omp parallel for schedule(dynamic)
<   for (int i = 0; i < n; i++) {
<     float Fx = 0.0f; float Fy = 0.0f; float Fz = 0.0f;
<
<     for (int j = 0; j < n; j++) {
<       float dy = p.y[j] - p.y[i];
<       float dz = p.z[j] - p.z[i];
<       float dx = p.x[j] - p.x[i];
<       float distSqr = dx*dx + dy*dy + dz*dz + SOFTENING;
<       float invDist = 1.0f / sqrtf(distSqr);
<       float invDist3 = invDist * invDist * invDist;
<
<       Fx += dx * invDist3; Fy += dy * invDist3; Fz += dz * invDist3;
<
---
> void bodyForce(BodySystem p, float dt, int n, int tileSize) {
>   for (int tile = 0; tile < n; tile += tileSize) {
>     int to = tile + tileSize;
>     if (to > n) to = n;
>
>     #pragma omp parallel for schedule(dynamic)
>     for (int i = 0; i < n; i++) {
>       float Fx = 0.0f; float Fy = 0.0f; float Fz = 0.0f;
>
>       for (int j = tile; j < to; j++) {
>         float dy = p.y[j] - p.y[i];
>         float dz = p.z[j] - p.z[i];
>         float dx = p.x[j] - p.x[i];
>         float distSqr = dx*dx + dy*dy + dz*dz + SOFTENING;
>         float invDist = 1.0f / sqrtf(distSqr);
>         float invDist3 = invDist * invDist * invDist;
>
>         Fx += dx * invDist3; Fy += dy * invDist3; Fz += dz * invDist3;
>       }
>
>     p.vx[i] += dt*Fx; p.vy[i] += dt*Fy; p.vz[i] += dt*Fz;
33,34d39
<

```

Test on the mic: nbody-align

24 threads

```
telzur@xphisrv2-mic0:~/nbody$ time ./nbody-align
```

Iteration 1: 0.170 seconds

Iteration 2: 0.125 seconds

Iteration 3: 0.126 seconds

Iteration 4: 0.125 seconds

Iteration 5: 0.125 seconds

Iteration 6: 0.125 seconds

Iteration 7: 0.126 seconds

Iteration 8: 0.125 seconds

Iteration 9: 0.125 seconds

Iteration 10: 0.126 seconds

30000 Bodies: average 7.175 Billion Interactions / second

real 0m1.338s

user 0m30.210s

sys 0m0.040s

240 threads

```
telzur@xphisrv2-mic0:~/nbody$ time ./nbody-align
```

Iteration 1: 0.275 seconds

Iteration 2: 0.023 seconds

Iteration 3: 0.023 seconds

Iteration 4: 0.024 seconds

Iteration 5: 0.023 seconds

Iteration 6: 0.023 seconds

Iteration 7: 0.024 seconds

Iteration 8: 0.023 seconds

Iteration 9: 0.023 seconds

Iteration 10: 0.024 seconds

30000 Bodies: average 38.823 Billion Interactions / second

real 0m0.568s

user 0m57.530s

sys 0m0.640s

```
XPHI> diff ./mic/nbody-block.c ./mic/nbody-align.c
```

```
26a27,28
```

```
>     #pragma vector aligned
```

```
>     #pragma simd
```

```
53a56,60
```

```
> if ( tileSize % (CACHELINE/sizeof(float)) ) {
```

```
>     printf("ERROR: blockSize not multiple of %d vector elements\n", CACHELINE/(int)sizeof(float));
```

```
>     exit(1);
```

```
> }
```

```
>
```

```
55c62
```

```
< float *buf = (float*)malloc(bytes);
```

```
---
```

```
> float *buf = (float*)_mm_malloc(bytes, CACHELINE);
```

```
92c99
```

```
< free(buf);
```

```
---
```

```
> _mm_free(buf);
```

```
XPHI> █
```




xphisrv2 x

Applications Places System

17 °C Mon Jan 4, 12:08 Guy Tel-Zu

Intel® Xeon Phi™ Coprocessor Platform Control Panel

Cards Advanced

Cards: 2 Memory: 15.03 GB Cores: 120

Utilization View (All Devices)

System [%]
User [%]

64

Average Core
Temperature °C

1.21 GB

Total
Memory Usage


217 Watts

Total
Power Usage

mic0: Core Histogram View



Individual Core Usage

System [%]
User [%]

Part 2:

Materials thanks to: **Colfax Research**

Optimization Techniques for the Intel MIC Architecture

- <http://colfaxresearch.com/optimization-techniques-for-the-intel-mic-architecture-part-1-of-3-multi-threading-and-parallel-reduction/>
- <http://colfaxresearch.com/optimization-techniques-for-the-intel-mic-architecture-part-2-of-3-strip-mining-for-vectorization/>
- <http://colfaxresearch.com/optimization-techniques-for-the-intel-mic-architecture-part-3-of-3-false-sharing-and-padding/>
- <http://colfaxresearch.com/download/1233/>
- <http://colfaxresearch.com/download/1241/>



September 08, 2015

GUY TEL-ZUR

has successfully completed

PARALLEL PROGRAMMING AND OPTIMIZATION FOR INTEL® ARCHITECTURE

The HOW (Hands-on Workshop) series, a 20-hour Web-based training program with hands-on component on best practices for programming Intel® Xeon® processors and Intel® Xeon Phi™ coprocessors.

A handwritten signature in black ink, reading "Gautam Shah".

Gautam Shah
President and CEO, Colfax International

- | | |
|-------------|---|
| Fundamental | - issued to students who attended training for courses that do not have post-training exams |
| Craftsman | - issued to students who attended training and successfully passed post-training exams |
| Trainer | - issued to students who attended special courses and passed tests that qualify the attendee to conduct the named Colfax Developer Training program |

* Intel, Intel Xeon, and Intel Xeon Phi are trademarks of Intel Corporation in the U.S. and other countries.

LEVEL OF ACCOMPLISHMENT

FUNDAMENTAL



Colfax International | www.colfax-intl.com | CA, USA

```
[telzur@xphisrv2 Colfax_Tutorial_Binning]$ ./app-CPU
```

Particle Binning Optimization Demo (single precision)

Additional information is available in accompanying papers at <http://colfaxresearch.com/>

(c) Colfax International, 2015.

Initialization... done in 0.661 seconds.

Computing reference result... done in 2.457 seconds.

Benchmarking...

=====

IMPLEMENTATION 0: serial

Trial	Time, s	Speedup	GP/s *
1	2.432e+00	n/a	5.52e-02 **
2	2.415e+00	n/a	5.56e-02 **
3	2.414e+00	n/a	5.56e-02
4	2.412e+00	n/a	5.57e-02
5	2.406e+00	n/a	5.58e-02
6	2.407e+00	n/a	5.58e-02
7	2.400e+00	n/a	5.59e-02
8	2.412e+00	n/a	5.57e-02
9	2.406e+00	n/a	5.58e-02
10	2.414e+00	n/a	5.56e-02

Average performance: 1.00 5.57e-02 +- 1.05e-04 GP/s

* - Performance unit 1 GP/s is 10^9 particles binned per second.

** - warm-up, not included in average

=====

=====

IMPLEMENTATION 1: parallel, atomics for reduction

Trial	Time, s	Speedup	GP/s *
1	3.213e+00	0.75	4.18e-02 **
2	3.016e+00	0.80	4.45e-02 **
3	2.984e+00	0.81	4.50e-02
4	2.983e+00	0.81	4.50e-02
5	2.984e+00	0.81	4.50e-02
6	2.984e+00	0.81	4.50e-02
7	2.986e+00	0.81	4.50e-02
8	2.985e+00	0.81	4.50e-02
9	2.984e+00	0.81	4.50e-02
10	2.985e+00	0.81	4.50e-02

Average performance: 0.81 4.50e-02 +- 1.40e-05 GP/s

* - Performance unit 1 GP/s is 10^9 particles binned per second.

** - warm-up, not included in average

=====

IMPLEMENTATION 2: parallel, thread-private containers for reduction

Trial	Time, s	Speedup	GP/s *
1	1.852e-01	13.01	7.25e-01 **
2	1.852e-01	13.01	7.25e-01 **
3	1.853e-01	13.00	7.24e-01
4	1.851e-01	13.01	7.25e-01
5	1.852e-01	13.01	7.25e-01
6	1.852e-01	13.01	7.25e-01
7	1.852e-01	13.01	7.25e-01
8	1.851e-01	13.01	7.25e-01
9	1.852e-01	13.00	7.25e-01
10	1.855e-01	12.98	7.23e-01

Average performance: 13.01 7.25e-01 +- 4.73e-04 GP/s

* - Performance unit 1 GP/s is 10^9 particles binned per second.

** - warm-up, not included in average

```
=====
=====
IMPLEMENTATION 3: vectorized

Trial    Time, s    Speedup    GP/s *
  1  8.954e-02    26.90    1.50e+00 **
  2  8.946e-02    26.93    1.50e+00 **
  3  8.947e-02    26.92    1.50e+00
  4  8.950e-02    26.91    1.50e+00
  5  8.953e-02    26.90    1.50e+00
  6  8.955e-02    26.90    1.50e+00
  7  8.949e-02    26.92    1.50e+00
  8  8.948e-02    26.92    1.50e+00
  9  8.948e-02    26.92    1.50e+00
 10  8.949e-02    26.92    1.50e+00

-----
Average performance:  26.91  1.50e+00 +- 4.15e-04 GP/s
-----

* - Performance unit 1 GP/s is 10^9 particles binned per second.
** - warm-up, not included in average

=====
=====
IMPLEMENTATION 4: vectorized, aligned

Trial    Time, s    Speedup    GP/s *
  1  5.861e-02    41.10    2.29e+00 **
  2  5.848e-02    41.19    2.30e+00 **
  3  5.851e-02    41.17    2.29e+00
  4  5.851e-02    41.17    2.29e+00
  5  5.852e-02    41.16    2.29e+00
  6  5.854e-02    41.15    2.29e+00
  7  5.867e-02    41.06    2.29e+00
  8  5.850e-02    41.18    2.29e+00
  9  5.850e-02    41.18    2.29e+00
 10  5.854e-02    41.15    2.29e+00

-----
Average performance:  41.15  2.29e+00 +- 2.07e-03 GP/s
-----

* - Performance unit 1 GP/s is 10^9 particles binned per second.
** - warm-up, not included in average
```

```
=====
IMPLEMENTATION 5: global container for reduction in threads-first layout
```

Trial	Time, s	Speedup	GP/s *
1	1.194e+00	2.02	1.12e-01 **
2	1.193e+00	2.02	1.12e-01 **
3	1.194e+00	2.02	1.12e-01
4	1.194e+00	2.02	1.12e-01
5	1.194e+00	2.02	1.12e-01
6	1.193e+00	2.02	1.12e-01
7	1.193e+00	2.02	1.12e-01
8	1.194e+00	2.02	1.12e-01
9	1.196e+00	2.01	1.12e-01
10	1.193e+00	2.02	1.12e-01

```
-----
Average performance: 2.02 1.12e-01 +- 6.35e-05 GP/s
-----
```

```
* - Performance unit 1 GP/s is 10^9 particles binned per second.
```

```
** - warm-up, not included in average
```

```
=====
IMPLEMENTATION 6: global container for reduction in threads-last layout
```

Trial	Time, s	Speedup	GP/s *
1	1.464e-01	16.45	9.17e-01 **
2	1.468e-01	16.41	9.14e-01 **
3	1.465e-01	16.44	9.16e-01
4	1.468e-01	16.41	9.15e-01
5	1.467e-01	16.42	9.15e-01
6	1.459e-01	16.51	9.20e-01
7	1.464e-01	16.45	9.17e-01
8	1.464e-01	16.46	9.17e-01
9	1.471e-01	16.38	9.12e-01
10	1.462e-01	16.47	9.18e-01

```
-----
Average performance: 16.44 9.16e-01 +- 2.11e-03 GP/s
-----
```

```
* - Performance unit 1 GP/s is 10^9 particles binned per second.
```

```
** - warm-up, not included in average
```


=====

IMPLEMENTATION 6: global container for reduction in threads-last layout

Trial	Time, s	Speedup	GP/s *
1	1.464e-01	16.45	9.17e-01 **
2	1.468e-01	16.41	9.14e-01 **
3	1.465e-01	16.44	9.16e-01
4	1.468e-01	16.41	9.15e-01
5	1.467e-01	16.42	9.15e-01
6	1.459e-01	16.51	9.20e-01
7	1.464e-01	16.45	9.17e-01
8	1.464e-01	16.46	9.17e-01
9	1.471e-01	16.38	9.12e-01
10	1.462e-01	16.47	9.18e-01

Average performance: 16.44 **9.16e-01 +- 2.11e-03 GP/s**

* - Performance unit 1 GP/s is 10^9 particles binned per second.

** - warm-up, not included in average

=====

=====

IMPLEMENTATION 7: global container for reduction in threads-last layout with padding

Trial	Time, s	Speedup	GP/s *
1	6.128e-02	39.31	2.19e+00 **
2	6.123e-02	39.34	2.19e+00 **
3	6.122e-02	39.35	2.19e+00
4	6.122e-02	39.35	2.19e+00
5	6.122e-02	39.35	2.19e+00
6	6.127e-02	39.32	2.19e+00
7	6.123e-02	39.34	2.19e+00
8	6.126e-02	39.32	2.19e+00
9	6.125e-02	39.33	2.19e+00
10	6.126e-02	39.32	2.19e+00

Average performance: 39.33 **2.19e+00 +- 7.39e-04 GP/s**

* - Performance unit 1 GP/s is 10^9 particles binned per second.

** - warm-up, not included in average

[telzur@xphisrv2 Colfax_Tutorial_Binning]\$


Same tests on the mic

cp the application to mic0:/home/telzur/colfax


Cp the libraries:

```
[telzur@xphisrv2 Colfax_Tutorial_Binning]$ scp /opt/intel/composer_xe_2015.3.187/mkl/lib/mic/*  
mic0:/home/telzur/lib/.
```

Implementation 0: serial



Implementations 1..7: parallel



```
telzur@xphisrv2:~/Teaching/PP/XeonPhi/mini-nbody-master
File Edit View Search Terminal Help
telzur@xphisrv2-mic0:~/colfax$ ./app-MIC

Particle Binning Optimization Demo (single precision)
Additional information is available in accompanying papers at http://colfaxresearch.com/

(c) Colfax International, 2015.

Initialization... done in 1.748 seconds.
Computing reference result... done in 28.269 seconds.
Benchmarking...

=====
IMPLEMENTATION 0: serial

Trial   Time, s   Speedup   GP/s *
  1  2.892e+01   n/a  4.64e-03 **
  2  2.892e+01   n/a  4.64e-03 **
  3  2.892e+01   n/a  4.64e-03
  4  2.892e+01   n/a  4.64e-03
  5  2.892e+01   n/a  4.64e-03
  6  2.892e+01   n/a  4.64e-03
  7  2.892e+01   n/a  4.64e-03
  8  2.892e+01   n/a  4.64e-03
  9  2.892e+01   n/a  4.64e-03
 10  2.892e+01   n/a  4.64e-03

-----
Average performance:   1.00  4.64e-03 +- 6.12e-08 GP/s
-----

* - Performance unit 1 GP/s is 10^9 particles binned per second.
** - warm-up, not included in average

=====
IMPLEMENTATION 1: parallel, atomics for reduction

Trial   Time, s   Speedup   GP/s *
  1  3.561e+00   8.12  3.77e-02 **
  2  3.558e+00   8.13  3.77e-02 **
  3  3.561e+00   8.12  3.77e-02
  4  3.558e+00   8.13  3.77e-02
  5  3.558e+00   8.13  3.77e-02
  6  3.558e+00   8.13  3.77e-02
  7  3.558e+00   8.13  3.77e-02
  8  3.561e+00   8.12  3.77e-02
  9  3.560e+00   8.12  3.77e-02
 10  3.560e+00   8.13  3.77e-02

-----
Average performance:   8.13  3.77e-02 +- 1.25e-05 GP/s
-----

* - Performance unit 1 GP/s is 10^9 particles binned per second.
```

=====

IMPLEMENTATION 2: parallel, thread-private containers for reduction

Trial	Time, s	Speedup	GP/s *
1	4.457e-01	64.90	3.01e-01 **
2	4.438e-01	65.18	3.02e-01 **
3	4.444e-01	65.09	3.02e-01
4	4.463e-01	64.80	3.01e-01
5	4.454e-01	64.94	3.01e-01
6	4.499e-01	64.28	2.98e-01
7	4.449e-01	65.01	3.02e-01
8	4.450e-01	65.00	3.02e-01
9	4.442e-01	65.11	3.02e-01
10	4.453e-01	64.95	3.01e-01

Average performance: 64.90 3.01e-01 +- 1.15e-03 GP/s

* - Performance unit 1 GP/s is 10^9 particles binned per second.

** - warm-up, not included in average

=====

IMPLEMENTATION 3: vectorized

Trial	Time, s	Speedup	GP/s *
1	4.975e-02	581.43	2.70e+00 **
2	5.737e-02	504.20	2.34e+00 **
3	3.303e-02	875.68	4.06e+00
4	3.262e-02	886.55	4.11e+00
5	3.243e-02	891.94	4.14e+00
6	3.306e-02	874.81	4.06e+00
7	3.243e-02	891.91	4.14e+00
8	3.267e-02	885.43	4.11e+00
9	3.249e-02	890.34	4.13e+00
10	3.302e-02	875.94	4.06e+00

Average performance: 884.08 4.10e+00 +- 3.25e-02 GP/s

* - Performance unit 1 GP/s is 10^9 particles binned per second.

** - warm-up, not included in average

```
telzur@xphisrv2:~/Teaching/PP/XeonPhi/mini-nbody-master
File Edit View Search Terminal Help

=====
IMPLEMENTATION 4: vectorized, aligned

Trial    Time, s    Speedup    GP/s *
  1  3.467e-02    834.19    3.87e+00 **
  2  2.911e-02    993.54    4.61e+00 **
  3  2.899e-02    997.71    4.63e+00
  4  2.911e-02    993.70    4.61e+00
  5  2.902e-02    996.79    4.63e+00
  6  2.911e-02    993.53    4.61e+00
  7  2.920e-02    990.37    4.60e+00
  8  2.917e-02    991.46    4.60e+00
  9  2.888e-02   1001.48    4.65e+00
 10  2.856e-02   1012.77    4.70e+00

-----
Average performance: 997.23  4.63e+00 +- 3.14e-02 GP/s
-----
* - Performance unit 1 GP/s is 10^9 particles binned per second.
** - warm-up, not included in average

=====
IMPLEMENTATION 5: global container for reduction in threads-first layout

Trial    Time, s    Speedup    GP/s *
  1  5.467e-01    52.91    2.46e-01 **
  2  5.282e-01    54.76    2.54e-01 **
  3  5.283e-01    54.74    2.54e-01
  4  5.281e-01    54.77    2.54e-01
  5  5.292e-01    54.66    2.54e-01
  6  5.276e-01    54.82    2.54e-01
  7  5.280e-01    54.78    2.54e-01
  8  5.279e-01    54.79    2.54e-01
  9  5.280e-01    54.78    2.54e-01
 10  5.294e-01    54.63    2.54e-01

-----
Average performance:  54.75  2.54e-01 +- 2.86e-04 GP/s
-----
* - Performance unit 1 GP/s is 10^9 particles binned per second.
** - warm-up, not included in average

=====
```

```
telzur@xphisrv2:~/Teaching/PP/XeonPhi/mini-nbody-master
File Edit View Search Terminal Help

=====
IMPLEMENTATION 6: global container for reduction in threads-last layout


Trial    Time, s    Speedup    GP/s *
  1  5.146e-02    562.07    2.61e+00 **
  2  3.995e-02    724.01    3.36e+00 **
  3  4.076e-02    709.54    3.29e+00
  4  4.049e-02    714.34    3.31e+00
  5  4.017e-02    720.00    3.34e+00
  6  4.080e-02    708.95    3.29e+00
  7  4.028e-02    718.03    3.33e+00
  8  4.070e-02    710.66    3.30e+00
  9  4.059e-02    712.51    3.31e+00
 10  4.053e-02    713.60    3.31e+00
-----
Average performance: 713.45  3.31e+00 +- 1.71e-02 GP/s
-----
* - Performance unit 1 GP/s is 10^9 particles binned per second.
** - warm-up, not included in average

=====
IMPLEMENTATION 7: global container for reduction in threads-last layout with padding

Trial    Time, s    Speedup    GP/s *
  1  2.884e-02   1002.80    4.65e+00 **
  2  2.913e-02    992.89    4.61e+00 **
  3  2.912e-02    993.19    4.61e+00
  4  2.915e-02    992.17    4.60e+00
  5  2.915e-02    992.20    4.60e+00
  6  2.915e-02    992.20    4.60e+00
  7  2.910e-02    993.94    4.61e+00
  8  4.611e-02    627.22    2.91e+00
  9  3.049e-02    948.62    4.40e+00
 10  2.916e-02    992.03    4.60e+00
-----
Average performance: 941.45  4.37e+00 +- 5.55e-01 GP/s
-----
* - Performance unit 1 GP/s is 10^9 particles binned per second.
** - warm-up, not included in average

telzur@xphisrv2-mic0:~/colfax$
```

KMP_AFFINITY



That's it