

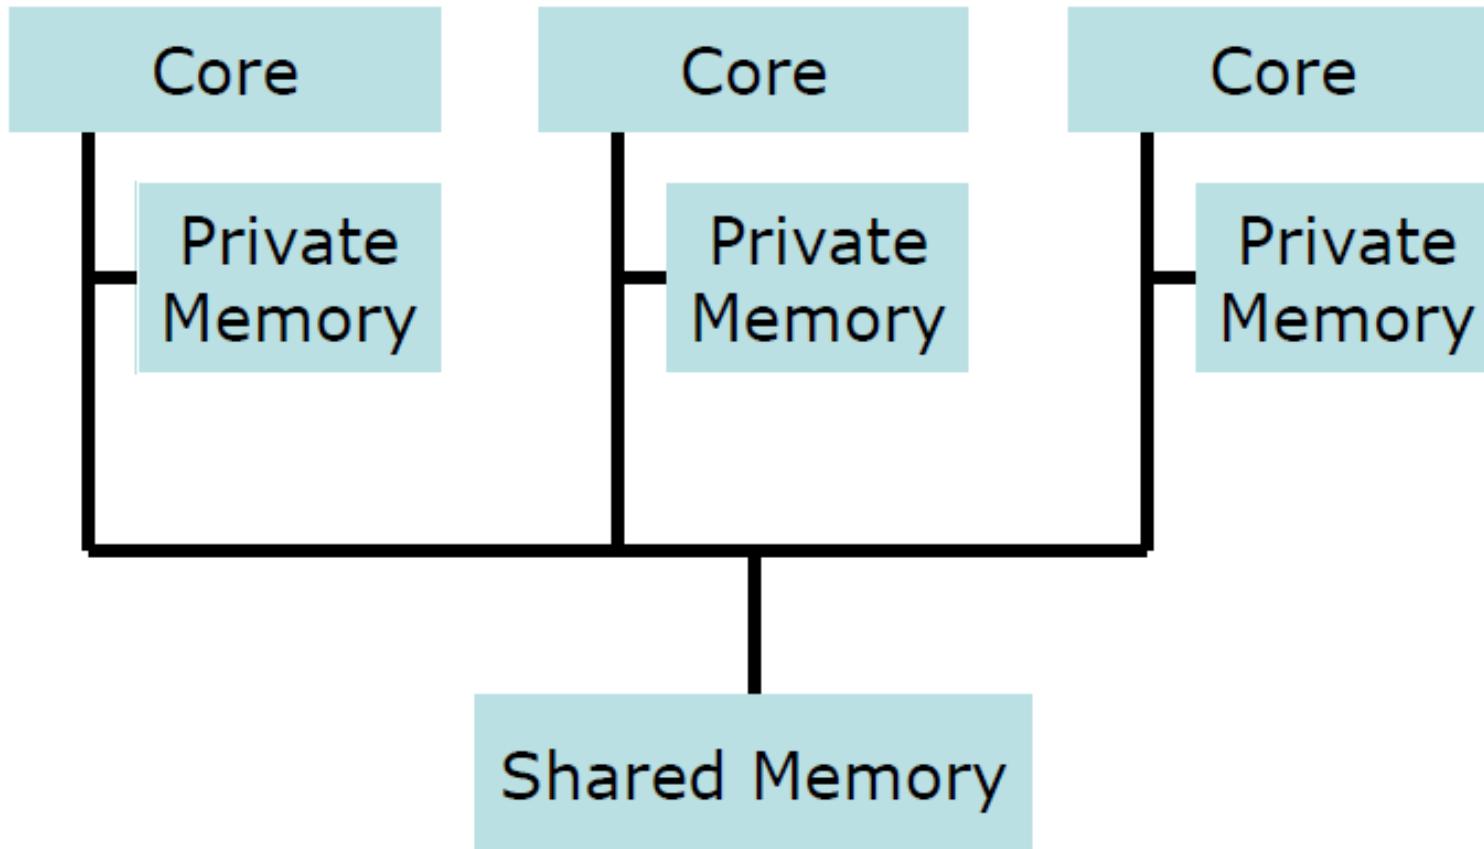
Shared Memory and OpenMP

Selected Slides from Intel's
INTEL® DEVELOPER ZONE:

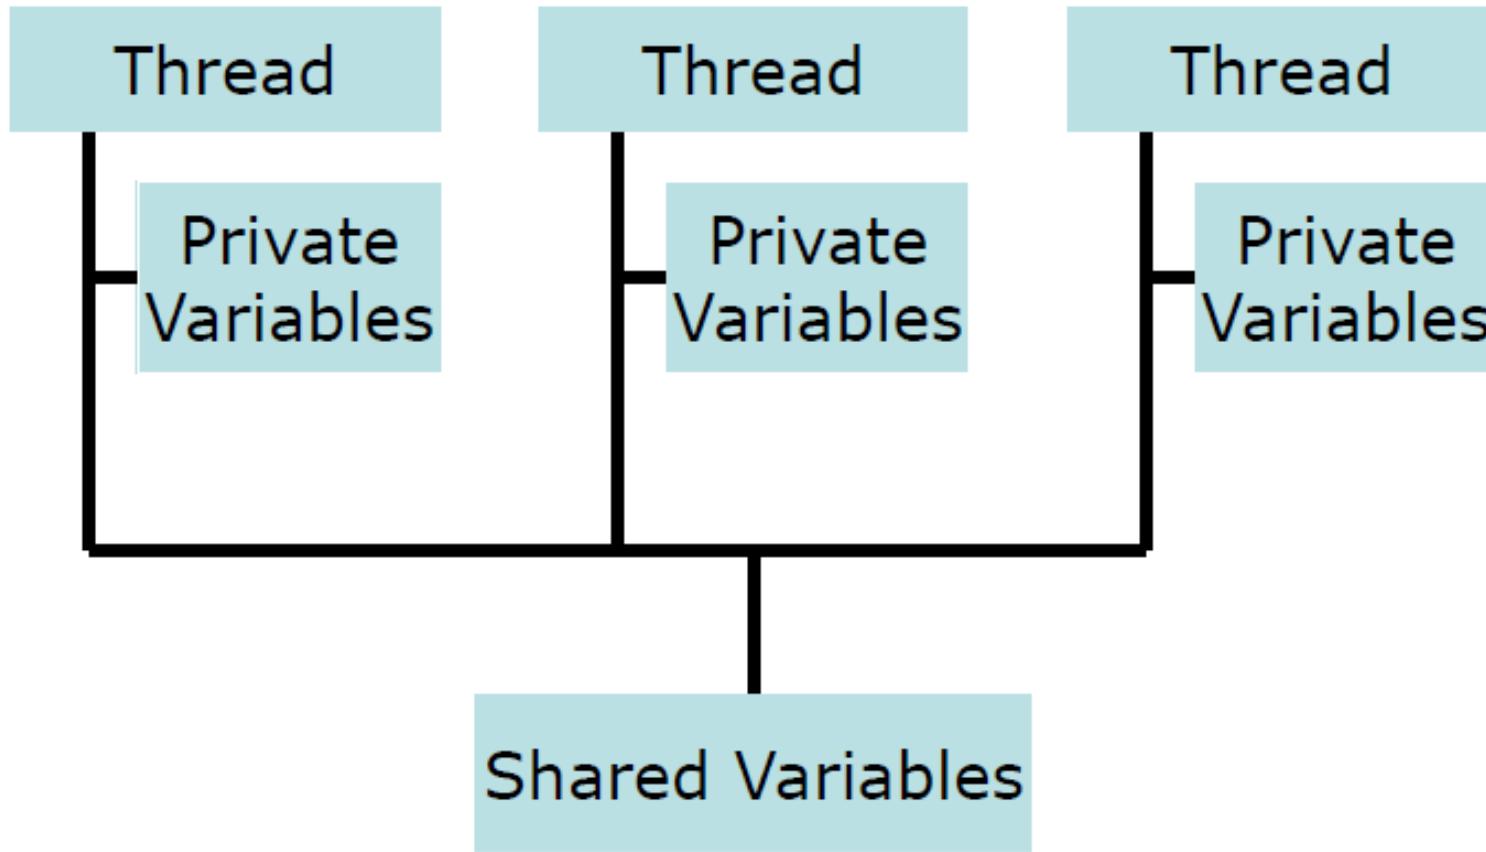
Courseware - Parallel Programming Basics

<http://software.intel.com/en-us/courseware/249633>

The Shared-Memory Model



The Threads Model



What Is a Thread?

“A unit of control within a process” — Carver and Tai

Main thread executes program’s “main” function

Main thread may create other threads to execute other functions

Threads have own program counter, copy of core registers, and stack of activation records

Threads share process’s data, code, address space, and other resources

Threads have lower overhead than processes

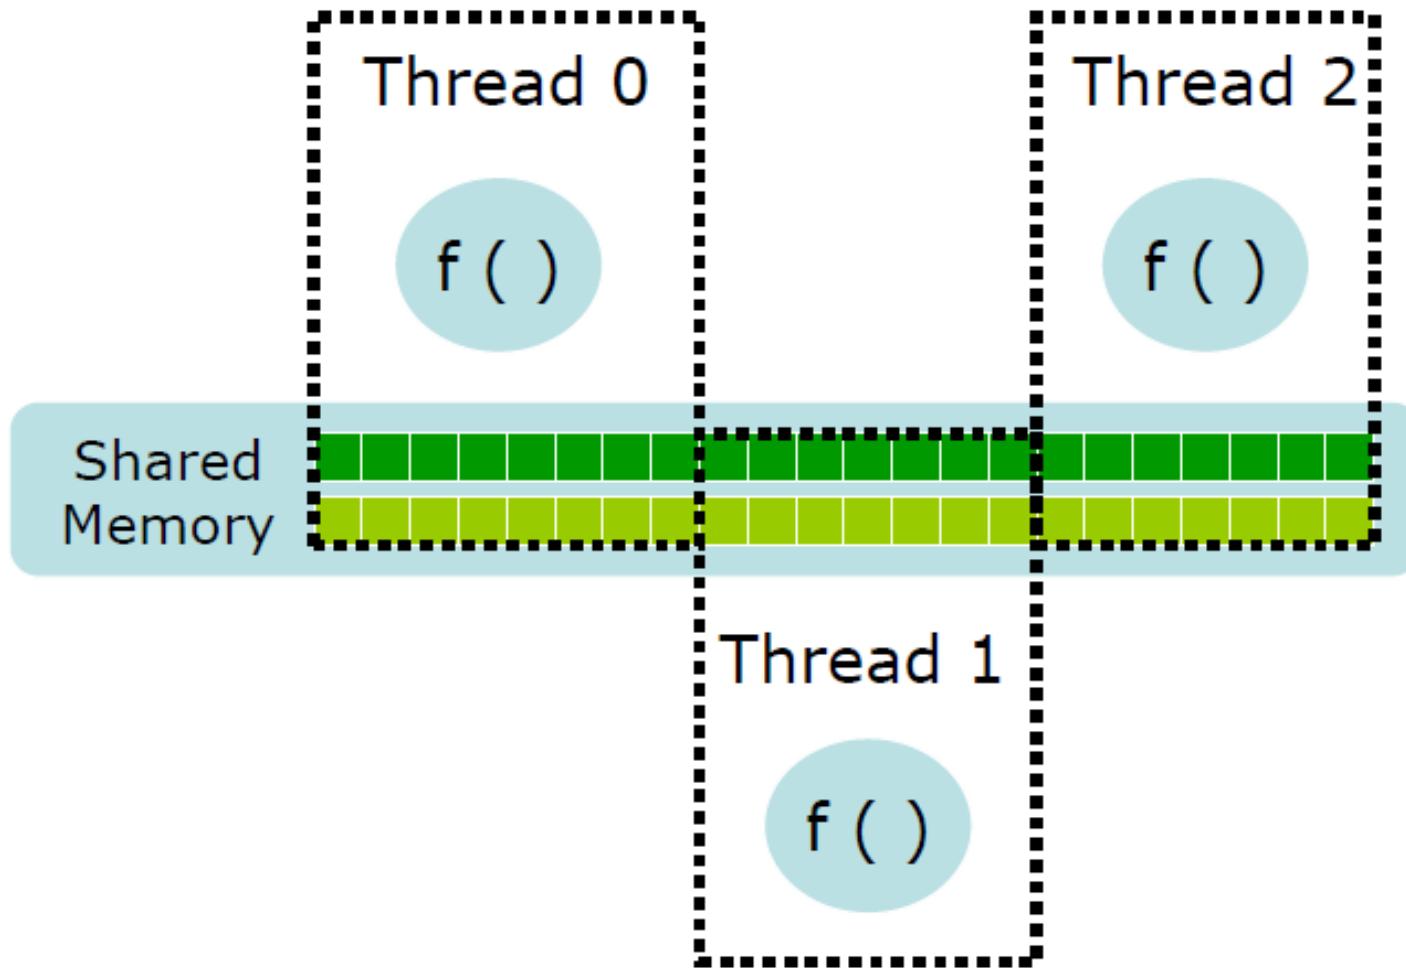


Copyright © 2009, Intel Corporation. All rights reserved.

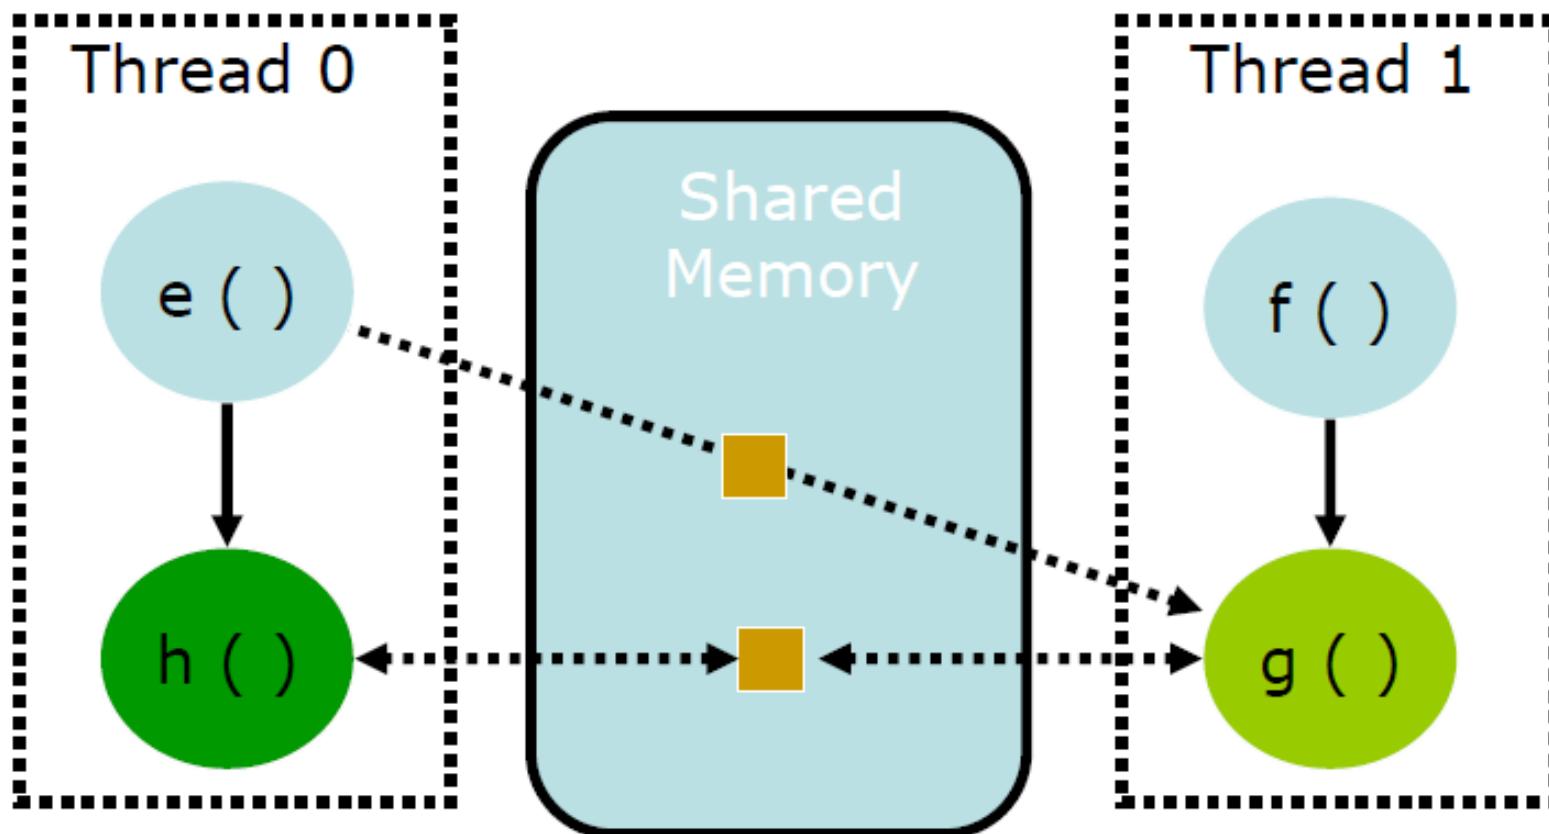
Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. * Other brands and names are the property of their respective owners.



Domain Decomposition Using Threads



Task Decomposition Using Threads



What are tasks?

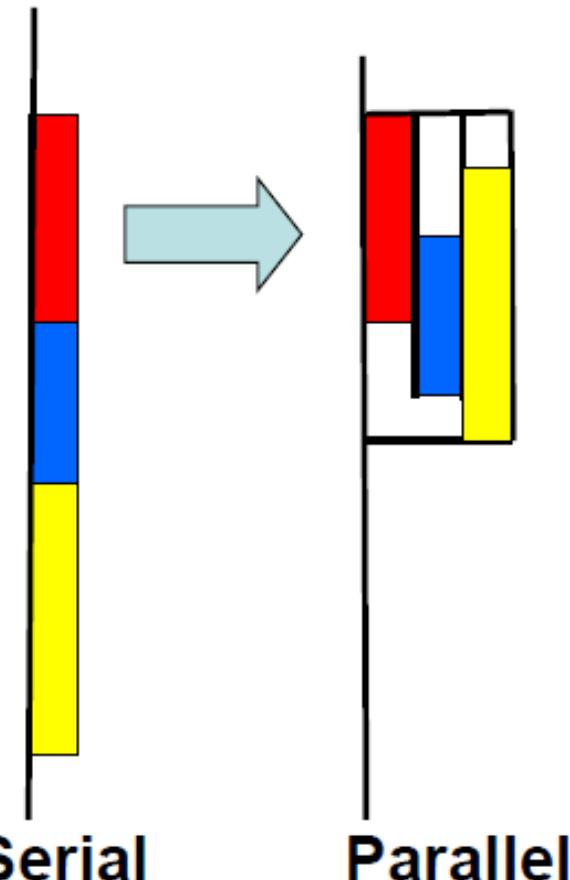
Tasks are independent units of work

- Threads are assigned to perform the work of each task
- Tasks may be deferred
- Tasks may be executed immediately

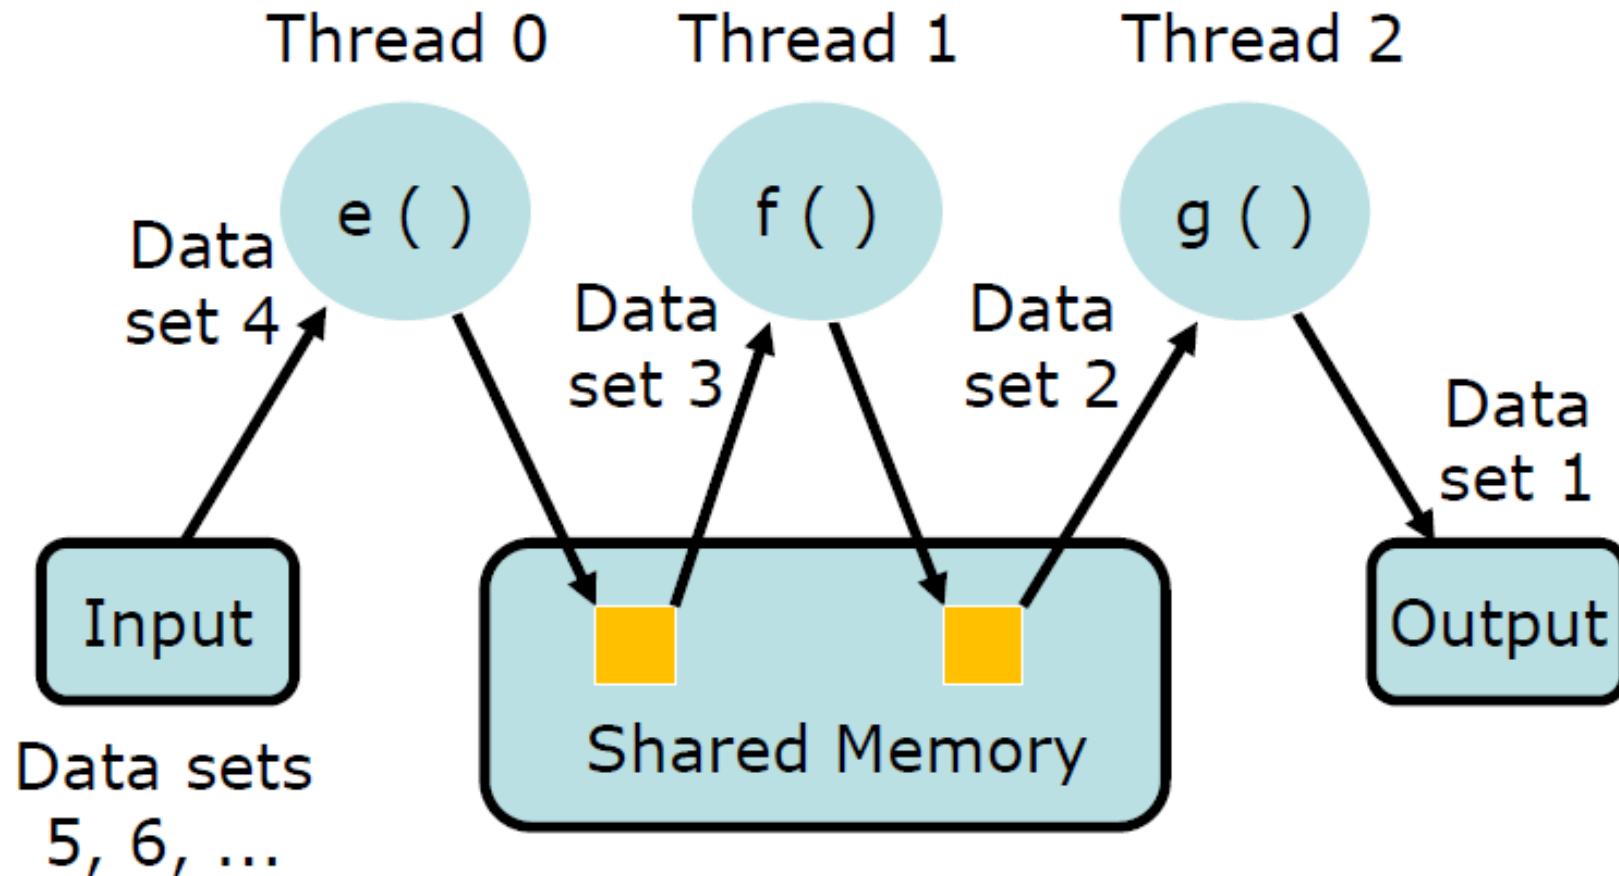
The runtime system decides which of the above

Tasks are composed of:

- **code** to execute
- **data** environment
- **internal control variables** (ICV)



Pipelining Using Threads



Domain Decomposition

Sequential Code:

```
int a[1000], i;  
for (i = 0; i < 1000; i++) a[i] = foo(i);
```

Thread 0:

```
for (i = 0; i < 500; i++) a[i] = foo(i);
```

Thread 1:

```
for (i = 500; i < 1000; i++) a[i] = foo(i);
```

Domain Decomposition

Sequential Code:

```
int a[1000], i;  
for (i = 0; i < 1000; i++) a[i] = foo(i);
```

Thread 0:

```
for (i = 0; i < 500; i++) a[i] = foo(i);
```

Thread 1:

```
for (i = 500; i < 1000; i++) a[i] = foo(i);
```

Private

Shared

Task Decomposition

```
int e;

main () {
    int x[10], j, k, m;    j = f(x, k); m = g(x, k); ...
}

int f(int *x, int k)
{
    int a;    a = e * x[k] * x[k];    return a;
}

int g(int *x, int k)
{
    int a;    k = k-1;    a = e / x[k];    return a;
}
```



Copyright © 2009, Intel Corporation. All rights reserved.

Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. * Other brands and names are the property of their respective owners.



Task Decomposition

```
int e;  
  
main () {  
    int x[10], j, k, m;    j = f(x, k);    m = g(x, k);  
}
```

```
int f(int *x, int k)          Thread 0  
{  
    int a;    a = e * x[k] * x[k];    return a;  
}  
  
int g(int *x, int k)          Thread 1  
{  
    int a;    k = k-1;    a = e / x[k];    return a;  
}
```

Task Decomposition

Volatile int **e**; Static variable: Shared

```
main () {  
    int x[10], j, k, m;    j = f(x, k);    m = g(x, k);  
}
```

```
int f(int *x, int k)                                Thread 0  
{  
    int a;    a = e * x[k] * x[k];    return a;  
}
```

```
int g(int *x, int k)                                Thread 1  
{  
    int a;    k = k-1;    a = e / x[k];    return a;  
}
```

Task Decomposition

```
Volatile int e;  
main () {  
    int x[10], j, k, m;    j = f(x, k);    m = g(x, k);  
}
```

Heap variable: Shared

```
int f(int *x, int k)          Thread 0  
{  
    int a;    a = e * x[k] * x[k];    return a;  
}  
  
int g(int *x, int k)          Thread 1  
{  
    int a;    k = k-1;    a = e / x[k];    return a;  
}
```

Task Decomposition

```
Volatile int e;
```

```
main () {  
    int x[10], j, k, m;    j = f(x, k);    m = g(x, k);  
}  
Function's local variables: Private
```

```
int f(int *x, int k)  
{  
    int a;    a = e * x[k] * x[k];    return a;  
}
```

Thread 0

```
int g(int *x, int k)  
{  
    int a;    k = k-1;    a = e / x[k];    return a;  
}
```

Thread 1

OpenMP

Fork/Join Programming Model

When program begins execution, only master thread active

Master thread executes sequential portions of program

For parallel portions of program, master thread forks (creates or awakens) additional threads

At join (end of parallel section of code), extra threads are suspended or die



Copyright © 2009, Intel Corporation. All rights reserved.

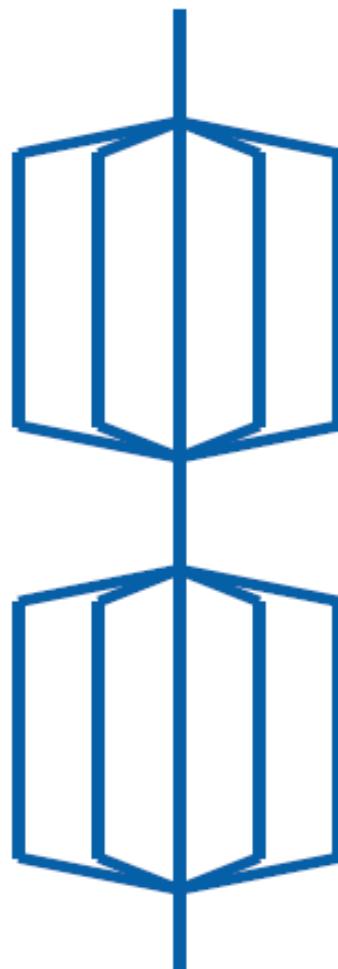
Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. * Other brands and names are the property of their respective owners.

4



Relating Fork/Join to Code

```
for {  
    [ ]  
    [ ]  
}  
[ ]  
[ ]  
  
for {  
    [ ]  
    [ ]  
}  
[ ]  
[ ]
```



Sequential code

Parallel code

Sequential code

Parallel code

Sequential code

Pragma: parallel for

The compiler directive

```
#pragma omp parallel for
```

tells the compiler that the `for` loop which immediately follows can be executed in parallel

The number of loop iterations must be computable at run time before loop executes

Loop must not contain a `break`, `return`, or `exit`

Loop must not contain a `goto` to a label outside loop

Example

```
int first, *marked, prime, size;  
...  
#pragma omp parallel for  
for (i = first; i < size; i += prime)  
marked[i] = 1;
```

Threads are assigned an independent set of iterations

Threads must wait at the end of construct

Pragma: parallel

Sometimes the code that should be executed in parallel goes beyond a single for loop

The parallel pragma is used when a block of code should be executed in parallel

```
#pragma omp parallel
{
    DoSomeWork(res, M);
    DoSomeOtherWork(res, M);
}
```



Copyright © 2009, Intel Corporation. All rights reserved.

Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. * Other brands and names are the property of their respective owners.

10



Pragma: for

The `for` pragma can be used inside a block of code already marked with the `parallel` pragma

Loop iterations should be divided among the active threads

There is a *barrier synchronization* at the end of the `for` loop

```
#pragma omp parallel
{
    DoSomeWork(res, M);

#pragma omp for
    for (i = 0; i < M; i++) {
        res[i] = huge();
    }
    DoSomeMoreWork(res, M);
}
```

Which Loop to Make Parallel?

```
main () {  
    int i, j, k;  
    float **a, **b;  
    ...  
    for (k = 0; k < N; k++)          Loop-carried dependences  
        for (i = 0; i < N; i++)      Can execute in parallel  
            for (j = 0; j < N; j++)  Can execute in parallel  
                a[i][j] = MIN(a[i][j], a[i][k] + a[k][j]);
```

Minimizing Threading Overhead

There is a fork/join for every instance of

```
#pragma omp parallel for  
for ( ) {  
    ...  
}
```

Since fork/join is a source of overhead, we want to maximize the amount of work done for each fork/join

Hence we choose to make the middle loop parallel

Almost Right, but Not Quite

```
main () {  
    int i, j, k;  
    float **a, **b;  
    ...  
    for (k = 0; k < N; k++)  
        #pragma omp parallel for  
        for (i = 0; i < N; i++)  
            for (j = 0; j < N; j++)  
                a[i][j] = MIN(a[i][j], a[i][k] + a[k][j]);  
}
```

Problem: j is a shared variable



Copyright © 2009, Intel Corporation. All rights reserved.

Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. * Other brands and names are the property of their respective owners.



Problem Solved with private Clause

```
main () {  
    int i, j, k;  
    float **a, **b;  
    ...  
    for (k = 0; k < N; k++)  
        #pragma omp parallel for private (j)  
        for (i = 0; i < N; i++)  
            for (j = 0; j < N; j++)  
                a[i][j] = MIN(a[i][j], a[i][k] + a[k][j]);
```

Tells compiler to make listed variables private



OpenMP reduction Clause

Reductions are so common that OpenMP provides a reduction clause for the parallel for pragma

```
reduction (op : list)
```

A PRIVATE copy of each list variable is created and initialized depending on the “op”

- The identity value “op” (e.g., 0 for addition)

These copies are updated locally by threads

At end of construct, local copies are combined through “op” into a single value and combined with the value in the original SHARED variable

Reduction Example

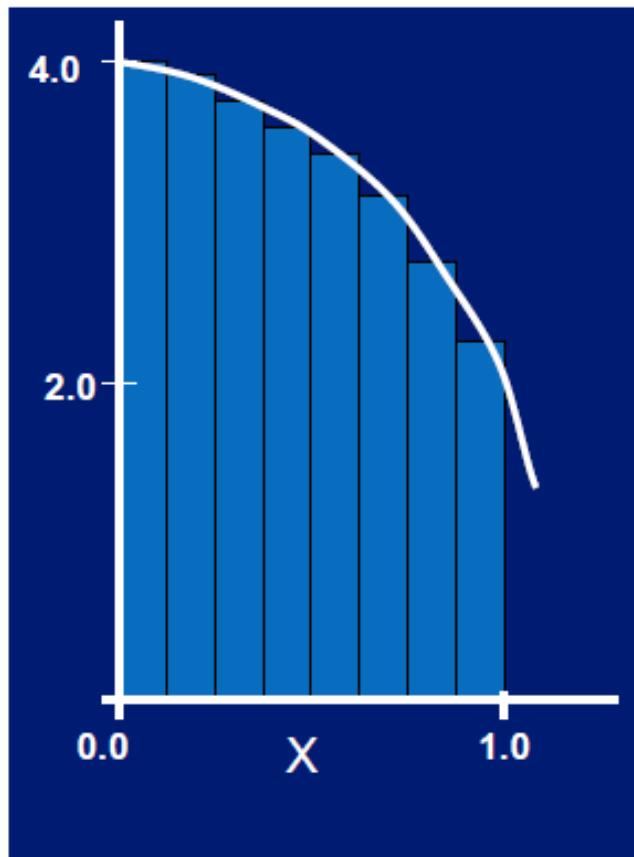
```
#pragma omp parallel for reduction(+:sum)
    for(i = 0; i < N; i++) {
        sum += a[i] * b[i];
    }
```

Local copy of sum for each thread

All local copies of sum added together and stored in shared copy

Numerical Integration Example

$$\int_0^1 \frac{4.0}{(1+x^2)} dx = \pi$$



```
static long num_rects=100000;
double width, pi;

void main()
{ int i;
  double x, sum = 0.0;

  width = 1.0/(double) num_rects;
  for (i = 0; i < num_rects; i++) {
    x = (i+0.5)*width;
    sum = sum + 4.0/(1.0 + x*x);
  }
  pi = width * sum;
  printf("Pi = %f\n",pi);
}
```

Numerical Integration: What's Shared?

```
static long num_rects=100000;
double width, pi;

void main()
{ int i;
  double x, sum = 0.0;

  width = 1.0/(double) num_rects;
  for (i = 0; i < num_steps; i++) {
    x = (i+0.5)*step;
    sum = sum + 4.0/(1.0 + x*x);
  }
  pi = step * sum;
  printf("Pi = %f\n",pi);
}
```

What variables can be shared?
width, num_rects

Numerical Integration: What's Private?

```
static long num_rects=100000;
double width, pi;

void main()
{ int i;
  double x, sum = 0.0;

  width = 1.0/(double) num_rects;
  for (i = 0; i < num_steps; i++) {
    x = (i+0.5)*step;
    sum = sum + 4.0/(1.0 + x*x);
  }
  pi = step * sum;
  printf("Pi = %f\n",pi);
}
```

What variables need to
be private?

x, i



Copyright © 2009, Intel Corporation. All rights reserved.

Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. * Other brands and names are the property of their respective owners.

23



Numerical Integration: Any Reductions?

```
static long num_rects=100000;
double width, pi;

void main()
{ int i;
  double x, sum = 0.0;

  width = 1.0/(double) num_rects;
  for (i = 0; i < num_steps; i++) {
    x = (i+0.5)*step;
    sum = sum + 4.0/(1.0 + x*x);
  }
  pi = step * sum;
  printf("Pi = %f\n",pi);
}
```

What variables should be set up for reduction?

sum



Copyright © 2009, Intel Corporation. All rights reserved.

Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. * Other brands and names are the property of their respective owners.

24



Solution to Computing Pi

```
static long num_rects=100000;
double width, pi;

void main()
{ int i;
  double x, sum = 0.0;
#pragma omp parallel for private(x) reduction(+:sum)
  width = 1.0/(double) num_rects;
  for (i = 0; i < num_steps; i++) {
    x = (i+0.5)*step;
    sum = sum + 4.0/(1.0 + x*x);
  }
  pi = step * sum;
  printf("Pi = %f\n",pi);
}
```



Copyright © 2009, Intel Corporation. All rights reserved.

Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. * Other brands and names are the property of their respective owners.



Locks Are Dangerous

Suppose a lock is used to guarantee mutually exclusive access to a shared variable

Imagine two threads, each with its own critical region

Thread 1

```
a += 5;  
b += 7;  
a += b;  
a += 11;
```

Thread 2

```
b += 5;  
a += 7;  
a += b;  
b += 11;
```



Copyright © 2009, Intel Corporation. All rights reserved.

Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. * Other brands and names are the property of their respective owners.



Faulty Implementation

Thread 1

```
lock (lock_a);  
a += 5;  
lock (lock_b);  
b += 7;  
a += b;  
unlock (lock_b);  
a += 11;  
unlock (lock_a);
```

Thread 2

```
lock (lock_b);  
b += 5;  
lock (lock_a);  
a += 7;  
a += b;  
unlock (lock_a);  
b += 11;  
unlock (lock_b);
```

Faulty Implementation

Thread 1

```
lock (lock_a);  
a += 5;  
lock (lock_b);  
b += 7;  
a += b;  
unlock (lock_b);  
a += 11;  
unlock (lock_a);
```

What happens if
threads are at
this point at the
same time?

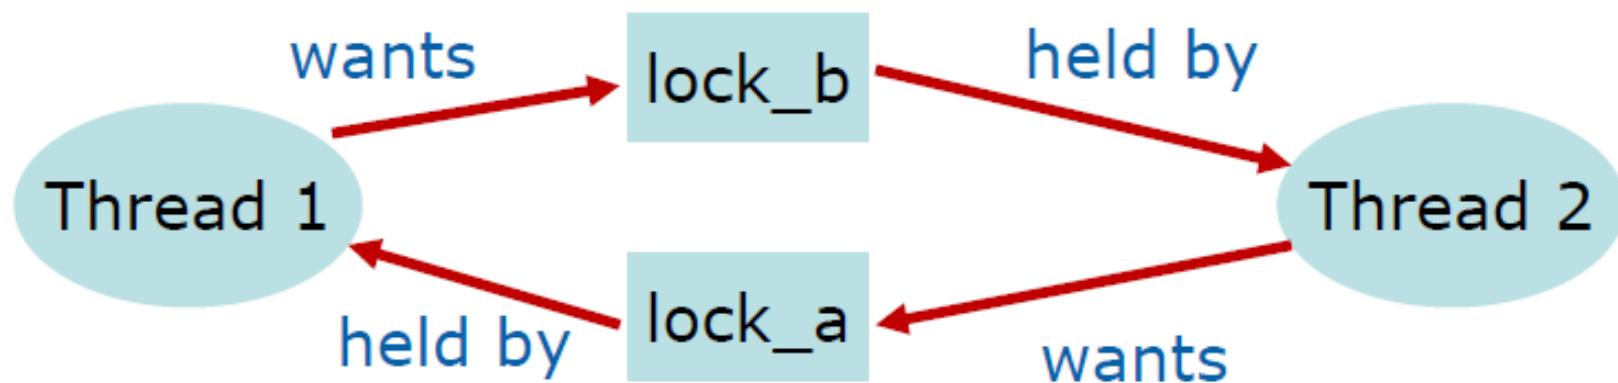
Thread 2

```
lock (lock_b);  
b += 5;  
lock (lock_a);  
a += 7;  
a += b;  
unlock (lock_a);  
b += 11;  
unlock (lock_b);
```

Deadlock

A situation involving two or more threads (processes) in which no thread may proceed because each is waiting for a resource held by another

Can be represented by a resource allocation graph



A graph of deadlock contains a cycle

Correct Implementation

Thread 1

```
lock (lock_a);  
a += 5;  
lock (lock_b);  
b += 7;  
a += b;  
unlock (lock_b);  
a += 11;  
unlock (lock_a);
```

Thread 2

```
lock (lock_a);  
lock (lock_b);  
b += 5;  
a += 7;  
a += b;  
unlock (lock_a);  
b += 11;  
unlock (lock_b);
```

Threads must lock
lock_a before lock_b

Pragma: single

Denotes block of code to be executed by only one thread

- First thread to arrive is chosen

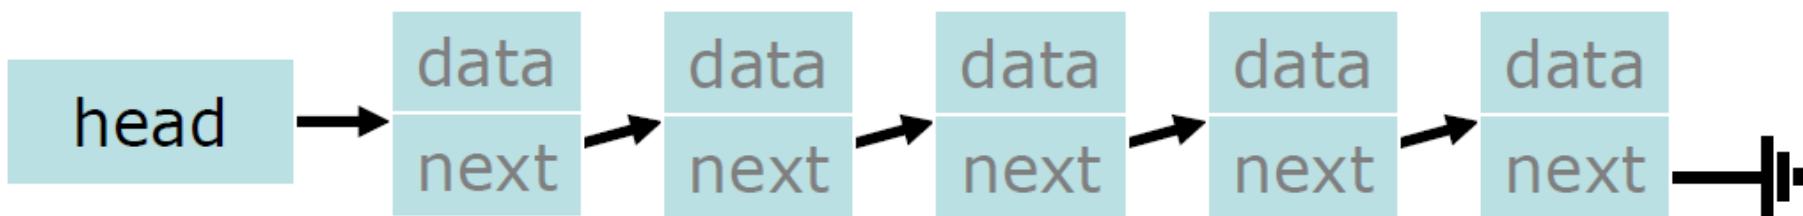
Implicit barrier at end

```
#pragma omp parallel
{
    DoManyThings();
#pragma omp single
    {
        printf("Many Things done\n");
    } // threads wait here for single
    DoManyMoreThings();
}
```

A Linked List Example

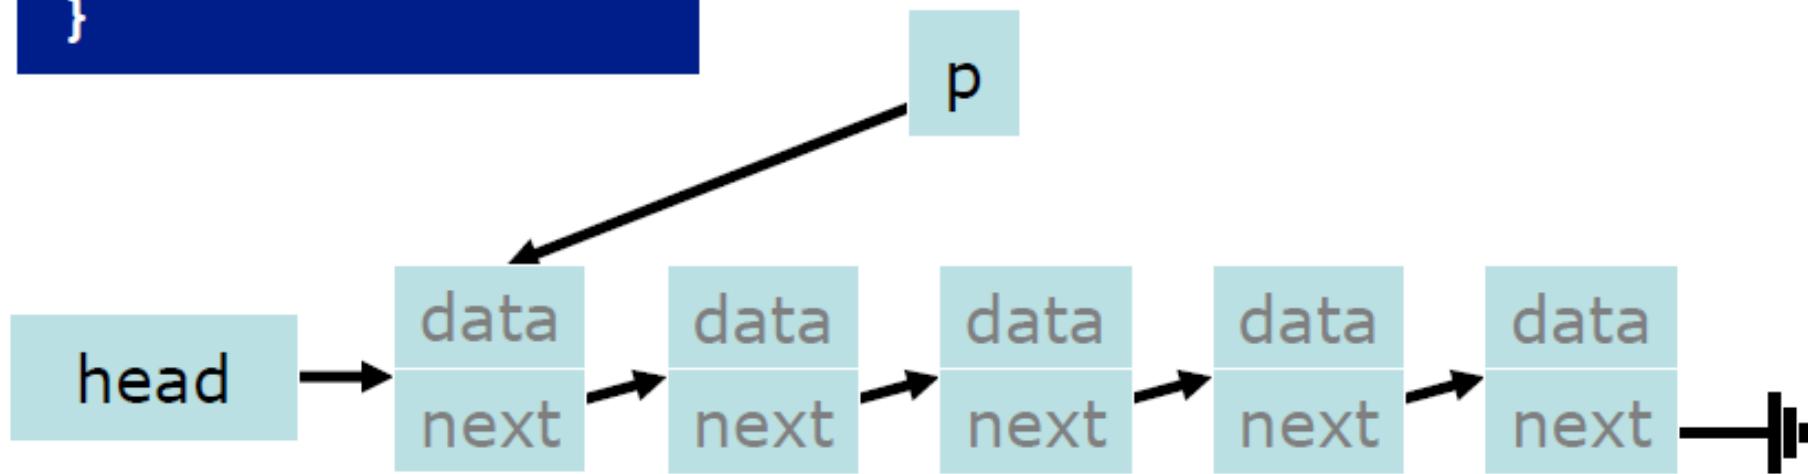
```
node *p = head;  
while (p) {  
    process(p);  
    p = p->next;  
}
```

Serial version



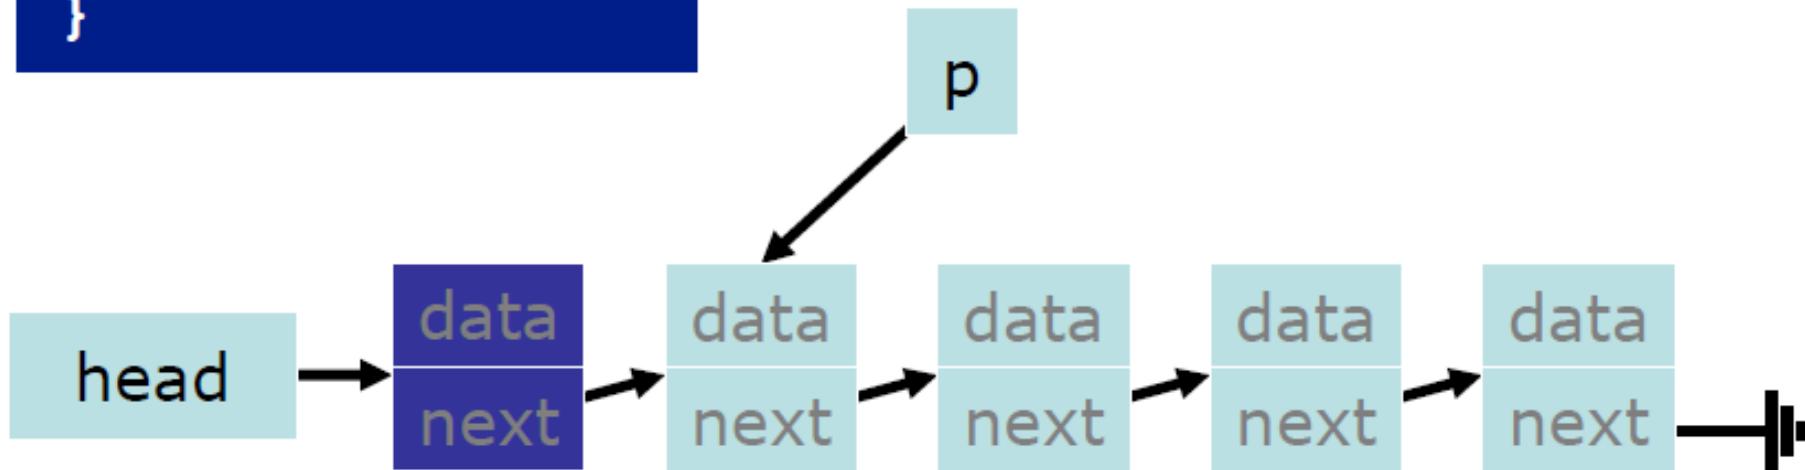
A Linked List Example

```
node *p = head;  
while (p) {  
    process(p);  
    p = p->next;  
}
```



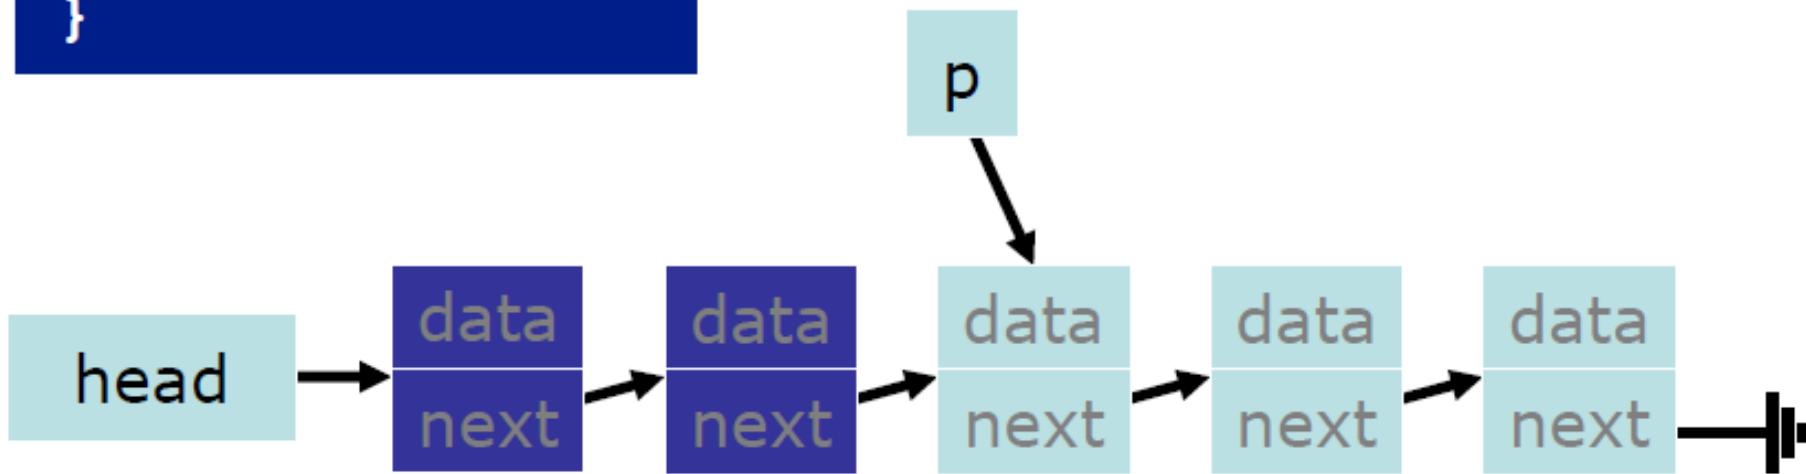
A Linked List Example

```
node *p = head;  
while (p) {  
    process(p);  
    p = p->next;  
}
```



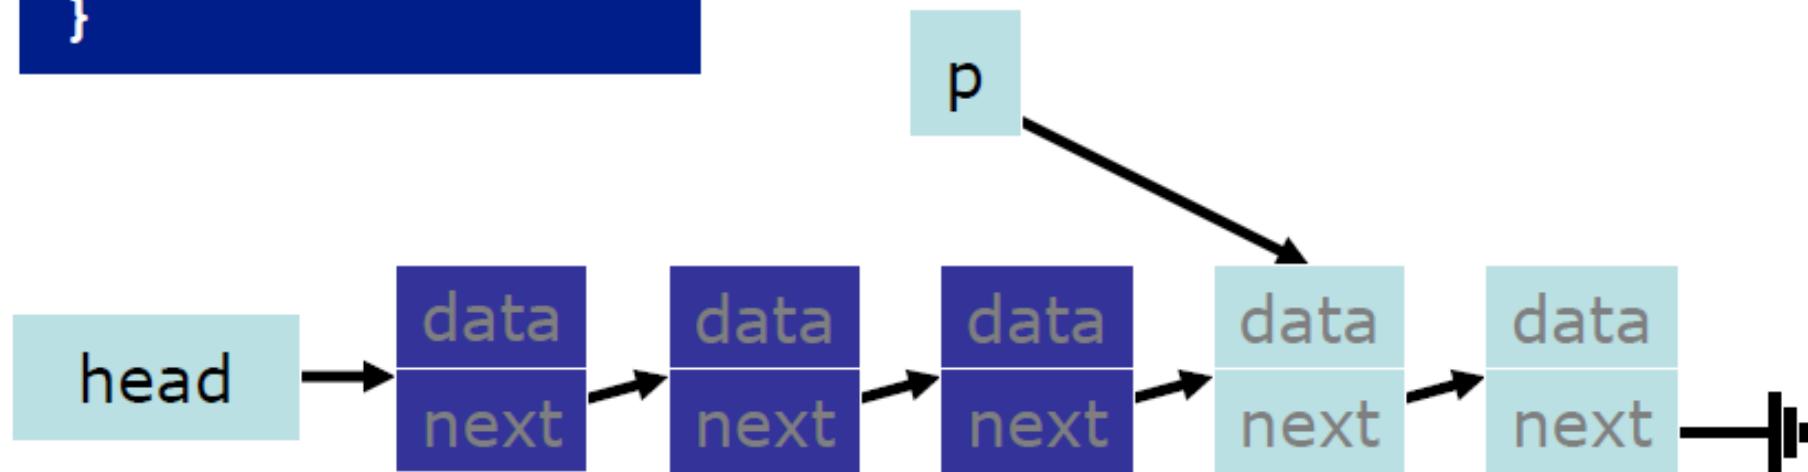
A Linked List Example

```
node *p = head;  
while (p) {  
    process(p);  
    p = p->next;  
}
```



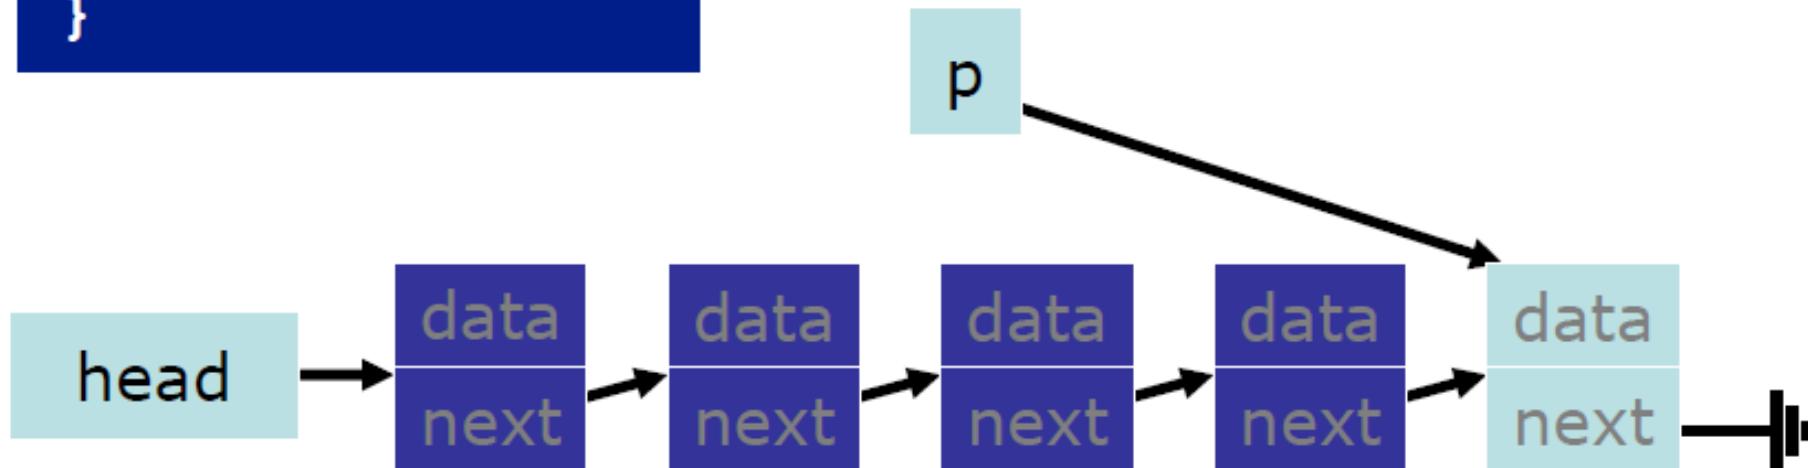
A Linked List Example

```
node *p = head;  
while (p) {  
    process(p);  
    p = p->next;  
}
```



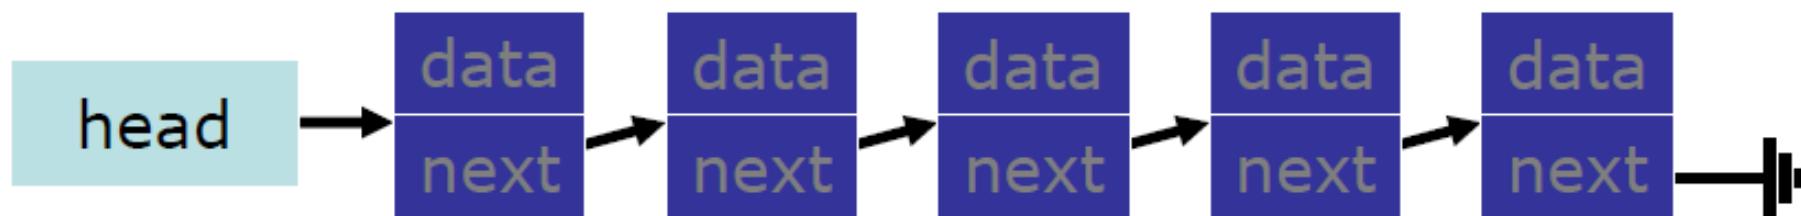
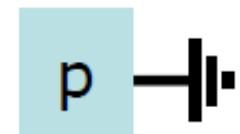
A Linked List Example

```
node *p = head;  
while (p) {  
    process(p);  
    p = p->next;  
}
```



A Linked List Example

```
node *p = head;  
while (p) {  
    process(p);  
    p = p->next;  
}
```



Task Construct – Explicit Task View

```
node *p = head;  
#pragma omp parallel  
{  
#pragma omp single  
while (p) {  
#pragma omp task  
process(p);  
p = p->next;  
}  
}
```

A team of threads is forked at the `omp parallel` construct

A single thread, T0, executes the `while` loop

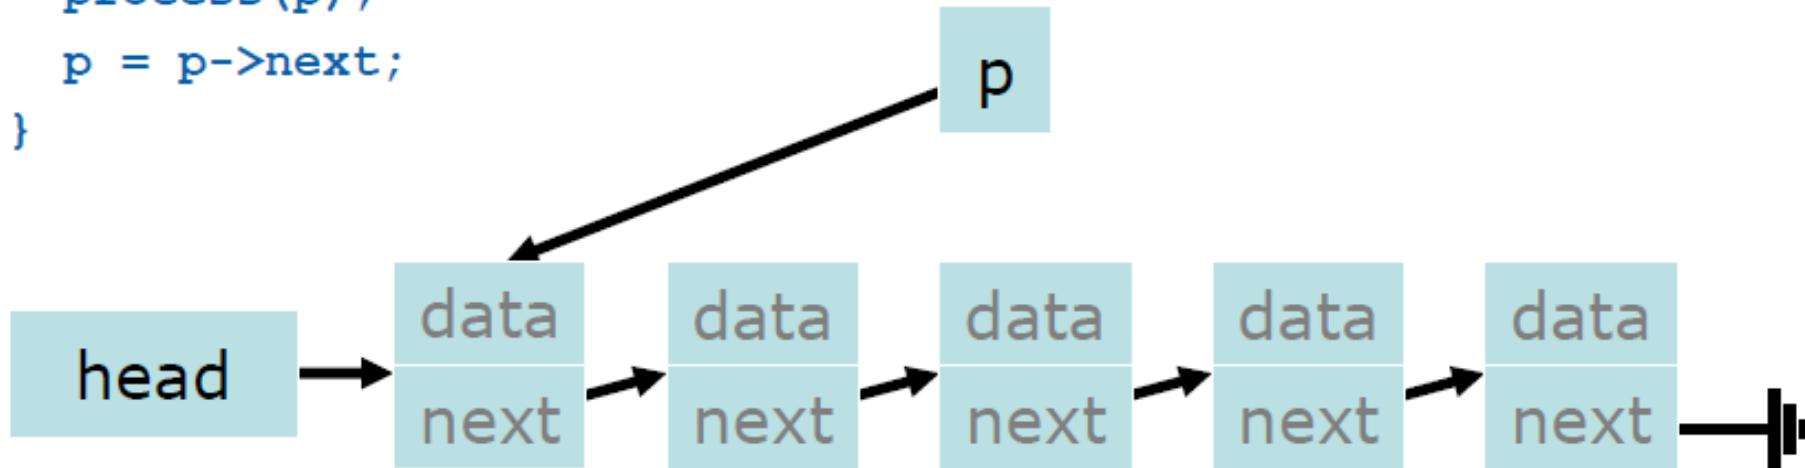
Each time T0 crosses the `omp task` construct it generates a new task

Each task runs in a thread

All tasks complete at the barrier at the end of the parallel region's single construct

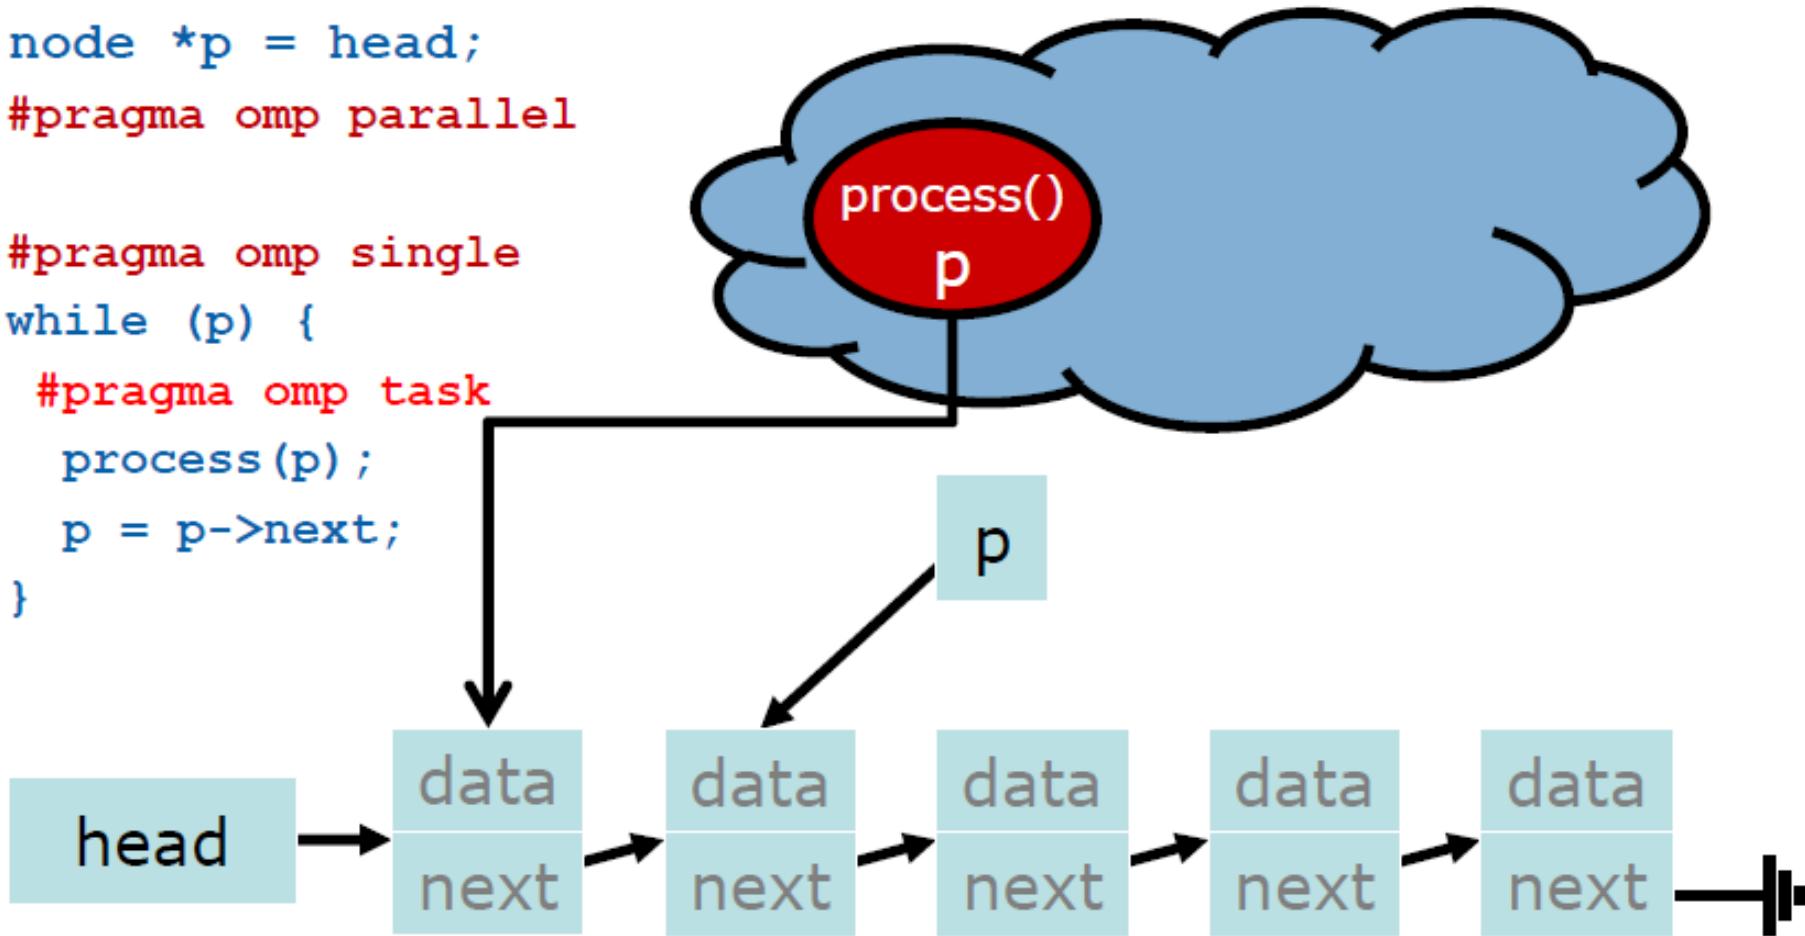
A Linked List Example

```
node *p = head;  
#pragma omp parallel  
{  
#pragma omp single  
while (p) {  
#pragma omp task  
process(p);  
p = p->next;  
}  
}
```



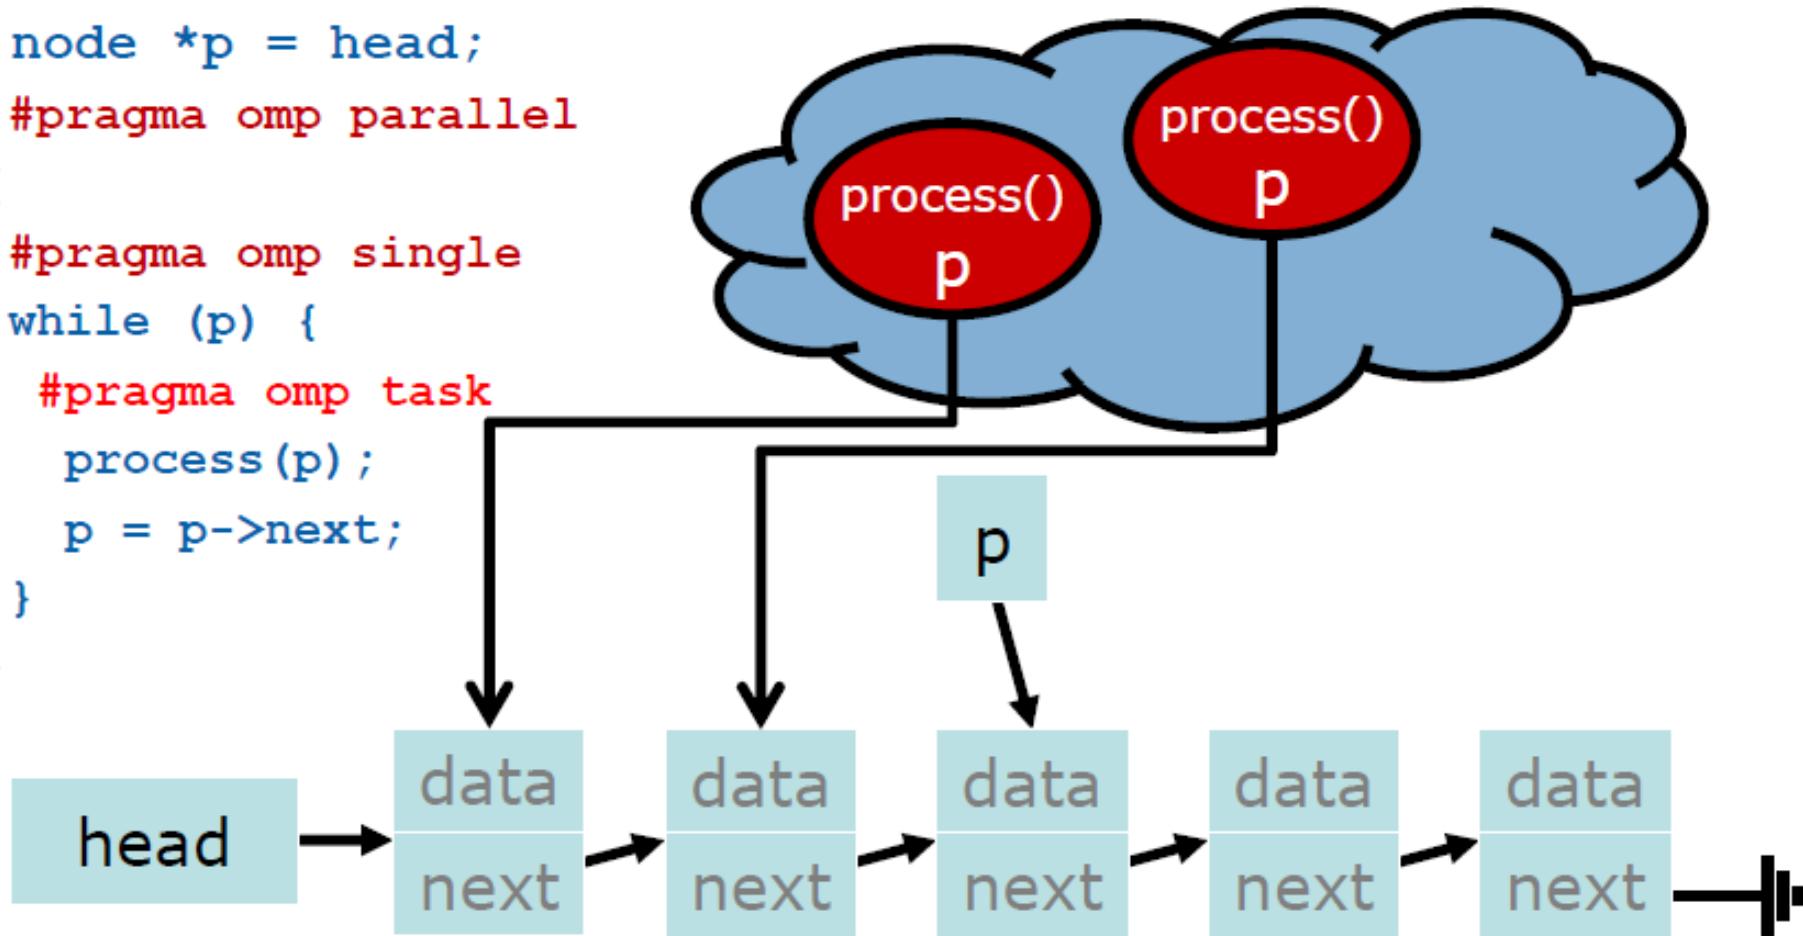
A Linked List Example

```
node *p = head;  
#pragma omp parallel  
{  
#pragma omp single  
while (p) {  
#pragma omp task  
process(p);  
p = p->next;  
}  
}
```



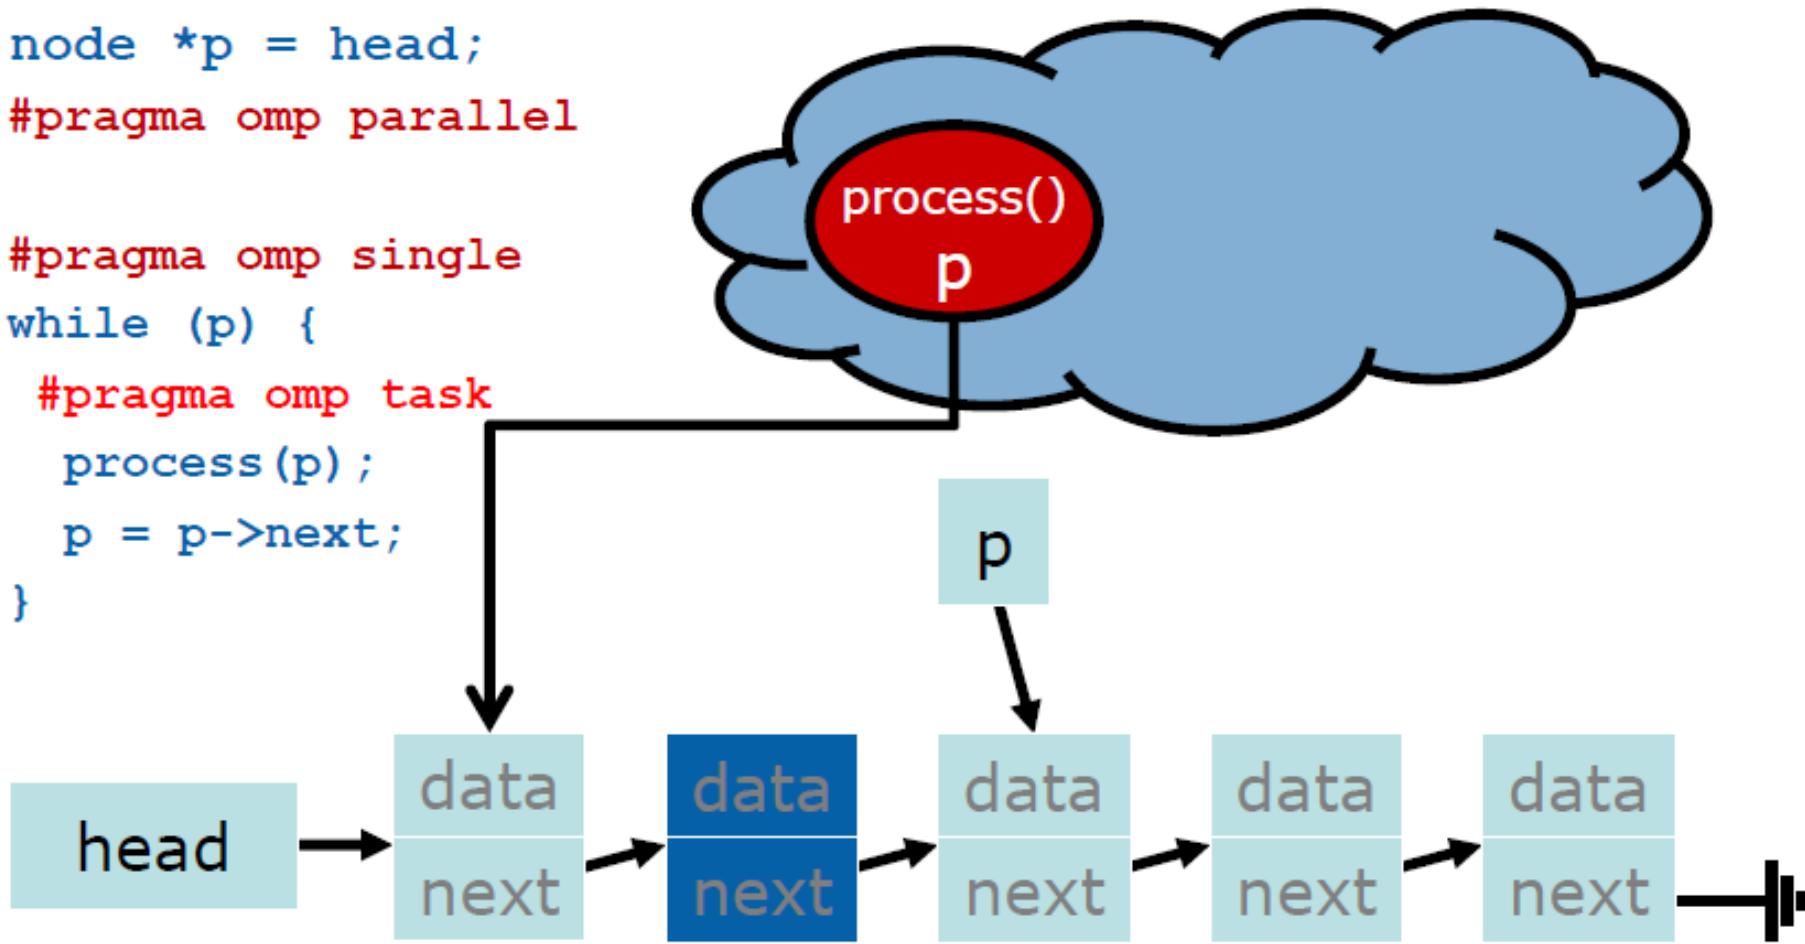
A Linked List Example

```
node *p = head;  
#pragma omp parallel  
{  
#pragma omp single  
while (p) {  
#pragma omp task  
    process(p);  
    p = p->next;  
}  
}
```



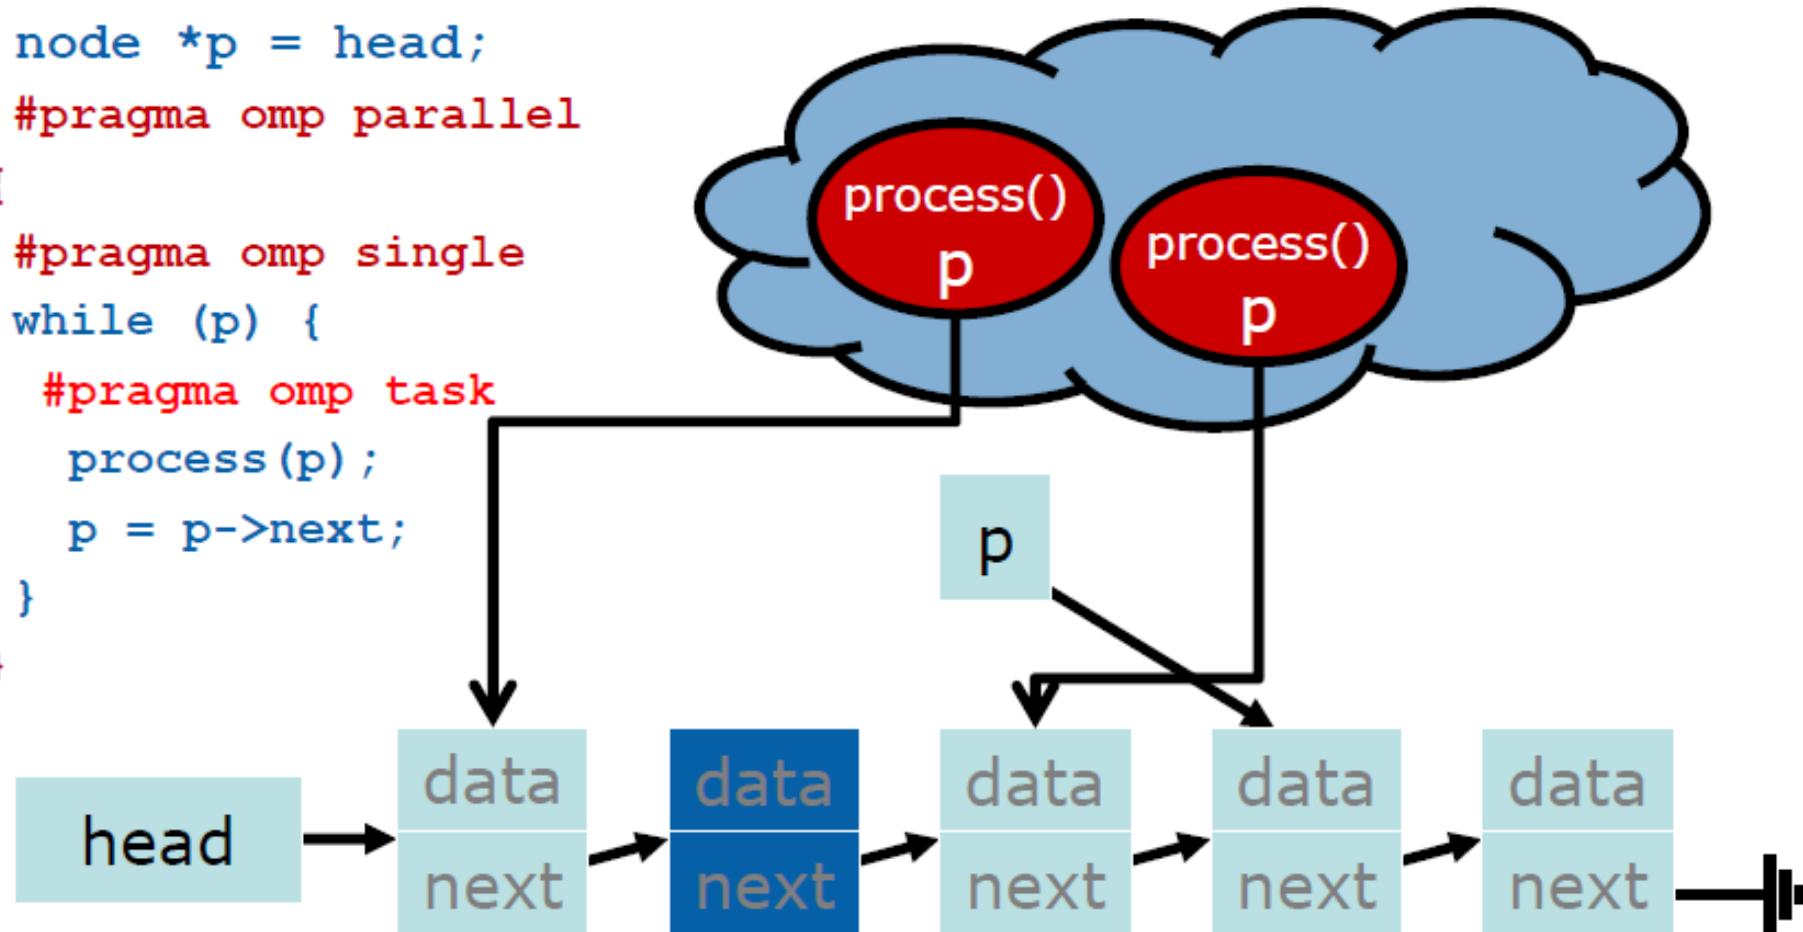
A Linked List Example

```
node *p = head;  
#pragma omp parallel  
{  
#pragma omp single  
while (p) {  
#pragma omp task  
process(p);  
p = p->next;  
}  
}
```



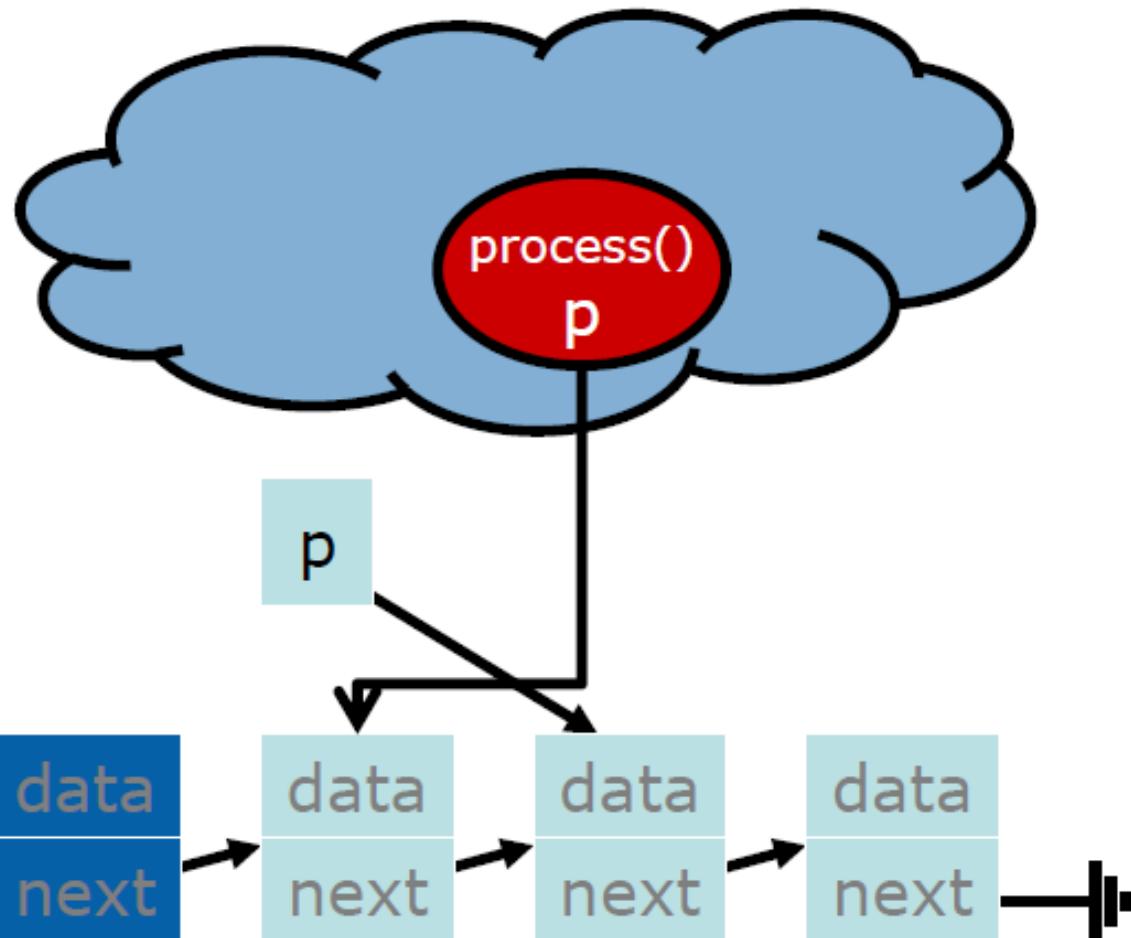
A Linked List Example

```
node *p = head;  
#pragma omp parallel  
{  
#pragma omp single  
while (p) {  
#pragma omp task  
process(p);  
p = p->next;  
}  
}
```



A Linked List Example

```
node *p = head;  
#pragma omp parallel  
{  
#pragma omp single  
while (p) {  
#pragma omp task  
process(p);  
p = p->next;  
}  
}
```



A Linked List Example

```
node *p = head;  
#pragma omp parallel  
{  
#pragma omp single  
while (p) {  
#pragma omp task  
    process(p);  
    p = p->next;  
}  
}
```



Copyright © 2009, Intel Corporation. All rights reserved.

Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. * Other brands and names are the property of their respective owners.

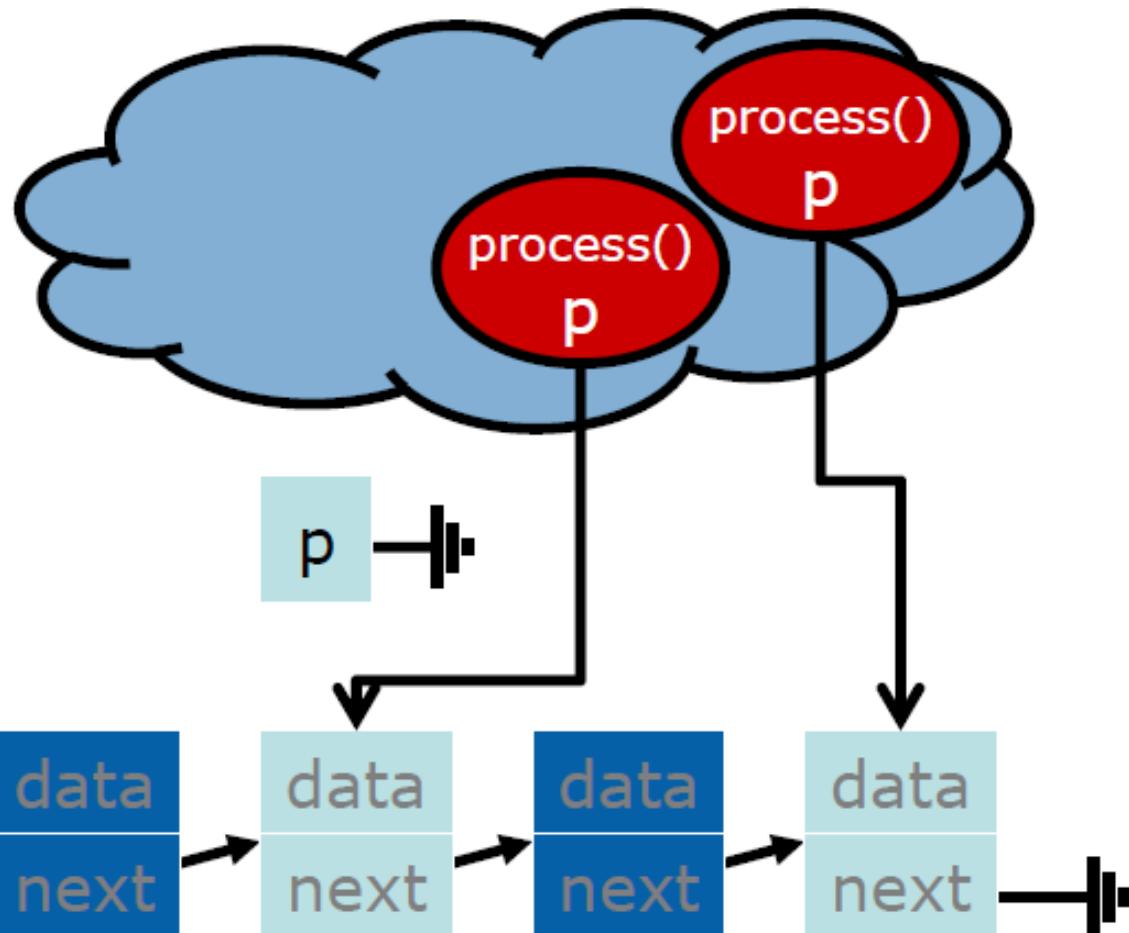
A Linked List Example

```
node *p = head;  
#pragma omp parallel  
{  
#pragma omp single  
while (p) {  
#pragma omp task  
process(p);  
p = p->next;  
}  
}
```



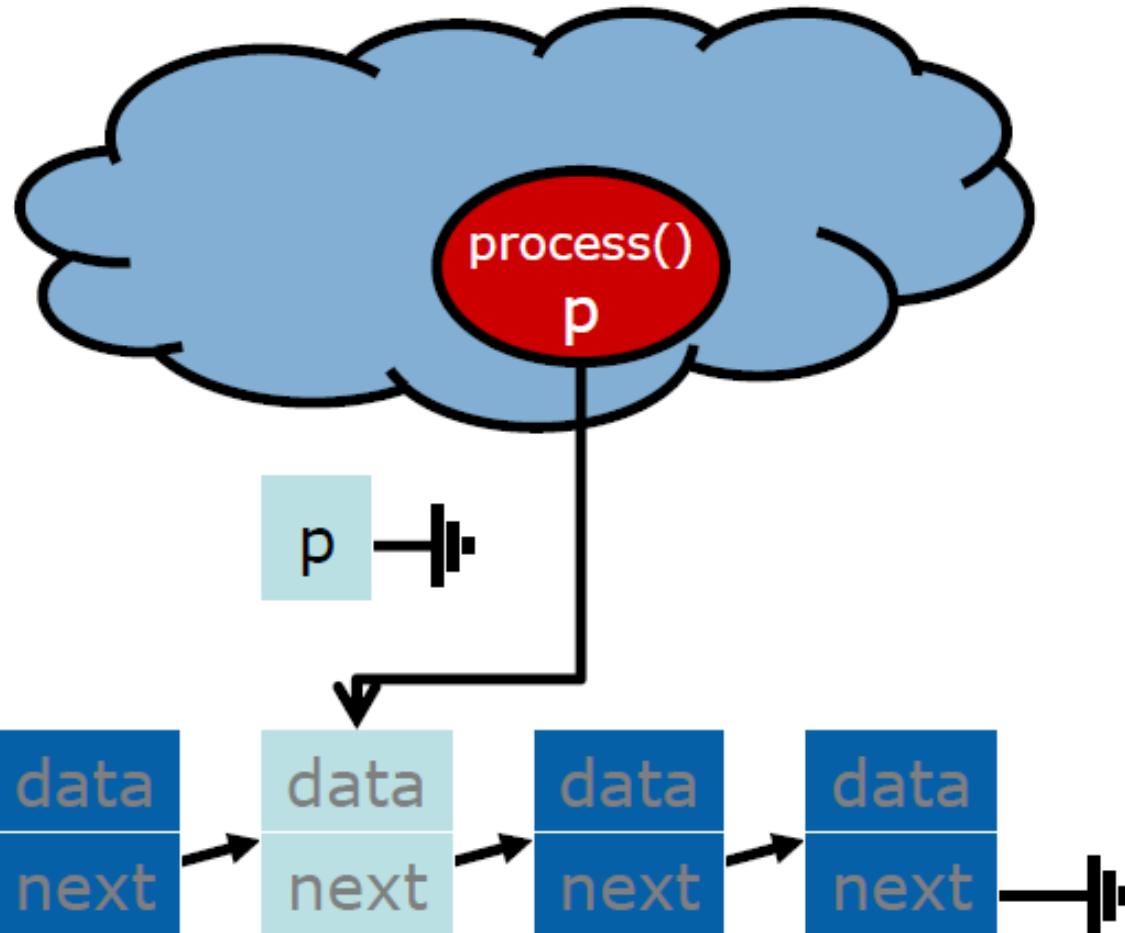
A Linked List Example

```
node *p = head;  
#pragma omp parallel  
{  
#pragma omp single  
while (p) {  
#pragma omp task  
process(p);  
p = p->next;  
}  
}
```



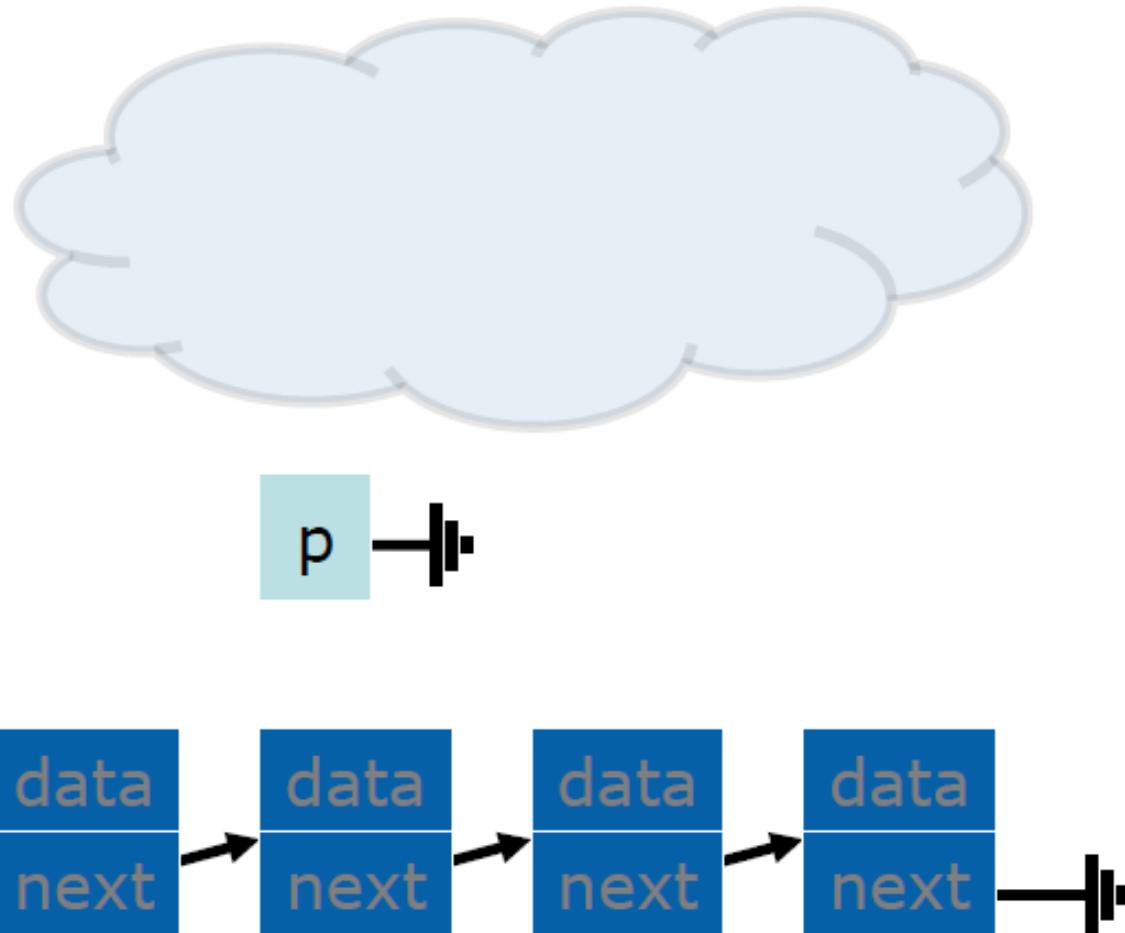
A Linked List Example

```
node *p = head;  
#pragma omp parallel  
{  
#pragma omp single  
while (p) {  
#pragma omp task  
process(p);  
p = p->next;  
}  
}
```



A Linked List Example

```
node *p = head;  
#pragma omp parallel  
{  
#pragma omp single  
while (p) {  
#pragma omp task  
process(p);  
p = p->next;  
}  
}
```



When are tasks guaranteed to be complete?

Tasks are guaranteed to be complete:

- At thread or task barriers
- At the directive: `#pragma omp barrier`
- At the directive: `#pragma omp taskwait`

Example: Naive Fibonacci Calculation

Recursion typically used to calculate Fibonacci number

Widely used as toy benchmark

- Easy to code
- Has unbalanced task graph

```
long SerialFib( long n ) {  
    if( n < 2 )  
        return n;  
    else  
        return SerialFib(n-1) + SerialFib(n-2);  
}
```

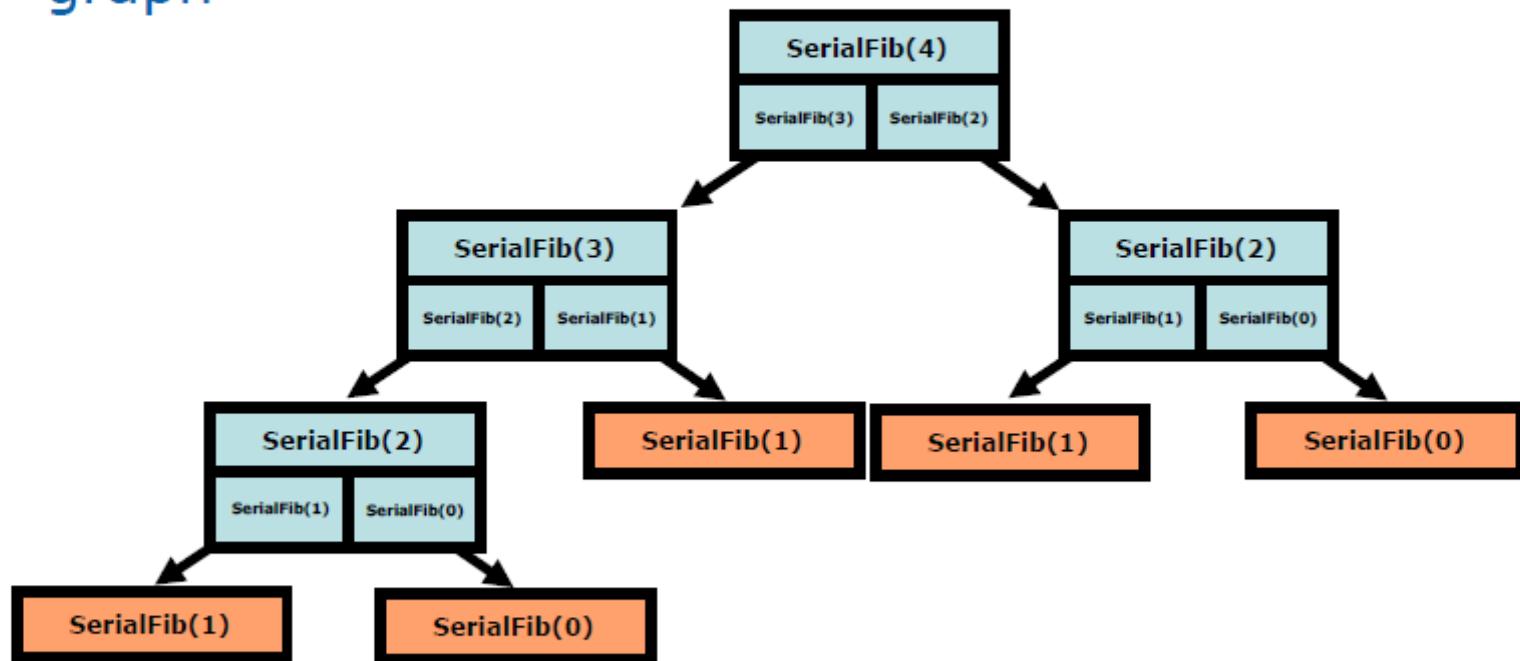


Copyright © 2009, Intel Corporation. All rights reserved.
Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. * Other brands and names are the property of their respective owners.



Example: Naive Fibonacci Calculation

We can envision Fibonacci computation as a task graph



Fibonacci - Task Spawning Solution

```
long ParallelFib(long n)
{ long sum;
#pragma omp parallel
{
#pragma omp single
    FibTask(n, &sum);
}
return sum;
}
```

Write a helper function to set up parallel region
Call FibTask() to do computation
Use sum return parameter in FibTask()



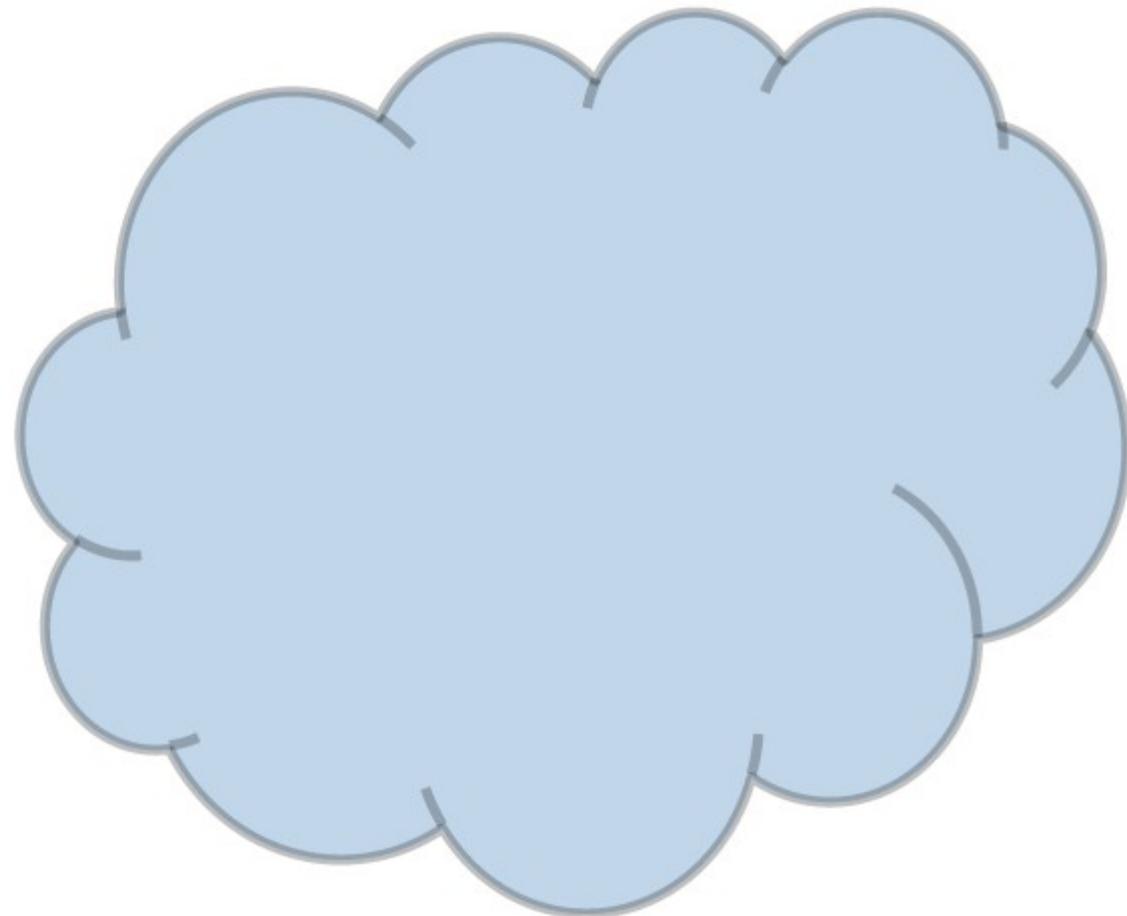
Copyright © 2009, Intel Corporation. All rights reserved.

Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. * Other brands and names are the property of their respective owners.



Fibonacci Task Solution Example

```
FibTask(8, *sum)  
long x, y;  
FibTask(7, &x);  
FibTask(6, &y);  
*sum = x + y;
```



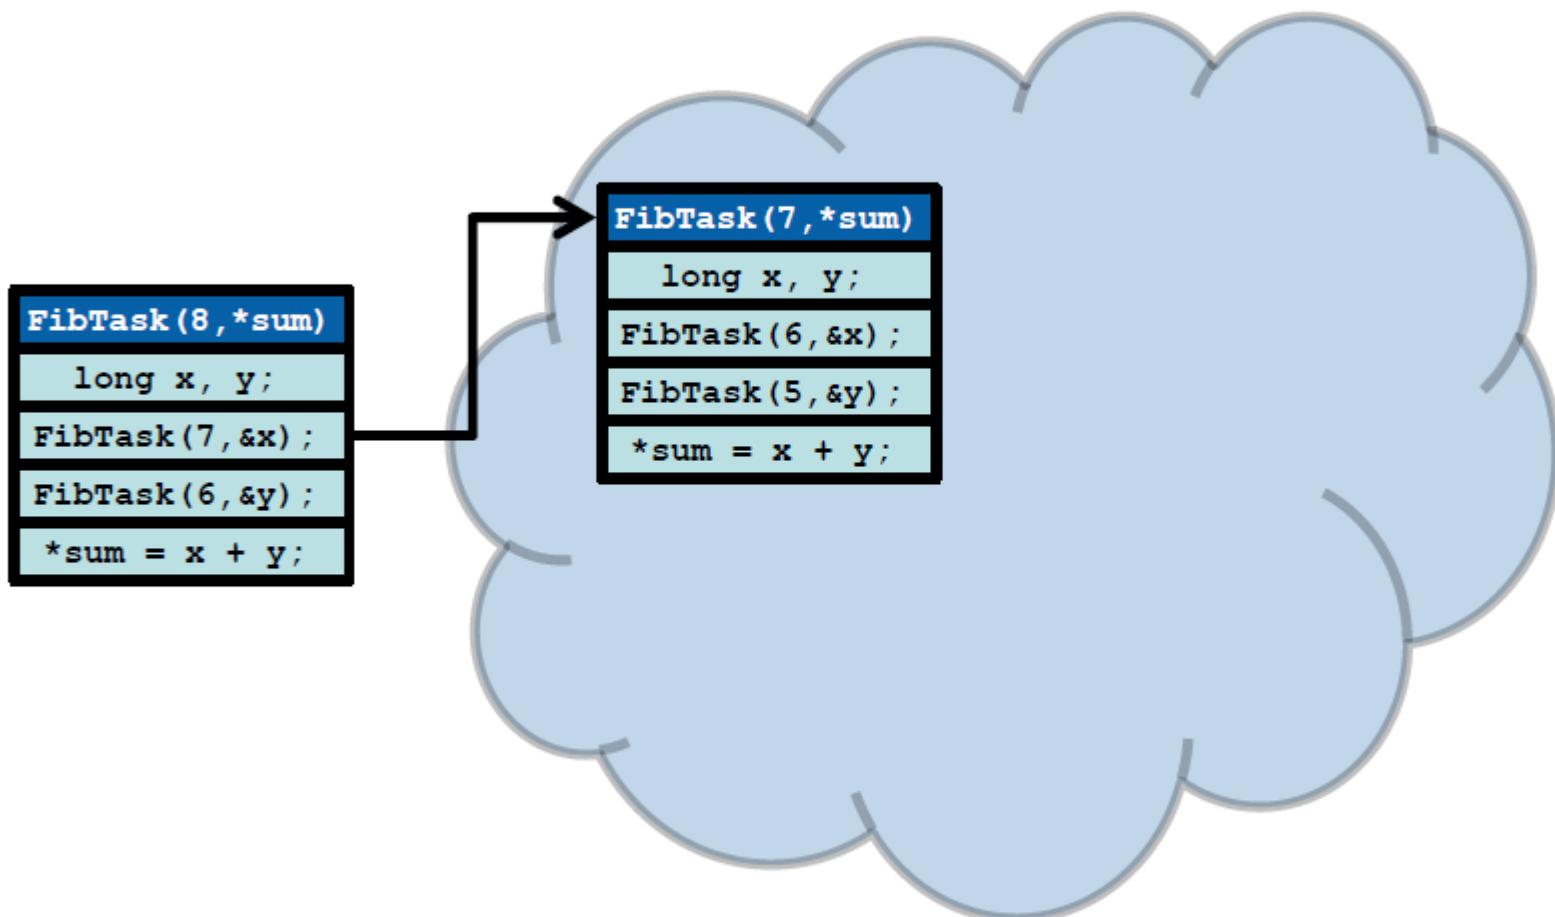
Copyright © 2009, Intel Corporation. All rights reserved.

Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. * Other brands and names are the property of their respective owners.

30



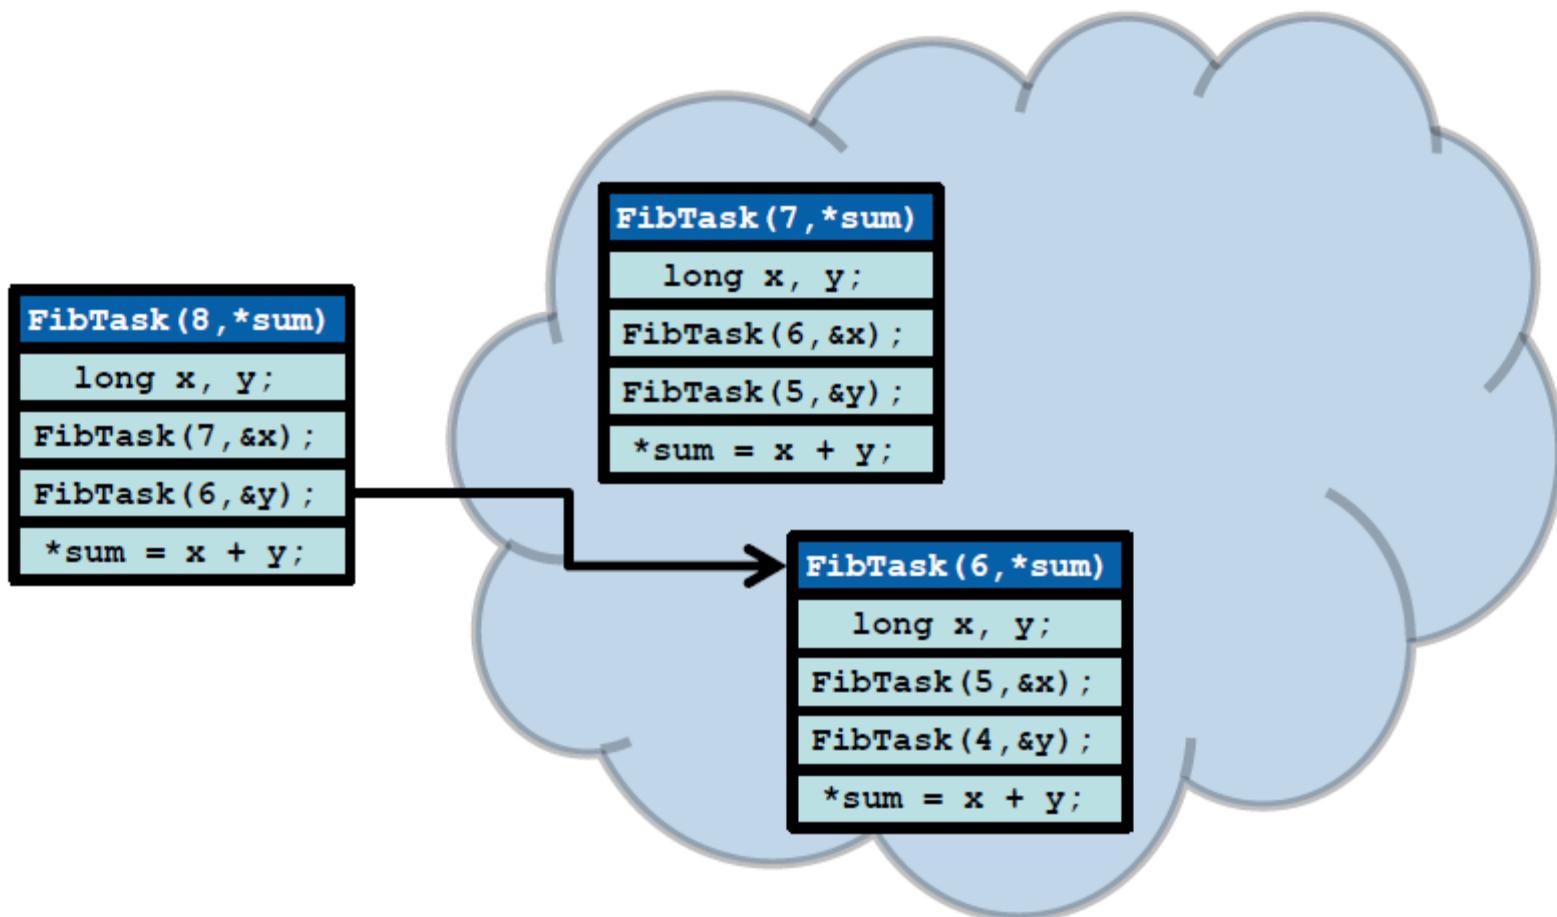
Fibonacci Task Solution Example



Copyright © 2009, Intel Corporation. All rights reserved.

Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. * Other brands and names are the property of their respective owners.

Fibonacci Task Solution Example



Copyright © 2009, Intel Corporation. All rights reserved.

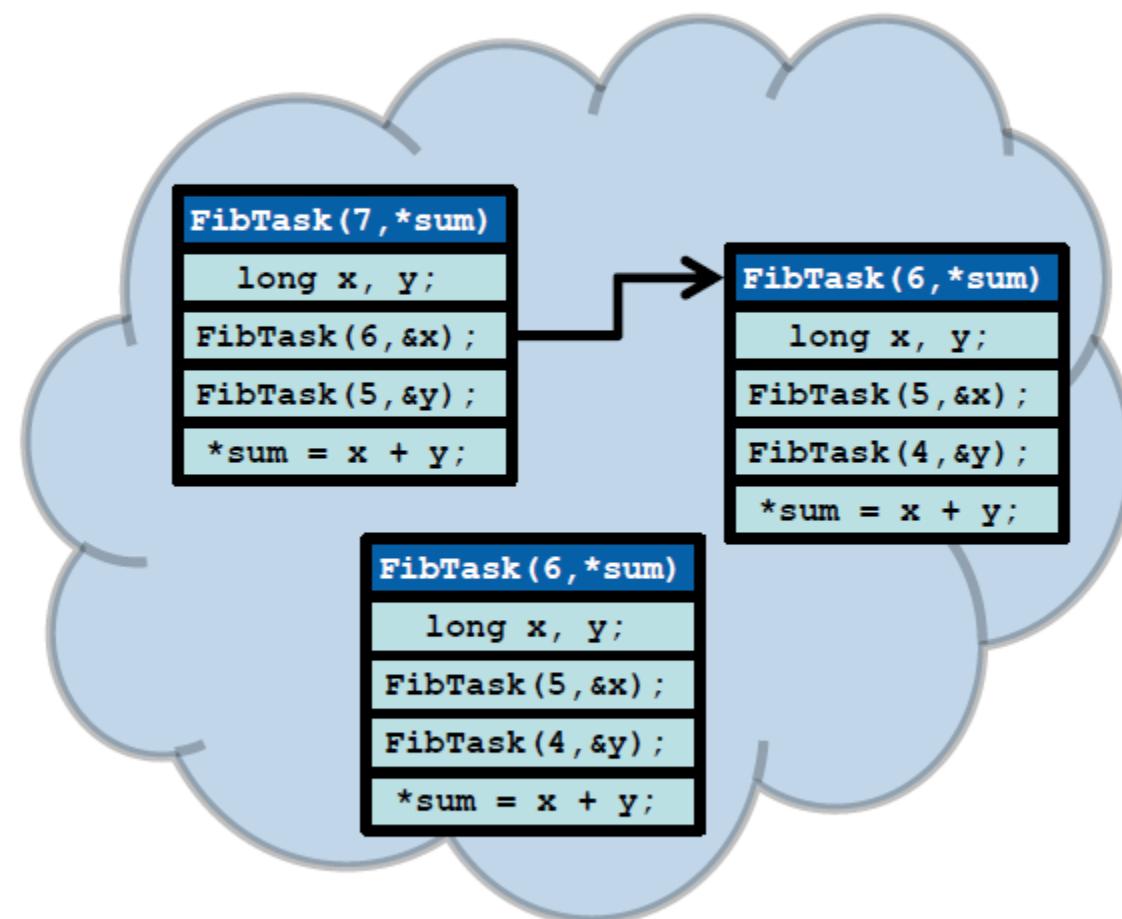
Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. * Other brands and names are the property of their respective owners.

32



Fibonacci Task Solution Example

```
FibTask(8, *sum)  
long x, y;  
FibTask(7, &x);  
FibTask(6, &y);  
*sum = x + y;
```



Copyright © 2009, Intel Corporation. All rights reserved.

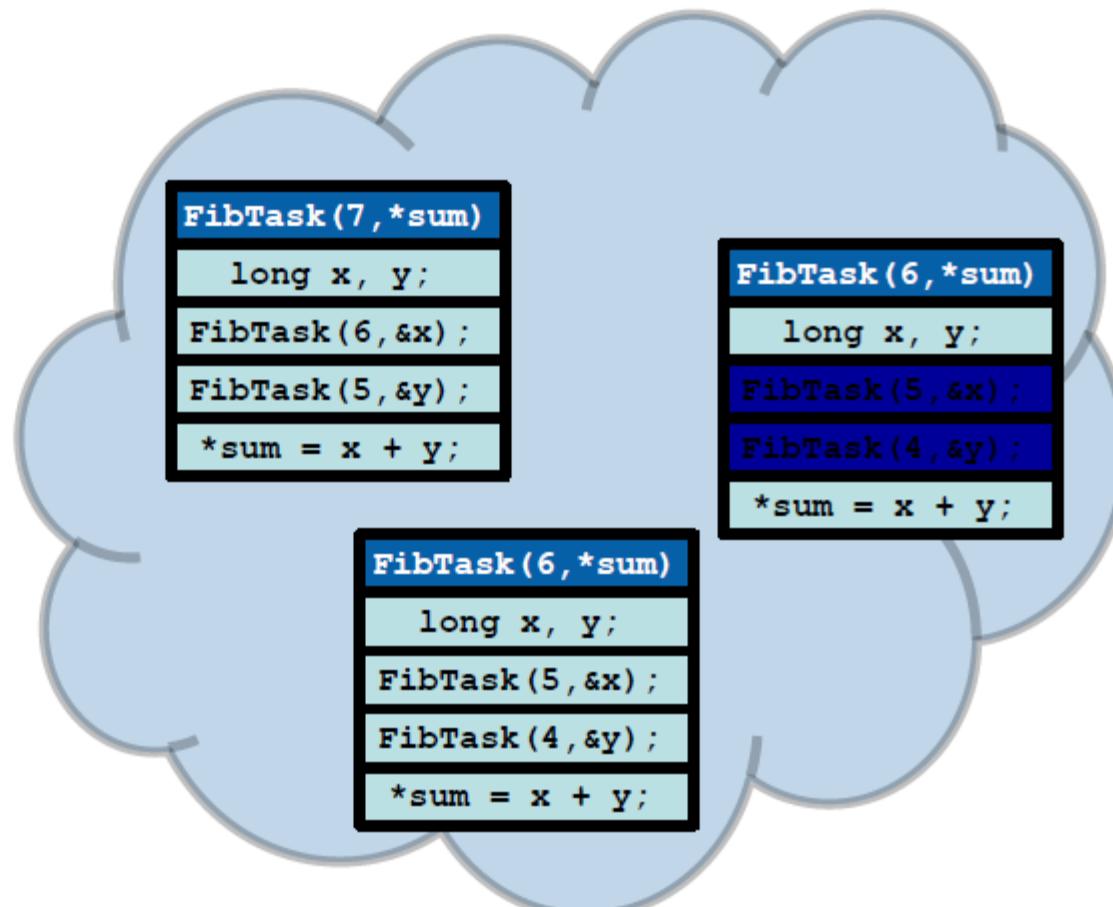
Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. * Other brands and names are the property of their respective owners.

33



Fibonacci Task Solution Example

```
FibTask(8,*sum)
long x, y;
FibTask(7,&x);
FibTask(6,&y);
*sum = x + y;
```



Copyright © 2009, Intel Corporation. All rights reserved.

Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. * Other brands and names are the property of their respective owners.

Fibonacci Task Solution Example

```
FibTask(8,*sum)
long x, y;
FibTask(7,&x);
FibTask(6,&y);
*sum = x + y;
```

```
FibTask(7,*sum)
long x, y;
FibTask(6,&x);
FibTask(5,&y);
*sum = x + y;
```

```
FibTask(6,*sum)
long x, y;
FibTask(5,&x);
FibTask(4,&y);
*sum = x + y;
```

```
FibTask(6,*sum)
long x, y;
FibTask(5,&x);
FibTask(4,&y);
*sum = x + y;
```

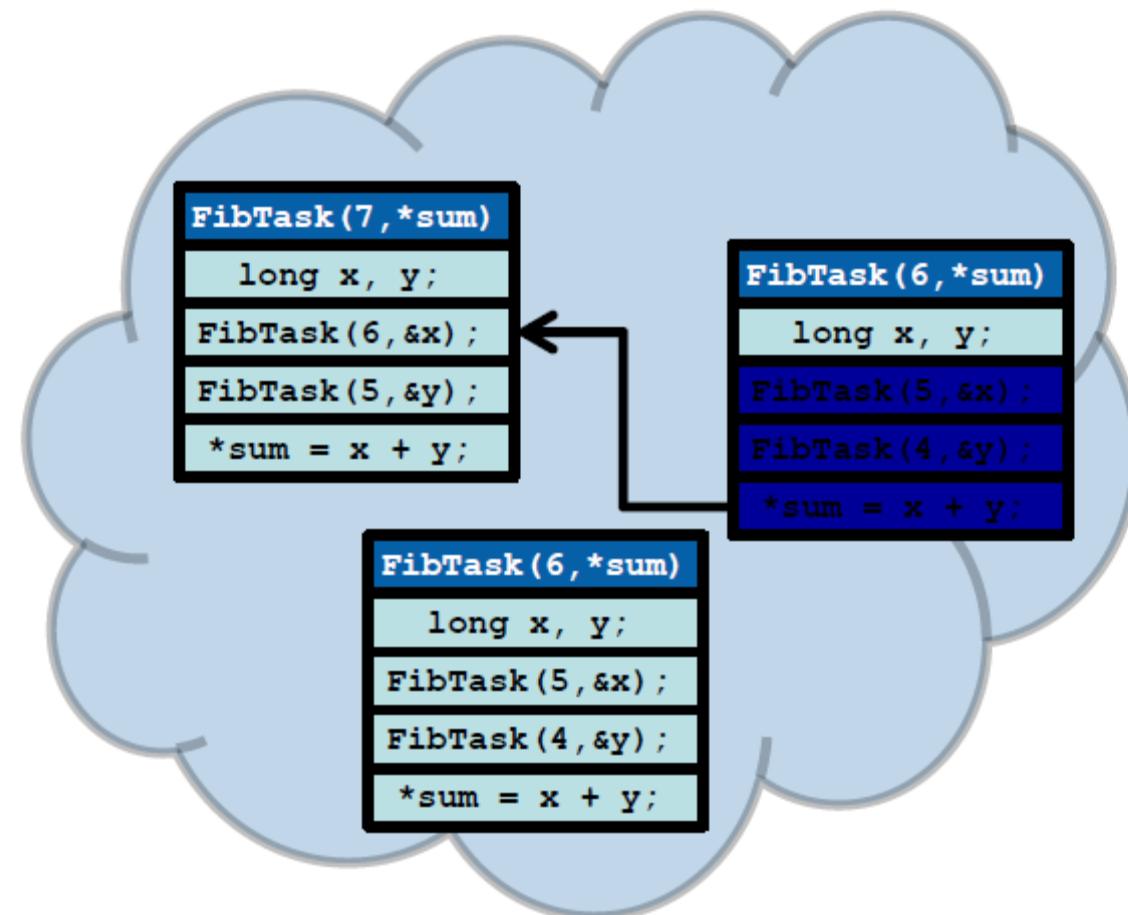


Copyright © 2009, Intel Corporation. All rights reserved.

Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. * Other brands and names are the property of their respective owners.

Fibonacci Task Solution Example

```
FibTask(8,*sum)
long x, y;
FibTask(7,&x);
FibTask(6,&y);
*sum = x + y;
```

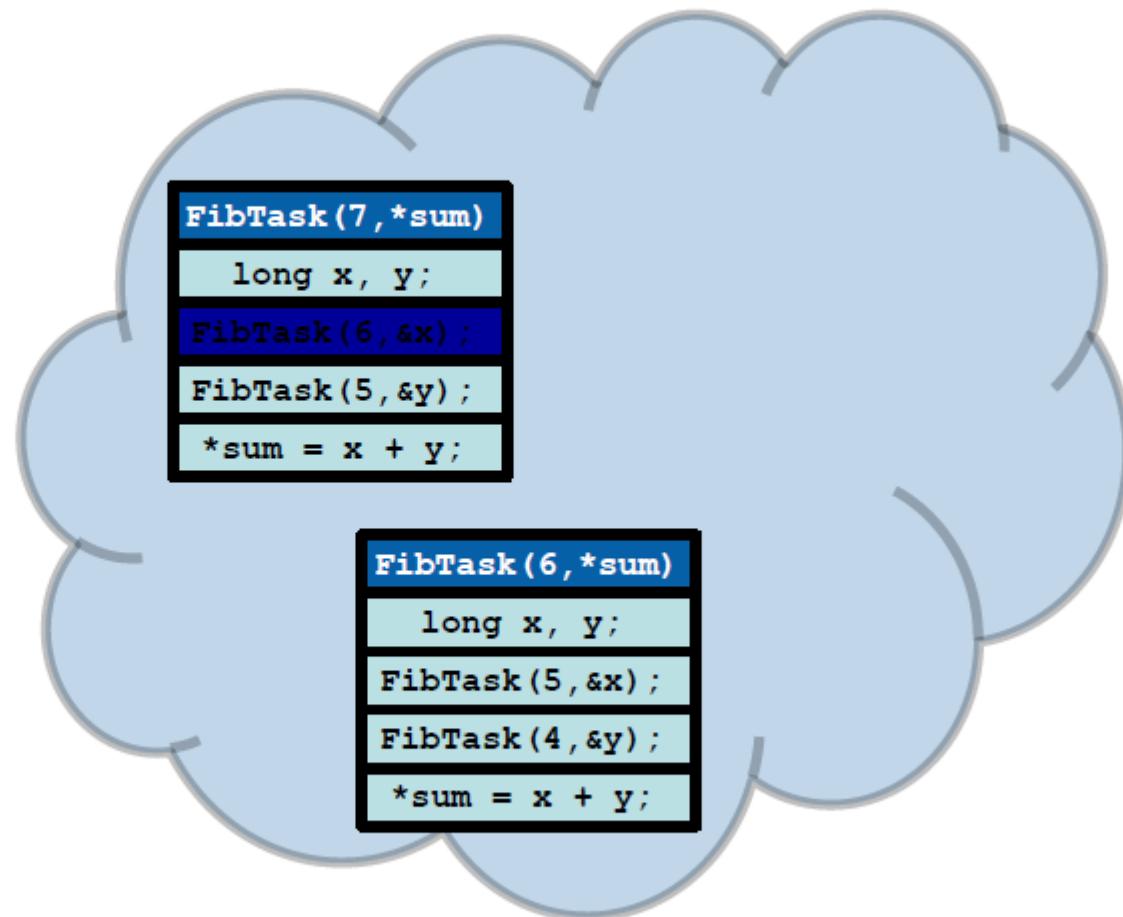


Copyright © 2009, Intel Corporation. All rights reserved.

Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. * Other brands and names are the property of their respective owners.

Fibonacci Task Solution Example

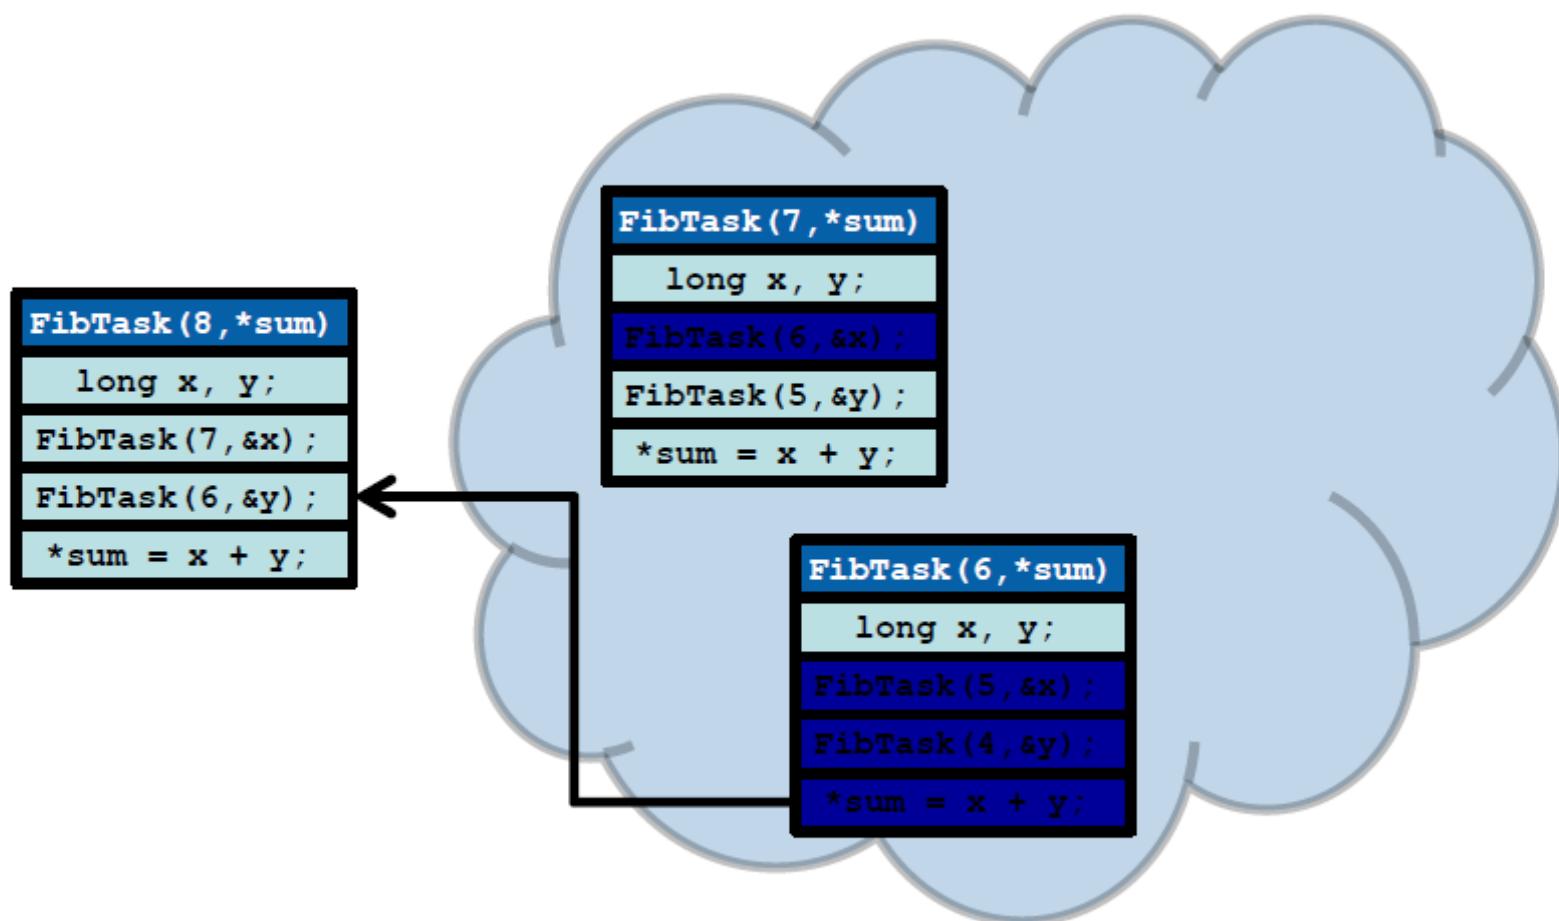
```
FibTask(8,*sum)
long x, y;
FibTask(7,&x);
FibTask(6,&y);
*sum = x + y;
```



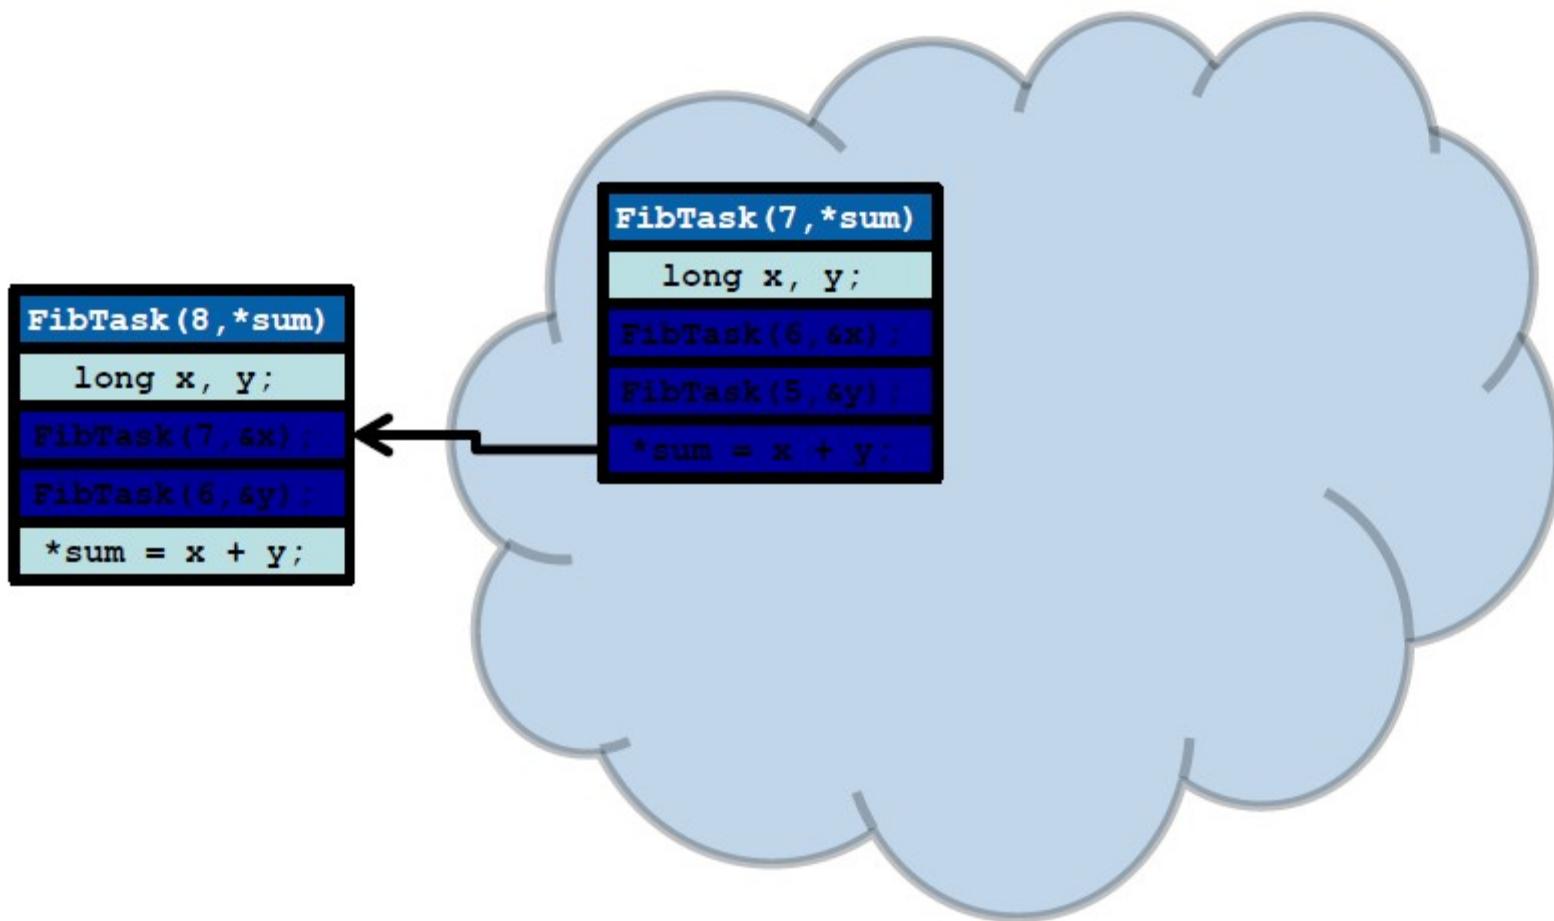
Copyright © 2009, Intel Corporation. All rights reserved.

Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. * Other brands and names are the property of their respective owners.

Fibonacci Task Solution Example



Fibonacci Task Solution Example



Copyright © 2009, Intel Corporation. All rights reserved.

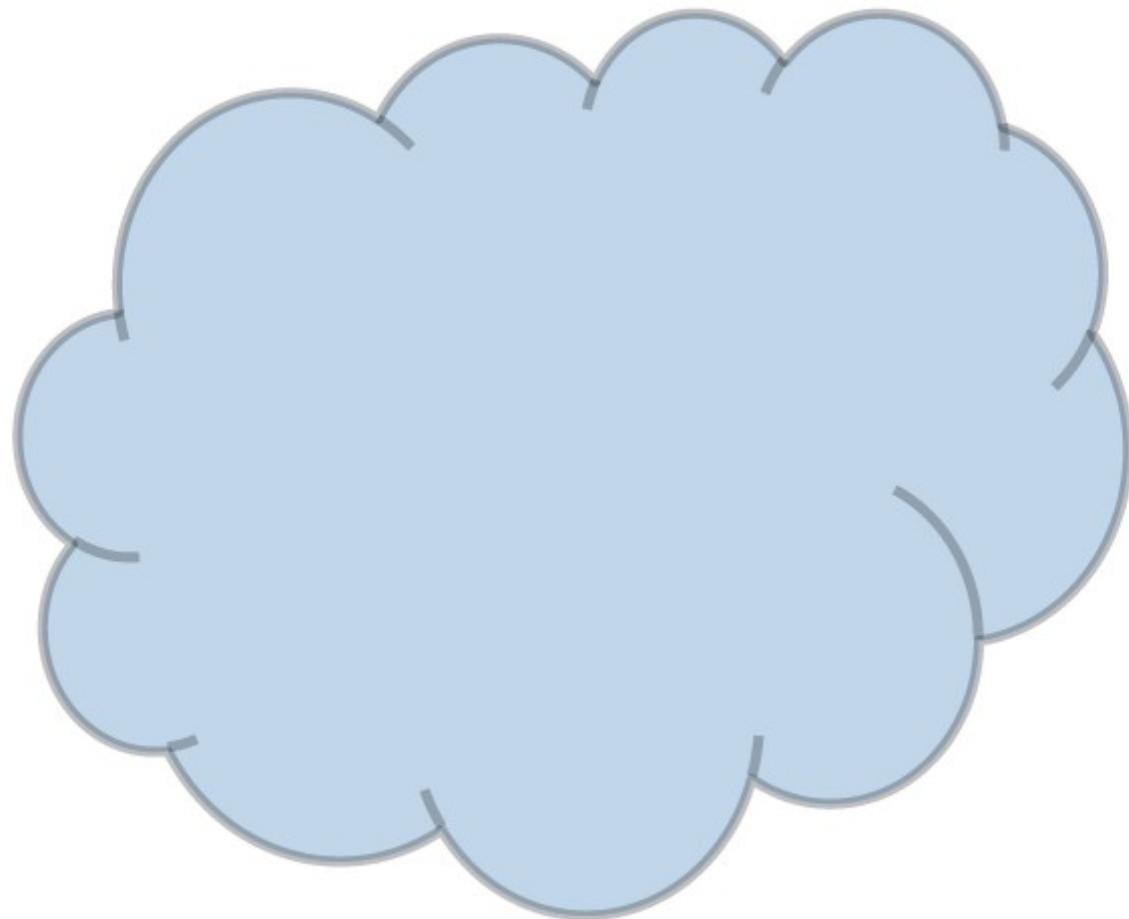
Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States or other countries. * Other brands and names are the property of their respective owners.

39



Fibonacci Task Solution Example

```
FibTask(8, *sum)  
long x, y;  
FibTask(7, &x);  
FibTask(6, &y);  
*sum = x + y;
```



Copyright © 2009, Intel Corporation. All rights reserved.
Intel and the Intel logo are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States
or other countries. * Other brands and names are the property of their respective owners.

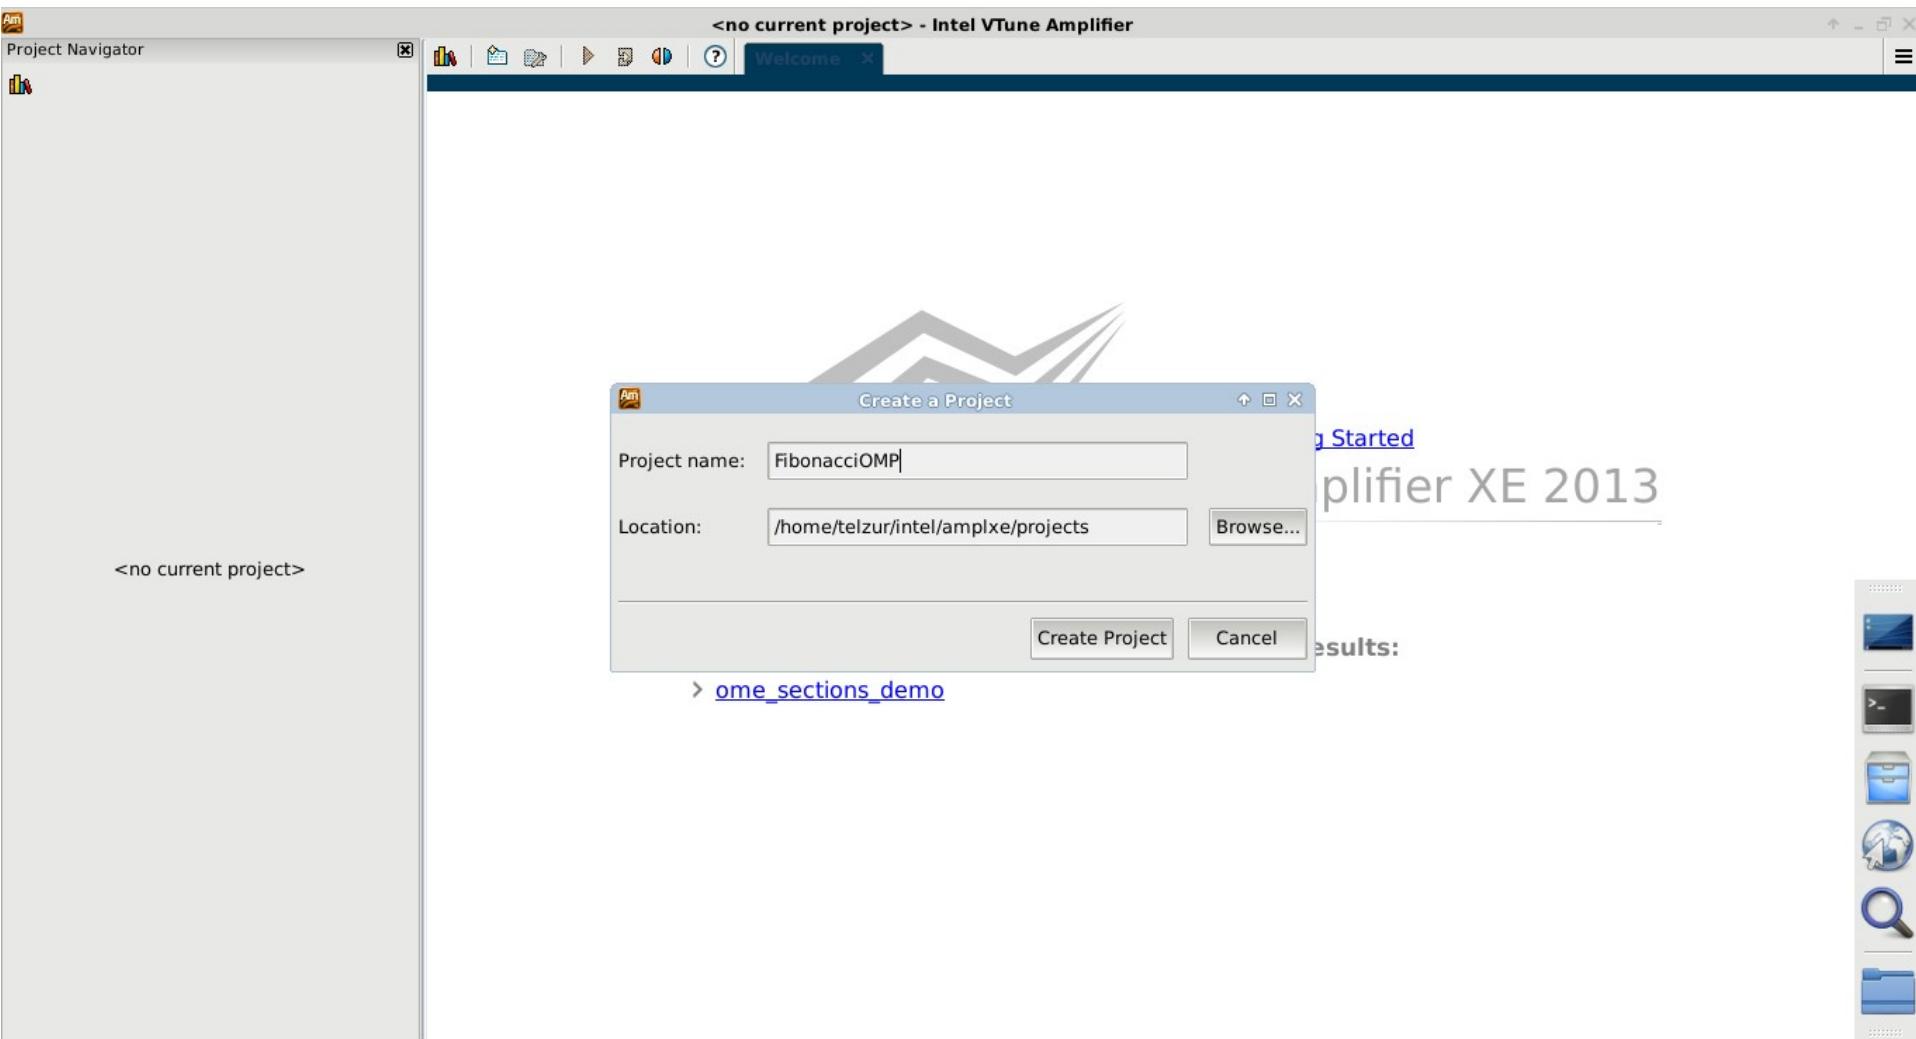
40

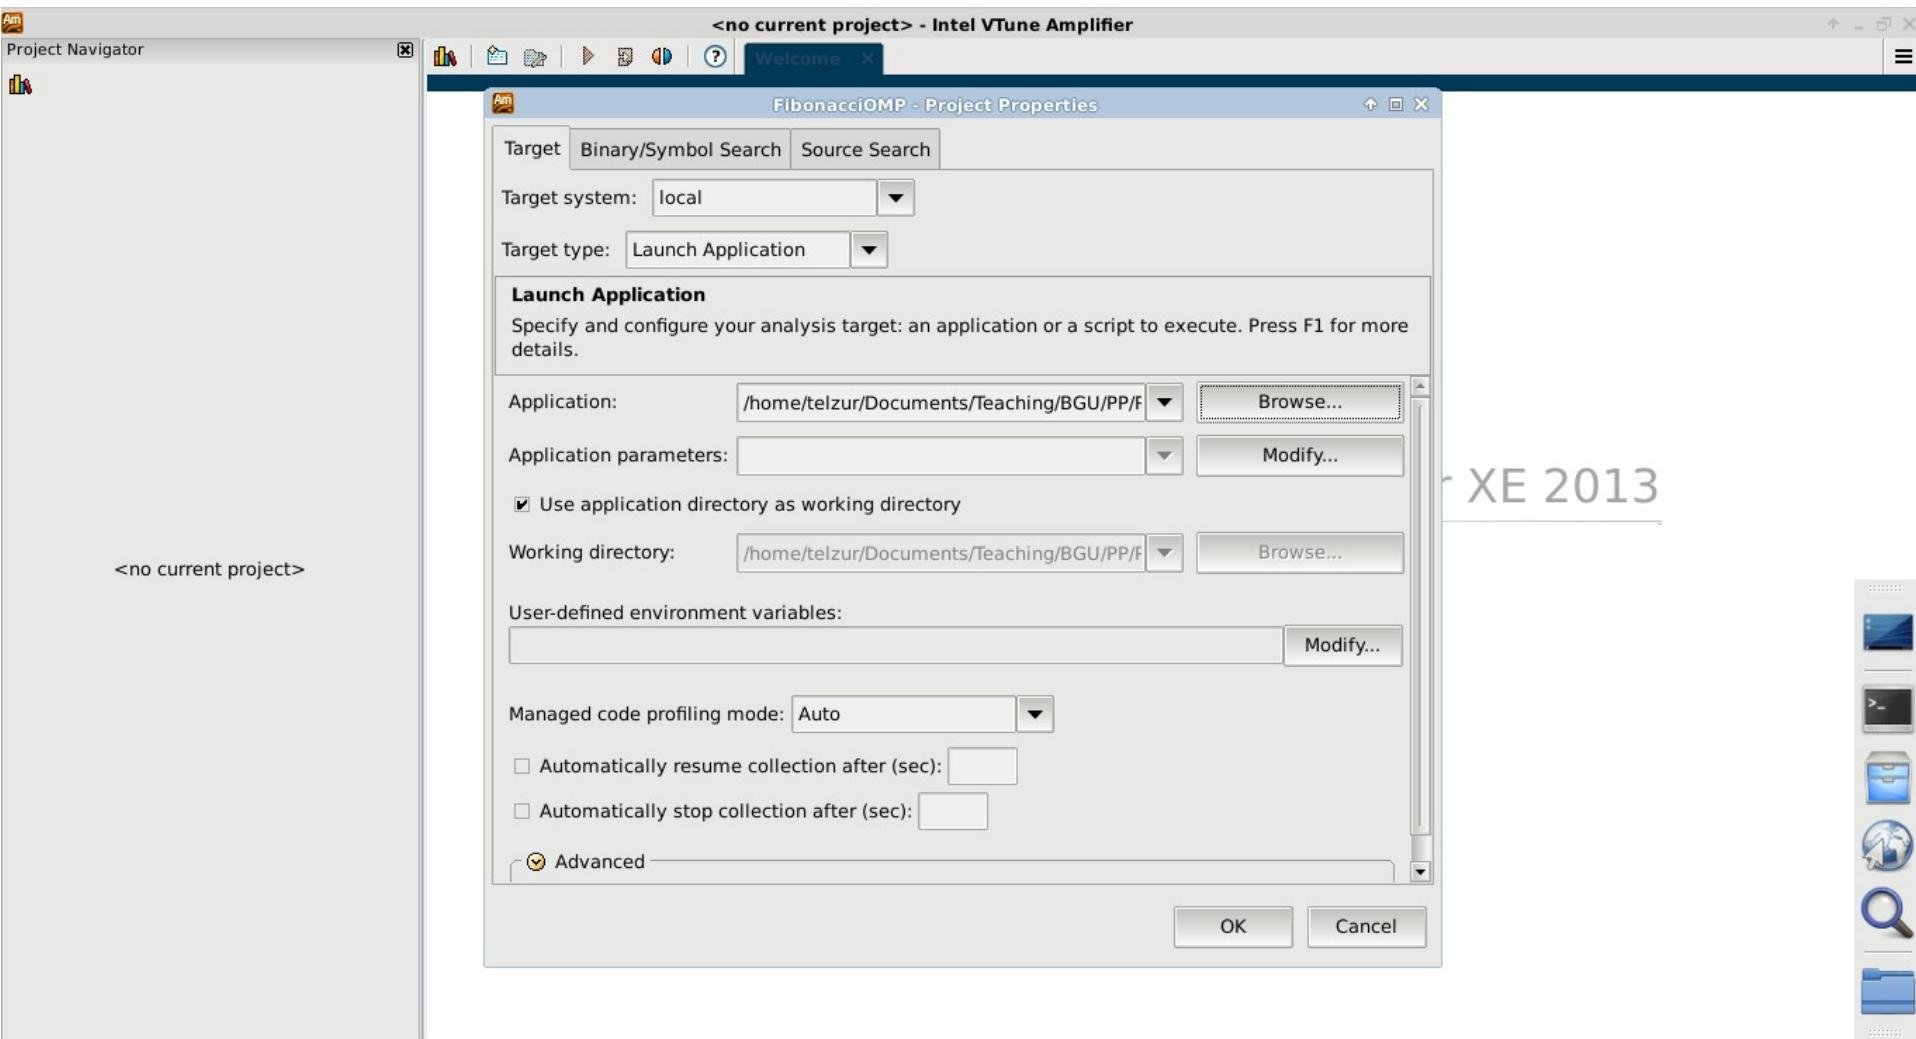


Back to lecture09 presentation

Next slides added by Guy

```
$ cp ./fibonacci.c fibonacci_intel.c  
  
$ /opt/intel/bin/icc -g -openmp -o ./fibonacci_intel  
./fibonacci_intel.c  
  
$ /opt/intel/vtune_amplifier_xe_2013/bin64/amplxe-gui &
```





Project Navigator

/home/telzur/intel/amplxe/projects/FibonacciOMP - Intel VTune Amplifier

Choose Analysis Type

Analysis Type

Algorithm Analysis

- Basic Hotspots
- Advanced Hotspots
- Concurrency
- Locks and Waits

Microarchitecture Analysis

- General Exploration
- Bandwidth

CPU Specific Analysis

- Intel Core 2 Processor
- Nehalem / Westmere
- Sandy Bridge Analysis
- Haswell Analysis

Knights Corner Platform

Custom Analysis

Basic Hotspots

Identify your most time-consuming source code. This analysis type cannot be used to profile the system but must either launch an application/process or attach to one. This analysis type uses user-mode sampling and tracing collection. Press F1 for more details.

CPU sampling interval, ms: 10

Analyze user tasks

Details

Start

Start Paused

Project Properties

Command Line...

This screenshot shows the Intel VTune Amplifier XE 2013 interface. The main window displays the 'Choose Analysis Type' dialog. On the left, there's a tree view of available analysis types under 'Analysis Type'. Under 'Algorithm Analysis', 'Basic Hotspots' is selected. Other options include 'Advanced Hotspots', 'Concurrency', and 'Locks and Waits'. Under 'Microarchitecture Analysis', there are 'General Exploration' and 'Bandwidth' options. A expanded section for 'CPU Specific Analysis' lists 'Intel Core 2 Processor', 'Nehalem / Westmere', 'Sandy Bridge Analysis', and 'Haswell Analysis'. Another section for 'Knights Corner Platform' and 'Custom Analysis' is also visible. The central area of the dialog provides a brief description of 'Basic Hotspots', stating it identifies time-consuming source code and requires launching an application or attaching to one. It uses user-mode sampling and tracing collection. A 'CPU sampling interval' dropdown is set to 10 ms. A checked checkbox allows for analyzing user tasks. A 'Details' button is present. To the right, there are large 'Start' and 'Start Paused' buttons, along with a 'Project Properties' button. The bottom right corner has a 'Command Line...' button. The overall interface is dark-themed with blue and white highlights.

