

Programming with Shared Memory

Part 1

Threads

Accessing shared data

Critical sections

Shared memory multiprocessor system

Any memory location can be accessible by any of the processors.

A *single address space* exists, meaning that each memory location is given a unique address within a single range of addresses.

Programming a shared memory multiprocessor system can take advantage of data stored in the shared memory, which is accessible by all processors without having to send the data to destination through message passing

Shared Memory Programming

Generally, shared memory programming more convenient although it does require access to shared data by different processors to be **carefully** controlled by the programmer.

Shared memory systems have been around for a long time but with the advent of multi-core systems, it has become very important to be able to program for them


Methods for Programming Shared Memory Multiprocessors

- Using heavyweight processes.
- Using threads. Example Pthreads, Java threads
- Using a completely new programming language for parallel programming - not popular. Example **Ada**.
- Modifying the syntax of an existing sequential programming language to create a parallel programming language. Example **UPC – נמצאת על המכונה הוירטואלית**
- Using an existing sequential programming language supplemented with compiler directives and libraries for specifying parallelism. Example **OpenMP**

We will look mostly at threads and OpenMP


Threads and Processes

Added by Guy:



Reference: <http://www.cs.cf.ac.uk/Dave/C/node29.html>

FORK-JOIN construct



UNIX System Calls

No join routine - use `exit()` and `wait()`

SPMD model

```

:
pid = fork();                /* fork */
    Code to be executed by both child and parent
if (pid == 0) exit(0); else wait(0)/* join */
:

```

Wait()

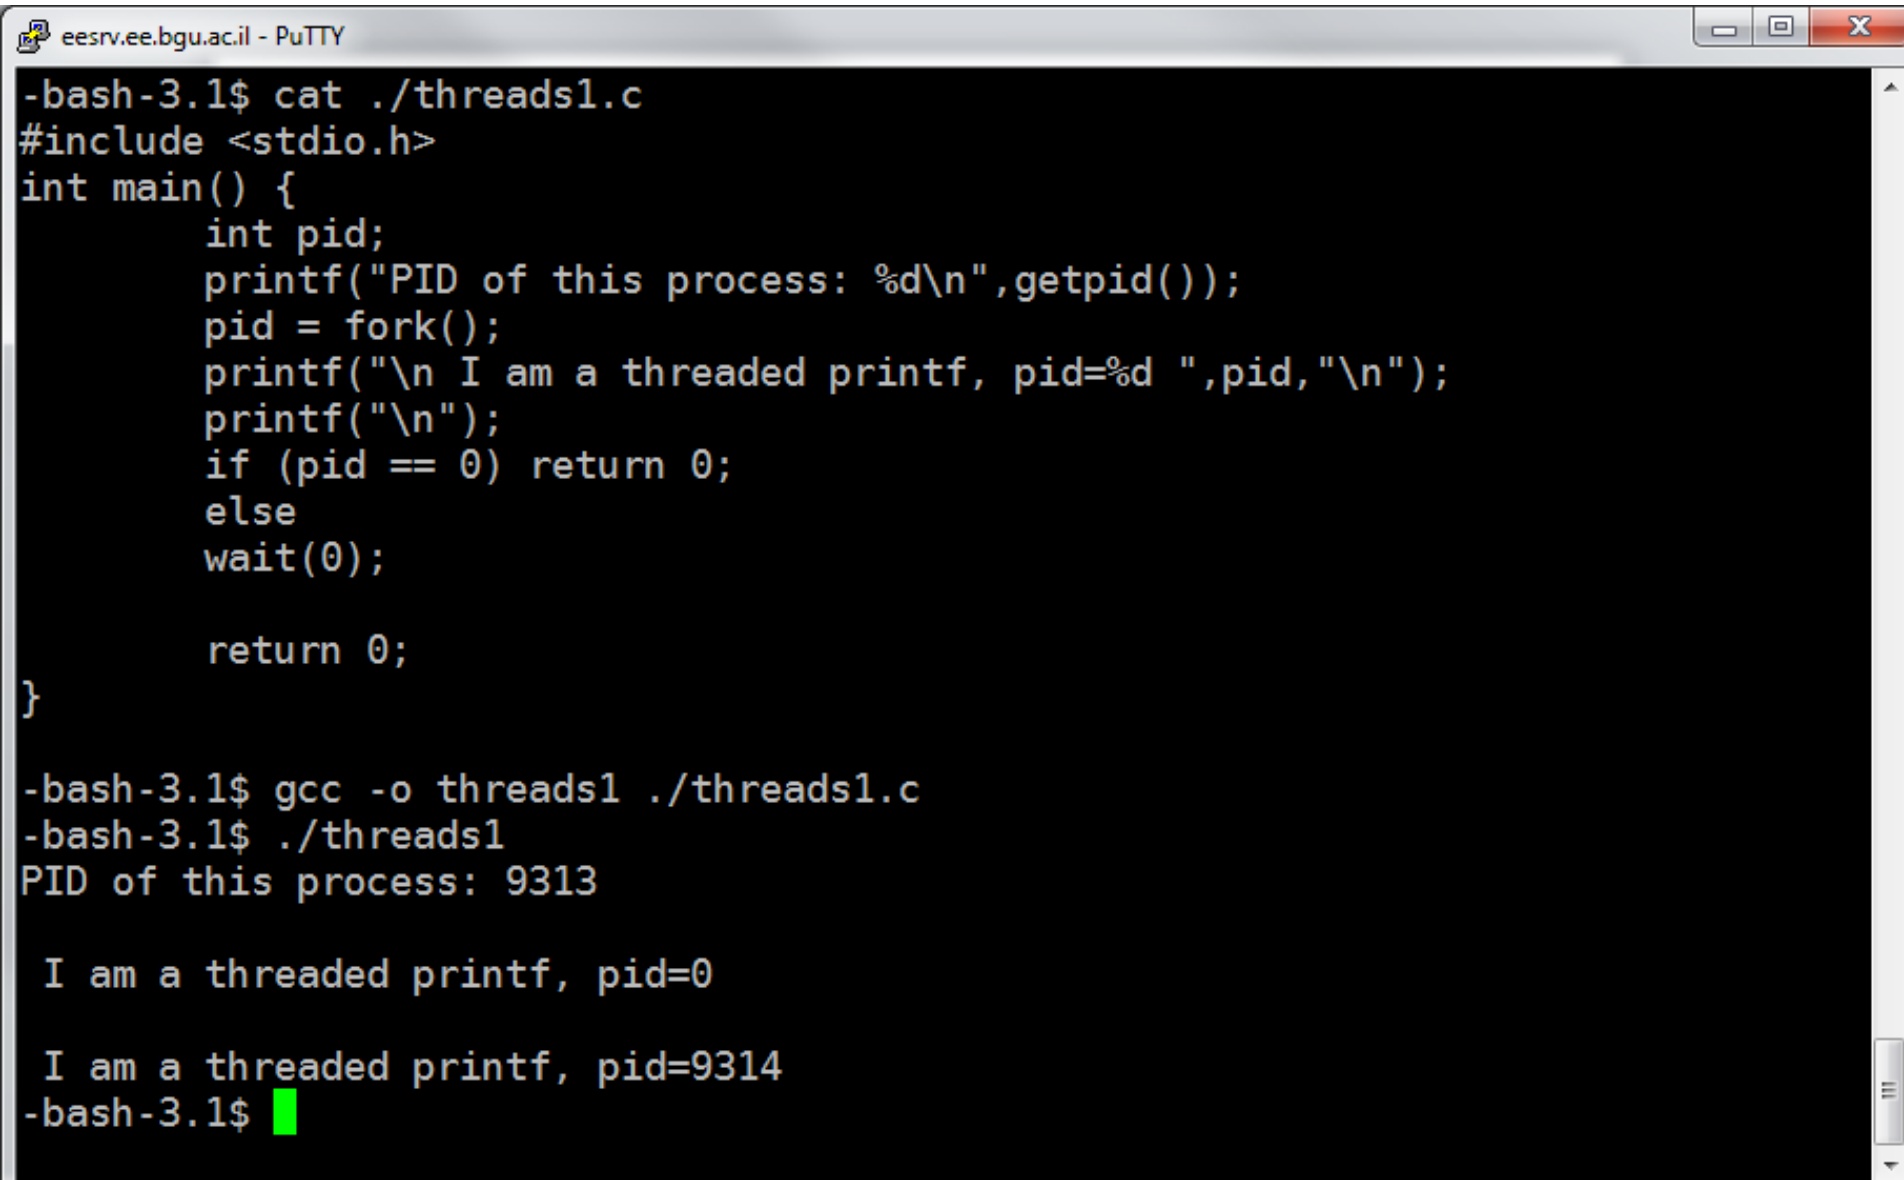
The parent process will often want to wait until all child processes have been completed. this can be implemented with the wait() function call.

Wait() Blocks calling process until the child process terminates. If child process has already terminated, the wait() call returns immediately. if the calling process has multiple child processes, the function returns when one returns.

Open "top":
top -u tel-zur

Guy - demo

Demos under
~/../08/code



```
eesrv.ee.bgu.ac.il - PuTTY
-bash-3.1$ cat ./threads1.c
#include <stdio.h>
int main() {
    int pid;
    printf("PID of this process: %d\n",getpid());
    pid = fork();
    printf("\n I am a threaded printf, pid=%d ",pid,"\n");
    printf("\n");
    if (pid == 0) return 0;
    else
    wait(0);

    return 0;
}

-bash-3.1$ gcc -o threads1 ./threads1.c
-bash-3.1$ ./threads1
PID of this process: 9313

I am a threaded printf, pid=0

I am a threaded printf, pid=9314
-bash-3.1$
```

UNIX System Calls

SPMD model with different code for master process and forked slave process.

```
pid = fork();  
if (pid == 0) {  
    code to be executed by slave  
} else {  
    Code to be executed by parent  
}  
if (pid == 0) exit(0); else wait(0);  
:  
:
```



Spawns a new process!

Guy - demo

eesrv.ee.bgu.ac.il - PuTTY

```
-bash-3.1$ cat ./threads2.c
#include <stdio.h>
int main() {
    int pid;
    printf("PID of this proces: %d\n",getpid());
    pid = fork();
    if (pid==0)
        printf("I am thread 0, pid=%d\n",pid);
    else
        printf("I am thread 1, pid=%d\n",pid);
    if (pid == 0) return 0; else wait(0);

    return 0;
}

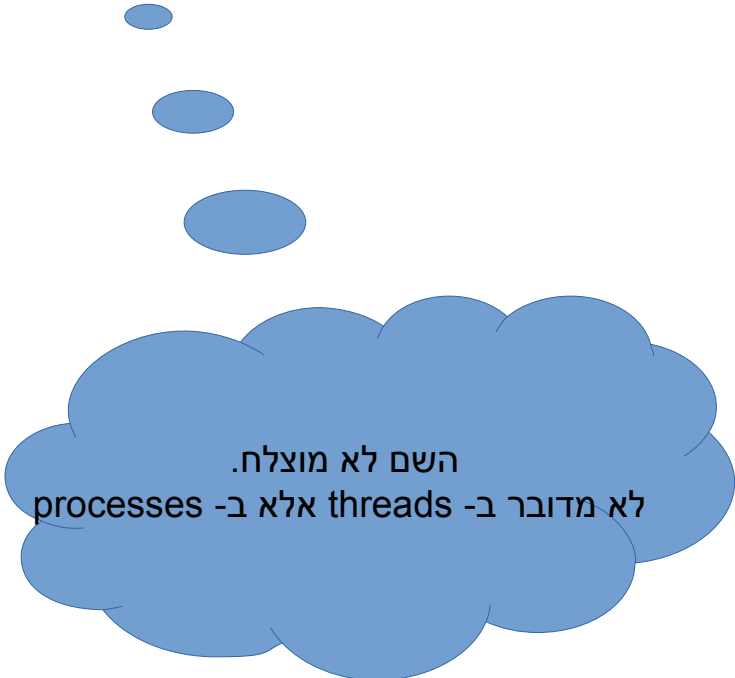
-bash-3.1$ gcc -o threads2 threads2.c
-bash-3.1$ ./threads2
PID of this proces: 9342
I am thread 0, pid=0
I am thread 1, pid=9343
-bash-3.1$
```

...same program with a delay added (cpu_burn): threads6.c

```
Threads> ./threads6  
PID of this process: 22498
```

```
I am a threaded printf,  
pid=22499
```

```
I am a threaded printf, pid=0  
^C
```



השם לא מוצלח.
לא מדובר ב- threads אלא ב- processes

fork() starts a new process (not a new thread!)

telzur@GL553VD ~/Documents/Teaching/PP2019A/lectures/08/code/Threads/solarisstudio

File Edit View Search Terminal Help

```
top - 14:18:48 up 7 days, 18:48, 1 user, load average: 1.74, 1.21, 0.68
Threads: 1589 total, 3 running, 1506 sleeping, 0 stopped, 2 zombie
%Cpu(s): 25.6 us, 0.4 sy, 0.0 ni, 73.9 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem : 16307588 total, 276600 free, 11052272 used, 4978716 buff/cache
KiB Swap: 16659452 total, 15657860 free, 1001592 used. 4745948 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
22498	telzur	20	0	4508	860	796	R	99.9	0.0	0:58.75	threads6
22499	telzur	20	0	4508	80	0	R	99.9	0.0	0:58.70	threads6
1472	root	20	0	1147328	265628	107512	S	2.6	1.6	52:09.70	Xorg
22504	telzur	20	0	43472	5248	3068	R	1.6	0.0	0:00.66	top
22515	telzur	20	0	495552	30780	24156	S	1.0	0.2	0:00.13	mate-screenshot
8201	telzur	20	0	1121324	25500	17068	S	0.7	0.2	0:56.75	mate-settings-d
8232	telzur	20	0	869488	44680	15208	S	0.7	0.3	6:37.10	marco
8266	telzur	20	0	526136	49768	14952	S	0.7	0.3	1:04.26	wnck-applet
1675	root	20	0	1147328	265628	107512	S	0.3	1.6	3:18.81	InputThread

This is the output of `top -h`



C threads6.c x



```
3 // This is needed in order to watch at the same time
4 // 'top' (task manager)
5 // Guy Tel-Zur, December 7, 2018
6 #include<stdio.h>
7 void cpu_burn() {
8     int iMAX = 1000000;
9     int i,j;
10    float fNORM, fMAX;
11    float x,y;
12    fMAX = iMAX; // convert to float
13    fNORM = fMAX * fMAX;
14    for (j=0;j<iMAX;j++)
15        for (i=0;i<iMAX;i++) {
16            y = i; // convert to float
17            x = y*y/fNORM;
18        }
19
20    return 0;
21 }
22 int main() {
23     int pid;
24
25     printf("PID of this process: %d\n",getpid());
26     pid = fork();
27     printf("\n I am a threaded printf, pid=%d ",pid);
28     printf("\n");
29     cpu_burn();
30     if (pid == 0) return 0;
31     else
32         wait(0);
33 }
```



```
Threads> ./threads6  
PID of this process: 23728
```


```
I am a threaded printf, pid=23729
```

```
I am a threaded printf, pid=23729
```

```
I am a threaded printf, pid=0
```

```
I am a threaded printf, pid=0
```

```
^C
```

```
Threads>   
File Edit View Search Terminal Help
```

```
top - 14:28:44 up 7 days, 18:58, 1 user, load average: 2.91, 1.59, 1.11  
threads: 1583 total, 5 running, 1502 sleeping, 0 stopped, 2 zombie  
Cpu(s): 51.2 us, 2.5 sy, 0.0 ni, 46.3 id, 0.0 wa, 0.0 hi, 0.0 si, 0.0 st  
MiB Mem : 16307588 total, 409376 free, 11073884 used, 4824328 buff/cache  
MiB Swap: 16659452 total, 15664772 free, 994680 used. 4727368 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
3729	telzur	20	0	4508	80	0	R	99.9	0.0	1:15.93	threads6
3730	telzur	20	0	4508	80	0	R	99.9	0.0	1:15.93	threads6
3731	telzur	20	0	4508	80	0	R	99.9	0.0	1:15.92	threads6
3728	telzur	20	0	4508	820	756	R	99.7	0.0	1:15.88	threads6
1472	root	20	0	1146048	272212	106700	S	5.9	1.7	52:28.87	Xorg
3726	telzur	20	0	43472	5232	3048	R	1.3	0.0	0:01.30	top
8568	telzur	20	0	5675736	1.788g	115904	S	0.7	11.5	67:12.63	Chrome_IOThread
3917	telzur	20	0	495532	30880	24268	S	0.7	0.2	0:00.12	mate-screenshot
568	root	-51	0	0	0	0	S	0.3	0.0	8:50.58	irq/128-iwlwifi
1213	telzur	20	0	2284352	494708	65468	S	0.3	3.0	9:24.36	chromium-browser
1629	root	-51	0	0	0	0	S	0.3	0.0	9:00.68	irq/130-nvidia
7057	telzur	20	0	1007324	107420	61722	S	0.3	1.2	0:26.22	chromium-browser

4 processes



C threads6.c x


```
11     float x,y;
12     fMAX = iMAX; // convert to float
13     fNORM = fMAX * fMAX;
14     for (j=0;j<iMAX;j++)
15     {
16         for (i=0;i<iMAX;i++) {
17             y = i; // convert to float
18             x = y*y/fNORM;
19         }
20     }
21     return 0;
22 }
23 int main() {
24     int pid,qid;
25     printf("PID of this process: %d\n",getpid());
26     pid = fork();
27     qid = fork();
28     printf("\n I am a threaded printf, pid=%d ",pid);
29     printf("\n");
30     cpu_burn();
31     if (pid == 0) return 0;
32     else
33         wait(0);
34
35     return 0;
36 }
37
38
```



Differences between a process and threads

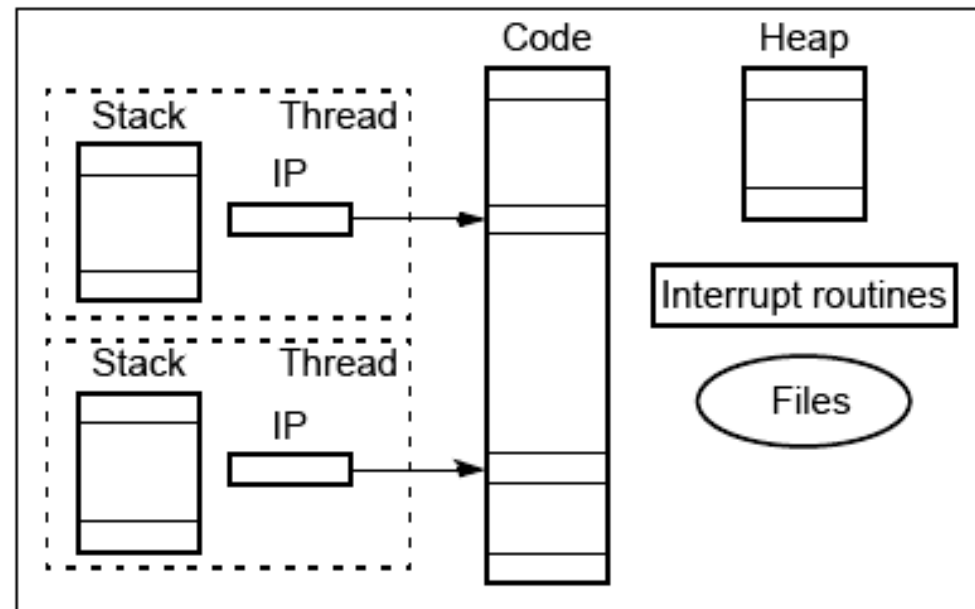
“heavyweight” process - completely separate program with its own variables, stack, and memory allocation.

(a) Process



Threads - shares the same memory space and global variables between routines.

(b) Threads



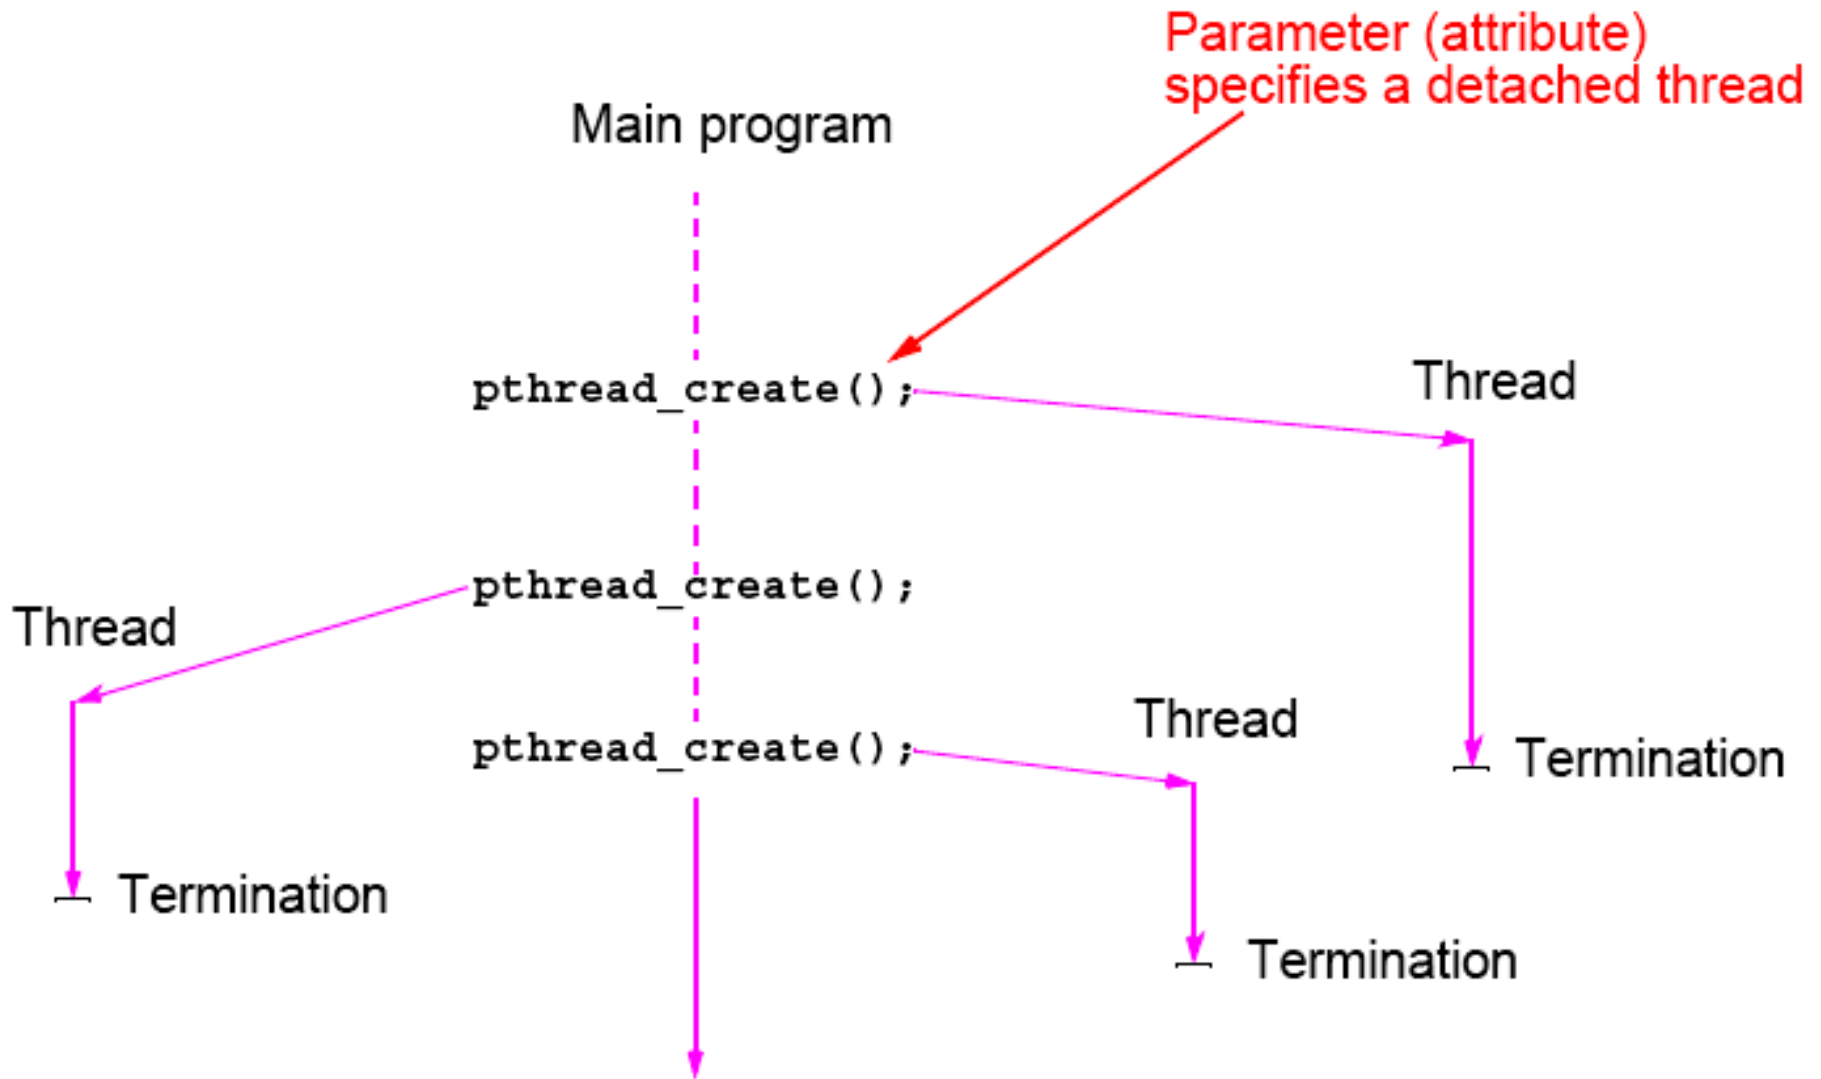
Detached Threads

It may be that thread are not bothered when a thread it creates terminates and then a join not needed.

Threads not joined are called *detached threads*.

When detached threads terminate, they are destroyed and their resource released.

Pthreads Detached Threads



Statement Execution Order

Single processor: Processes/threads typically executed until blocked.

Multiprocessor: Instructions of processes/threads interleaved in time.

Example

Process 1

Instruction 1.1

Instruction 1.2

Instruction 1.3

Process 2

Instruction 2.1

Instruction 2.2

Instruction 2.3

Many possible orderings, e.g.:

Instruction 1.1

Instruction 1.2

Instruction 2.1

Instruction 1.3

Instruction 2.2

Instruction 2.3

assuming instructions cannot be divided into smaller steps.

A sequence of instructions might perform a specific task, for example print out a message.

If two processes were to print messages, using interleaved instructions, the messages could appear garbled - the individual characters of each message could be interleaved if special care is not taken in coding the print routines.

Thread-Safe Routines

Thread safe if they can be called from multiple threads simultaneously and always produce correct results.

Standard I/O thread safe (prints messages without interleaving the characters).

System routines that return time *may not be thread safe*.

Routines that access shared data may require special care to be made thread safe.

Re-ordering code also done intentionally by:

- Compilers to optimize performance
- Processors internally, again to optimize performance

Compilers will re-order code prior to execution while processors will re-order the code during execution.

In both cases the objective is to best utilize the available computer resources and minimize any waiting time.

Compiler/Processor Optimizations

Compiler and processor reorder instructions to improve performance (execution speed).

Example

Suppose one had the code:

```
·  
·  
a = b + 5;  
x = y * 4;  
p = x + 9;  
·  
·
```

and the processor can perform, as is usual, multiple arithmetic operations at the same time.

Compiler/Processor Optimizations

Can reorder to:

```
·  
·  
x = y * 4;  
a = b + 5;  
p = x + 9;  
·  
·
```

and still be logically correct. This gives the multiply operation longer time to complete before the result (x) is needed in the last statement.

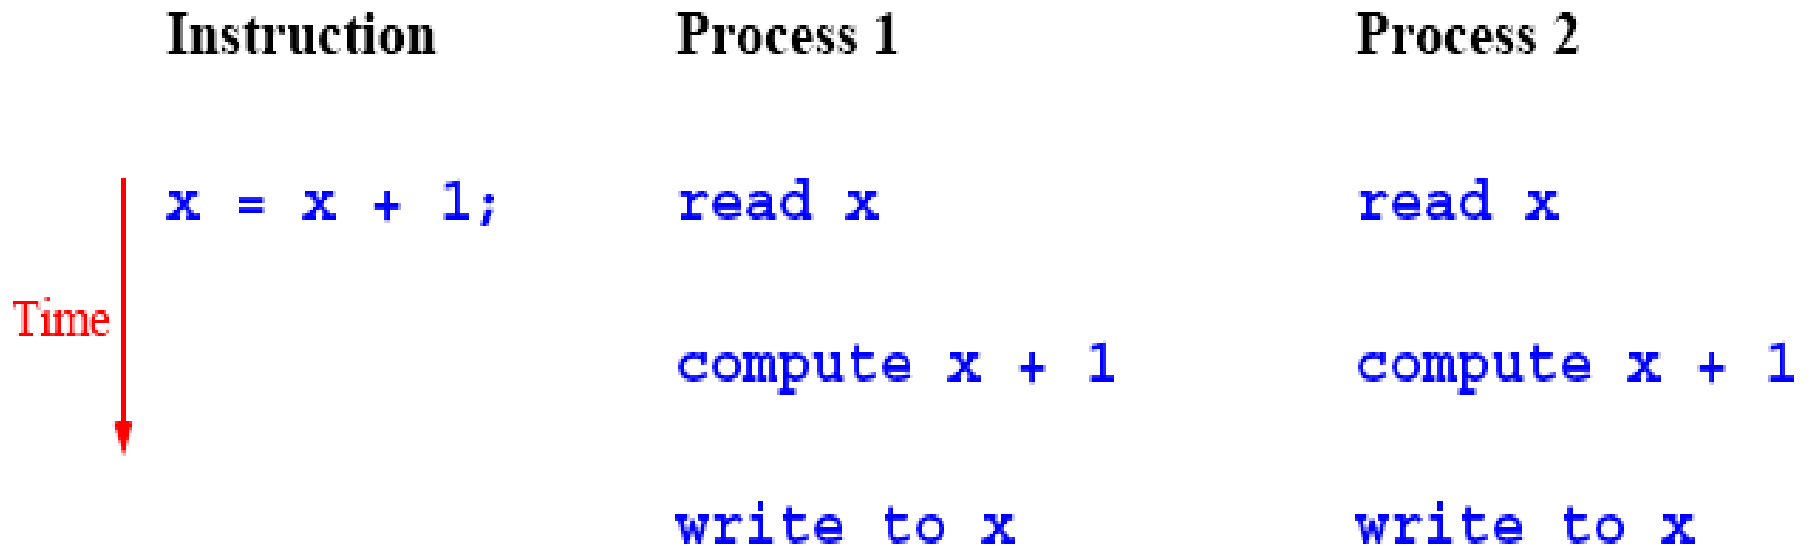
Very common for processors to execute machines instructions out of order for increased speed.

Accessing Shared Data


Accessing shared data needs careful control.

Consider two processes each of which is to add one to a shared data item, x .

Location x is read, $x + 1$ computed, and the result written back to the location:



Conflict in accessing shared variable



Guy – demo

vdwarf2.ee.bgu.ac.il - PuTTY

```
vdwarf2.ee.bgu.ac.il> cat threads3.c
```

```
int main() {  
    int pid;  
    int x = 0;  
    pid = fork();  
    if (pid==0) {  
        x = x+1;  
        printf("I am thread 0, x=%d\n",x);  
    }  
    else {  
        x = x+1;  
        printf("I am thread 1, x=%d\n",x);  
    }  
    if (pid == 0) exit (0); else wait(0);  
    printf("Finally x=%d\n",x);  
    return 0;  
}
```

```
vdwarf2.ee.bgu.ac.il> ./threads3
```

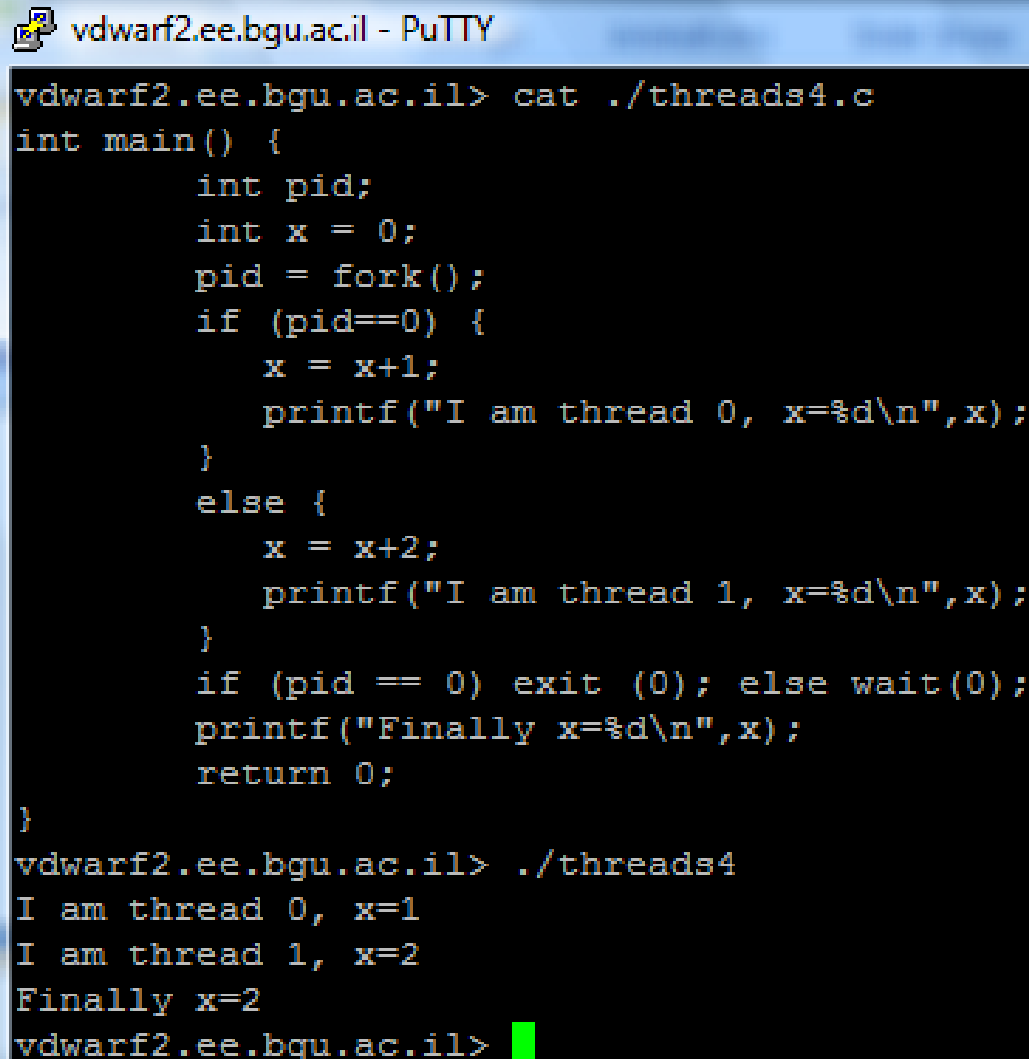
```
I am thread 0, x=1
```

```
I am thread 1, x=1
```

```
Finally x=1
```

```
vdwarf2.ee.bgu.ac.il> █
```

Guy – demo; x is following the master thread



The screenshot shows a PuTTY terminal window with the title 'vdwarf2.ee.bgu.ac.il - PuTTY'. The terminal displays the source code of a C program named 'threads4.c' and its execution output. The program uses 'fork()' to create two threads. Thread 0 increments 'x' by 1, and Thread 1 increments 'x' by 2. Both threads then wait for each other using 'wait(0)'. The final output shows 'x=2', indicating that the master thread's updates to 'x' are visible to the child threads.

```
vdwarf2.ee.bgu.ac.il> cat ./threads4.c
int main() {
    int pid;
    int x = 0;
    pid = fork();
    if (pid==0) {
        x = x+1;
        printf("I am thread 0, x=%d\n",x);
    }
    else {
        x = x+2;
        printf("I am thread 1, x=%d\n",x);
    }
    if (pid == 0) exit (0); else wait(0);
    printf("Finally x=%d\n",x);
    return 0;
}
vdwarf2.ee.bgu.ac.il> ./threads4
I am thread 0, x=1
I am thread 1, x=2
Finally x=2
vdwarf2.ee.bgu.ac.il> █
```

Critical Section

A mechanism for ensuring that only one process accesses a particular resource at a time.

critical section – a section of code for accessing resource
Arrange that only one such critical section is executed at a time.

This mechanism is known as *mutual exclusion*.

Concept also appears in an operating systems.

Locks

Simplest mechanism for ensuring mutual exclusion of critical sections.

A lock - a 1-bit variable that is a 1 to indicate that a process has entered the critical section and a 0 to indicate that no process is in the critical section.

Operates much like that of a door lock:

A process coming to the “door” of a critical section and finding it open may enter the critical section, locking the door behind it to prevent other processes from entering. Once the process has finished the critical section, it unlocks the door and leaves.

Control of critical sections through busy waiting

Must be indivisible

Process 1

```
while (lock == 1) do_nothing;  
lock = 1;
```

Critical section

lock = 0;

Process 2

```
while (lock == 1) do_nothing;
```

Better to deschedule
process

lock = 1;

Critical section

lock = 0;

Critical Sections Serializing Code

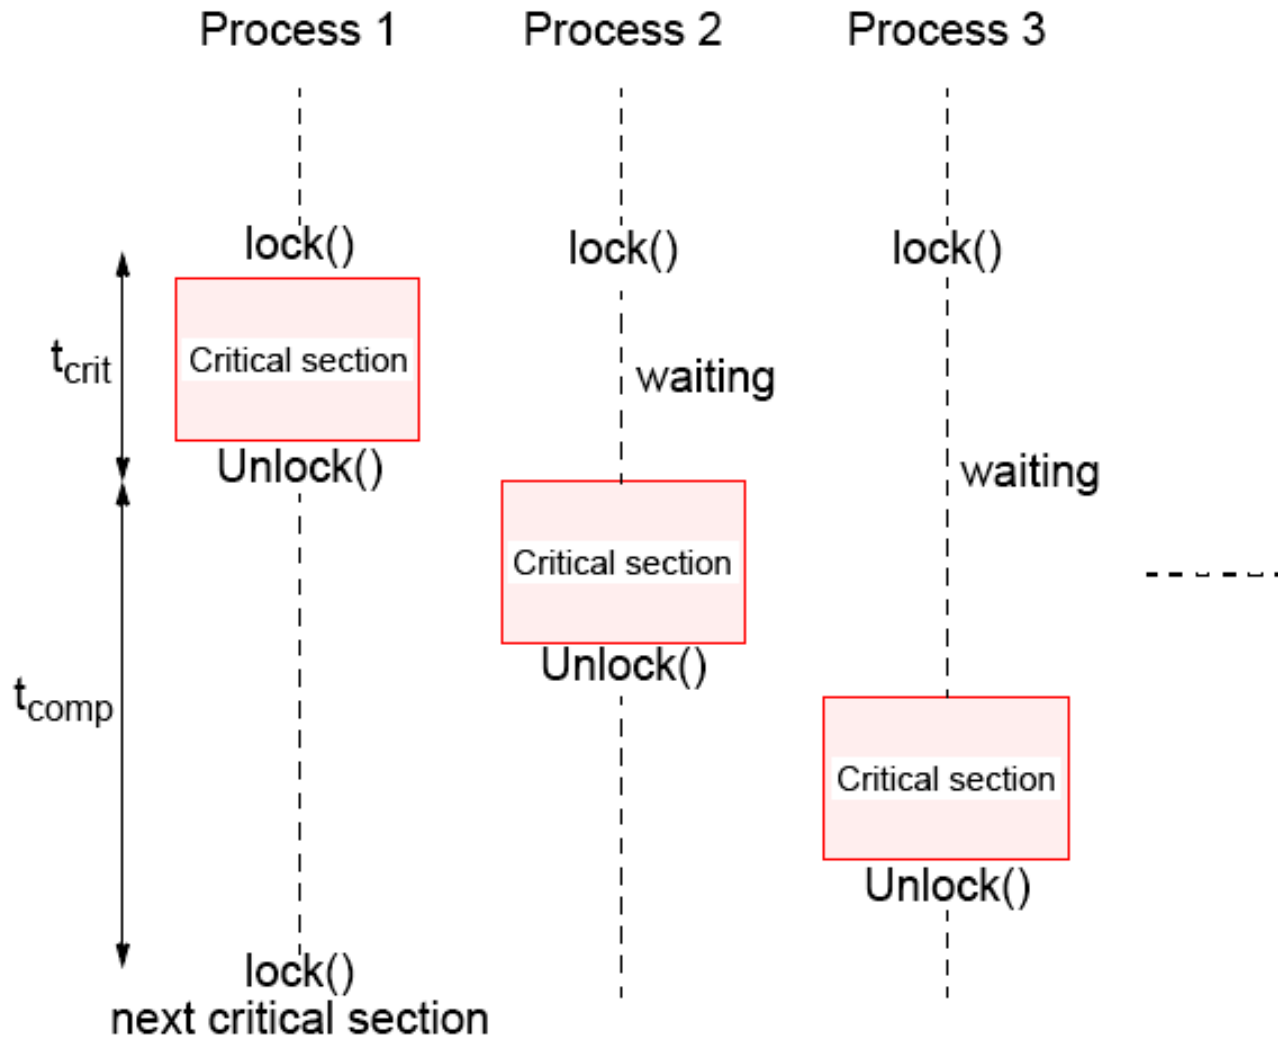
High performance programs should have as few as possible critical sections as their use can serialize the code.

Suppose, all processes happen to come to their critical section together.

They will execute their critical sections one after the other.

In that situation, the execution time becomes almost that of a single processor.

Illustration




When $t_{comp} < pt_{crit}$, less than p processor will be active

Deadlock


Can occur with two processes when one requires a resource held by the other, and this process requires a resource held by the first process.

Two-process deadlock



Deadlock (deadly embrace)

Deadlock can also occur in a circular fashion with several processes having a resource wanted by another.



Semaphores

A positive integer (including zero) operated upon by two operations:

P operation on semaphore s

Waits until s is greater than zero and then decrements s by one and allows the process to continue.

V operation on semaphore s

Increments s by one and releases one of the waiting processes (if any).

P and **V** operations are performed indivisibly.

Mechanism for activating waiting processes implicit in **P** and **V** operations.

Though exact algorithm not specified, algorithm expected to be fair. Processes delayed by **P**(s) are kept in abeyance(*) until released by a **V**(s) on the same semaphore.

השהייה, אי-הפעלה (*)

Devised by Dijkstra in 1968.

Letter **P** from Dutch word *passeren*, meaning “to pass”

Letter **V** from Dutch word *vrijgeven*, meaning “to release”

Mutual exclusion of critical sections can be achieved with one semaphore having the value 0 or 1 (a binary semaphore), which acts as a lock variable, but the P and V operations include a process scheduling mechanism:

Process 1

Noncritical section

.

P(s)

Critical section

V(s)

.

Noncritical section

Process 2

Noncritical section

.

P(s)

Critical section

V(s)

.

Noncritical section

Process 3

Noncritical section

.

P(s)

Critical section

V(s)

.


Noncritical section

Monitor (=condition variable + lock)

Suite of procedures that provides only way to access shared resource. Only one process can use a monitor procedure at any instant.

Could be implemented using a semaphore or lock to protect entry, i.e.:

```
monitor_proc1() {  
lock(x);  
  
    ▪  
monitor body  
  
    ▪  
unlock(x);  
return;  
}
```

Source: [http://en.wikipedia.org/wiki/Monitor_\(synchronization\)](http://en.wikipedia.org/wiki/Monitor_(synchronization))

Program example

To sum the elements of an array, **a[1000]**:

```
int sum, a[1000];  
    sum = 0;  
    for (i = 0; i < 1000; i++)  
        sum = sum + a[i];
```

UNIX Processes

Calculation will be divided into two parts, one doing even i and one doing odd i ; i.e.,

Process 1

```
sum1 = 0;  
for (i = 0; i < 1000; i = i + 2)  
    sum1 = sum1 + a[i];
```

Process 2

```
sum2 = 0;  
for (i = 1; i < 1000; i = i + 2)  
    sum2 = sum2 + a[i];
```

Each process will add its result (sum1 or sum2) to an accumulating result, sum :

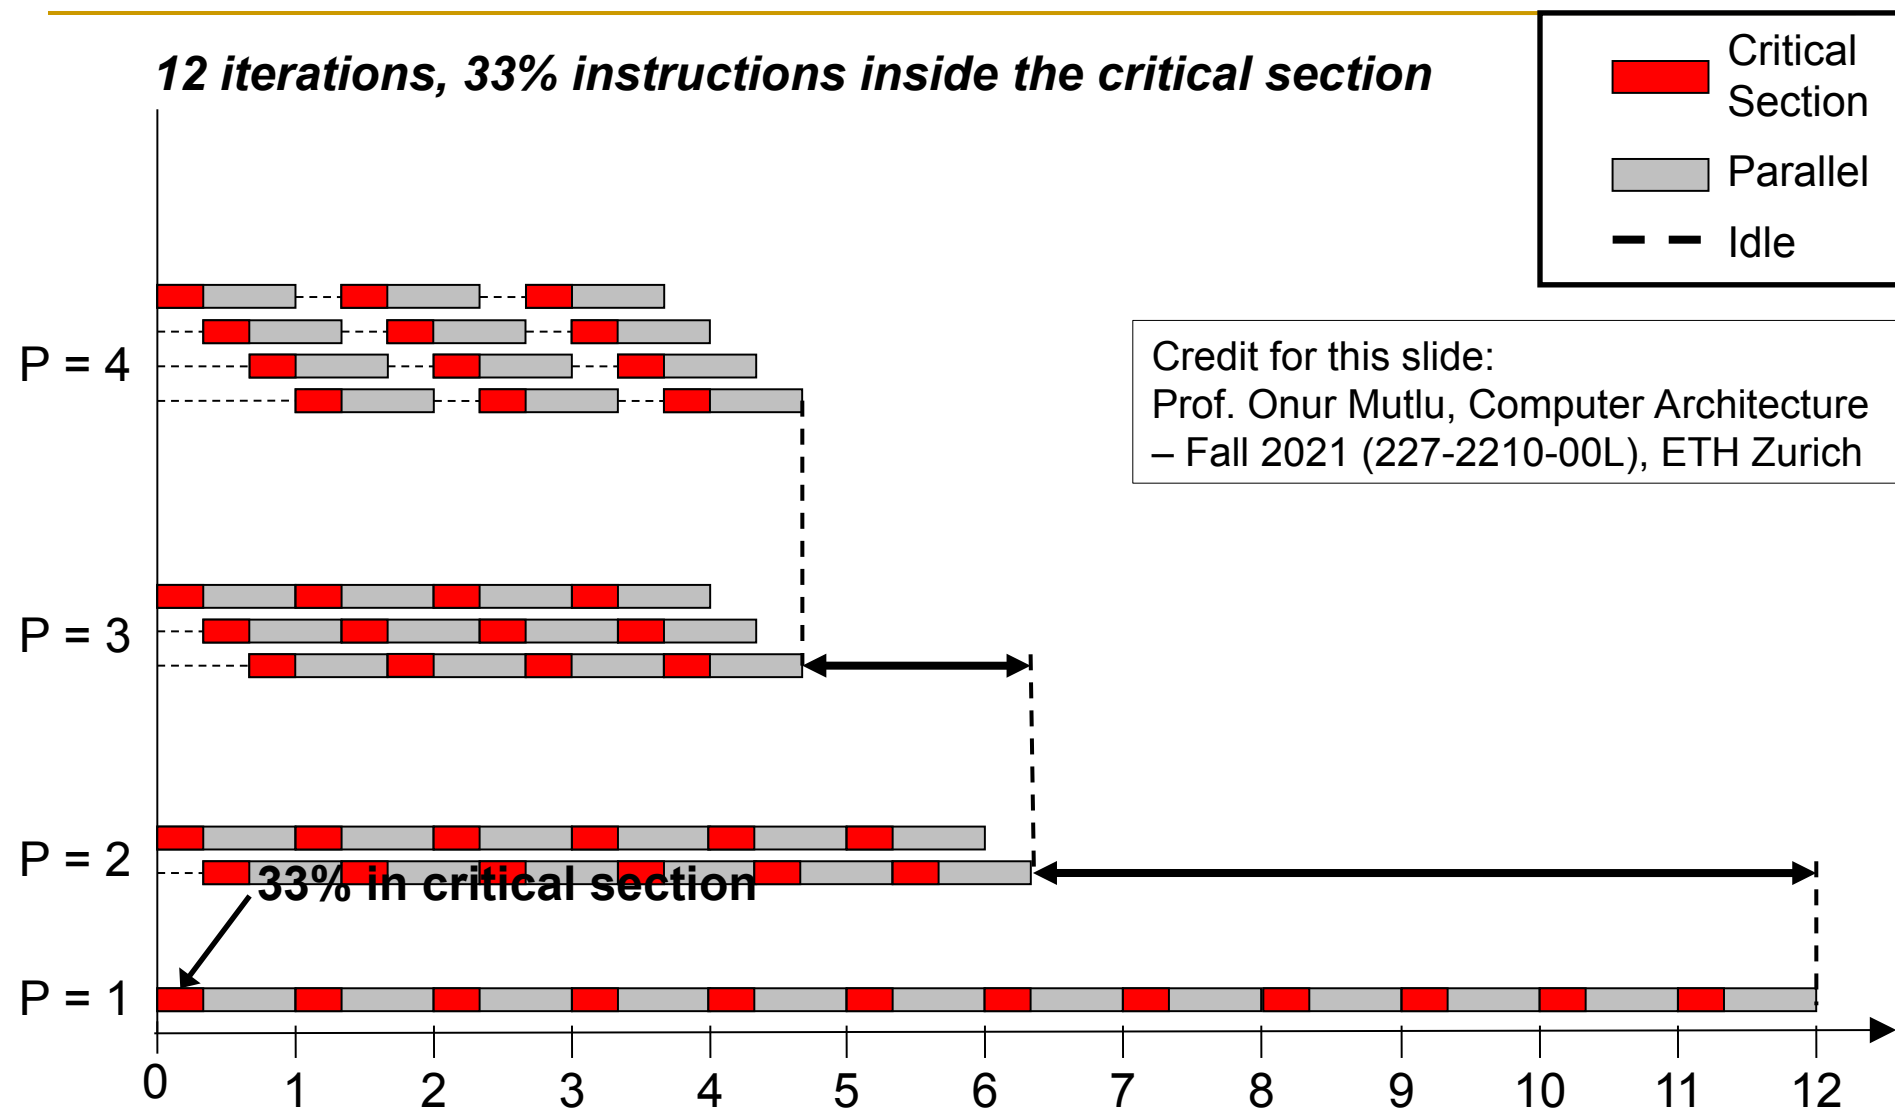
```
sum = sum + sum1;
```

```
sum = sum + sum2;
```

Sum will need to be shared and protected by a lock. Shared data structure is created:

Contention for Critical Sections

12 iterations, 33% instructions inside the critical section



Contention for Critical Sections

12 iterations, 33% instructions inside the critical section

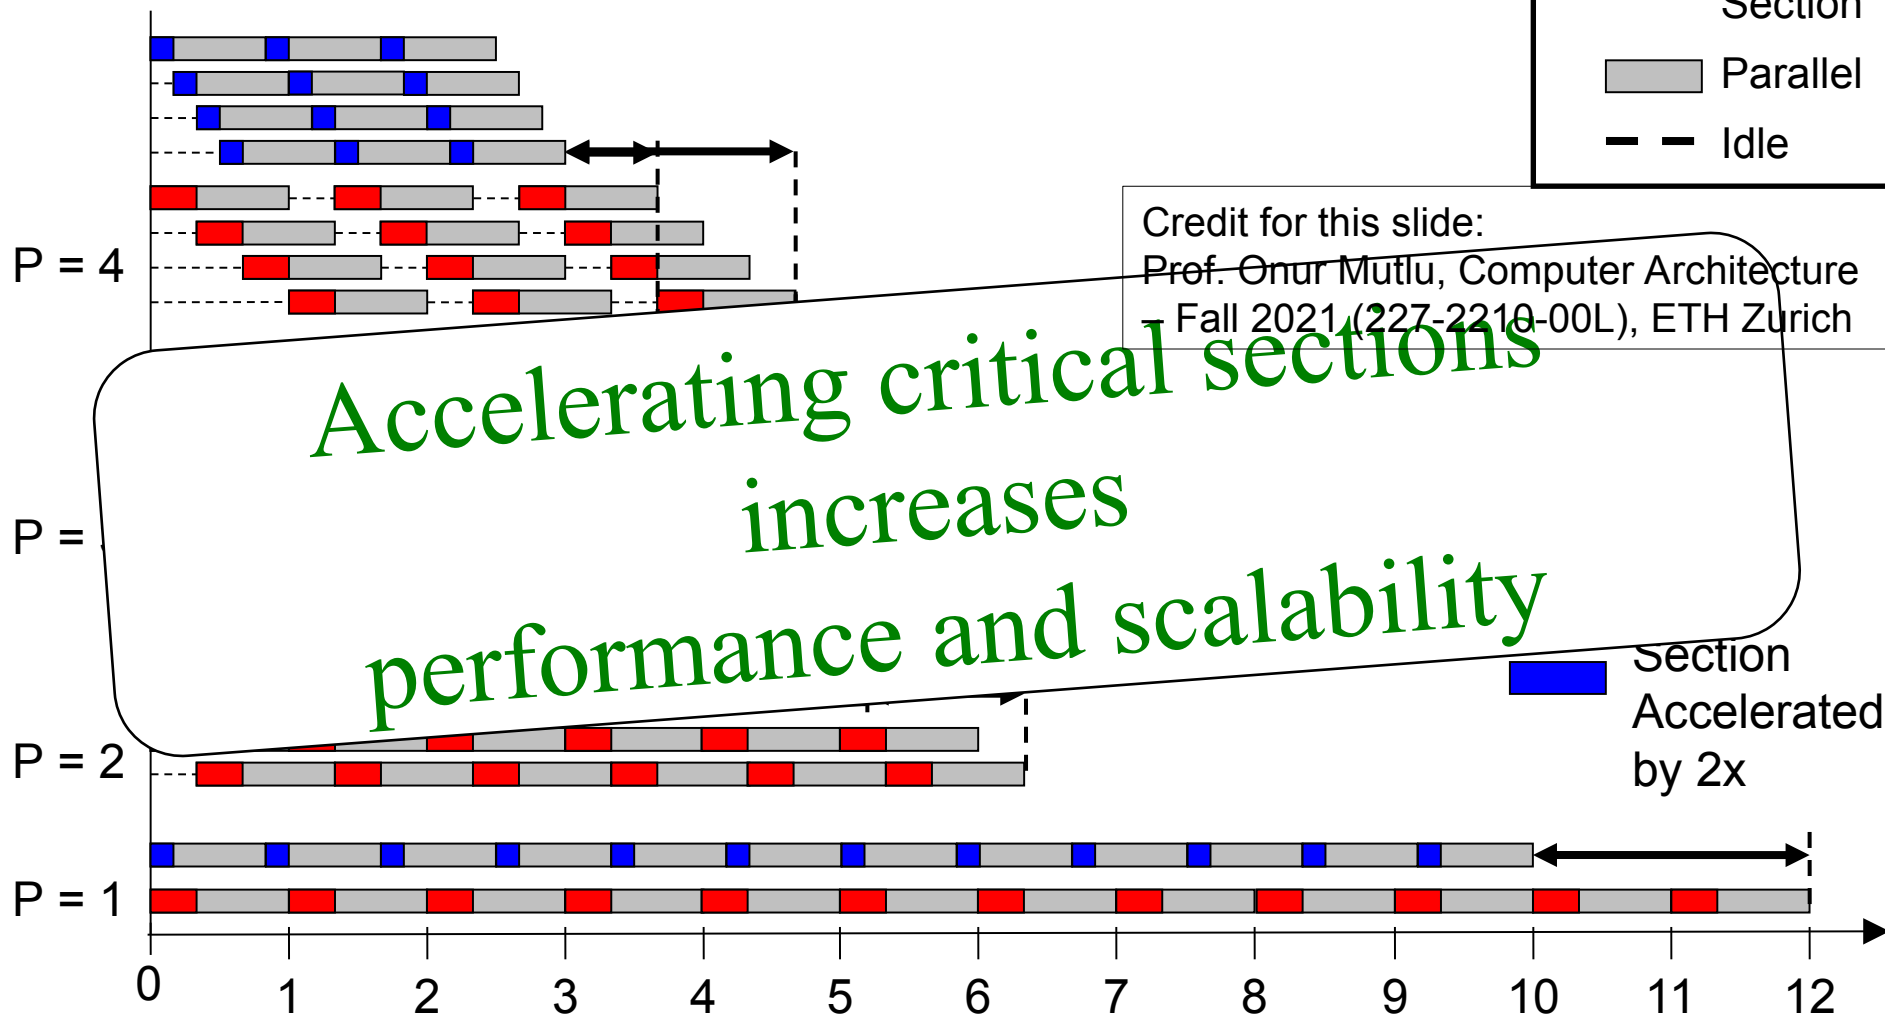
Critical Section
Parallel
Idle

Credit for this slide:

Prof. Onur Mutlu, Computer Architecture
Fall 2021 (227-2210-00L), ETH Zurich

Accelerating critical sections
increases
performance and scalability

Section Accelerated by 2x



עד כאן מצגת זו

Programming with Shared Memory

Part 2

Introduction to OpenMP

OpenMP

An accepted standard developed in the late 1990s by a group of industry specialists.

Consists of a small set of compiler directives, augmented with a small set of library routines and environment variables using the base language Fortran and C/C++.

Several OpenMP compilers available.

OpenMP

- Uses a thread-based shared memory programming model
- OpenMP programs will create multiple threads
- All threads have access to global memory
- Data can be shared among all threads or private to one thread
- Data transfer hidden from programmer
- Synchronization occurs but often implicit


OpenMP uses “**fork-join**” model but thread-based.

Initially, a single thread is executed by a **master thread**. **Parallel regions** (sections of code) can be executed by multiple threads (a team of threads).

parallel directive creates a team of threads with a specified block of code executed by the multiple threads in parallel. The exact number of threads in the team determined by one of several ways.

Other directives used within a **parallel** construct to specify parallel for loops and different blocks of code for threads.

Fork/join model



For C/C++, the OpenMP directives are contained in `#pragma` statements. The OpenMP `#pragma` statements have the format:

`#pragma omp directive_name ...`

where `omp` is an OpenMP keyword.

May be additional parameters (clauses) after the directive name for different options.

Some directives require code to specified in a structured block that follows the directive and then the directive and structured block form a “construct”.

Parallel Directive

```
#pragma omp parallel  
    structured_block
```

creates multiple threads, each one executing the specified structured_block, either a single statement or a compound statement created with { ... } with a single entry point and a single exit point.

There is an implicit barrier at the end of the construct.
The directive corresponds to forall construct.

Hello world example

OpenMP
directive for a
parallel region



```
#pragma omp parallel {
```

```
    printf("Hello World from thread = %d\n", omp_get_thread_num(),  
          omp_get_num_threads());  
}
```

From an 8-processor/core machine:

Hello World from thread 0 of 8
Hello World from thread 4 of 8
Hello World from thread 3 of 8
Hello World from thread 2 of 8
Hello World from thread 7 of 8
Hello World from thread 1 of 8
Hello World from thread 6 of 8
Hello World from thread 5 of 8

Private thread variables and shared variables

Could be declared within each parallel region but OpenMP provides **private** and **shared** clauses

```
int tid;
```

```
...
```

```
#pragma omp parallel private(tid) {  
    tid = omp_get_thread_num();  
    printf("Hello World from thread = %d\n", tid);  
}
```

Guy: demo

```
bash% export OMP_NUM_THREADS=8
```

```
csh% setenv OMP_NUM_THREADS 8
```

Folder:

```
.../lectures/08/code/HelloOpenMP
```

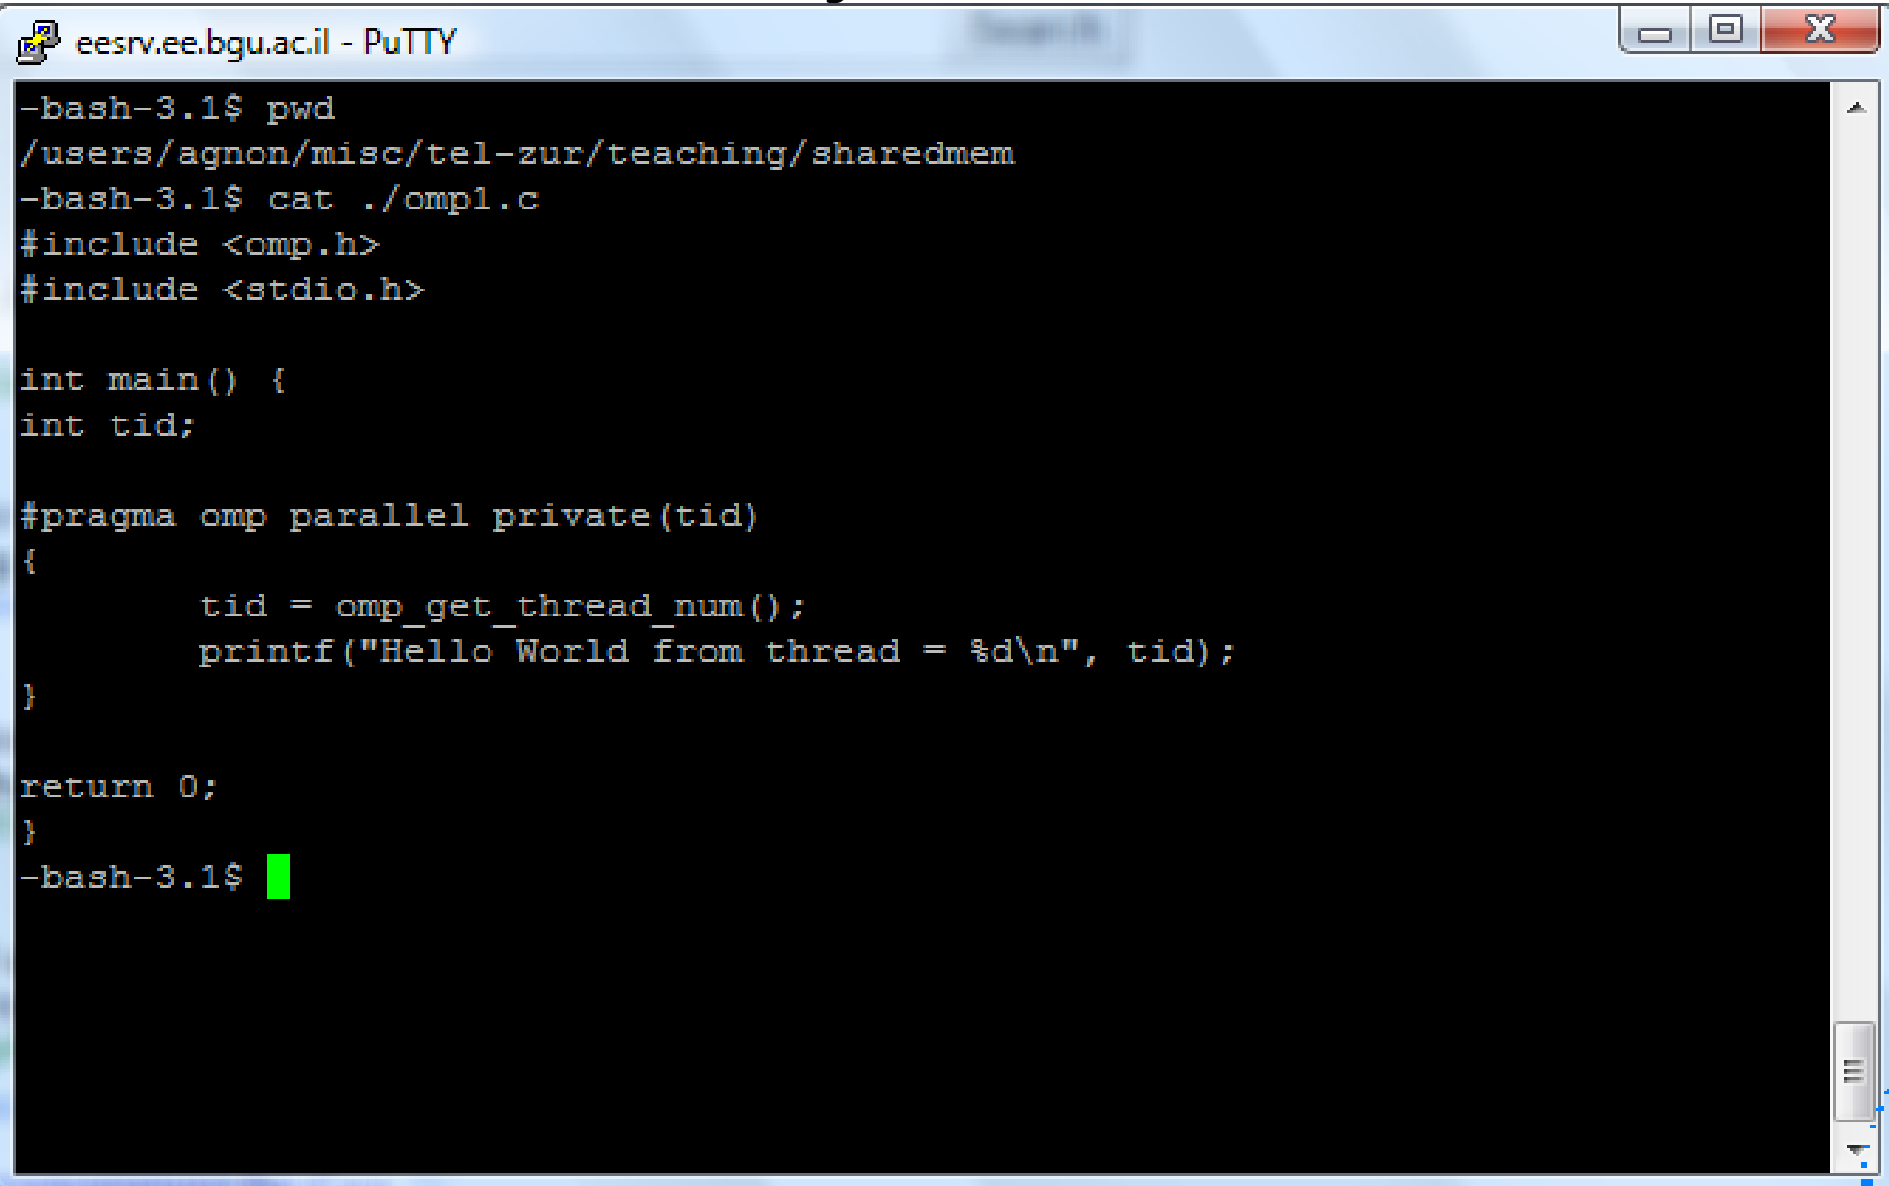
Program:

```
omp1.c
```

Compile:

```
gcc -fopenmp -o omp1 ./omp1.c
```


Guy: Demo



The image shows a PuTTY terminal window titled "eesrv.ee.bgu.ac.il - PuTTY". The terminal displays the following C code being cat'd from a file named "omp1.c":

```
-bash-3.1$ pwd
/users/agnon/misc/tel-zur/teaching/sharedmem
-bash-3.1$ cat ./omp1.c
#include <omp.h>
#include <stdio.h>

int main() {
int tid;

#pragma omp parallel private(tid)
{
    tid = omp_get_thread_num();
    printf("Hello World from thread = %d\n", tid);
}

return 0;
}
-bash-3.1$
```

The code is a C program that uses OpenMP to create a parallel region. Inside the parallel region, each thread prints its thread ID. The program is currently at the prompt, ready for execution.

Guy: demo

```
File Edit View Search Terminal Help
telzur@GL553VD ~/science/Teaching/PP/lectures/08/code/HelloOpenMP $ export OMP_NUM_THREADS=8
telzur@GL553VD ~/science/Teaching/PP/lectures/08/code/HelloOpenMP $ gcc -fopenmp -o omp1 ./omp1.c
telzur@GL553VD ~/science/Teaching/PP/lectures/08/code/HelloOpenMP $ ./omp1
Hi, now I am serial
Hello World from thread 0=
Hello World from thread 7=
Hello World from thread 5=
Hello World from thread 3=
Hello World from thread 1=
Hello World from thread 6=
Hello World from thread 2=
Hello World from thread 4=
Now I am serial again
telzur@GL553VD ~/science/Teaching/PP/lectures/08/code/HelloOpenMP $
```

Example

```
#pragma omp parallel private(x, num_threads)
{
    x = omp_get_thread_num();
    num_threads = omp_get_num_threads();
    a[x] = 10*num_threads;
}
```

Two library routines

`omp_get_num_threads()` returns number of threads that are currently being used in parallel directive

`omp_get_thread_num()` returns thread number (an integer from 0 to `omp_get_num_threads() - 1` where thread 0 is the master thread).

Array `a[]` is a global array, and `x` and `num_threads` are declared as private to the threads.

Number of threads in a team

Established by either:

1. `num_threads` clause after the `parallel` directive, or
 2. `omp_set_num_threads()` library routine being previously called,
or
 3. the environment variable `OMP_NUM_THREADS` is defined
- in the order given or is system dependent if none of the above.

Number of threads available can also be altered automatically to achieve best use of system resources by a “dynamic adjustment” mechanism.

Work-Sharing

Three constructs in this classification:

sections
for
single

In all cases, there is **an implicit barrier** at the end of the construct unless a **nowait** clause is included.

Sections

The construct

```
#pragma omp sections
{
    #pragma omp section
    structured_block
    #pragma omp section
    structured_block
    .
    .
    .
}
```

cause the structured blocks to be shared among threads in team.

`#pragma omp sections` precedes the set of structured blocks.

`#pragma omp section` prefixes each structured block.

The first `section` directive is optional.

Example

```
#pragma omp parallel shared(a,b,c,d,nthreads) private(i,tid) {  
    tid = omp_get_thread_num();
```

One
thread
does this

```
#pragma omp sections nowait {  
    #pragma omp section {  
        printf("Thread %d doing section 1\n",tid);  
        for (i=0; i<N; i++) {  
            c[i] = a[i] + b[i];  
            printf("Thread %d: c[%d]= %f\n",tid,i,c[i]);  
        }  
    }  
}
```

Another
thread
does this

```
#pragma omp section {  
    printf("Thread %d doing section 2\n",tid);  
    for (i=0; i<N; i++) {  
        d[i] = a[i] * b[i];  
        printf("Thread %d: d[%d]= %f\n",tid,i,d[i]);  
    }  
}  
} /* end of sections */  
} /* end of parallel section */
```

```

#include<stdio.h>
#include<omp.h>
int main() {
    int id;
    #pragma omp parallel
    {
        #pragma omp sections
        {
            #pragma omp section
            {
                printf ("Section 1.1 id = %d, \n", omp_get_thread_num());
            }
            #pragma omp section
            {
                printf ("Section 1.2 id = %d, \n", omp_get_thread_num());
            }
        }
        #pragma omp sections
        {
            #pragma omp section
            {
                printf ("Section 2.1 id = %d, \n", omp_get_thread_num());
            }
            #pragma omp section
            {
                printf ("Section 2.2 id = %d, \n", omp_get_thread_num());
            }
        }
    } // end pragma omp parallel
    return 0;
}

```

Demo: /home/telzur/Documents/Teaching/BGU/PP/PP2015A/lectures/08/code

Terminal



```
telzur@LIFEBOOK ~/Documents/Teaching/BGU/PP/PP2015A/lectures/08/code $ gcc -fopenmp -o omp_
sections_demo ./omp_sections_demo.c
telzur@LIFEBOOK ~/Documents/Teaching/BGU/PP/PP2015A/lectures/08/code $ ./omp_sections_demo
Section 1.1 id = 5,
Section 1.2 id = 7,
Section 2.1 id = 1,
Section 2.2 id = 6,
telzur@LIFEBOOK ~/Documents/Teaching/BGU/PP/PP2015A/lectures/08/code $ ./omp_sections_demo
Section 1.1 id = 6,
Section 1.2 id = 3,
Section 2.1 id = 1,
Section 2.2 id = 0,
telzur@LIFEBOOK ~/Documents/Teaching/BGU/PP/PP2015A/lectures/08/code $ ./omp_sections_demo
Section 1.2 id = 5,
Section 1.1 id = 4,
Section 2.2 id = 2,
Section 2.1 id = 0,
telzur@LIFEBOOK ~/Documents/Teaching/BGU/PP/PP2015A/lectures/08/code $
```

Terminal

```
$ export OMP_NUM_THREADS=2
$ ./omp_sections_demo
Section 1.1 id = 0,
Section 1.2 id = 1,
Section 2.1 id = 0,
Section 2.2 id = 1,
$ export OMP_NUM_THREADS=3
$ ./omp_sections_demo
Section 1.1 id = 0,
Section 1.2 id = 1,
Section 2.2 id = 1,
Section 2.1 id = 2,
$ export OMP_NUM_THREADS=4
$ ./omp_sections_demo
Section 1.1 id = 0,
Section 1.2 id = 3,
Section 2.1 id = 0,
Section 2.2 id = 1,
$ █
```

Intel Parallel Studio

- export
LD_LIBRARY_PATH=\$LD_LIBRARY_PATH:/opt/intel/composer_xe_2013_sp1.3.174/compiler/lib/intel64
- echo 0 | sudo tee
/proc/sys/kernel/yama/ptrace_scope
- /opt/intel/bin/icc -openmp -g -o
./omp_sections_demo_intel ./omp_sections_demo.c
- **Intel Inspector:**
/opt/intel/vtune_amplifier_xe/bin64/amplxe-gui &

Intel

/home/telzur/intel/inspxe/projects/OMP_sections_demo - Intel Inspector

File View Help

Project Navigator

/home/telzur/intel/i...

OMP_sections_demo

ome_sections_demo

pi_wrong_race

race

safecode

Welcome

New Inspector Result

Configure Analysis Type

Analysis Type

Memory Error Analysis

2x-20x

10x-40x

20x-80x

Detect Leaks

Detect Memory Problems

Locate Memory Problems

Memory Overhead

Detect Memory Problems

Copy

Medium scope memory error analysis type. Increases the load on the system and the time and resources required to perform analysis. Press F1 for more details.

☒ Detect memory leaks upon application exit

☒ Enable interactive memory growth detection

☒ Enable on-demand memory leak detection

☒ Report still-allocated memory at application exit

Stack frame depth: 8

☒ Analyze without debugger

Run an analysis and report all detected problems. Use to view correctness issues without stopping in the debugger to examine them.

☐ Enable debugger when problem detected

Run an analysis under the debugger and stop every time a problem is detected. Use to allow investigation of every problem detected.

Start

Stop

Close

Reset Growth Tracking

Measure Growth

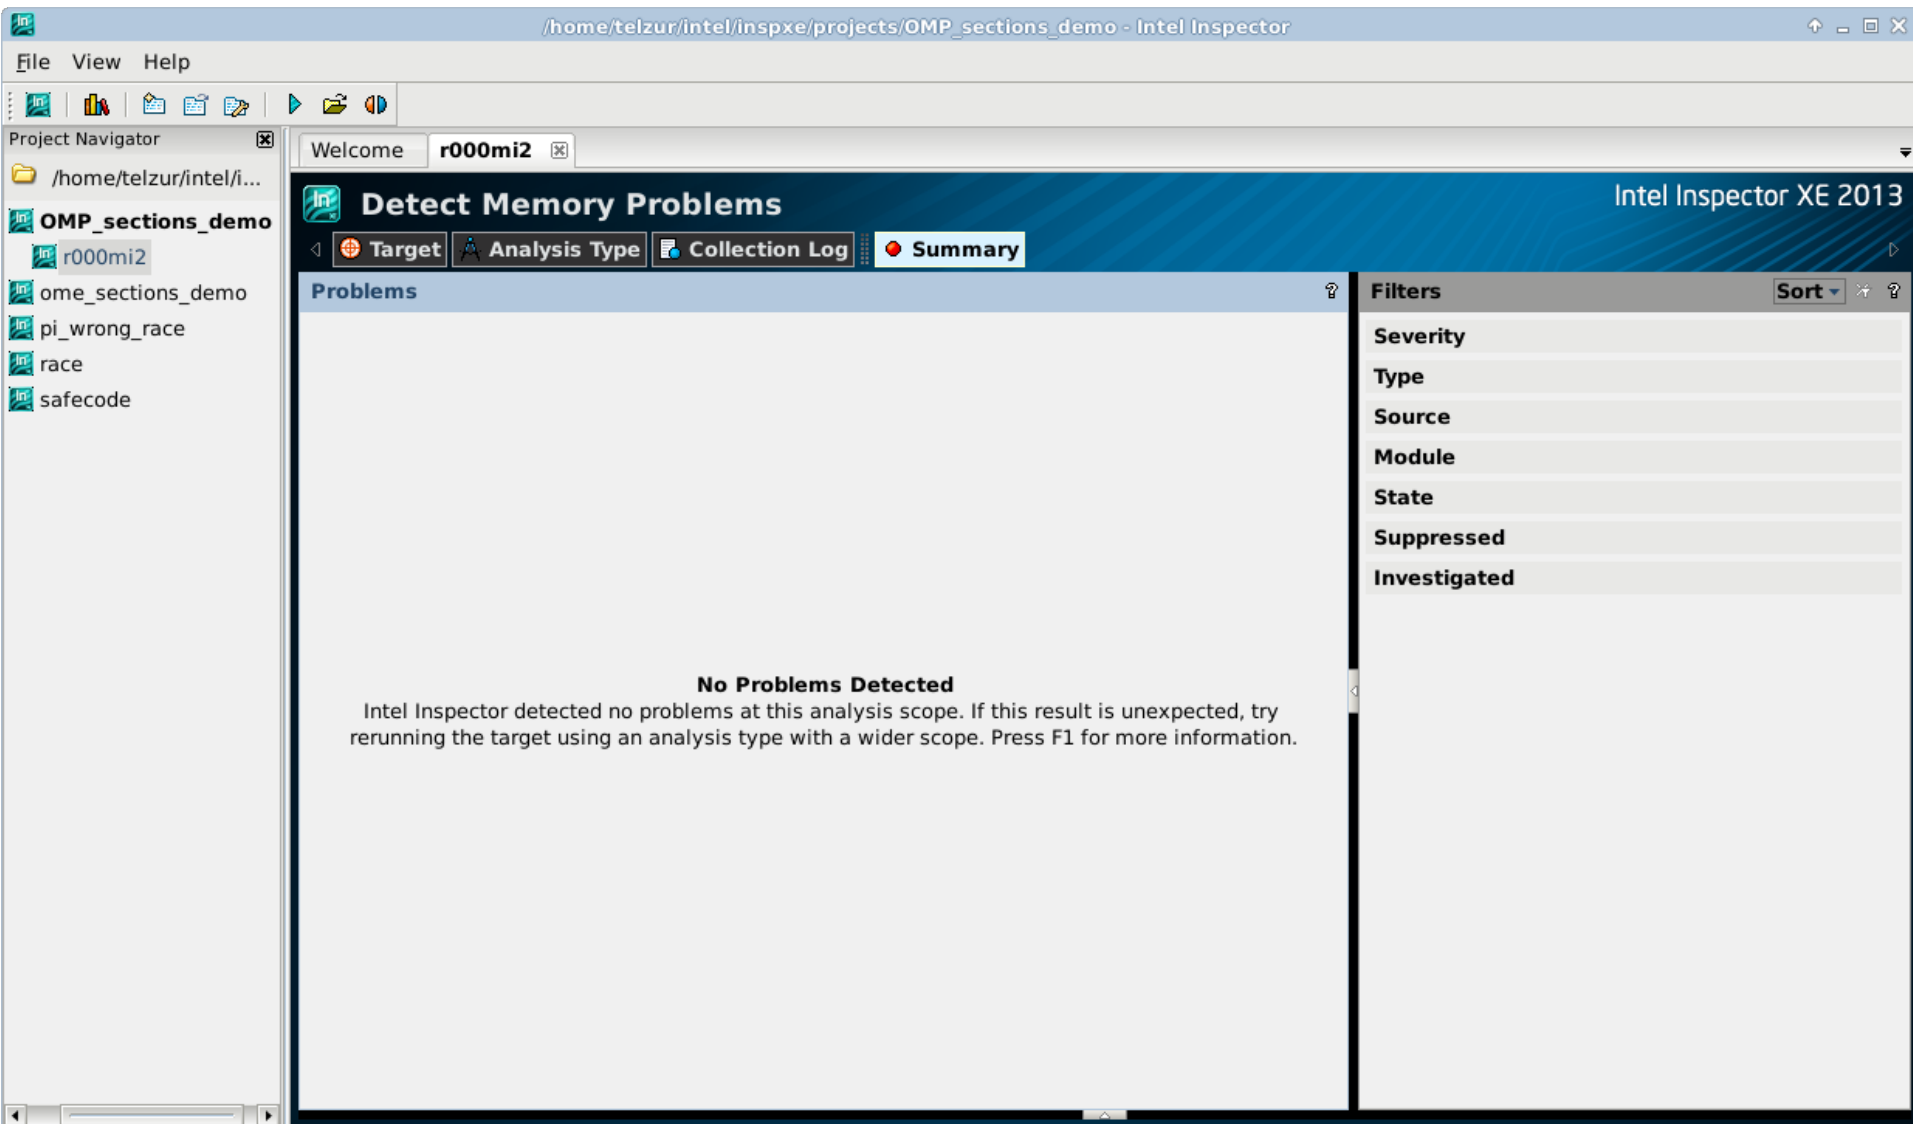
Reset Leak Tracking

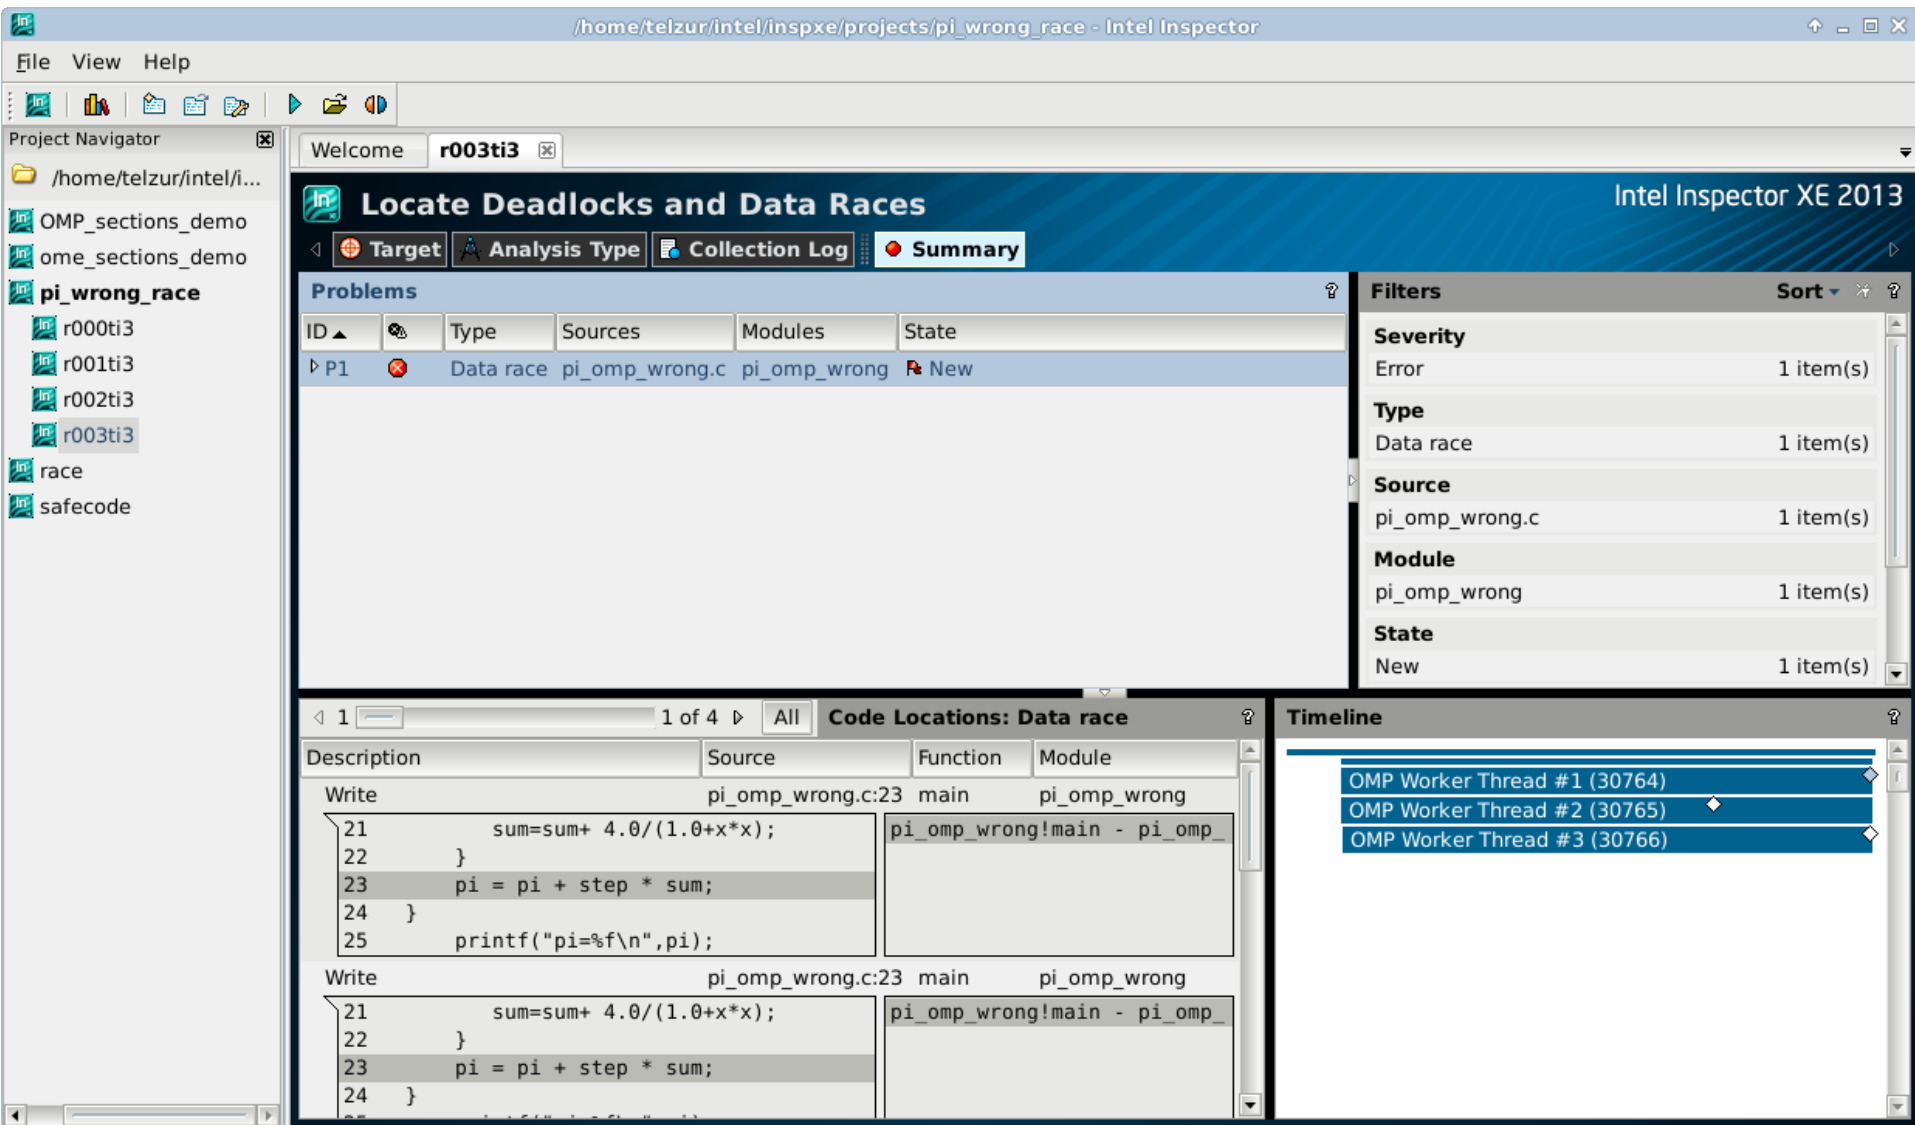
Find Leaks

Project Properties...

Command Line...

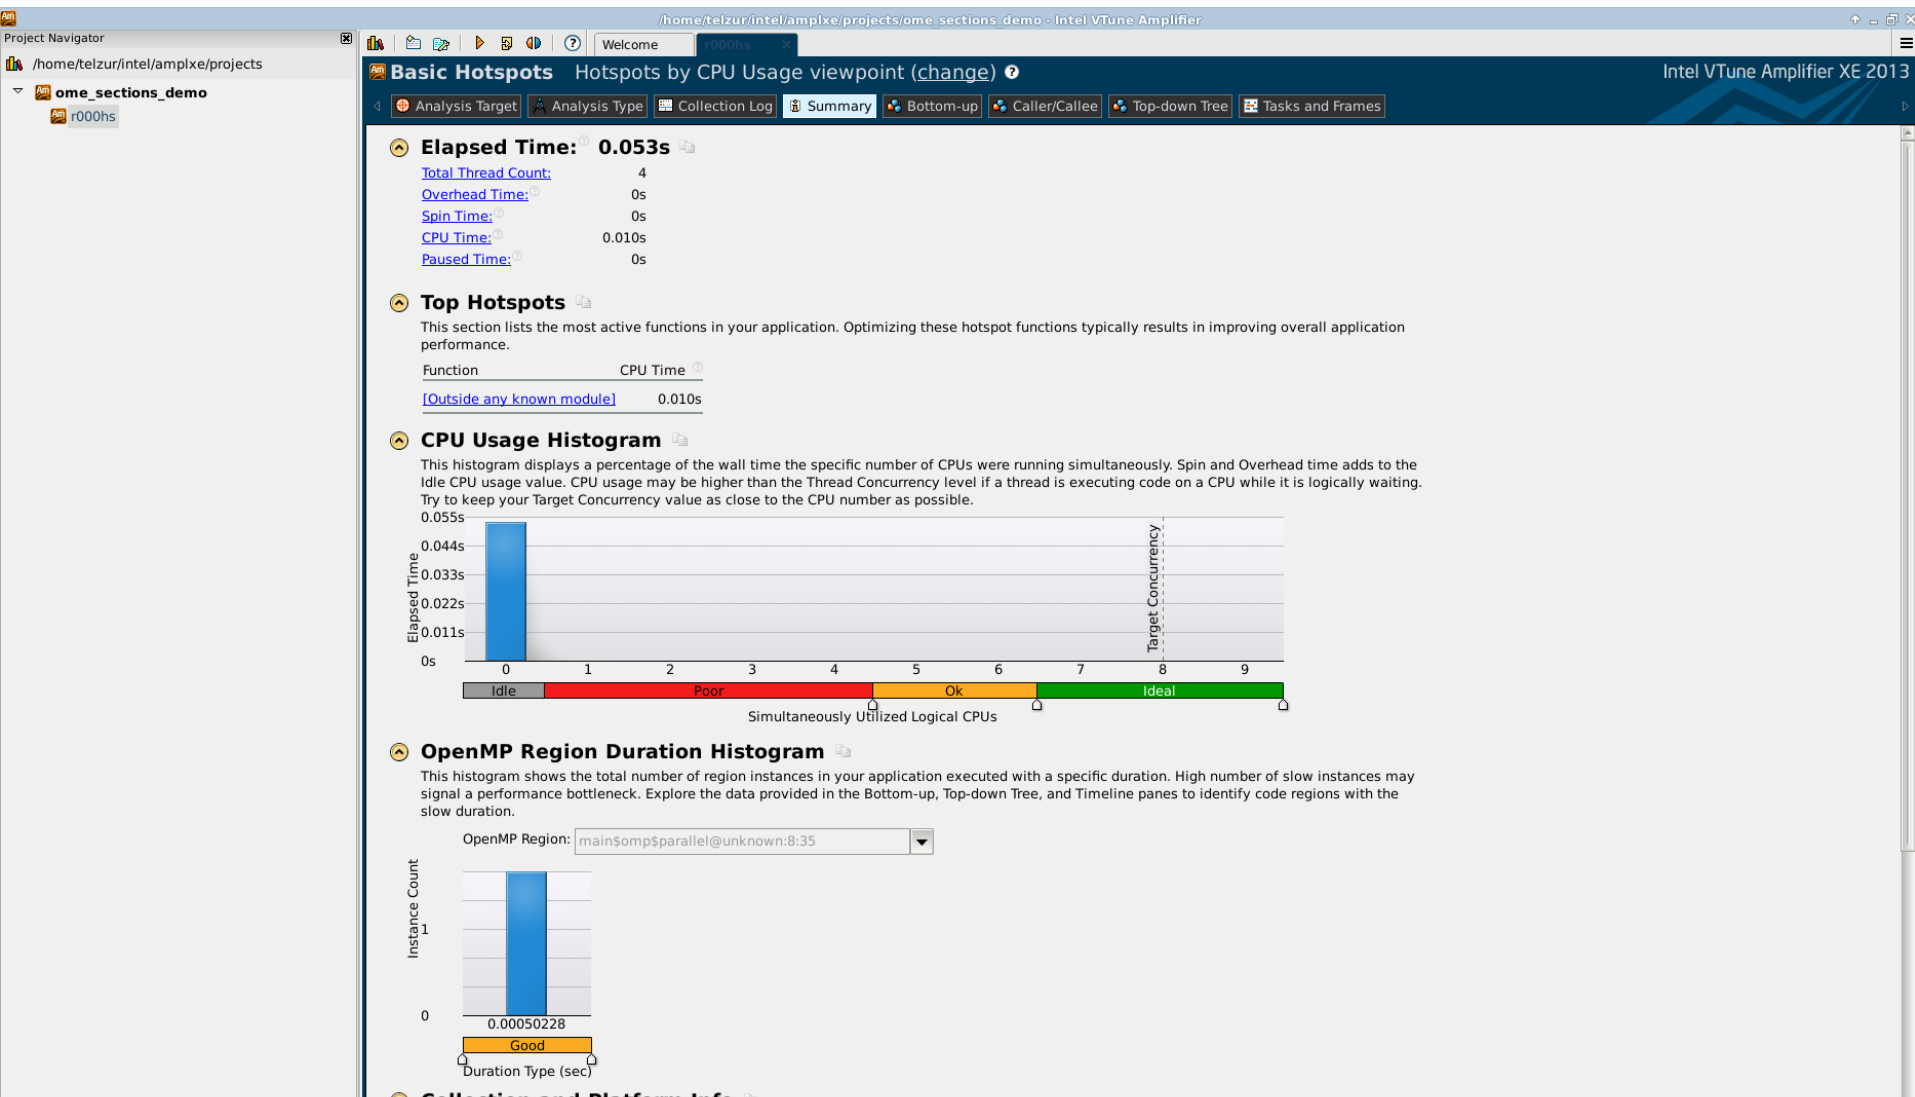
No race conditions were detected





Intel's Vtune

```
$ /opt/intel/vtune_amplifier_xe/bin64/amplxe-gui
```



The screenshot displays the Intel VTune Amplifier XE 2013.1.0.100 interface. The main window is titled 'Basic Hotspots' and shows 'Hotspots by CPU Usage viewpoint (change)'. The x-axis represents time in milliseconds (5ms to 50ms). The y-axis lists threads: omp_sections_de (TID: 32211), OMP Worker Thread #1 (TID: 32226), OMP Worker Thread #2 (TID: 32227), and OMP Worker Thread #3 (TID: 32228). The chart shows that omp_sections_de is the most active thread, running for approximately 45ms. The other threads show shorter durations. The right sidebar shows the 'Data Of Interest (CPU Metrics)' section with 'CPU Usage' selected, displaying '100.0% (0.010s of 0.010s)'. The bottom status bar indicates 'No filters are applied.' and 'Any Process'.

Single

The directive

```
#pragma omp single  
structured block
```

cause the structured block to be executed by one thread only.

Master Directive

The `master` directive:

```
#pragma omp master  
    structured_block
```

causes the master thread to execute the structured block.

Different to those in the work sharing group in that there is no implied barrier at the end of the construct (nor the beginning). Other threads encountering this directive will ignore it and the associated structured block, and will move on.

If a `parallel` directive is followed by a single `sections` directive, it can be combined into:

```
#pragma omp parallel sections {  
    #pragma omp section  
        structured_block  
    #pragma omp section  
        structured_block  
        .  
        .  
        .  
}
```

with similar effect.

(In both cases, the `nowait` clause is not allowed.)

For Loop

```
#pragma omp for  
for_loop
```

causes the for loop to be divided into parts and parts shared among threads in the team. The for loop must be of a simple form.

Way that for loop divided can be specified by an additional “schedule” clause.

Example: the clause `schedule (static, chunk_size)` cause for loop be divided into sizes specified by `chunk_size` and allocated to threads in a round robin fashion.

Example

```
#pragma omp parallel shared(a,b,c,nthreads,chunk) private(i,tid) {  
    tid = omp_get_thread_num();  
  
    if (tid == 0) {  
        nthreads = omp_get_num_threads();  
        printf("Number of threads = %d\n", nthreads);  
    }  
    printf("Thread %d starting...\n",tid);  
  
    #pragma omp for schedule(dynamic,chunk)  
    for (i=0; i<N; i++) {  
        c[i] = a[i] + b[i];  
        printf("Thread %d: c[%d]= %f\n",tid,i,c[i]);  
    }  
  
} /* end of parallel section */
```

Combined Parallel Work-sharing Constructs

If a `parallel` directive is followed by a single `for` directive, it can be combined into:

```
#pragma omp parallel for  
for_loop
```

with similar effects, i.e. it has the effect of each thread executing the same for loop.

Loop Scheduling and Partitioning

OpenMP offers scheduling clauses to add to for construct:

- Static

`#pragma omp parallel for schedule (static,chunk_size)`

Partitions loop iterations into equal sized chunks specified by `chunk_size`. Chunks assigned to threads in round robin fashion.

- Dynamic

`#pragma omp parallel for schedule (dynamic,chunk_size)`

Uses internal work queue. Chunk-sized block of loop assigned to threads as they become available.

- Guided

`#pragma omp parallel for schedule (guided,chunk_size)`

Similar to dynamic but chunk size starts large and gets smaller to reduce time threads have to go to work queue.

$$\text{chunk size} = \left\lceil \frac{\text{number of iterations remaining}}{2 * \text{number of threads}} \right\rceil$$

- Runtime

`#pragma omp parallel for schedule (runtime)`

Uses OMP_SCHEDULE environment variable to specify which of static, dynamic or guided should be used.

SCHEDULE: Describes how iterations of the loop are divided among the threads in the team. The default schedule is implementation dependent.

STATIC

Loop iterations are divided into pieces of size chunk and then statically assigned to threads. If chunk is not specified, the iterations are evenly (if possible) divided contiguously among the threads.

DYNAMIC

Loop iterations are divided into pieces of size chunk, and dynamically scheduled among the threads; when a thread finishes one chunk, it is dynamically assigned another. The default chunk size is 1.

GUIDED

Iterations are dynamically assigned to threads in blocks as threads request them until no blocks remain to be assigned. Similar to DYNAMIC except that the block size decreases each time a parcel of work is given to a thread. The size of the initial block is proportional to: $\text{number_of_iterations} / \text{number_of_threads}$ Subsequent blocks are proportional to $\text{number_of_iterations_remaining} / \text{number_of_threads}$ The chunk parameter defines the minimum block size. The default chunk size is 1.

RUNTIME

The scheduling decision is deferred until runtime by the environment variable OMP_SCHEDULE. It is illegal to specify a chunk size for this clause.

AUTO

The scheduling decision is delegated to the compiler and/or runtime system.

<https://computing.llnl.gov/tutorials/openMP/>

Example

```
#pragma omp parallel for schedule(kind [,chunk size])
```

specified, must be a positive integer.

Four different loop scheduling types (kinds) can be provided to OpenMP, as shown in the following table. The optional parameter (chunk), when

Kind	Description
static	Divide the loop into equal-sized chunks or as equal as possible in the case where the number of loop iterations is not evenly divisible by the number of threads multiplied by the chunk size. By default, chunk size is loop_count/number_of_threads.Set chunk to 1 to interleave the iterations.
dynamic	Use the internal work queue to give a chunk-sized block of loop iterations to each thread. When a thread is finished, it retrieves the next block of loop iterations from the top of the work queue. By default, the chunk size is 1. Be careful when using this scheduling type because of the extra overhead involved.
guided	Similar to dynamic scheduling, but the chunk size starts off large and decreases to better handle load imbalance between iterations. The optional chunk parameter specifies them minimum size chunk to use. By default the chunk size is approximately loop_count/number_of_threads.
auto	When schedule (auto) is specified, the decision regarding scheduling is delegated to the compiler. The programmer gives the compiler the freedom to choose any possible mapping of iterations to threads in the team.
runtime	Uses the OMP_schedule environment variable to specify which one of the three loop-scheduling types should be used. OMP_SCHEDULE is a string formatted exactly the same as would appear on the parallel construct.



Example: OpenMP

```
#pragma isat tuning
variable(@omp_schedule_type, [static, dynamic, guided])
variable(@omp_schedule_chunk, range(0, 1000, 100))
search(dependent)

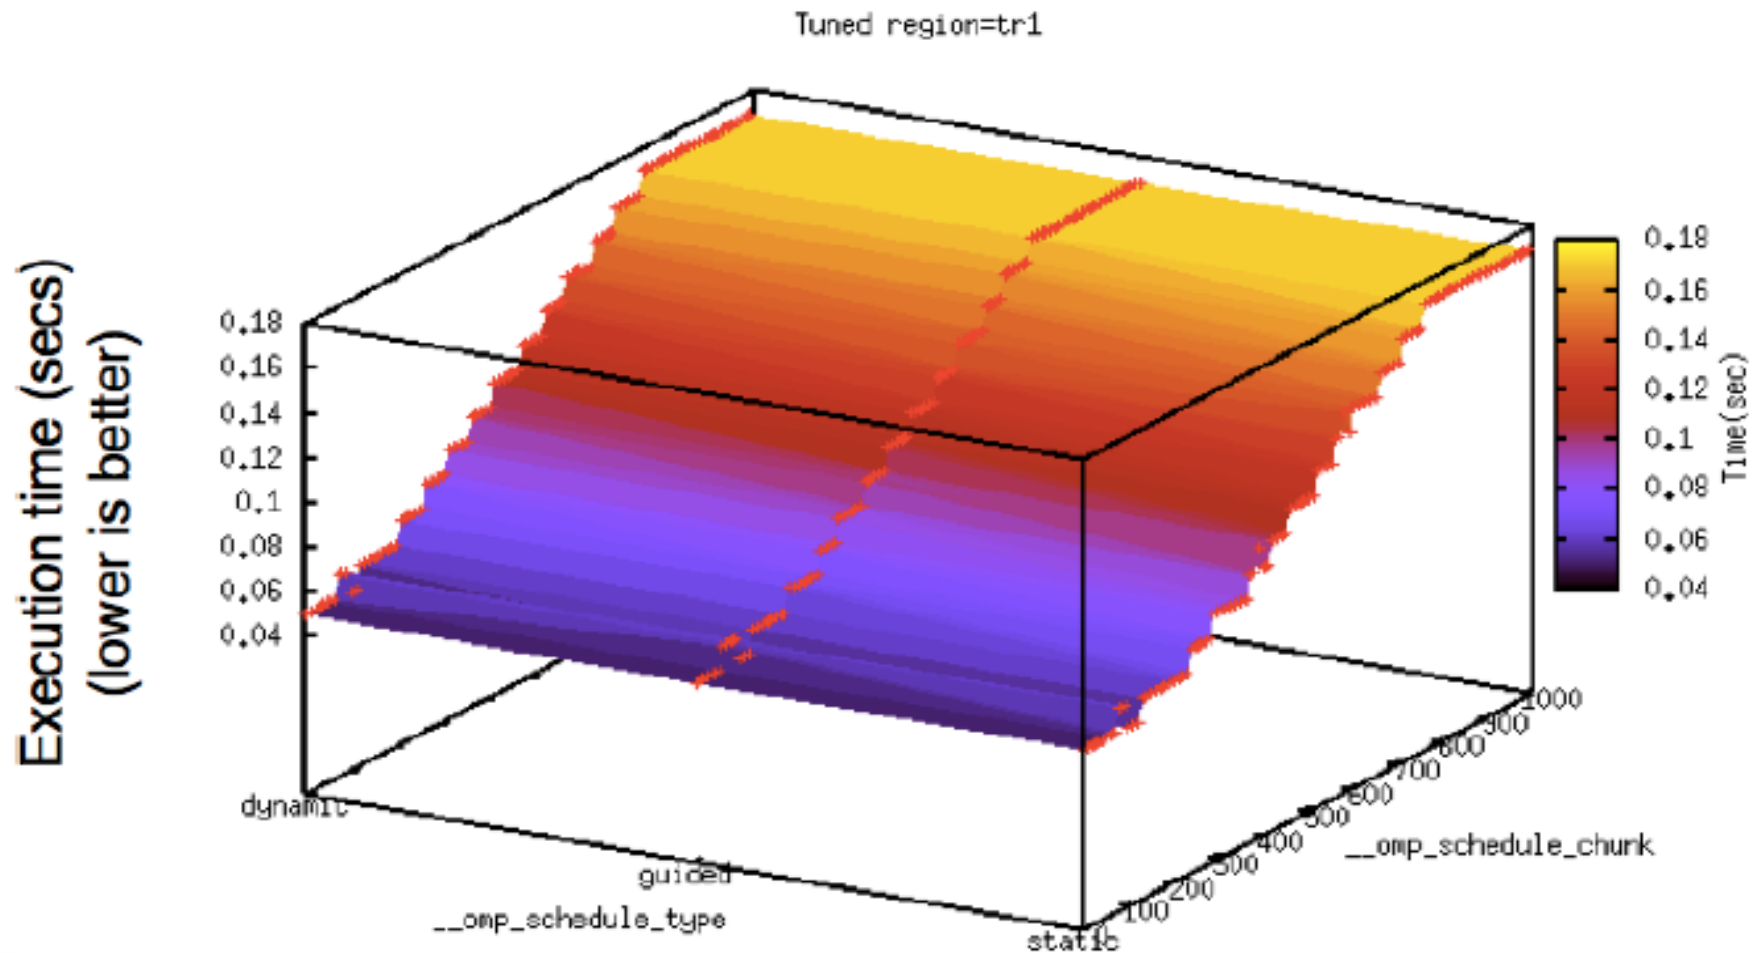
#pragma omp parallel for
for(i = 0; i < N; i++)
    C[i] = A[i] * B[i];
```

❖ Two tunable parameters: schedule type and chunk size



Visualization of Result

Tuning the schedule for `parallel_for`



Reduction clause

Used combined the result of the iterations into a single value c.f. with MPI _Reduce().


Can be used with parallel, for, and sections,

Example

```
sum = 0  
#pragma omp parallel for reduction(+:sum)  
for (k = 0; k < 100; k++ ) {  
    sum = sum + funct(k);  
}
```

Operation

Variable



Private copy of sum created for each thread by compiler.
Private copy will be added to sum at end.
Eliminates here the need for critical sections.

Private variables

private clause – creates private copies of variables for each thread

firstprivate clause - as private clause but initializes each copy to the values given immediately prior to parallel construct.

lastprivate clause – as private but “the value of each lastprivate variable from the sequentially last iteration of the associated loop, or the lexically last section directive, is assigned to the variable’s original object.”

Synchronization Constructs

Critical

The `critical` directive will only allow one thread execute the associated structured block. When one or more threads reach the `critical` directive:

```
#pragma omp critical name  
    structured_block
```

they will wait until no other thread is executing the same critical section (one with the same name), and then one thread will proceed to execute the structured block. name is optional. All critical sections with no name map to one undefined name.

Barrier

When a thread reaches the barrier

`#pragma omp barrier`

it waits until all threads have reached the barrier and then they all proceed together.

There are restrictions on the placement of barrier directive in a program. In particular, all threads must be able to reach the barrier.

Atomic

The atomic directive

```
#pragma omp atomic  
expression_statement
```

implements a critical section efficiently when the critical section simply updates a variable (adds one, subtracts one, or does some other simple arithmetic operation as defined by `expression_statement`).

ensures the serialisation of a particular operation and its much faster (less overhead)

More information

Full information on OpenMP at

<http://openmp.org/wp/>

Guy: **Dependency Analysis**

Next two slides taken from

“slides8d.ppt”

Bernstein's Conditions

Set of conditions sufficient to determine whether two processes can be executed simultaneously. Given:

I_i is the set of memory locations read (input) by process P_i .

O_j is the set of memory locations written (output) by process P_j .

For two processes P_1 and P_2 to be executed simultaneously, inputs to process P_1 must not be part of outputs of P_2 , and inputs of P_2 must not be part of outputs of P_1 ; i.e.,

$$I_1 \cap O_2 = \phi$$

$$I_2 \cap O_1 = \phi$$

where ϕ is an empty set. Set of outputs of each process must also be different; i.e.,

$$O_1 \cap O_2 = \phi$$

If the three conditions are all satisfied, the two processes can be executed concurrently.

Example

Suppose the two statements are (in C)

`a = x + y;`

`b = x + z;`

We have

$I_1 = (x, y)$

$O_1 = (a)$

$I_2 = (x, z)$

$O_2 = (b)$

and the conditions

$I_1 \cap O_2 = \phi$

$I_2 \cap O_1 = \phi$

$O_1 \cap O_2 = \phi$

are satisfied. Hence, the statements `a = x + y` and `b = x + z` can be executed simultaneously.